

# Installation Guide



## **Installation Guide**

Publication Date: 01/10/2018

SUSE LLC 10 Canal Park Drive Suite 200 Cambridge MA 02141 USA

https://www.suse.com/documentation ▶

# Contents

## Installation Overview xiii

1 For More Information xiv

4.1

4.2

Introduction 18

Install Phase Configuration 18

- 1	PRE-INSTALLATION 1
1	Overview 2
2	Pre-Installation Checklist 3
2.1	BIOS and iLO Settings 3
2.2	Network Setup and Configuration 4
2.3	Lifecycle Manager 7
2.4	Information for the nic_mappings.yml Input File 8
2.5	Control Plane 9
2.6	Compute Hosts 10
2.7	VSA Hosts 11
2.8	Additional Comments 12
2.9	SLES Pre-Installation Checks 12 Introduction 12 • Sample iptables rules must be removed on SLES 12
3	Using Git for Configuration Management 15
3.1	Initialization on a new deployment 15
3.2	Updating any configuration, including the default configuration 16
4	Boot from SAN and Multipath Configuration 18

iii Installation Guide

4.3	QLogic FCoE restrictions and additional configurations 19	
4.4	Red Hat Compute Host for FCoE 22 Installing the HPE Helion OpenStack 5.0 iso for nodes that support Boot from SAN 24	
5	Installing the L2 Gateway Agent for the Networking Service 32	
5.1	Introduction 32	
5.2	Sample network topology (for illustration purposes) 33	
5.3	Networks 33	
5.4	HPE 5930 switch configuration 35	
5.5	Configuring the provider data path network 36	
5.6	Enabling and configuring the L2 gateway agent 40	
5.7	Routing between software and hardware 42	
5.8	Connecting a baremetal server to the HPE 5930 switch 44	
5.9	Configuration on a baremetal server 45 NIC bonding and IRF configuration 45	
5.10	Scale numbers tested 46	
5.11	L2 gateway commands 46	
Ш	CLOUD INSTALLATION 48	
6	Overview 49	
7	Installing via the GUI 51	
7.1	Before you begin 52	
7.2	Creating a CSV file for import 53	
7.3	Run the GUI to install your cloud 54	
7 4	Use the installer securely 55	

**iv** Installation Guide

8	DNS Service Installation Overview 57		
8.1	Getting Started 57		
8.2	Install the DNS Service with PowerDNS 57 Install DNS Service with PowerDNS 57 • Configure the Backend 58		
8.3	Install the DNS Service with BIND 59 Install DNS Service with BIND 59 • Configure the Backend 59		
8.4	Install the DNS Service with InfoBlox 60 Prerequisites 61 • Configure the Backend 61		
8.5	Configure DNS Domain and NS Records 65		
9	Magnum Overview 67		
9.1	Magnum Architecture 67		
9.2	Install the Magnum Service 83 Installing Magnum as part of new HPE Helion OpenStack 5.0 environment 84 • Adding Magnum to existing HPE Helion OpenStack 5.0 environment 84		
9.3	Integrate Magnum with the DNS Service 86		
10	Installing Mid-scale and Entry-scale KVM 89		
10.1	Important Notes 89		
10.2	Before You Start 90		
10.3	Setting Up the Lifecycle Manager 90		
10.4	Configuring Your Environment 93		
10.5	Provisioning Your Baremetal Nodes 95		
10.6	Running the Configuration Processor 98		
10.7	Configuring TLS 98		
10.8	Deploying the Cloud 99		

7.5

Final Steps 56

Installation Guide

11	Installation for Helion Entry Scale ESX (with OVSvApp, KVM with VSA Model) 101
11.1	Important Notes 101
11.2	Before You Start 102
11.3	Prerequisites 102
11.4	Deploy ESX Cloud with OVSvApp 104
11.5	Procedure to Deploy ESX Cloud with OVSvApp 104  Setting Up the Lifecycle Manager 104 • Prepare and Deploy Cloud  Controllers 107 • Provisioning Your Baremetal Nodes 109 • Running the  Configuration Processor 111 • Deploying the Cloud 112
11.6	Prepare and Deploy ESX Computes and OVSvAPPs 113  Validate the block storage 119 • Validate the compute 120 • Validate the neutron 121
12	Sample activationtemplate.json File for ESX Compute 123
13	Installing Baremetal (Ironic) 140
13.1	Installation for HPE Helion Entry-scale Cloud with Ironic Flat Network 140
	Before You Start 140 • Setting Up the Lifecycle Manager 140 • Configure Your Environment 144 • Provisioning Your Baremetal Nodes 145 • Running the Configuration Processor 147 • Deploying the Cloud 148 • Ironic configuration 149 • Node Configuration 150 • TLS Certificates with Ironic Python Agent (IPA) Images 151
13.2	Ironic in Multiple Control Plane 154  Networking for Baremetal in Multiple Control Plane 154 • Handling Optional  Swift Service 155 • Instance Provisioning 155

Configuring a Block Storage Backend (Optional) 100

Post-Installation Verification and Administration 100

10.9

10.10

vi Installation Guide

- 13.3 Provisioning Baremetal Nodes with Flat Network Model 156
  Supplied Images 157 Provisioning a node 157 Creating a node using agent\_ilo 158 Creating a node using agent\_ipmi 160 Create
  Flavor 163 Create network port 164 Create Glance
  Image 164 Generate Key Pair 165 Determine the neuron network
  ID 165 Boot the node 166 Create Glance Images using diskimage-builder (DIB) 168 Creating BIOS images for RHEL 168 Creating BIOS images for Ubuntu 168 Creating UEFI images 168
- 13.4 Provisioning Baremetal Nodes with Multi-Tenancy 170
- 13.5 View Ironic System Details 178

View details about the server using nova show <nova-node-id> 178 · View detailed information about a node using ironic node-show <ironic-node-id> 180 · View detailed information about a port using ironic port-show <ironic-port-id> 182 · View detailed information about a hypervisor using nova hypervisor-list and nova hypervisor-show 182 · View a list of all running services using nova service-list 183

13.6 Troubleshooting Ironic Installation 184

No valid host was found. There are not enough hosts available. 184 • Deployment to a node fails and in "ironic nodelist" command, the power\_state column for the node is shown as "None" 191 • Error Downloading Image 192 • Using node-inspection can cause temporary claim of IP addresses 192 • Node permanently stuck in deploying state 193 • The NICs in the baremetal node should come first in boot order 193 • Increase in the number of nodes can cause power commands to fail 193 • DHCP succeeds with PXE but times out with iPXE 194

13.7 Node Cleaning 196

Setup 196 • In use 197 • Troubleshooting 199 • Disabling Node Cleaning 199

13.8 Ironic and OneView 200

Enabling Ironic OneView driver in HPE Helion OpenStack 200 • Adding OneView Appliance Credentials 200 • Encrypting the OneView Password 200 • Decrypting the OneView Password 201 • Registering

vii Installation Guide

Properties 201 • Creating Port for Driver 201 • Creating a Node 201 • Getting Data using REST API 202 • Ironic OneView CLI 202 13.9 RAID Configuration for Ironic 203 13.10 Audit Support for Ironic 211 API Audit Logging 211 • Enabling API Audit Logging 212 • Sample Audit Event 212 Installation for HPE Helion Entry-scale Cloud with Swift 14 **Only 214** 14.1 Important Notes 214 14.2 Before You Start 215 14.3 Setting Up the Lifecycle Manager 215 14.4 Configure Your Environment 218 14.5 Running the Configuration Processor 220 14.6 Provisioning Your Baremetal Nodes 221 14.7 Deploying the Cloud 223 14.8 Post-Installation Verification and Administration 224 **Installing SLES Compute 225** 15 15.1 SLES Compute Node Installation Overview 225 15.2 SLES Support 225 15.3 Using the Lifecycle Manager to Deploy SLES Compute Nodes 226 Deploying legacy BIOS SLES compute nodes 226 15.4 Provisioning SLES Yourself 227 Introduction 227 • Configure Lifecycle Manager to Enable SLES 227 • Install SLES 12 SP2 228 • Assign a static IP 228 • Add stack user and home directory 229 • Allow user stack to sudo without password 229 • Add zypper repository 229 • Add Required Packages 230 • Set up passwordless SSH access 230

Baremetal Node for OneView Driver 201 • Updating Node

viii Installation Guide

HLM-Hypervisor instructions 232
Introduction 232
Model changes 232 passthrough-network-groups 232 • hlm-hypervisor 234
Bootstrap Instructions 241
Customizing for VCP 241 Installing a fully input model managed Virtual Control Plane based HPE Helion OpenStack Cloud 242 • Provisioning just the Virtual Control Plane VMs 242 • Manually running each phase of the Virtual Control Plane provisioning 242
NIC Mapping for HLM Virtual-Controllers 243 Recommended Mappings 244
Notification of actions to be taken post upgrade 245
Reboot of a HLM-Hypervisor node 245
Staged install on HLM-Hypervisor individual play books. 246
Virtual Controller Replacement 247
Ansible host-specific variables for network-group access 248  Example 3rd-Party Ansible for Network Configuration Access 252
Ansible host-specific variables for gluster 253
"Empty" HLM Hypervisor Node 253
Monasca Monitoring of HLM Hypervisors 254  Monitored Metrics 254 • Configured Alarms 255 • HLM Hypervisor  Monitoring Configuration 255
Integrations 257
Ceph Overview 257  Overview 257 • Deployment Architecture 259 • Ceph  Networking 259 • Placement of service component 261 • Ceph

Using SLES as a Ceph Client 231

15.5

ix Installation Guide

Deployment Architecture 262 • Alternative supported architecture 263 • Core networking 264 • Placement of service components 264 • Hardware recommendations 265 • Ceph Deployment and Configurations 265 • Alternative Supported Choices 281 • Usage of Ceph Storage 298 • Managing Ceph Clusters After Deployment 318

- 17.2 Configuring for VSA Block Storage Backend 319
   Prerequisites 319 Notes 319 Configure HPE StoreVirtual
   VSA 320 Using Ansible 321 Using the CMC Utility 322 Configure
   VSA as the Backend 337 Post-Installation Tasks 341
- 17.3 Configuring for 3PAR Block Storage Backend 341
   Prerequisites 342 Notes 342 Multipath Support 343 Configure
   3PAR FC as a Cinder Backend 344 Configure 3PAR iSCSI as Cinder backend 346 Post-Installation Tasks 348
- 17.4 Ironic OneView Integration 349
   Prerequisites 349 Integrating with OneView 349 Registering Node in OneView 351 Provisioning Ironic Node 352
- 18 Troubleshooting the Installation 361
- 18.1 Issues during Lifecycle-manager Setup 361
- 18.2 Issues while Provisioning your Baremetal Nodes 361
- 18.3 Issues while Updating Configuration Files 368
- 18.4 Issues while Deploying the Cloud 370
  - 19 Troubleshooting the Block Storage Backend Configuration 376
- 20 Troubleshooting the ESX 378
- 20.1 Issue: If hlm-ux-services.service is not running, the EON resource-activate and resource-deactivate commands fails 378
- 20.2 Issue: ESX onboard Compute 378
- 20.3 Issue: Migrate eon-conductor 379

x Installation Guide

20.4 Issue: ESX Cluster shows UNKNOWN in OpsConsole 382 20.5 Issue: Unable to view the VM console in Horizon UI 383 Ш POST-INSTALLATION 385 21 Overview 386 Cloud Verification 387 22 22.1 API Verification 387 Prerequisites 387 • Tempest Integration Tests 388 • Running the Tests 388 • Viewing Test Results 389 • Customizing the Test Run 389 • Run Tests for Specific Services and Exclude Specific Features 390 • Run Tests Matching a Series of White and Blacklists 390 23 UI Verification 393 23.1 Verifying Your Block Storage Backend 393 Create a Volume 393 • Attach Volume to an Instance 395 • Detach Volume from Instance 395 • Delete Volume 396 • Verifying Your Object Storage (Swift) 396 23.2 Verify the Object Storage (Swift) Operations 398 23.3 Uploading an Image for Use 398 Running the Playbook 399 • How to Curate Your Own Images 399 • Using the GlanceClient CLI to Create Images 399 23.4 Creating an External Network 400 Notes 400 • Using the Ansible Playbook 400 • Using the NeutronClient CLI 401 • Next Steps 402 **Installing OpenStack Clients 403** 24 25 Configuring Transport Layer Security (TLS) 406 25.1 Configuring TLS in the input model 407 25.2 User-provided certificates and trust chains 410 25.3 Edit the input model to include your certificate files 410

xi Installation Guide

25.4	Generate a self-signed CA 411		
25.5	Generate a certificate signing request 414		
25.6	Generate a server certificate 415		
25.7	Upload to the lifecycle manager 418		
25.8	Configuring the cipher suite 419		
25.9	Testing 419		
25.10	Verifying that the trust chain is correctly deployed 420		
25.11	Turning TLS on or off 420		
26	Configuring Availability Zones 422		
27	Configuring Load Balancer as a Service 423		
27.1	Prerequisites 424		
27.2	Octavia Load Balancing Provider 424		
27.3	Setup of prerequisites 425		
27.4	Create Load Balancers 440		
27.5	Create Floating IPs for Load Balancer 444		
27.6	Testing the Octavia Load Balancer 445		
28	Other Common Post-Installation Tasks 448		
28.1	Determining Your User Credentials 448		
28.2	Configure your lifecycle manager to use the command-line tools 448		
28.3	Protect home directory 449		
28.4	Back up Your SSH Keys 449		
28.5	Retrieving Service Endpoints 449		
28.6	Other Common Post-Installation Tasks 450		

**xii** Installation Guide

# **Installation Overview**

Before jumping into your installation, we recommend taking the time to read through our documentation to get an overview of the sample configurations HPE Helion OpenStack 5.0 offers. We have highly tuned example configurations for each of these Cloud models:

Name	Location
	~/helion/examples/entry-scale-kvm-
	vsa
	~/helion/examples/entry-scale-kvm-
	<u>vsa-mml</u>
	~/helion/examples/entry-scale-kvm-
	<u>ceph</u>
	~/helion/examples/mid-scale-kvm-vsa
	~/helion/examples/entry-scale-esx-
	kvm-vsa
	~/helion/examples/entry-scale-esx-
	kvm-vsa-mml
	~/helion/examples/entry-scale-swift
	~/helion/examples/entry-scale-
	<u>ironic-flat-network</u>
	~/helion/examples/entry-scale-
	<u>ironic-multi-tenancy</u>

- •
- •
- •
- •
- •

## **Using the Command-line**

#### You should use the command-line if:

- You are installing a more complex or large-scale cloud.
- You have more than 15 nodes.
- You need to use availability zones or the server groups functionality of the cloud model.
   See the for more information.
- You want to customize the cloud configuration beyond the tuned defaults that HPE provides out of the box.
- You need deeper customizations than are possible to express using the GUI.

## Instructions for installing via the command-line are here:

- Chapter 10, Installing Mid-scale and Entry-scale KVM
- Chapter 11, Installation for Helion Entry Scale ESX (with OVSvApp, KVM with VSA Model)

## 1 For More Information

Other useful documents that will help you understand, plan, configure, and install your cloud are listed below.

- Chapter 3, Using Git for Configuration Management
- Chapter 18, Troubleshooting the Installation
- Chapter 22, Cloud Verification
- Chapter 28, Other Common Post-Installation Tasks

xiv For More Information

# I Pre-Installation

- 1 Overview 2
- 2 Pre-Installation Checklist 3
- 3 Using Git for Configuration Management 15
- 4 Boot from SAN and Multipath Configuration 18
- 5 Installing the L2 Gateway Agent for the Networking Service 32

## 1 Overview

We have a *Chapter 2, Pre-Installation Checklist* we recommend going through to ensure your environment meets the requirements of the Cloud model you choose.

Once you have an idea of the configuration you want to choose for your cloud and you have gone through the pre-installation steps, you have two options for installation. You can use a GUI that runs in your web browser or you can install via the command-line that exposes the full power and flexibility of HPE Helion OpenStack 5.0.

## Using the GUI

You should use the GUI if:

- You are not planning to deploy availability zones or use L3 segmentation in your initial deployment.
- You are happy with the tuned HPE default OpenStack configurations.

Instructions for installing via the GUI are here.

• Chapter 7, Installing via the GUI

Note that reconfiguration of your cloud, at the moment, can only be done via the command-line. The GUI installer is for initial installation only.

# 2 Pre-Installation Checklist



## **Important**

The formatting of this page facilitates printing it out and using it to record details of your setup.

This checklist is focused on the Entry-scale KVM with VSA model but you can alter it to fit the example configuration you chose for your cloud.

# 2.1 BIOS and iLO Settings

Ensure that the following BIOS and iLO settings are applied to each baremetal server:

#	Item		
	Setup the iLO Advanced license in the iLO configuration		
	Choose either UEFI or Legacy BIOS in the BIOS settings		
	Verify the Date and Time settings in the BIOS.		
	Note: Verify Time and Date		
	It is important that all of your systems have the correct date/time because the		
	HPE StoreVirtual VSA license has a start date. If the start date is later than the system time then the installation will fail. Note also that Helion installs and runs		
	with UTC, not local time.		
	Ensure that Wake-on-LAN is disabled in the BIOS		
	Ensure that the NIC port to be used for PXE installation has PXE enabled in the BIOS		
	Ensure that all other NIC ports have PXE disabled in the BIOS		
	Ensure all hardware in the server not directly used by Helion is disabled		

BIOS and iLO Settings

# 2.2 Network Setup and Configuration

Before installing HPE Helion OpenStack, the following networks must be provisioned and tested. The networks are not installed or managed by the Cloud. You must install and manage the networks as documented in the .

Note that if you want a pluggable IPAM driver, it must be specified at install time. Only with a clean install of HPE Helion OpenStack 5.0 can you specify a different IPAM driver. If upgrading, you must use the default driver. More information can be found in .

Use these checklists to confirm and record your network configuration information.

#### **Router**

The IP router used with HPE Helion OpenStack must support the updated of its ARP table through gratuitous ARP packets.

## **PXE Installation Network**

When provisioning the IP range, allocate sufficient IP addresses to cover both the current number of servers and any planned expansion. Use the following table to help calculate the requirements:

Instance	Description	IPs
Deployer O/S		1
Controller server O/S (x3)		3
Compute servers (2nd thru 100th)	single IP per server	
VSA host servers	single IP per server	

#	Item	Value
	Network is untagged	
	No DHCP servers other than Helion are on the network	
	Switch PVID used to map any "internal" VLANs to untagged	
	Routable to the IPMI network	
	IP CIDR	

#	Item	Value
	IP Range (Usable IPs)	begin:
		end:
	Default IP Gateway	

## **Management Network**

The management network is the backbone used for the majority of HPE Helion OpenStack management communications. Control messages are exchanged between the Controllers, Compute hosts, and Cinder backends through this network. In addition to the control flows, the management network is also used to transport Swift and iSCSI based Cinder block storage traffic between servers.

When provisioning the IP Range, allocate sufficient IP addresses to cover both the current number of servers and any planned expansion. Use the following table to help calculate the requirements:

Instance	Description	IPs
Controller server O/S (x3)		3
Controller VIP		1
Compute servers (2nd thru 100th)	single IP per server	
VSA VM servers	single IP per server	
VSA VIP per cluster		

#	Item	Value
	Network is untagged	
	No DHCP servers other than Helion are on the network	
	Switch PVID used to map any "internal" VLANs to untagged	
	IP CIDR	
	IP Range (Usable IPs)	begin:
		end:

#	Item	Value
	Default IP Gateway	
	VLAN ID	

## **IPMI Network**

The IPMI network is used to connect the IPMI interfaces on the servers that are assigned for use with implementing the cloud. For HPE ProLiant servers, the IPMI network connects to the HPE iLO management device port for the server. This network is used by Cobbler to control the state of the servers during baremetal deployments.

#	Item	Value
	Network is untagged	
	Routable to the Management Network	
	IP Subnet	
	Default IP Gateway	

### **External API Network**

The External network is used to connect OpenStack endpoints to an external public network such as a company's intranet or the public internet in the case of a public cloud provider.

When provisioning the IP Range, allocate sufficient IP addresses to cover both the current number of servers and any planned expansion. Use the following table to help calculate the requirements.

Instance	Description	IPs
Controller server O/S (x3)		3
Controller VIP		1

#	Item	Value
	VLAN Tag assigned:	
	IP CIDR	

#	Item	Value
	IP Range (Usable IPs)	begin: end:
	Default IP Gateway	
	VLAN ID	

### **External VM Network**

The External VM network is used to connect cloud instances to an external public network such as a company's intranet or the public internet in the case of a public cloud provider. The external network has a predefined range of Floating IPs which are assigned to individual instances to enable communications to and from the instance to the assigned corporate intranet/internet. There should be a route between the External VM and External API networks so that instances provisioned in the cloud, may access the Cloud API endpoints, using the instance floating IPs.

#	Item	Value
	VLAN Tag assigned:	
	IP CIDR	
	IP Range (Usable IPs)	begin: end:
	Default IP Gateway	
	VLAN ID	

# 2.3 Lifecycle Manager

This server contains the HPE Helion OpenStack installer, which is based on Git, Ansible, and Cobbler.

#	<u>.</u>	Item	Value
		Disk Requirement: Single 8GB disk needed per the Book "Planning",	
		Chapter 2 "Hardware and Software Support Matrix"	

7 Lifecycle Manager

#	Item	Value
	Chapter 10, Installing Mid-scale and Entry-scale KVM	
	Ensure your local DNS nameserver is placed into your <a href="//etc/resolv.conf">/etc/resolv.conf</a> file	
	Install and configure NTP for your environment	
	Ensure your NTP server(s) is placed into your <a href="/&gt;/etc/ntp.conf">/etc/ntp.conf</a> file	
	NTP time source:	

# 2.4 Information for the nic\_mappings.yml Input File

Log onto each type of physical server you have and issue platform-appropriate commands to identify the <u>bus-address</u> and <u>port-num</u> values that may be required. For example, run the following command:

```
sudo lspci -D | grep -i net
```

and use this information for the bus-address value in your nic\_mappings.yml file.

# NIC Adapter PCI Bus Address Output

To find the port-num use:

```
cat /sys/class/net/<device name>/dev_port
```

where the 'device-name' is the name of the device **currently mapped** to this address, not necessarily the name of the device **to be mapped**.

Network Device Port Number Output	

## 2.5 Control Plane

The Control Plane consists of three servers, in a highly available cluster, that host the core HPE Helion OpenStack services including Nova, Keystone, Glance, Cinder, Heat, Neutron, Swift, Ceilometer, and Horizon. Additional services include mysql, ip-cluster, apache2, rabbitmq, memcached, zookeeper, kafka, influxDB, storm, monasca, logging, and cmc.



## Note

To mitigate the "split-brain" situation described in it is recommended that you have HA network configuration with Multi-Chassis Link Aggregation (MLAG) and NIC bonding configured for all the controllers to deliver system-level redundancy as well network-level resiliency. Also reducing the ARP timeout on the TOR switches will help.

#### **Control Plane 1**

#	Item	Value
	Disk Requirement: 3x 512 GB disks (or enough space to create three logical drives with that amount of space)	
	Ensure the disks are wiped	
	MAC address of first NIC	
	A second NIC, or a set of bonded NICs are required	
	iLO IP address	
	iLO Username/Password	

## **Control Plane 2**

#	Item	Value
	Disk Requirement: 3x 512 GB disks (or enough space to create three logical drives with that amount of space)	
	Ensure the disks are wiped	

9 Control Plane

#	Item	Value
	MAC address of first NIC	
	A second NIC, or a set of bonded NICs are required	
	iLO IP address	
	iLO Username/Password	

## **Control Plane 3**

#	Item	Value
	Disk Requirement: 3x 512 GB disks (or enough space to create three logical drives with that amount of space)	
	Ensure the disks are wiped	
	MAC address of first NIC	
	A second NIC, or a set of bonded NICs are required	
	iLO IP address	
	iLO Username/Password	

# 2.6 Compute Hosts

One or more KVM Compute servers will be used as the compute host targets for instances.

#	Item
	Disk Requirement: 2x 512 GB disks (or enough space to create three logical drives with that amount of space)
	A NIC for PXE boot and a second NIC, or a NIC for PXE and a set of bonded NICs are required
	Ensure the disks are wiped

10 Compute Hosts

Table to record your Compute host details:

ID	NIC MAC Address	iLO User- name/Pass- word	iLO IP Address	CPU/Mem/Disk

# 2.7 VSA Hosts

Three or more servers with local disk volumes running HPE StoreVirtual VSA to provide Cinder block storage resources.



## Note

The cluster created from VSA nodes should be quorum. In other words, the cluster must be formed with either 3,5,7 and so on.

#	Item
	Disk Requirement: 3x 512 GB disks (or enough space to create three logical drives with that amount of space)
	- The VSA appliance deployed on a host is expected to consume $\sim\!40$ GB of disk space from the host root disk for ephemeral storage to run the VSA virtual machine.
	A NIC for PXE boot and a second NIC, or a NIC for PXE and a set of bonded NICs are required
	Ensure the disks are wiped

11 VSA Hosts

Table to record your VSA host details:

ID	NIC MAC Address	iLO User- name/Pass- word	CPU/Mem/ Disk	VSA Data Volume

## 2.8 Additional Comments

This section is so you can add any additional information that you deem necessary.

## 2.9 SLES Pre-Installation Checks

## 2.9.1 Introduction

## This document needs review for SLES content.

As part of the preparation to provision a SLES compute node you need to be aware of the issues raised in this article and take any remedial action required.

## 2.9.2 Sample iptables rules must be removed on SLES

HPE Helion OpenStack 5.0 uses iptables to secure access to lifecycle manager network interfaces and on Red Hat this requires the package <u>iptables-services</u> to be installed. The package will provide sample iptables configurations for IPv4 and IPv6 if none existed before. This sample configuration is inappropriate for HPE Helion OpenStack operation and the node will not be able to run HOS with these rules installed.

12 Additional Comments

The files installed are:

- /etc/sysconfig/iptables
- /etc/sysconfig/ip6tables

If these files **do not exist** on the candidate Red Hat systems the rest of this note may be skipped as the HPE Helion OpenStack 5.0 install will prevent the installation of the sample files. However, if these files do exist, there are a number of steps that you must follow before you install HPE Helion OpenStack 5.0.

The contents of these two files are displayed here for reference:

### /etc/sysconfig/iptables

```
# sample configuration for iptables service
# you can edit this manually or use system-config-firewall
# please do not ask us to add additional ports/services to this default configuration
*filter
:INPUT ACCEPT [0:0]
:FORWARD ACCEPT [0:0]
:OUTPUT ACCEPT [0:0]
-A INPUT -m state --state RELATED,ESTABLISHED -j ACCEPT
-A INPUT -p icmp -j ACCEPT
-A INPUT -i lo -j ACCEPT
-A INPUT -p tcp -m state --state NEW -m tcp --dport 22 -j ACCEPT
-A INPUT -j REJECT --reject-with icmp-host-prohibited
-A FORWARD -j REJECT --reject-with icmp-host-prohibited
COMMIT
```

#### /etc/sysconfig/ip6tables

```
# sample configuration for ip6tables service
# you can edit this manually or use system-config-firewall
# please do not ask us to add additional ports/services to this default configuration
*filter
:INPUT ACCEPT [0:0]
:FORWARD ACCEPT [0:0]
:OUTPUT ACCEPT [0:0]
-A INPUT -m state --state RELATED,ESTABLISHED -j ACCEPT
-A INPUT -p ipv6-icmp -j ACCEPT
-A INPUT -i lo -j ACCEPT
-A INPUT -i lo -j ACCEPT
-A INPUT -b tcp -m state --state NEW -m tcp --dport 22 -j ACCEPT
-A INPUT -d fe80::/64 -p udp -m udp --dport 546 -m state --state NEW -j ACCEPT
-A INPUT -j REJECT --reject-with icmp6-adm-prohibited
-A FORWARD -j REJECT --reject-with icmp6-adm-prohibited
```

#### COMMIT

These rules are applied to network traffic on *any interface* on the system whereas HPE Helion OpenStack components and OpenStack components manage *interface specific* rules. With reference to the security policies that you may require for your environment, it is essential that you either:

- Delete the files <u>/etc/sysconfig/iptables</u> and <u>/etc/sysconfig/ip6tables</u> as none of the contained rules are required.
- Ensure that any remaining rules are limited to interfaces not used by HPE Helion OpenStack.

# Important

If these files existed on your system, and did contain content, the corresponding rules are currently installed and active on the system. Once you delete these two files (or edit them to limit the rules to interfaces not used by HPE Helion OpenStack), you will need to reboot the system to activate the new settings.

# 3 Using Git for Configuration Management

In HPE Helion OpenStack 5.0, a local git repository is used to track configuration changes and the configuration processor (CP) uses this repository. The introduction of a git workflow also means that your configuration history is maintained, making rollbacks easier as well as keeping a record of previous configuration settings.

The git repository is installed by the lifecycle manager on the lifecycle manager node.

The git repository provides a way for you to merge changes that you pull down as "upstream" (that is, from HP) updates, and allows you to manage your own configuration changes.

# 3.1 Initialization on a new deployment

#### hos

This holds the upstream source code corresponding to the contents of the  $\[ \]$ -/helion directory on a pristine fresh installation. Every source code release that is downloaded from HPE is applied as a fresh commit to this branch. This branch contains no customization by the end user.

#### site

This branch begins life as a copy of the first 'hos' drop. It is onto this branch that you commit your configuration changes. It's the branch most visible to the end user.

#### ansible

This branch contains the variable definitions generated by the CP that our main ansible playbooks need. This includes the <a href="verb\_hosts">verb\_hosts</a> file that describes to ansible what servers are playing what roles. The <a href="ready-deployment">ready-deployment</a> playbook takes this output and assembles a <a href="ready-deployment">~/scratch</a> directory containing the ansible playbooks together with the variable definitions in this branch. The result is a working ansible directory <a href="ready-deployment">~/scratch/ansible/next/</a> hos/ansible from which the main deployment playbooks may be successfully run.

#### cp-persistent

This branch contains the persistent state that the CP needs to maintain. That state is mostly the assignment of IP addresses and roles to particular servers. Some operational procedures may involve editing the contents of this branch: for example, retiring a machine from service or repurposing it.

Two temporary branches are created and populated at run time:

#### staging-ansible

This branch hosts the most recent commit that will be appended to the ansible branch.

#### staging-cp-persistent

This branch hosts the most recent commit that will be appended to the cp-persistent branch.



## Note

The information above provides insight into the workings of the configuration processor and the git repository. However, in practice you can simply follow the steps below to make configuration changes.

# 3.2 Updating any configuration, including the default configuration

When you need to make updates to a configuration you must:

1. Check out the <u>site</u> branch. You may already be on that branch. If so, git will tell you that and the command will leave you there.

```
git checkout site
```

- 2. Edit the YAML file or files that contain the configuration you want to change.
- 3. Commit the changes to the site branch.

```
git add -A
git commit -m "your commit message goes here in quotes"
```

If you want to add a single file to your git repository, you can use the command below, as opposed to using **git add -A**:

```
git add PATH_TO_FILE
```

For example, if you made a change to your <u>servers.yml</u> file and wanted to only commit that change, you would use this command:

```
git add ~/hestepon/my_cloud/definition/data/servers.yml
```

4. To produce the required configuration processor output from those changes. Review the output files manually if required, run the configuration processor:

```
cd ~/hestepon/hos/ansible ansible-playbook -i hosts/localhost config-processor-run.yml
```

5. Ready the deployment area

```
ansible-playbook -i hosts/localhost ready-deployment.yml
```

6. Run the deployment playbooks from the resulting scratch directory.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts site.yml
```

# 4 Boot from SAN and Multipath Configuration

## 4.1 Introduction

For information about supported hardware for multipathing, see .



When exporting a LUN to a node for boot from san, you should ensure that "LUN 0" is assigned to the LUN and configure any setup dialog that is necessary in the firmware to consume this LUN0 for OS boot.

# Important

Any hosts that are connected to 3PAR storage must have a host persona of 2-generic-alua set on the 3PAR. Refer to the 3PAR documentation for the steps necessary to check this and change if necessary.

iSCSI boot from SAN is not supported. For more information on the use of Cinder with multipath, see *Section 17.3.3, "Multipath Support"*.

In order to allow HPE Helion OpenStack 5.0 use volumes from a SAN, a number of configuration options need to be specified for both the install phases and for the osconfig phases. In all cases the devices that are utilized are devices for which multipath is configured.

# 4.2 Install Phase Configuration

For FC connected nodes and for FCoE nodes where the network processor used is from the Emulex family such as for the 650FLB, the following changes need to be made.

1. In each stanza of the servers.yml insert a line stating boot-from-san: true

- id: controller2
ip-addr: 192.168.10.4
role: CONTROLLER-ROLE
server-group: RACK2

18 Introduction

```
nic-mapping: HP-DL360-4PORT
mac-addr: "8a:8e:64:55:43:76"
ilo-ip: 192.168.9.4
ilo-password: password
ilo-user: admin
boot-from-san: true
```

This uses the disk /dev/mapper/mpatha as the default device on which to install the OS.

2. In the disk input models you also need to specify the devices that will be used via their multipath names (which will be of the form <a href="//dev/mapper/mpatha">/dev/mapper/mpatha</a>, <a href="//dev/mapper/mpatha">/dev/mapper/mpatha</a>, <a href="/dev/mapper/mpatha">/dev/mapper/mpatha</a>) etc).

```
volume-groups:
    - name: hlm-vg
    physical-volumes:

# NOTE: 'sda_root' is a templated value. This value is checked in
    # os-config and replaced by the partition actually used on sda
    #e.g. sda1 or sda5
    - /dev/mapper/mpatha_root

...

- name: vg-comp
    physical-volumes:
    - /dev/mapper/mpathb
```

# 4.3 QLogic FCoE restrictions and additional configurations

If you are using network cards such as Qlogic Flex Fabric 536 and 630 series then there are additional OS configuration steps to support the importation of LUNs as well as some restrictions on supported configurations.

The restrictions are:

- 1. Only one network card can be enabled in the system.
- 2. The FCoE interfaces on this card are dedicated to FCoE traffic these cannot have ip addresses associated with them.
- 3. Nic mapping cannot be used.

In addition to the configuration options above, you also need to specify the FCoE interfaces for install and for os configuration. There are 3 places where you need to add additional configuration options for fcoe-support:

• In <u>servers.yml</u>, which is used for configuration of the system during OS install, FCoE interfaces need to be specified for each server. In particular, the mac addresses of the FCoE interfaces need to be given, **not** the symbolic name (e.g. eth2).

```
- id: compute1
ip-addr: 10.245.224.201
role: COMPUTE-ROLE
server-group: RACK2
mac-addr: 6c:c2:17:33:4c:a0
ilo-ip: 10.1.66.26
ilo-user: hlinux
ilo-password: hlinux123
boot-from-san: True
fcoe-interfaces:
    - 6c:c2:17:33:4c:a1
    - 6c:c2:17:33:4c:a9
```

# Important

Nic mapping can not be used.

• For the osconfig phase, you will need to specify the <u>fcoe-interfaces</u> as a peer of <u>net-work-interfaces</u> in the net interfaces.yml file:

# Important

The mac addresses specified in the  $\underline{\texttt{fcoe-interfaces}}$  stanza in  $\underline{\texttt{servers.yml}}$  must correspond to the symbolic names used in the  $\underline{\texttt{fcoe-interfaces}}$  stanza in  $\underline{\texttt{net\_in-terfaces.yml}}$ .

Also, to satisfy the restriction outlined in #multipath\_boot\_from\_san/restriction2 nabove, there can be no overlap between the devices in <a href="fcoe-interfaces">fcoe-interfaces</a> and those in <a href="mailto:network-interfaces">network-interfaces</a> in the <a href="net\_interfaces.yml">network-interfaces</a> in the <a href="mailto:network-interfaces">network-interfaces</a>.

and eth3 are fcoe-interfaces while eth0 is in network-interfaces.

• As part of the initial install from an iso, additional parameters need to be supplied on the kernel command line:

```
multipath=true partman-fcoe/interfaces=<mac address1>,<mac address2> disk-detect/
fcoe/enable=true --- quiet
```

See the sections of installing from an ISO using UEFI (Section 4.4.1.2, "UEFI Boot Mode with QLogic FCoE") and BIOS (Section 4.4.1.4, "Legacy BIOS Boot Mode with QLogic FCoE").

Since nic mapping is not used to guarantee order of the networks across the system the installer will remap the network interfaces in a deterministic fashion as part of the install. As part of the installer dialogue, if DHCP is not configured for the interface, it is necessary to confirm that the appropriate interface is assigned the ip address. The network interfaces may not be at the names expected when installing via an iso. When you are asked to apply an ip address to an interface, enter Alt-F2 and in the console window issue an ip a command to examine the interfaces and their associated mac addresses. Take a note of the interface name with the expected mac address and use this in the subsequent dialogue. Enter Alt-F1 to return to the installation screen. You should note that the names of the interfaces may have changed after the installation completes. These names are used consistently in any subsequent operations.

Therefore, even if FCoE is not used for boot from SAN (e.g. for cinder), then it is recommended that <u>fcoe-interfaces</u> be specified as part of install (without the multipath or disk detect options). Alternatively, you need to run osconfig-fcoe-reorder.yml before site.yml or

osconfig-run.yml is invoked to reorder the networks in a similar manner to the installer. In this case, the nodes will need to be manually rebooted for the network reorder to take effect. You will run osconfig-fcoe-reorder.yml in the following scenarios:

- If you have used a third-party installer to provision your baremetal nodes
- If you are booting from a local disk (i.e. one that is not presented from the SAN) but you want to use FCoE later, for example, for cinder.

#### To run the command:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts osconfig-fcoe-reorder.yml
```

If you do not run <u>osconfig-fcoe-reorder.yml</u>, you will encounter a failure in <u>osconfig-run.yml</u>.

If you are booting from a local disk, the LUNs that will be imported over FCoE will not be visible before site.yml or osconfig-run.yml has been run. However, if you need to import the LUNs before this, for instance, in scenarios where you need to run wipe\_disks.yml, then you can run the fcoe-enable playbook across the nodes in question. This will configure FCoE and import the LUNs presented to the nodes.

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/verb_hosts fcoe-enable.yml
```

## 4.4 Red Hat Compute Host for FCoE

When installing a Red Hat compute host, the names of the network interfaces will have names like <a href="mailto:ens1f2">ens1f2</a> rather than <a href="mailto:eth2">eth2</a> therefore a separate role and associated <a href="mailto:network-interfaces">network-interfaces</a> and <a href="mailto:foce-interfaces">fcoe-interfaces</a> descriptions will have to be provided in the input model for the Red Hat compute hosts. Here are some excerpts highlighting the changes required:

### net\_interfaces.yml

```
    name: RHEL-COMPUTE-INTERFACES
    fcoe-interfaces:
    name: fcoe
    devices:
    ens1f2
    ens1f3
```

### 

device:

name: ens1f0 network-groups: - EXTERNAL-VM - GUEST

- MANAGEMENT

### control\_plane.yml

- name: rhel-compute

resource-prefix: rhcomp

server-role: RHEL-COMPUTE-ROLE

allocation-policy: any

min-count: 0

service-components:

ntp-clientnova-compute

- nova-compute-kvm

- neutron-l3-agent

- neutron-metadata-agent

- neutron-openvswitch-agent

- neutron-lbaasv2-agent

#### server\_roles.yml

- name: RHEL-COMPUTE-ROLE

interface-model: RHEL-COMPUTE-INTERFACES

disk-model: COMPUTE-DISKS

#### servers.yml

- id: QlogicFCoE-Cmp2

ip-addr: 10.245.224.204
role: RHEL-COMPUTE-ROLE
server-group: RACK2

mac-addr: "6c:c2:17:33:4c:a0"

ilo-ip: 10.1.66.26
ilo-password: hlinux123

ilo-user: hlinux
boot-from-san: True

distro-id: rhel72-x86\_64-multipath

fcoe-interfaces:

- 6c:c2:17:33:4c:a1 - 6c:c2:17:33:4c:a9

# Important

Note that you need to add a <u>-multipath</u> suffix to the <u>distro-id</u> value when using multipath with RHEL.

If you are installing a Red Hat compute node with QLogic FCoE boot from SAN, you will need to edit the <a href="https://www.var/lib/cobbler/kickstarts/rhel72-anaconda-ks-multipath.cfg">/var/lib/cobbler/kickstarts/rhel72-anaconda-ks-multipath.cfg</a> file. and uncomment the **two** sections that are prefaced by:

```
#Qlogic-FCOE: Uncomment the below lines if using qlogic fcoe boot from san
```

The values will be overwritten on running the cobbler-deploy play.



## Note

It is highly recommended as part of a Red Hat install that only one disk is presented to the node during the install phase. You should ensure that any additional LUNs that are required are attached and visible on the compute hosts before running the site.yml playbook.

# 4.4.1 Installing the HPE Helion OpenStack 5.0 iso for nodes that support Boot from SAN

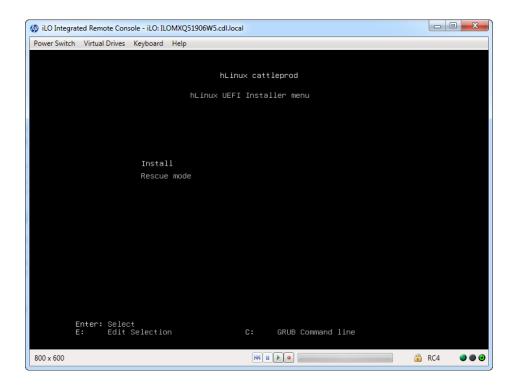
This sections describes how to install the iso on the Lifecycle Manager to support the following configurations:

- Section 4.4.1.1, "UEFI Boot Mode"
- Section 4.4.1.2, "UEFI Boot Mode with QLogic FCoE"
- Section 4.4.1.3, "Legacy BIOS Boot Mode"
- Section 4.4.1.4, "Legacy BIOS Boot Mode with QLogic FCoE"

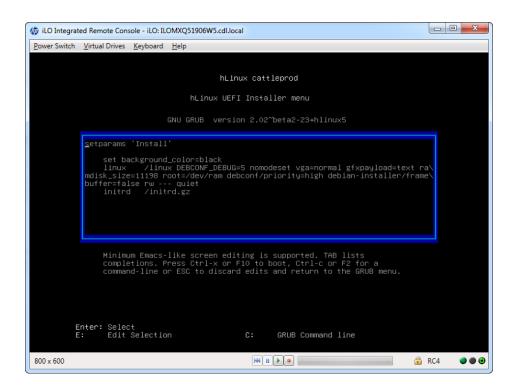
The lifecycle manager will then automatically handle the installation on nodes that supports Boot from SAN based on the configuration information in <u>servers.yml</u> and the disk models, as described in the preceding section.

#### 4.4.1.1 UEFI Boot Mode

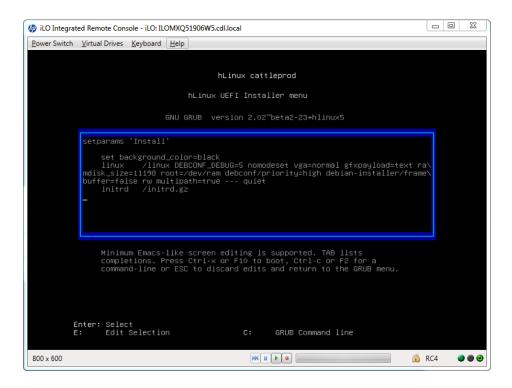
On the installer boot menu, press  $extbf{E}$  to enter the Edit Selection screen.



This brings up the Edit Selection screen.



Enter the text multipath=true before --- quiet:

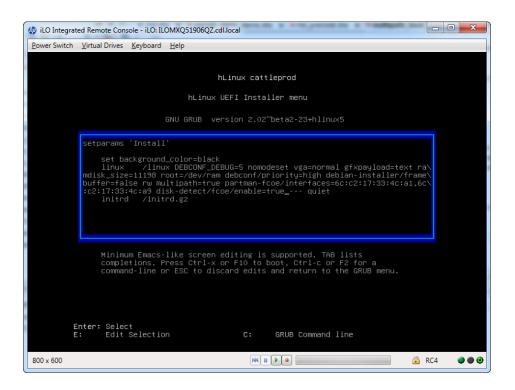


Press Ctrl-X or F10 to proceed with the install.

## 4.4.1.2 UEFI Boot Mode with QLogic FCoE

At the installer boot menu, press **E** to enter the <u>Edit Selection</u> screen as described in the preceding section. In addition to inserting <u>multipath=true</u>, it is necessary to supply details of the FCoE network interfaces. In the example below, the interfaces are specified as:

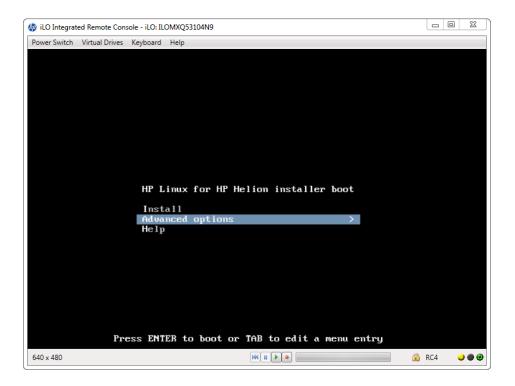
partman-fcoe/interfaces=6c:c2:17:33:4c:a1,6c:c2:17:33:4c:a9 disk-detect/fcoe/enable=true



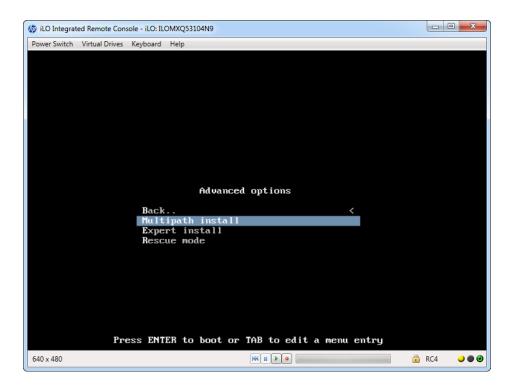
Press Ctrl—X or F10 to proceed with the install.

## 4.4.1.3 Legacy BIOS Boot Mode

On the installer boot menu, navigate (using 1) to the Advanced options entry and then press Enter.



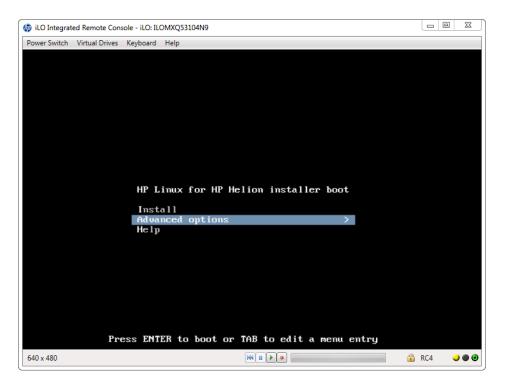
This will bring up the Advanced options menu:



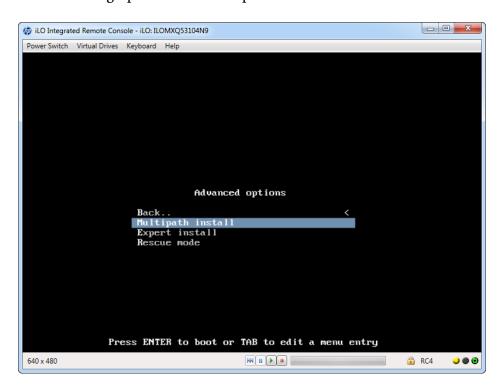
Navigate to the Multipath install entry and press <a>Enter</a> to start the installation.

## 4.4.1.4 Legacy BIOS Boot Mode with QLogic FCoE

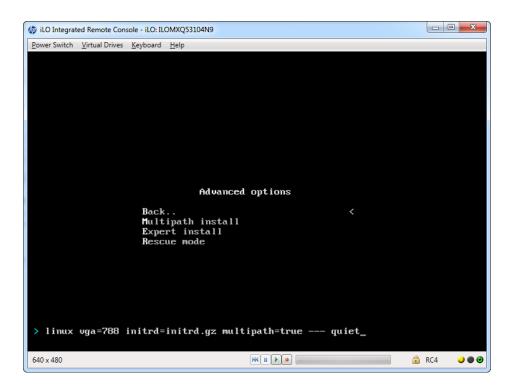
On the installer boot menu, navigate (using 1) to the Advanced options entry and then press Enter.



This will bring up the Advanced options menu:

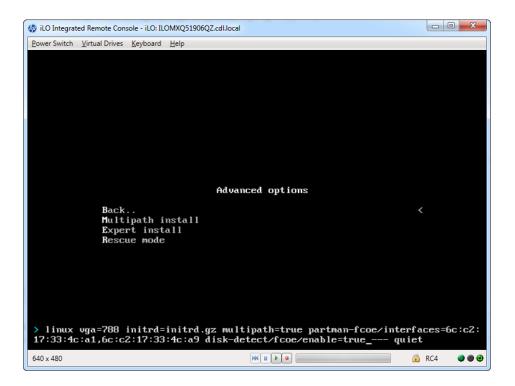


Navigate to the <u>Multipath install</u> entry and press <u>I</u> to edit the entry details. Notice that the kernel command line is now displayed at the bottom of the screen and that <u>multipath=true</u> is already specified.



Edit the kernel command line to add the FCoE network interfaces before the  $\frac{---}{}$ . In the example below, the interfaces are specified as:

 $\verb|partman-fcoe| interfaces=6c:c2:17:33:4c:a1,6c:c2:17:33:4c:a9 | disk-detect/fcoe/enable=true| | disk-detect/fcoe/enable=tru$ 



Now press Enter to start the installation.

# 5 Installing the L2 Gateway Agent for the Networking Service

## 5.1 Introduction

The L2 gateway is a service plug-in to the Neutron networking service that allows two L2 networks to be seamlessly connected to create a single L2 broadcast domain. The initial implementation provides for the ability to connect a virtual Neutron VxLAN network to a physical VLAN using a VTEP-capable HPE 5930 switch. The L2 gateway is to be enabled only for VxLAN deployments.

To begin L2 gateway agent setup, you need to configure your switch. These instructions use an HPE FlexFabric 5930 Switch Series (http://www8.hp.com/us/en/products/networking-switches/product-detail.html?oid=6604154#!tab=models) switch.

- Use case
- Sample network topology (#l2gw5930/topology) 

  ✓
- Networks (#l2gw5930/networks)
- HPE 5930 Switch Configuration (#l2gw5930/switchConfig)
- Enabling and configuring L2 gateway agent (#l2gw5930/enabling)
- Routing between software and hardware VTEP networks (#I2gw5930/routing)
- Connecting baremetal server to HPE 5930 switch (#l2gw5930/connecting)
- Configuration on baremetal server to receive tagged packets from the switch (#l2gw5930/baremetalConfig)
- NIC bonding and IRF configuration (#l2gw5930/NICbonding)
- Scale numbers tested (#l2gw5930/scale)
- L2 gateway commands (#l2gw5930/commands)

32 Introduction

# 5.2 Sample network topology (for illustration purposes)

When viewing the following network diagram, assume that the blue VNET has been created by the tenant and has been assigned a segmentation ID of 1000 (VNI 1000). The Cloud Admin is now connecting physical servers to this VNET.

Assume also that the blue L2 Gateway is created and points to the HW VTEPS, the physical ports, the VLAN, and if it is an access or trunk port (tagged or untagged)

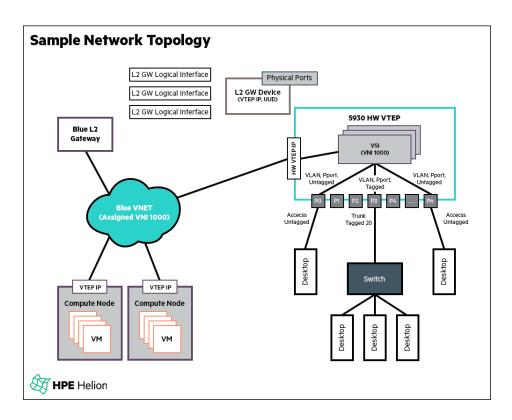


FIGURE 5.1: FIGURE 1

## 5.3 Networks

The following diagram illustrates an example network configuration.

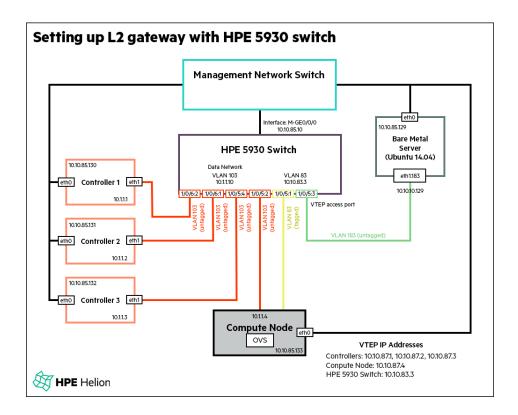


FIGURE 5.2: FIGURE 2

An L2 gateway is useful in extending virtual networks (VxLAN) in a cloud onto physical VLAN networks. The L2 gateway switch converts VxLAN packets into VLAN packets and back again, as shown in the following diagram. This topic assumes a VxLAN deployment.



#### FIGURE 5.3: FIGURE 3

Management Network: 10.10.85.0/24

• Data Network: 10.1.1.0/24

• Tenant VM Network: 10.10.10.0/24

Note: These IP ranges are used in the topology shown in the diagram for illustration only.

34 Networks

# 5.4 HPE 5930 switch configuration

- 1. Telnet to the 5930 switch and provide your username and password.
- 2. Go into system view:

```
system-view
```

3. Create the required VLANs and VLAN ranges:

```
vlan 103
vlan 83
vlan 183
vlan 1261 to 1270
```

4. Assign an IP address to VLAN 103. This is used as a data path network for VxLAN traffic.

```
interface vlan 103
ip address 10.1.1.10 255.255.255.0
```

5. Assign an IP address to VLAN 83. This is used as a hardware VTEP network.

```
interface vlan 83
ip address 10.10.83.3 255.255.255.0
```

The 5930 switch has a fortygigE1/0/5 interface to which a splitter cable is connected that splits the network it into four tengigEthernet (tengigEthernet1/0/5:1 to tengigEthernet1/0/5:4) interfaces:

- engigEthernet1/0/5:1 and tengigEthernet1/0/5:2 are connected to the compute node. This is required just to bring the interface up. In other words, in order to have the HPE 5930 switch work as a router, there should be at least one interface of that particular VLAN up. Alternatively, the interface can be connected to any host or network element.
- tengigEthernet1/0/5:3 is connected to a baremetal server.
- tengigEthernet1/0/5:4 is connected to controller 3, as shown in Figure 2 (#l2g-w5930/fig2)

The switch's fortygigE1/0/6 interface to which the splitter cable is connected splits it into four tengigEthernet (tengigEthernet1/0/6:1 to tengigEthernet1/0/6:4) interfaces:

- tengigEthernet1/0/6:1 is connected to controller 2
- tengigEthernet1/0/6:2 is connected to controller 1
   Note: 6:3 and 6:4 are not used although they are available.
- 6. Split the fortygigE 1/0/5 interface into tengig interfaces:

```
interface fortygigE 1/0/5
using tengig
The interface FortyGigE1/0/5 will be deleted. Continue? [Y/N]: y
```

7. Configure the Ten-GigabitEthernet1/0/5:1 interface:

```
interface Ten-GigabitEthernet1/0/5:1
port link-type trunk
port trunk permit vlan 83
```

# 5.5 Configuring the provider data path network

1. Configure the Ten-GigabitEthernet1/0/5:2 interface:

```
interface Ten-GigabitEthernet1/0/5:2

port link-type trunk

port trunk permit vlan 103

port trunk permit vlan 1261 to 1270

port trunk pvid vlan 103
```

2. Configure the Ten-GigabitEthernet1/0/5:4 interface:

```
interface Ten-GigabitEthernet1/0/5:4
port link-type trunk
port trunk permit vlan 103
port trunk permit vlan 1261 to 1270
port trunk pvid vlan 103
```

3. Configure the Ten-GigabitEthernet1/0/5:3 interface:

```
interface Ten-GigabitEthernet1/0/5:3
port link-type trunk
```

```
port trunk permit vlan 183
vtep access port
```

4. Split the fortygigE 1/0/6 interface into tengig interfaces:

```
interface fortygigE 1/0/6
using tengig
The interface FortyGigE1/0/6 will be deleted. Continue? [Y/N]: y
```

5. Configure the Ten-GigabitEthernet1/0/6:1 interface:

```
interface Ten-GigabitEthernet1/0/6:1
port link-type trunk
port trunk permit vlan 103
port trunk permit vlan 1261 to 1270
port trunk pvid vlan 103
```

6. Configure the Ten-GigabitEthernet1/0/6:2 interface:

```
interface Ten-GigabitEthernet1/0/6:2

port link-type trunk

port trunk permit vlan 103

port trunk permit vlan 1261 to 1270

port trunk pvid vlan 103
```

7. Enable l2vpn:

```
l2vpn enable
```

8. Configure a passive TCP connection for OVSDB on port 6632:

```
ovsdb server ptcp port 6632
```

9. Enable OVSDB server:

```
ovsdb server enable
```

10. Enable a VTEP process:

```
vtep enable
```

11. Configure 10.10.83.3 as the VTEP source IP. This acts as a hardware VTEP IP.

```
tunnel global source-address 10.10.83.3
```

12. Configure the VTEP access port:

```
interface Ten-GigabitEthernet1/0/5:3
vtep access port
```

13. Disable VxLAN tunnel mac-learning:

```
vxlan tunnel mac-learning disable
```

14. Display the current configuration of the 5930 switch and verify the configuration:

```
display current-configuration
```

After switch configuration is complete, you can dump OVSDB to see the entries.

1. Run the ovsdb-client from any Linux machine reachable by the switch:

```
ovsdb-client dump --pretty tcp:10.10.85.10:6632
sdn@small-hLinux:~$ ovsdb-client dump --pretty tcp:10.10.85.10:6632
Arp_Sources_Local table
_uuid locator src_mac
-----
Arp_Sources_Remote table
_uuid locator src_mac
----
Global table
_uuid
                              managers switches
2c891edc-439b-4144-84d9-fa4cd88092bf []
                                      [f5f4b43b-40bc-4640-b580-d4b12011f688]
Logical Binding Stats table
_uuid bytes_from_local bytes_to_local packets_from_local packets_to_local
Logical Router table
_uuid description name static_routes switch_binding
-----
Logical Switch table
_uuid description name tunnel_key
-----
Manager table
```

```
_uuid inactivity_probe is_connected max_backoff other_config status target
Mcast Macs Local table
MAC _uuid ipaddr locator_set logical_switch
--- -----
Mcast_Macs_Remote table
MAC _uuid ipaddr locator_set logical_switch
--- ----- ----- ------
Physical_Locator table
_uuid dst_ip encapsulation_type
_____
Physical_Locator_Set table
_uuid locators
-----
Physical_Port table
_uuid
                        description name
port_fault_status vlan_bindings vlan_stats
------
fda90870-656c-479b-91ee-852b70e2007e "" "Ten-GigabitEthernet1/0/5:3" [UP]
      {}
Physical_Switch table
                      description management_ips name
_uuid
                   switch_fault_status tunnel_ips tunnels
f5f4b43b-40bc-4640-b580-d4b12011f688 ""
"L2GTWY02" [fda90870-656c-479b-91ee-852b70e2007e] []
["10.10.83.3"] []
Tunnel table
_uuid bfd_config_local bfd_config_remote bfd_params bfd_status local remote
Ucast_Macs_Local table
MAC _uuid ipaddr locator logical_switch
--- ----- ----- ------
Ucast_Macs_Remote table
MAC _uuid ipaddr locator logical_switch
```

--- ----- ------

# 5.6 Enabling and configuring the L2 gateway agent

1. Update the input model (in <u>control\_plane.yml</u>) to specify where you want to run the neutron-l2gateway-agent. For example, see the line in bold in the following yml:

```
product:
version: 2
control-planes:
- name: cp
  region-name: region1
failure-zones:
- AZ1
 common-service-components:
    - logging-producer
    - monasca-agent
    - freezer-agent
    - stunnel
    - lifecycle-manager-target
clusters:
- name: cluster1
 cluster-prefix: c1
 server-role: ROLE-CONTROLLER
 member-count: 2
 allocation-policy: strict
service-components:
- neutron-l2gateway-agent
```

2. Update <a href="lagateway\_agent.ini.j2">lagateway\_agent.ini.j2</a>. For example, here the IP address (10.10.85.10) must be the management IP address of your 5930 switch. Open the file in vi:

```
$ vi /home/stack/my_cloud/config/neutron/l2gateway_agent.ini.j2
```

3. Then make the changes:

```
[ovsdb]
# (StrOpt) OVSDB server tuples in the format
# <ovsdb_name>:<ip address>:<port>[,<ovsdb_name>:<ip address>:<port>]
```

```
# - ovsdb_name: a symbolic name that helps identifies keys and certificate files
# - ip address: the address or dns name for the ovsdb server
# - port: the port (ssl is supported)
ovsdb_hosts = hardware_vtep:10.10.85.10:6632
```

4. By default, the L2 gateway agent initiates a connection to OVSDB servers running on the L2 gateway switches. Set the attribute <a href="enable\_manager">enable\_manager</a> to <a href="True">True</a> if you want to change this behavior (to make L2 gateway switches initiate a connection to the L2 gateway agent). In this case, it is assumed that the Manager table in the OVSDB hardware\_vtep schema on the L2 gateway switch has been populated with the management IP address of the L2 gateway agent and the port.

```
#enable_manager = False
#connection can be initiated by the ovsdb server.
#By default 'enable_manager' value is False, turn on the variable to True
#to initiate the connection from ovsdb server to l2gw agent.
```

5. If the port that is configured with <a href="enable\_manager">enable\_manager</a> = True is any port other than 6632, update the <a href="enable-manager">2.0/services/neutron/l2gateway-agent.yml</a> input model file with that port number:

```
endpoints:
- port: '6632'
roles:
- ovsdb-server
```

6. Note: The following command can be used to set the Manager table on the switch from a remote system:

```
sudo vtep-ctl --db=tcp:10.10.85.10:6632 set-manager tcp:10.10.85.130:6632
```

7. For SSL communication, the command is:

```
sudo vtep-ctl --db=tcp:10.10.85.10:6632 set-manager ssl:10.10.85.130:6632
```

where **10.10.85.10** is the management IP address of the L2 gateway switch and **10.10.85.130** is the management IP of the host on which the L2 gateway agent runs. Therefore, in the above topology, this command has to be repeated for **10.10.85.131** and **10.10.85.132**.

8. If you are not using SSL, comment out the following:

```
#l2_gw_agent_priv_key_base_path={{    neutron_l2gateway_agent_creds_dir }}/keys
```

```
#l2_gw_agent_cert_base_path={{ neutron_l2gateway_agent_creds_dir }}/certs
#l2_gw_agent_ca_cert_base_path={{ neutron_l2gateway_agent_creds_dir }}/ca_certs
```

9. If you are using SSL, then rather than commenting out the attributes, specify the directory path of the private key, the certificate, and the CA cert that the agent should use to communicate with the L2 gateway switch which has the OVSDB server enabled for SSL communication.

Make sure that the directory path of the files is given permissions 755, and the files' owner is root and the group is root with 644 permissions.

**Private key:** The name should be the same as the symbolic name used above in ovsdb\_hosts attribute. The extension of the file should be ".key". With respect to the above example, the filename will be hardware\_vtep.key

**Certificate** The name should be the same as the symbolic name used above in ovsdb\_hosts attribute. The extension of the file should be ".cert". With respect to the above example, the filename will be hardware\_vtep.cert

**CA certificate** The name should be the same as the symbolic name used above in ovsd-b\_hosts attribute. The extension of the file should be ".ca\_cert". With respect to the above example, the filename will be hardware\_vtep.ca\_cert

10. To enable the HPE 5930 switch for SSL communication, execute the following commands:

```
undo ovsdb server ptcp
undo ovsdb server enable
ovsdb server ca-certificate flash:/cacert.pem bootstrap
ovsdb server certificate flash:/sc-cert.pem
ovsdb server private-key flash:/sc-privkey.pem
ovsdb server pssl port 6632
ovsdb server enable
```

11. Data from the OVSDB sever with SSL can be viewed using the following command:

```
ovsdb-client -C <ca-cert.pem> -p <client-private-key.pem> -c <client-cert.pem> dump ssl:10.10.85.10:6632
```

# 5.7 Routing between software and hardware

#### VTEP networks

In order to allow L2 gateway switches to send VxLAN packets over the correct tunnels destined for the compute node and controller node VTEPs, you must ensure that the cloud VTEP (compute and controller) IP addresses are in different network/subnet from that of the L2 gateway switches. You must also create a route between these two networks. This is explained below.

1. In the following example of the input model file <a href="networks.yml">networks.yml</a>, the GUEST-NET represents the cloud data VxLAN network. REMOTE-NET is the network that represents the hardware VTEP network.

```
# networks.yml
networks:
    - name: GUEST-NET
    vlanid: 103
    tagged-vlan: false
    cidr: 10.1.1.0/24
    gateway-ip: 10.1.1.10
    network-group: GUEST

- name: REMOTE-NET
    vlanid: 183
    tagged-vlan: false
    cidr: 10.10.83.0/24
    gateway-ip: 10.10.83.3
    network-group: REMOTE
```

2. The route must be configured between the two networks in the <a href="network-groups.yml">network-groups.yml</a> input model file:

Note that the IP route is configured on the compute node. Per this route, the HPE 5930 acts as a gateway that routes between the two networks.

3. On the compute node, it looks like this:

```
stack@padawan-cp1-comp0001-mgmt:~$ sudo ip route 10.10.83.0/24 via 10.1.1.10 dev eth4
```

4. Run the following Ansible playbooks to apply the changes. config-processor-run.yml:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

ready-deployment.yml:

```
ansible-playbook -i hosts/localhost ready-deployment.yml
```

hlm-reconfigure.yml:

```
cd ~/scratch/ansible/next/hos/ansible/
ansible-playbook -i hosts/verb_hosts hlm-reconfigure.yml
```

#### Notes:

- 1. Make sure that the controller cluster is able to reach the management IP address of the L2 gateway switch. Otherwise, the L2 gateway agents running on the controllers will not be able to reach the gateway switches.
- 2. Make sure that the interface on the baremetal server connected to the 5930 switch is tagged (this is explained shortly).

# 5.8 Connecting a baremetal server to the HPE 5930 switch

As Ubuntu (baremetal server) is not aware of tagged packets, it is not possible for a virtual machine to communicate with a baremetal box.

You (the administrator) must manually perform configuration changes to the interface in the HPE 5930 switch to which the baremetal server is connected so that the switch can send untagged packets. Either one of the following command sets can be used to do so:

```
Interface <interface number>
service-instance <service-instance id>
```

encapsulation untagged
xconnect vsi <vsi-name>

#### Or:

```
Interface <interface number>
service-instance <service-instance id>
encapsulation s-vid <vlan-id>
xconnect vsi <vsi-name> access-mode ethernet
```

There are two ways of configuring the baremetal server to communicate with virtual machines. If the switch sends tagged traffic, then the baremetal server should be able to receive the tagged traffic.

# 5.9 Configuration on a baremetal server

## To receive tagged packets from the switch

If the configuration changes mentioned previously are not made on the switch to send untagged traffic to the baremetal server, then the following changes are required on the baremetal server so that it can receive tagged traffic from the switch.

Bare metal Management ip: 10.10.85.129 on interface em1 Switch 5930 is connected to baremetal on eth1 Need to set the IP address into tagged interface of eth1

1. create tagged (VLAN 183) interface

```
vconfig add ethl 183
```

2. Assign the IP address (10.10.10.129) to eth1.183 tagged interface (IP from the subnet 10.10.10.0/24 as VM (10.10.10.4) spawned in Compute node belongs to this subnet).

```
ifconfig eth1.183 10.10.10.129/24
```

## 5.9.1 NIC bonding and IRF configuration

With L2 gateway in actiondeployment, NIC bonding can be enabled on compute nodes. For more details on NIC bonding, please refer to . In order to achieve high availability, HPE 5930 switches can be configured to form a cluster using Intelligent Resilient Framework (IRF). Please refer to the HPE FlexFabric 5930 Switch Series configuration guide (http://h20564.www2.hpe.com/hpsc/doc/public/display?docId=c04567141) for details.

## 5.10 Scale numbers tested

- Number of neutron port MACs tested on a single switch: 4000
- Number of HPE 5930 switches tested: 2
- Number of baremetal connected to a single HPE 5930 switch: 100
- Number of L2 gateway connections to different networks: 800

# 5.11 L2 gateway commands

These commands are not part of the L2 gateway deployment. They are to be executed after L2 gateway is deployed.

1. Create Network

```
neutron net-create net1
```

2. Create Subnet

```
neutron subnet-create net1 10.10.10.0/24
```

3. Boot a tenant VM ( nova boot –image < Image-id > --flavor 2 –nic net-id = < net\_id > < VM name > )

```
nova boot --image 1f3cd49d-9239-49cf-8736-76bac5360489 --flavor 2 --nic net-
id=4f6c58b6-0acc-4e93-bb4c-439b38c27a23 VM
```

Assume the VM was assigned the IP address 10.10.10.4

4. Create the L2 gateway filling in your information here:

```
neutron l2-gateway-create --device name="<switch
name>",interface_names="<interface_name>" <gateway-name>
```

For this example:

```
neutron l2-gateway-create --device name="L2GTWY02",interface_names="Ten-
GigabitEthernet1/0/5:3" gw1
```

46 Scale numbers tested

Ping from the VM (10.10.10.4) to baremetal server and from baremetal server (10.10.10.129) to the VM Ping should not work as there is no gateway connection created yet.

5. Create 12 gateway Connection

```
neutron l2-gateway-connection-create gw1 net1 --segmentation-id 183
```

- 6. Ping from VM (10.10.10.4) to baremetal and from baremetal (10.10.10.129) to VM Ping should work.
- 7. Delete 12 gateway Connection

```
neutron l2-gateway-connection-delete <gateway id/gateway_name>
```

8. Ping from the VM (10.10.10.4) to baremetal and from baremetal (10.10.10.129) to the VM. Ping should not work as 12 gateway connection was deleted.

47

# **II** Cloud Installation

- 6 Overview 49
- 7 Installing via the GUI 51
- 8 DNS Service Installation Overview 57
- 9 Magnum Overview 67
- 10 Installing Mid-scale and Entry-scale KVM 89
- 11 Installation for Helion Entry Scale ESX (with OVSvApp, KVM with VSA Model) 101
- 12 Sample activationtemplate.json File for ESX Compute 123
- 13 Installing Baremetal (Ironic) 140
- 14 Installation for HPE Helion Entry-scale Cloud with Swift Only 214
- 15 Installing SLES Compute 225
- 16 HLM-Hypervisor instructions 232
- 17 Integrations 257
- 18 Troubleshooting the Installation **361**
- 19 Troubleshooting the Block Storage Backend Configuration 376
- 20 Troubleshooting the ESX 378

# 6 Overview

Before jumping into your installation, we recommend taking the time to read through our *Book "Planning"*, *Chapter 10 "Example Configurations"* documentation to get an overview of the sample configurations HPE Helion OpenStack 5.0 offers. We have highly tuned example configurations for each of these Cloud models:

- •
- •
- •
- •
- •

Name	Location
	~/helion/examples/entry-scale-kvm- vsa
	~/helion/examples/entry-scale-kvm-vsa-mml
	~/helion/examples/entry-scale-kvm-ceph
	~/helion/examples/mid-scale-kvm-vsa
	~/helion/examples/entry-scale-esx- kvm-vsa
	~/helion/examples/entry-scale-esx- kvm-vsa-mml
	~/helion/examples/entry-scale-swift
	~/helion/examples/entry-scale- ironic-flat-network
	~/helion/examples/entry-scale- ironic-multi-tenancy

## **Using the Command-line**

You should use the command-line if:

- You are installing a more complex or large-scale cloud.
- You have more than 15 nodes.
- You need to use availability zones or the server groups functionality of the cloud model. See the for more information.
- You want to customize the cloud configuration beyond the tuned defaults that HPE provides out of the box.
- You need deeper customizations than are possible to express using the GUI.

## Instructions for installing via the command-line are here:

- Chapter 10, Installing Mid-scale and Entry-scale KVM
- Chapter 11, Installation for Helion Entry Scale ESX (with OVSvApp, KVM with VSA Model)

# 7 Installing via the GUI

HPE Helion OpenStack 5.0 comes with a UI installer for installing your cloud the first time. Rather than searching through and editing a number of YAML files to input your specific configuration data, finding the correct location, and maintaining the format of the files, the UI presents fields in which you may enter your information. In addition, this latest version allows you to populate server information via a .csv file import.

Using the installer, the series Cobbler and Ansible scripts that initiate the installation of HPE Helion OpenStack are run for you behind the scenes.

The UI app populates the configuration files that the configuration processor uses as input to its processes. Therefore you can inspect the files whenever you like, or before reconfiguring your cloud.

Note, the GUI installer will periodically save the configuration files during the process and when you click the Deploy button at the end of the process. This also means that if you were to try to change anything in the files themselves mid process, the values would get overwritten by values you have entered in the GUI.

These YAML configuration files are stored in a local git repository as well, as explained in *Chapter 3, Using Git for Configuration Management*.

When installing via the GUI, you begin with a model (or a cloud configuration template) and edit the contents to reflect your environment.

However, the full range of customizations is not possible via the GUI as some configuration items are not exposed. Below is the list of what can and cannot be changed:

#### Changes to the following items may be made:

- servers (including Linux for HPE Helion installation configuration)
- networks
- disk models
- interface models
- NIC mappings
- NTP servers
- Name servers
- tags in network groups

### Changes to the following items cannot be made:

- server groups
- server roles
- network groups
- firewall rules
- cloudConfig.yml (i.e. DNS, SMTP, firewall settings)
- control planes

# 7.1 Before you begin

Before you run the GUI installer to install your cloud, there are a number of things you need to do:

- 1. Prepare your servers
- 2. Gather your information
  - Server names
  - IP addresses
  - Server Roles
  - PXE MAC addresses
  - PXE IP addresses
  - PXE interfaces
  - IPMI/iLO IP address, username, password
- 3. Choose your *Book "Planning"*, *Chapter 10 "Example Configurations"* template. No action other than an understanding of your needs is necessary at this point. You will indicate which model you wish to deploy in a GUI screen. Note, however, that the only VSA options for

52 Before you begin

the *Entry-scale KVM with VSA* model is to have 0 or 3 VSA nodes. Note also that you cannot run the lifecycle manager on a dedicated node using the GUI installer; it can only run in the control plane using the models exposed in the GUI.

4. Before you use the GUI to install your cloud, you may install the operating system, Linux for HPE Helion, on your nodes (servers) if you prefer. Otherwise, the installer will install it for you.

If you are installing the operating system on all your nodes yourself, note that you must do so using the Linux for HPE Helion image that is included in the HPE Helion OpenStack 5.0 package. You may find the instructions on the following pages helpful even when using your own tools to install the operating systems:

- Chapter 10, Installing Mid-scale and Entry-scale KVM
- Chapter 11, Installation for Helion Entry Scale ESX (with OVSvApp, KVM with VSA Model)



## Note

When following the guidance for operating system installation on those pages, stop before the sections for **Running the configuration processor**. The GUI installer will run those steps, and in the case where you have the GUI also install the operating systems, it will run all the previous steps (playbooks) as well.

# 7.2 Creating a CSV file for import

You may create a CSV file for import into the GUI, indicating your server information rather than entering it directly into the GUI. The following table shows the fields you need and those that are optional.

Field	Optional	OS Install Only?	Aliases
Server ID	NO	NO	server_id, id
IP Address	NO	NO	ip, ip_address, address
MAC Address	NO	YES	mac, mac_address, mac_addr
IPMI IP Address	NO	YES	ipmi_ip, ilo_ip

IPMI User	NO	YES	ipmi_user, ilo_user, user
IPMI Password	NO	YES	Ipmi_password, ilo_password, pass- word
Server Role	NO	NO	server_role, role
Server Group	YES	NO	server_group, group
NIC Mapping	NO	NO	server_nic_map, nic_map, nic_mapping

Here "OS Install Only?" indicates that this field is ONLY needed if you have indicated that you want the installer to install HPE Linux on the servers.

The aliases are all the valid names that can be used in the CSV file for the column header for a given field. Note field names are not case sensitive and that you can use either '' (space) or '-' (hyphen) in place of underscore for a field name.

So, an example header line in the CSV file could be:

```
Server ID, ip address, mac-address, IPMI_ip, password, user, server-role, server-GROUP, nic mapping
```

The order of the fields does NOT have to match the order in the table above.

# 7.3 Run the GUI to install your cloud

The GUI is found on the lifecycle manager node at: http://ecycle manager\_IP>:79/dayze-ro

To create a secure tunnel/port forwarding into Apache from the lifecycle manager to be able to use the GUI in a browser, SSH tunnel into the lifecycle manager with the username/password set up when installing the lifecycle manager. The username and password should be what was set in "Set up the lifecycle manager" in the installation instructions. For example:

```
ssh -v -N -L 8080:<lifecycle manager_IP>:79 <Username>@<lifecycle manager_IP>
```

Then on your local machine, the one you are tunneling from, point your web browser to: http://localhost:8080/dayzero

As you proceed through the GUI, you can either enter your server names and IP addresses manually or import them from the CSV file using the CSV fields described above.

In the Assign network groups: External VM network fields, **do not** add entries for CIDR, Gateway, or IP Range, only VLAN ID. Neutron will assign these.

When the installation is complete, the Ansible log will be displayed on the final page.

# 7.4 Use the installer securely

The GUI installer is a web-based application that runs on the lifecycle manager node and is implemented in node.js. Currently this application runs behind Apache 2.0 where requests are proxied to the node.js Express Web Server. The node.js Express Server runs on port 3000 and Apache runs on port 79.

You should access the installer through Apache and not directly from the Express server. An Example URL looks like this: http://
lifecycle manager-host-ip>:79/dayzero When you finish installing, the Installer will be disabled via Ansible. It is also disabled once deployment is completed following the CLI instructions..

The GUI is used only once: during the initial installation of your cloud.

For VSA deployments, once the GUI competes its installation, you must follow the steps outlined in Section 17.2, "Configuring for VSA Block Storage Backend".

After deployment, you may also continue to Chapter 22, Cloud Verification and Chapter 28, Other Common Post-Installation Tasks.

To understand a cloud configuration more thoroughly and to learn how to make any changes later, visit these documentation pages:

•

• Chapter 3, Using Git for Configuration Management



#### Note

Because the GUI installer is disabled after deployment is complete, whether you installed your cloud via the GUI or via the command line, if you want to re-enable the GUI later but are not re-installing from scratch, run these commands:

sudo a2ensite dayzero-apache.conf

# 7.5 Final Steps

• The public Horizon Dashboard connection is over TLS, thus please use HTTPS:// to navigate to Horizon with the IP address the installer presents once install is complete.

• Visit the for Ceph post-install steps.

56 Final Steps

## 8 DNS Service Installation Overview

The HPE Helion OpenStack DNS Service supports several different backends for domain name service. The choice of backend must be included in the deployment model before the HPE Helion OpenStack install is completed.

The backends that are available within the DNS Service are separated into two categories, selfcontained and external.

**TABLE 8.1:** 

Category	Backend	Description	Recommended For
Self-con- tained	PowerDNS 3.4.1, BIND 9.9.5	All components necessary will be installed and configured as part of the HPE Helion OpenStack install.	POCs and customers who wish to keep cloud and traditional DNS separated.
External	InfoBlox	The authoritative DNS server itself is external to HPE Helion OpenStack. Management and configuration is out of scope for the lifecycle manager but remains the responsibility of the customer.	Customers who wish to integrate with their existing DNS infrastructure.

# 8.1 Getting Started

## 8.2 Install the DNS Service with PowerDNS

## 8.2.1 Install DNS Service with PowerDNS

HPE Helion OpenStack DNS Service defaults to the PowerDNS Backend if another backend is not configured for domain name service. PowerDNS will be deployed to one or more control planes clusters. The following configuration example denotes where the PowerDNS service is installed.

57 Getting Started

## 8.2.2 Configure the Backend

To configure the backend for PowerDNS, follow these steps.

1. Ensure the DNS Service components, and the PowerDNS component have been placed on a cluster. PowerDNS may be placed on a separate cluster to the other DNS Service components.

```
control-planes:
          - name: control-plane-1
          region-name: region1
          clusters:
          name: cluster1
          service-components:
          - lifecycle-manager
          - mysql
          - ip-cluster
          - apache2
          - designate-api
          - designate-central
          - designate-pool-manager
          - designate-zone-manager
          - designate-mdns
          - designate-client
          - powerdns
```

2. Edit the \_~/helion/my\_cloud/definitions/data/network\_groups.yml file to include the powerdns-ext.

```
- name: EXTERNAL-API
hostname-suffix: extapi
component-endpoints:
- powerdns-ext
load-balancers:
- provider: ip-cluster
```

3. Edit the ~/helion/my\_cloud/definitions/data/firewall\_rules.yml to allow UDP/ TCP access.

```
    name: DNSudp
        # network-groups is a list of all the network group names that the rules apply
        to
        network-groups:
```

58 Configure the Backend

```
- EXTERNAL-API
     rules:
     - type: allow
       # range of remote addresses in CIDR format that this rule applies to
       remote-ip-prefix: 0.0.0.0/0
       port-range-min: 53
       port-range-max: 53
       protocol: udp
   - name: DNStcp
     # network-groups is a list of all the network group names that the rules apply
to
     network-groups:
     - EXTERNAL-API
     rules:
     - type: allow
       # range of remote addresses in CIDR format that this rule applies to
       remote-ip-prefix: 0.0.0.0/0
       port-range-min: 53
       port-range-max: 53
       protocol: tcp
```

Please see for post installation DNS Service configuration.

## 8.3 Install the DNS Service with BIND

#### 8.3.1 Install DNS Service with BIND

HPE Helion OpenStack DNS Service and BIND can be installed together instead of the default **PowerDNS** backend. BIND will be deployed to one or more control planes clusters. The following configuration example denotes where the BIND service is installed.

## 8.3.2 Configure the Backend

Ensure the DNS Service components, and the BIND component have been placed on a cluster. BIND may be placed on a separate cluster to the other DNS Service components. Additionally, ensure the default **powerdns** component has been removed.

```
control-planes:
- name: control-plane-1
```

```
region-name: region1

clusters:
    name: cluster1
service-components:
    lifecycle-manager
    mysql
    ip-cluster
    apache2
    ...
    designate-api
    designate-central
    designate-zone-manager
    designate-zone-manager
    designate-designate-client
    bind
```

#### **Update Input Model**

Once the backend is configured, you will need to add <u>bind-ext</u> to the <u>network\_groups.yml</u> file.

1. Edit <a href="helion/my\_cloud/definition/data/network\_groups.yml">helion/my\_cloud/definition/data/network\_groups.yml</a>, add <a href="bind-ext">bind-ext</a> to compnent-endpoints.

```
name: EXTERNAL-API
hostname-suffix: extapi
component-endpoints:
- bind-ext
```

2. Save file.

## 8.4 Install the DNS Service with InfoBlox

HPE Helion OpenStack DNS Service can be installed with the **InfoBlox** backend instead of the default **PowerDNS** backend. In order to use InfoBlox as the backend, all prerequisites must be satisfied and the configuration of InfoBlox complete.



#### Note

No DNS server will be deployed onto the HPE Helion OpenStack nodes. Instead, zones will be hosted on the **InfoBlox** servers.

# Important

The InfoBlox integration was enabled and validated in HPE Helion OpenStack 3.x. To use InfoBlox with version [4.x / 5.x] we recommend that you review the InfoBlox Deployment guide which can be found here: https://www.infoblox.com/wp-content/uploads/infoblox-deployment-guide-infoblox-openstack-driver.pdf  $\mathbb{Z}$  and address any questions with your InfoBlox support contact.

## 8.4.1 Prerequisites

- 1. An existing InfoBlox system deployed within your organization.
- 2. IPs and other credentials required to access the InfoBlox WS API.
- 3. Network connectivity to and from the cluster containing designate and the InfoBlox WS API.

This information is available from your InfoBlox system administrator.

## 8.4.2 Configure the Backend

If not already present, DNS Service components must be placed on a cluster. Additionally, ensure the default **powerdns** component has been removed, and **bind** has not been added. Replace the designate-mdns component with the designate-mdns-external component.

**61** Prerequisites

```
designate-pool-managerdesignate-zone-managerdesignate-mdns-externaldesignate-client
```

You will need to provide DNS Service information details on your InfoBlox deployment. Open the designate pool-manager configuration template:

```
$ cd ~/helion/my_cloud
$ nano config/designate/pool-manager.conf.j2
```

In the config/designate/pool-manager.conf.j2, find the following code block:

```
nameservers:
    - host: <infoblox-server-ip>
        port: <infoblox-port-number>
    ...
    also_notifies:
    - host: <infoblox-server-ip>
        port: <infoblox-port-number>
```

Make the following changes:

- 1. Uncomment the block for infoblox and also\_notifies. In Jinja2, this means replacing the {# and #} markers with {{ and }}.
- 2. Fill in the API URL, username, password, and all remaining fields.
- 3. Save the file.

Once complete, the block should look like this:

```
- type: infoblox
   description: infoblox Cluster

masters:
{% if DES_PMG.consumes_DES_MDN_EXT is defined %}
{% for mdn_member in DES_PMG.consumes_DES_MDN_EXT.members.private %}
        - host: {{ mdn_member.ip_address }}
        port: {{ mdn_member.port }}

{% endfor %}
{% endif %}

options:
    wapi_url: https://<infoblox-server-ip>/wapi/v2.2.2/
    username: hos-designate
    password: MySecretPassword
```

62 Configure the Backend

```
ns_group: designate
sslverify: False
dns_view: default
network_view: default
```

You will need to inspect and commit the changes before proceeding with the deployment:

```
$diff --git a/hos/ansible/roles/designate-pool-manager/templates/pools.yaml.j2 b/hos/
ansible/roles/designate-pool-manager/templates/pools.yaml.j2
index 291c6c9..b7fb39c 100644
--- a/hos/ansible/roles/designate-pool-manager/templates/pools.yaml.j2
+++ b/hos/ansible/roles/designate-pool-manager/templates/pools.yaml.j2
@ -28,6 +28,8 @@
      priority: 2
  nameservers:
    - host: <infoblox-server-ip>
      port: <infoblox-port-number>
{% if DES_PMG.consumes_FND_PDN is defined %}
 {% for pdn_member in DES_PMG.consumes_FND_PDN.members.private %}
     - host: {{ pdn_member.ip_address }}
@ -40,7 +42,9 @@
      port: {{ bnd_member.port }}
{% endfor %}
{% endif %}
-# also_notifies:
+ also_notifies:
   - host: <infoblox-server-ip>
      port: <infoblox-port-number>
  targets:
{% if DES_PMG.consumes_FND_PDN is defined %}
     - type: powerdns
@ -89,27 +93,27 @
{% endfor %}
{% endif %}
-#
-#
      - type: infoblox
        description: infoblox Cluster
-#
-#
-#
       masters:
     - type: infoblox
      description: infoblox Cluster
      masters:
{% if DES_PMG.consumes_DES_MDN_EXT is defined %}
{% for mdn_member in DES_PMG.consumes_DES_MDN_EXT.members.private %}
```

63 Configure the Backend

```
-#
          - host: {{ mdn_member.ip_address }}
            port: {{ mdn_member.port }}
-#
         - host: {{ mdn_member.ip_address }}
           port: {{ mdn_member.port }}
 {% endfor %}
 {% endif %}
-#
-#
        options:
-#
          wapi_url: https://127.0.0.1/wapi/v2.2.2/
-#
          username: admin
-#
          password: infoblox
-#
          ns_group: designate
-#
          sslverify: False
-#
          dns_view: default
-#
          network_view: default
-#
       options:
         wapi_url: https://127.0.0.1/wapi/v2.2.2/
         username: admin
         password: infoblox
         ns_group: designate
         sslverify: False
         dns_view: default
```

#### SSL and CA Certificate

To enable SSL Verify, edit the following file:

```
$ cd ~/helion/my_cloud
$ nano config/designate/pools.yaml.j2
```

In the infoblox section, set sslverify to true:

```
sslverify: True
```

To generate a CA certificate for InfoBlox, follow these steps:

- 1. Login to the InfoBlox User Interface
- 2. Generate a self signed certificate by selecting: System->Certificate->HTTP Cert->Generate Self Signed Certificate
- 3. Provide the hostname as the InfoBlox server-ip
- 4. Reload the InfoBlox User Interface. The certificate will be loaded and can be verified through the browser.

64 Configure the Backend

- 5. If you want to download the certificate, select: System-> Certificate-> HTTP Cert-> Download the Certificate
- Copy Certificate file to ~/helion/my\_cloud/config/tls/cacerts/

Commit changes.

```
$ git commit -a -m "Configure Designate InfoBlox Backend"
```

# 8.5 Configure DNS Domain and NS Records

To configure the default DNS domain and Name Server records for the default pool, follow these steps.

- 1. Ensure that <a href="designate\_config.yml">designate\_config.yml</a> file is present in the <a href="https://www.neuron.org.yml">~/helion/my\_cloud/definition/data/designate</a> folder. If the file or folder is not present, create the folder and copy <a href="designate\_config.yml">designate\_config.yml</a> file from one of the example input models (for example, <a href="https://www.neuron.org/neuron.org/">~/helion/examples/entry-scale-esx-kvm-vsa/data/designate/designate\_config.yml</a>).
- 2. Modify the **dns\_domain** and/or **ns\_records** entries in the designate\_config.yml file.

```
data:
dns_domain: example.org.
ns_records:
   hostname: ns1.example.org.
   priority: 1
   hostname: ns2.example.org.
   priority: 2
```

3. Edit your input model's <u>control\_plane.yml</u> file to include **DESIGNATE-CONFIG-CP1** in **configuration-data** section.

```
control-planes:
    - name: control-plane-1
    region-name: region1
    lifecycle-manager-target
    configuration-data:
        - DESIGNATE-CONFIG-CP1
        - NEUTRON-CONFIG-CP1
```

4. Continue your cloud deployment by reviewing and committing your changes.

\$ git add ~/helion/my\_cloud/definition/data/designate/designate\_config.yml
\$ git commit -m "Adding DNS Domain and NS Records"



## Note

In an , you will have 3 ns\_records since the DNS service runs on all three control planes. In a or the above example would be correct since there are only two controller nodes.

# 9 Magnum Overview

The HPE Helion OpenStack Magnum Service provides container orchestration engines such as Docker Swarm, Kubernetes, and Apache Mesos available as first class resources. HPE Helion OpenStack Magnum uses Heat to orchestrate an OS image which contains Docker and Kubernetes and runs that image in either virtual machines or bare metal in a cluster configuration.

# 9.1 Magnum Architecture

As an OpenStack API service, Magnum provides Container as a Service (CaaS) functionality. Magnum is capable of working with container orchestration engines (COE) such as Kubernetes, Docker Swarm, and Apache Mesos.

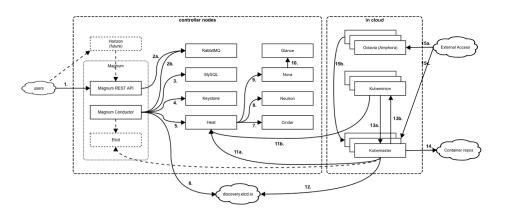
#### Components

- Magnum API: RESTful API for cluster and cluster template operations.
- Magnum Conductor: Performs operations on clusters requested by Magnum API in an asynchronous manner.
- Magnum CLI: Command-line interface to the Magnum API.
- **Etcd (planned, currently using public service)**: Remote key/value storage for distributed cluster bootstrap and discovery.
- **Kubemaster (in case of Kubernetes COE)**: One or more VM(s) or baremetal server(s), representing a control plane for Kubernetes cluster.
- **Kubeminion (in case of Kubernetes COE)**: One or more VM(s) or baremetal server(s), representing a workload node for Kubernetes cluster.
- Octavia VM aka Amphora (in case of Kubernetes COE with enabled load balancer functionality): One or more VM(s), created by HOS LBaaS v2, performing request load balancing for Kubemasters.

#### Data

Data Name	Confiden- tiality	In- tegri- ty	Avail- ability	Back- up?	Description
Session To- kens	Confiden- tial	High	Medium	No	Session tokens not stored.
System Request	Confiden- tial	High	Medium	No	Data in motion or in MQ not stored.
MySQL Data- base "Mag- num"	Confiden- tial	High	High	Yes	Contains user preferences. Backed up to Swift daily.
etcd data	Confiden- tial	High	Low	No	Kubemaster IPs and cluster info. Only used during cluster bootstrap.

## Service Architecture Diagram



#### **Interfaces**

**TABLE 9.1:** 

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Request Creden- tials	Request Au- tho- riza- tion	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
1	Exter- nal-API	HTTPS	User	Man- age clus- ters	Key- stone token	Man- age ob- jects which belong to cur- rent project	Mag- num API	Operation status with or without data	TLS Cer- tifi- cate	CRUD operations on cluster templates and clusters
2a	Inter- nal-API	AMQP over HTTPS	Mag- num API	En- queue mes- sages	Rab- bitMQ user- name, pass- word	Rab- bitMQ queue read/ write opera- tions	Rab- bitMQ	Operation status	TLS Cer- tifi- cate	Noti- fica- tions issued when cluster CRUD oper- ations re- quest- ed
2b	Inter- nal-API	AMQP over HTTPS	Mag- num Con- ductor	Read queued mes- sages	Rab- bitMQ user- name,	Rab- bitMQ queue read/ write	Rab- bitMQ	Operation status	TLS Cer- tifi- cate	Noti- fica- tions issued when

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Request Credentials	Request Authorization	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
					pass- word	opera- tions				cluster CRUD oper- ations re- quest- ed
3	Inter- nal-API	MySQL over HTTPS	Mag- num Con- ductor	Persist da- ta in MySQL	MySQL user- name, pass- word	Mag- num data- base	MySQL	Operation status with or without data	TLS Cer- tifi- cate	Persist clus- ter/clus- ter tem- plate data, read per- sisted data
4	Inter- nal-API	HTTPS	Mag- num Con- ductor	Create per- cluster user in ded- icat- ed do- main, no	Trustee do- main admin user- name, pass- word	Manage users in dedicated magnum domain	Key- stone	Operation status with or without data	TLS Cer- tifi- cate	Mag- num gern- erates user record in a dedi- cated Key-

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Request Credentials	Request Authorization	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
				role as- sign- ments initial- ly						stone do- main for each cluster
5	Inter- nal-API	HTTPS	Mag- num Con- ductor	Create per- cluster stack	Key- stone token	Limited to scope of authorized user	Heat	Operation status with or without data	TLS Cer- tifi- cate	Mag- num cre- ates Heat stack for each cluster
6	Exter- nal Net- work	HTTPS	Mag- num Con- ductor	Boot- strap a clus- ter in pub- lic dis- covery https:// discov- ery.etcd	over HTTPS. URL is only avail- able	and up- date	Pub- lic dis- covery ser- vice	Clus- ter dis- covery URL	TLS Cer- tifi- cate	Create key/ value reg- istry of spec- ified size in public stor- age. This is

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Re- quest Cre- den- tials	Re- quest Au- tho- riza- tion	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
					es need- ing it.					used to stand up a clus- ter of kuber- netes master nodes (re- fer to inter- face call #12)
7	Inter- nal-API	HTTPS	Heat En- gine	Create Cinder vol- umes	Key- stone token	Limited to scope of authorized user	Cinder API	Operation status with or without data	TLS Cer- tifi- cate	Heat cre- ates Cinder vol- umes as part of stack
8	Inter- nal-API	HTTPS	Heat En- gine	Create net-works,	Key- stone token	Limit- ed to scope	Neu- tron API	Operation status	TLS Cer-	Heat cre- ates

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Re- quest Cre- den- tials	Re- quest Au- tho- riza- tion	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
				routers, load bal- ancers		of au- tho- rized user		with or with- out data	tifi- cate	net- works, routers, load bal- ancers as part of the stack
9	Inter- nal-API	HTTPS	Heat En- gine	Create Nova VNs, attach vol- umes	Key- stone token	Limited to scope of authorized user	Nova API	Operation status with or without data	TLS Cer- tifi- cate	Heat cre- ates Nova VMs as part of the stack
10	Inter- nal-API	HTTPS	Nova	Read pre- con- fig- ured Glance image	Key- stone token	Limited to scope of authorized user	Glance API	Operation status with or without data	TLS Cer- tifi- cate	No- va is using pre- con- fig- ured im- age in

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Re- quest Cre- den- tials	Request Authorization	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion	
										Glance to boot- strap VMs	
11a	External-API	HTTPS	Cluster member (VM or Ironic node)	Heat notification	Key- stone token	Limited to scope of authorized user	Heat	Operation status with or without data	TLS Certificate	Heat is us- ing OS::Hea Con- dition re- source. VM is ex- pected to call Heat notifi- cation URL upon com- ple- tion of cer- tain boot-	t::Wait-

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Re- quest Cre- den- tials	Re- quest Au- tho- riza- tion	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
										strap opera- tion.
11b	Exter- nal-API	HTTPS	Cluster member (VM or ironic node)	Heat notifi- cation	Key-stone token	Limited to scope of authorized user	Heat	Operation status with or without data	TLS Cer- tifi- cate	Heat is us- ing OS::Hea Con- dition re- source. VM is ex- pected to call Heat notifi- cation URL upon com- ple- tion of cer- tain boot-

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Re- quest Cre- den- tials	Re- quest Au- tho- riza- tion	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
										strap opera- tion.
12	Exter- nal-API	HTTPS	Cluster member (VM or Ironic node)	Up- date cluster mem- ber state in a public reg- istry at https:// dis- cov- ery.etcd	es need-	and up- date	Pub- lic dis- covery ser- vicve	Operation status	TLS Cer- tifi- cate	Up- date key/ value pair in a reg- istry creat- ed by inter- face call #6
13a	A VxLAN encap- sulat- ed pri- vate net- work on the Guest	HTTPS	Cluster member (VM or Ironic node)	Various communications inside Kubernetes cluster	Ten- ant specif- ic	Ten- ant specif- ic	Cluster member (VM or Ironic node)	Ten- ant specif- ic	TLS Cer- tifi- cate	Various calls performed to build Kubernetes clus-

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Re- quest Cre- den- tials	Re- quest Au- tho- riza- tion	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
	net- work									ters, deploy ap- plica- tions and put work- load
13b	A VxLAN encapsulated private network on the Guest network	HTTPS	Cluster member (VM or Ironic node)	Various communications inside Kubernetes cluster	Ten- ant specif- ic	Ten- ant specif- ic	Cluster member (VM or Ironic node)	Ten- ant specif- ic	TLS Cer- tifi- cate	Various calls performed to build Kubernetes clusters, deploy applications and put work- load

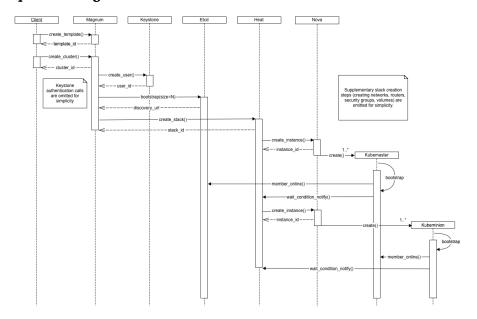
Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Re- quest Cre- den- tials	Re- quest Au- tho- riza- tion	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
14	Guest/ Exter- nal	HTTPS	Cluster member (VM or Ironic node)	Down-load con-tain-er images	None	None	Exter- nal	Con- tainer image data	TLS Cer- tifi- cate	Kubernetes is making calls to external repositories to download prepacked container images
15a	Exter- nal/EX- T_VM (Float- ing IP)	HTTPS	Ten- ant specif- ic	Ten- ant specif- ic	Ten- ant specif- ic	Ten- ant specif- ic	Oc- tavia load bal- ancer	Ten- ant specif- ic	Ten- ant specif- ic	External work-load han-dled by con-tain-

Inter- face	Net- work Name	Net- work Proto- col	Re- questor	Re- quest	Re- quest Cre- den- tials	Request Authorization	Lis- tener	Re- sponse	Re- sponse Cre- den- tials	De- scrip- tion of Op- era- tion
										er ap- plica- tions
15b	Guest	HTTPS	Ten- ant specif- ic	Ten- ant specif- ic	Ten- ant specif- ic	Ten- ant specif- ic	Cluster member (VM or Ironic node)	Ten- ant specif- ic	Ten- ant specif- ic	External work- load han- dled by con- tain- er ap- plica- tions
15c	Exter- nal/EX- T_VM (Float- ing IP)	HTTPS	Ten- ant specif- ic	Ten- ant specif- ic	Ten- ant specif- ic	Ten- ant specif- ic	Cluster member (VM or Ironic node)	Ten- ant specif- ic	Ten- ant specif- ic	External work- load han- dled by con- tain- er ap- plica- tions

## **Dependencies**

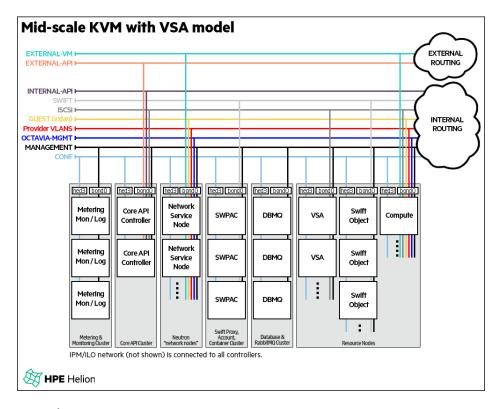
- Keystone
- RabbitMQ
- MySQL
- Heat
- Glance
- Nova
- Cinder
- Neutron
- Barbican
- Swift

### **Sequence Diagram**



## Implementation

Magnum API and Magnum Conductor are run on the HPE Helion OpenStack controllers (or core nodes in case of mid-scale deployments).



### **Security Groups**

Source CIDR/Security Group	Port/Range	Protocol	Notes
Any IP	22	SSH	Tenant Admin access
Any IP/Kubernetes Security Group	2379-2380	HTTPS	Etcd Traffic
Any IP/Kubernetes Security Group	6443	HTTPS	kube-apiserver
Any IP/Kubernetes Security Group	7080	HTTPS	kube-apiserver
Any IP/Kubernetes Security Group	8080	HTTPS	kube-apiserver
Any IP/Kubernetes Security Group	30000-32767	HTTPS	kube-apiserver

Source CIDR/Security Group	Port/Range	Protocol	Notes
Any IP/Kubernetes Security Group	any	tenant app specific	tenant app specific

#### **Network Ports**

Port/Range	Protocol	Notes
22	SSH	Admin Access
9511	HTTPS	Magnum API Access
2379-2380	HTTPS	Etcd (planned)

#### **Summary of Controls**

Summary of controls spanning multiple components and interfaces:

- Audit: Magnum performs logging. Logs are collected by the centralized logging service.
- Authentication: Authentication via Keystone tokens at APIs. Password authentication to MQ and DB using specific users with randomly-generated passwords.
- Authorization: OpenStack provides admin and non-admin roles that are indicated in session tokens. Processes run at minimum privilege. Processes run as unique user/group definitions (magnum/magnum). Appropriate filesystem controls prevent other processes from accessing service's files. Magnum config file is mode 600. Logs written using group adm, user magnum, mode 640. IPtables ensure that no unneeded ports are open. Security Groups provide authorization controls between in-cloud components.
- Availability: Redundant hosts, clustered DB, and fail-over provide high availability.
- **Confidentiality**: Network connections over TLS. Network separation via VLANs. Data and config files protected via filesystem controls. Unencrypted local traffic is bound to local-host. Separation of customer traffic on the TUL network via Open Flow (VxLANs).
- Integrity: Network connections over TLS. Network separation via VLANs. DB API integrity
  protected by SQL Alchemy. Data and config files are protected by filesystem controls.
  Unencrypted traffic is bound to localhost.

# 9.2 Install the Magnum Service

Installing the Magnum Service can be performed as part of a new HPE Helion OpenStack 5.0 environment or can be added to an existing HPE Helion OpenStack 5.0 environment. Both installations require container management services, running in Magnum cluster VMs with access to specific Openstack API endpoints. The following TCP ports need to be open in your firewall to allow access from VMs to external (public) HPE Helion OpenStack endpoints.

TCP Port	Service
5000	Idenity
8004	Heat
9511	Magnum

Additionally, Magnum is dependent on the following OpenStack services.

- Keystone
- Heat
- Nova KVM
- Neutron
- Glance
- Cinder
- Swift
- Barbican
- LBaaS v2 (Octavia) optional

# 9.2.1 Installing Magnum as part of new HPE Helion OpenStack 5.0 environment

Magnum components are already included in example HPE Helion OpenStack models based on Nova KVM, such as **entry-scale-kvm-vsa**, **entry-scale-kvm-vsa-mml** and **mid-scale**. These models contain the Magnum dependencies (see above). You can follow generic installation instruction for Mid-Scale and Entry-Scale KM model by using this guide:



#### Note

- 1. If you modify the cloud model to utilize a dedicated lifecycle manager, add <a href="mag-num-client">mag-num-client</a> item to the list of service components for the lifecycle manager cluster.
- 2. Magnum needs a properly configured external endpoint. While preparing the cloud model, ensure that <a href="external-name">external-name</a> setting in <a href="eata-network\_groups.yml">data/network\_groups.yml</a> is set to valid hostname, which can be resolved on DNS server, and a valid TLS certificate is installed for your external endpoint. For non-production test installations, you can omit <a href="external-name">external-name</a>. In test installations, the HPE Helion OpenStack installer will use an IP address as a public endpoint hostname, and automatically generate a new certificate, signed by the internal CA. Please refer to <a href="#">Chapter 25</a>, <a href="#">Configuring Transport Layer Security (TLS)</a> for more details.
- 3. To use LBaaS v2 (Octavia) for container management and container applications, follow the additional steps to configure LBaaS v2 in the guide .

# 9.2.2 Adding Magnum to existing HPE Helion OpenStack 5.0 environment

Adding Magnum to an already deployed HPE Helion OpenStack 5.0 installation or during and upgrade can be achieved by performing the following steps.

1. Add items listed below to the list of service components in <a href="https://home/stack/he-lion/my\_cloud/definition/data/control\_plane.yml">https://home/stack/he-lion/my\_cloud/definition/data/control\_plane.yml</a>. Add them to clusters which have <a href="mailto:server-role">server-role</a> set to <a href="mailto:CONTROLLER-ROLE">CONTROLLER-ROLE</a> (entry-scale models) or <a href="mailto:CORE\_ROLE">CORE\_ROLE</a> (midscale model).

- magnum-api

```
- magnum-conductor
```

- 2. If your environment utilizes a dedicated lifecycle manager, add <u>magnum-client</u> to the list of service components for the lifecycle manager.
- 3. Commit your changes to the local git repository. Run the following playbooks as described in *Chapter 3, Using Git for Configuration Management* for your Installation.
  - config-processor-run.yml
  - ready-deployment.yml
  - site.yml
- 4. Ensure that your external endpoint is configured correctly. The current public endpoint configuration can be verified by running the following commands on the lifecycle manager.

Ensure that the endpoint URL is using either an IP address, or a valid hostname, which can be resolved on the DNS server. If the URL is using an invalid hostname (for example, 'myhelion.test'), follow the steps described in *Chapter 25, Configuring Transport Layer Security (TLS)* to configure a valid external endpoint. You will need to update the <a href="external-name">external-name</a> setting in the <a href="eata/network\_groups.yml">data/network\_groups.yml</a> to a valid hostname, which can be resolved on DNS server, and provide a valid TLS certificate for the external endpoint. For non-production test installations, you can omit the <a href="external-name">external-name</a>. The HPE Helion OpenStack installer will use an IP address as public endpoint hostname, and automatically generate a new certificate, signed by the internal CA. Please refer to *Chapter 25, Configuring Transport Layer Security (TLS)* for more information.

5. Ensure that LBaaS v2 (Octavia) is correctly configured. For more information, see *Chapter 27, Configuring Load Balancer as a Service*.

# 9.3 Integrate Magnum with the DNS Service

Integration with DNSaaS may be needed if:

- 1. External endpoint is configured to use 'myhelion.test' as hostname and HPE Helion OpenStack frontend certificate is issued for this hostname.
- 2. Minions are registered using Nova VM names as hostnames Kubernetes API server. Most kubectl command won't work if the VM name (e.g. 'cl-mu3eevqizh-1-b3vifun6qtuh-kubeminion-ff4cqjgsuzhy') is not getting resolved at provided DNS server/

Follow these steps to integrate the Magnum Service with the DNS Service.

1. Allow connections from VMs to EXT-API

```
sudo modprobe 8021q
sudo ip link add link virbr5 name vlan108 type vlan id 108
sudo ip link set dev vlan108 up
sudo ip addr add 192.168.14.200/24 dev vlan108
sudo iptables -t nat -A POSTROUTING -o vlan108 -j MASQUERADE
```

2. Workaround for ????

```
$ grep enable_host_header hos/ansible/roles/designate-api/templates/api.conf.j2
enable_host_header = False
```

3. Run the designate reconfigure playbook.

```
$ cd ~/scratch/ansible/next/hos/ansible/
$ ansible-playbook -i hosts/verb_hosts designate-reconfigure.yml
```

4. Set up Designate to resolve myhelion.test correctly.

```
$ openstack zone create --email hostmaster@myhelion.test myhelion.test.
# wait for status to become active
$ EXTERNAL_VIP=$(grep HZN-WEB-extapi /etc/hosts | awk '{ print $1 }')
$ openstack recordset create --records $EXTERNAL_VIP --type A myhelion.test.
    myhelion.test.
# wait for status to become active
$ LOCAL_MGMT_IP=$(grep `hostname` /etc/hosts | awk '{ print $1 }')
$ nslookup myhelion.test $LOCAL_MGMT_IP
Server: 192.168.14.2
Address: 192.168.14.2#53
Name: myhelion.test
```

```
Address: 192.168.14.5
```

5. If you need to add/override a top level domain record:

```
$ openstack tld create --name net
$ openstack zone create --email hostmaster@proxy.houston.hpecorp.net
proxy.houston.hpecorp.net.
$ openstack recordset create --records 16.85.88.10 --type A
proxy.houston.hpecorp.net. proxy.houston.hpecorp.net.
$ nslookup proxy.houston.hpecorp.net 192.168.14.2
Server: 192.168.14.2
Address: 192.168.14.2#53
Name: proxy.houston.hpecorp.net
Address: 16.85.88.10
```

6. Enable propagation of dns\_assignment and dns\_name attributes to neutron ports, as per https://docs.openstack.org/draft/networking-guide/config-dns-int.html 🗗

7. Enable DNSaaS registration of created VMs by editing the \_~/helion/hos/ansible/roles/neutron-common/templates/neutron.conf.j2 file. You will need to add \_external\_dns\_driver = designate to the [DEFAULT] section and create a new [designate] section for the Designate specific configurations.

```
...
advertise_mtu = False
dns_domain = ksperf.
external_dns_driver = designate
{{ neutron_api_extensions_path|trim }}
```

```
{{ neutron_vlan_transparent|trim }}

# Add additional options here

[designate]
url = https://10.240.48.45:9001
admin_auth_url = https://10.240.48.45:35357/v3
admin_username = designate
admin_password = P8lZ9FdHuoW
admin_tenant_name = services
allow_reverse_dns_lookup = True
ipv4_ptr_zone_prefix_size = 24
ipv6_ptr_zone_prefix_size = 116
ca_cert = /etc/ssl/certs/ca-certificates.crt
```

#### 8. Commit your changes.

```
$ git commit -a -m "Enable DNSaaS registration of Nova VMs"
[site f4755c0] Enable DNSaaS registration of Nova VMs
1 file changed, 11 insertions(+)
```

# 10 Installing Mid-scale and Entry-scale KVM

•

•

•

•

•

•

•

•

## 10.1 Important Notes

- If you are looking for information about when to use the GUI installer and when to use the CLI, see the *Installation Overview*.
- Review the Book "Planning", Chapter 2 "Hardware and Software Support Matrix" that we have listed.
- Review the release notes to make yourself aware of any known issues and limitations.
- The installation process can occur in different phases. For example, you can install the control plane only and then add Compute nodes afterwards if you would like.
- If you run into issues during installation, we have put together a list of *Chapter 18, Troubleshooting the Installation* you can reference.
- Make sure all disks on the system(s) are wiped before you begin the install. (For Swift, refer to )
- There is no requirement to have a dedicated network for OS-install and system deployment, this can be shared with the management network. More information can be found in *Book "Planning"*, *Chapter 10 "Example Configurations"*.
- You may see the terms deployer and lifecycle manager used interchangeably. These are referring to the same nodes in your environment.

89 Important Notes

- When running the Ansible playbook in this installation guide, if a runbook fails you will see in the error response to use the <a href="e--limit">--limit</a> switch when retrying a playbook. This should be avoided. You can simply re-run any playbook without this switch.
- DVR is not supported with ESX compute.
- When you attach a Cinder volume to the VM running on the ESXi host, the volume will
  not get detected automatically. Make sure to set the image metadata vmware\_adaptertype=lsiLogicsas for image before launching the instance. This will help to discover the
  volume change appropriately.

## 10.2 Before You Start

We have put together a *Chapter 2, Pre-Installation Checklist* that should help with the recommended pre-installation tasks.

# 10.3 Setting Up the Lifecycle Manager

#### Installing the Lifecycle Manager

The lifecycle manager will contain the installation scripts and configuration files to deploy your cloud. You can set up the lifecycle manager on a dedicated node or you do so on your first controller node. The default choice is to use the first controller node as the lifecycle manager.

- 1. Sign in and download the product and signature files at the link below:
  - a. Software Entitlement Portal (http://www.hpe.com/software/entitlements)
- 2. You can verify the download was complete via the signature verification process outlined here (https://h20392.www2.hpe.com/portal/swdepot/displayProductInfo.do?product-Number=HPLinuxCodeSigning) .
- 3. Boot your lifecycle manager from the ISO contained in the download.
- 4. Enter "install" to start installation.



"install" is all lower case

90 Before You Start

- 5. Select the language. Note that only the English language selection is currently supported.
- 6. Select the location.
- 7. Select the keyboard layout.
- 8. Select the primary network interface, if prompted:
  - a. Assign IP address, subnet mask, and default gateway
- 9. Create new account:
  - a. Enter a username.
  - b. Enter a password.
  - c. Enter time zone.

Once the initial installation is finished, complete the lifecycle manager setup with these steps:

- 1. Ensure your lifecycle manager has a valid DNS nameserver specified in <a href="//etc/re-solv.conf">/etc/re-solv.conf</a>.
- 2. Set the environment variable LC\_ALL:

export LC\_ALL=C



#### Note

This can be added to ~stack/.bashrc or /etc/bash.bashrc.

At the end of this section you should have a node set up with Linux for HPE Helion on it.

#### Configure and Run the Lifecycle Manager



#### **Important**

It is critical that you don't run any of the commands below as the <u>root</u> user or use <u>sudo</u>, unless it is stated explicitly in the steps. Run then as the user you just created (or <u>stack</u> if you left the default of "stack").

 Log into your lifecycle manager as the user you created and mount the install media at / media/cdrom. It may be necessary to use wget or another file transfer method to transfer the install media to the lifecycle manager before completing this step. Here is the command to mount the media:

Example for HPE Helion OpenStack® 5.0:

```
sudo mount HelionOpenStack-5.0.iso /media/cdrom
```

2. Unpack the tarball that is in the <a href="media/cdrom/hos/">media/cdrom/hos/</a> directory: Example for HPE Helion OpenStack 5.0

```
tar xvf /media/cdrom/hos/hos-5.0.0-<timestamp>.tar
```

3. Run the hos-init.bash script which is included in the build: Example for HPE Helion OpenStack® 5.0

```
~/hos-5.0.0/hos-init.bash
```

You will be prompted to enter an optional SSH passphrase when running <a href="https://hos-init.bash">hos-init.bash</a>. This passphrase is used to protect the key used by Ansible when connecting to its client nodes. If you do not want to use a passphrase then just press <a href="mailto:Enter">Enter</a> at the prompt. For automated installation (e.g. CI) it is possible to disable SSH passphrase prompting by setting the <a href="https://hos-init.bash">HOS\_INIT\_AUTO</a> environment variable before running <a href="https://hos-init.bash">hos-init.bash</a>, like this:

```
export HOS_INIT_AUTO=y
```

If you have protected the SSH key with a passphrase then execute the following commands to avoid having to enter the passphrase on every attempt by Ansible to connect to its client nodes:

```
eval $(ssh-agent)
ssh-add ~/.ssh/id_rsa
```

At the end of this section you should have a local directory structure, as described below:

```
~/helion/
~/helion/examples/
clouds
~/helion/my_cloud/definition/
~/helion/my_cloud/config/
~/helion/hos/ansible directory
~/helion/hos/
Directory contains the config input files
Directory contains the config input files
Directory contains .j2 files which are symlinks to the
Directory contains files used by the installer
```



#### Warning

It is important that you do not add any extra files in the <u>~/helion</u> directory on your lifecycle manager. This includes temporary save files. Doing so will cause errors during the installation.

# •

#### **Important**

Be aware that the files in the \_-/helion/my\_cloud/config/ directory are symlinks to the \_-/helion/hos/ansible/ directory. For example, \_-/helion/my\_cloud/config/keepalived/defaults.yml is a symbolic link to \_-/helion/hos/ansible/roles/keepalived/defaults/main.yml.

```
ls -al ~/helion/my_cloud/config/keepalived/defaults.yml
lrwxrwxrwx 1 stack stack 55 May 24 20:38 /home/stack/helion/my_cloud/config/
keepalived/defaults.yml -> ../../hos/ansible/roles/keepalived/defaults/main.yml
```

If you are using a tool like <u>sed</u> to make edits to files in this directory, you might break the symbolic link and create a new copy of the file. To maintain the link, you will need to force **sed** to follow the link:

```
sed -i --follow-symlinks \
's$keepalived_vrrp_offset: 0$keepalived_vrrp_offset: 2$' \
~/helion/my_cloud/config/keepalived/defaults.yml
```

Alternatively, directly edit the target of the link <u>~/helion/hos/ansible/roles/</u>keepalived/defaults/main.yml.

For any troubleshooting information regarding these steps, see Section 18.1, "Issues during Lifecy-cle-manager Setup".

# 10.4 Configuring Your Environment

During the configuration phase of the installation you will be making modifications to the example configuration input files to match your cloud environment. You should use the *Book "Planning", Chapter 10 "Example Configurations"* documentation for detailed information on how to do this. There is also a <a href="README.md">README.md</a> file included in each of the example directories on the lifecycle manager that has useful information about the models.

In the steps below we show how to set up the directory structure with the example input files as well as use the optional encryption methods for your sensitive data.

- 1. Setup your configuration files, as follows:
  - a. Copy the example configuration files into the required setup directory and edit them to contain the details of your environment.
    - For example, if you want to use the Helion Mid-scale KVM with VSA model, you can use this command to copy the files to your cloud definition directory:

```
cp -r ~/helion/examples/mid-scale-kvm-vsa/* ~/helion/my_cloud/definition/
```

If you want to use the Helion Entry-scale KVM with VSA model, you can use this command to copy the files to your cloud definition directory:

```
cp -r ~/helion/examples/entry-scale-kvm-vsa/* ~/helion/my_cloud/definition/
```

- b. Begin inputting your environment information into the configuration files in the  $\sim$ /helion/my cloud/definition directory.
  - If you are using the Mid-scale or Entry-scale KVM with VSA model, see *Book "Planning"*, *Chapter 11 "Modifying the Entry-scale KVM with VSA Model for Your Environment"* for details. If you are using the Entry-scale KVM with Ceph model, see for details.
- 2. [OPTIONAL] You can use the <a href="hosencrypt.py">hosencrypt.py</a> script to encrypt your iLo passwords. This script uses OpenSSL.
  - a. Change to the Ansible directory:

```
cd ~/helion/hos/ansible
```

b. Put the encryption key into the following environment variable:

```
export HOS_USER_PASSWORD_ENCRYPT_KEY=<encryption key>
```

c. Run the python script below and follow the instructions. Enter a password that you want to encrypt.

```
./hosencrypt.py
```

- d. Take the string generated and place it in the <a href=""">"ilo-password"</a> field in your <a href="">~/he-lion/my\_cloud/definition/data/servers.yml"</a> file, remembering to enclose it in quotes.
- e. Repeat the above for each server.



#### Note

Before you run any playbooks, remember that you need to export the encryption key in the following environment variable: <a href="mailto:export HOS\_USER\_PASS-WORD\_ENCRYPT\_KEY=<encryption">encryption</a> key>

3. Commit your configuration to the local git repo (*Chapter 3, Using Git for Configuration Management*), as follows:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```



#### **Important**

This step needs to be repeated any time you make changes to your configuration files before you move onto the following steps. See *Chapter 3, Using Git for Configuration Management* for more information.

## 10.5 Provisioning Your Baremetal Nodes

To provision the baremetal nodes in your cloud deployment you can either use the automated operating system installation process provided by HPE Helion OpenStack or you can use the 3rd party installation tooling of your choice. We will outline both methods below:

**Using 3rd Party Baremetal Installers** 

If you do not wish to use the automated operating system installation tooling included with 5.0 then the requirements that have to be met using the installation tooling of your choice are:

- The operating system must be installed via the HPE Linux for HPE Helion OpenStack ISO
  provided on the Software Entitlement Portal (http://www.hpe.com/software/entitlements)
- Each node must have SSH keys in place that allows the same user from the lifecycle manager node who will be doing the deployment to SSH to each node without a password.
- Passwordless sudo needs to be enabled for the user.
- There should be a LVM logical volume as /root on each node.
- If the LVM volume group name for the volume group holding the "root" LVM logical volume is hlm-vg then it will align with the disk input models in the examples.
- Ensure that openssh-server, python, python-apt, and rsync are installed.

If you chose this method for installing your baremetal hardware, skip forward to the Chapter 10, Installing Mid-scale and Entry-scale KVM step.

If you would like to use the automated operating system installation tools provided by 5.0 then complete all of the steps below.

# Using the Automated Operating System Installation Provided by HPE Helion OpenStack Part One: Deploy Cobbler

This phase of the install process takes the baremetal information that was provided in <a href="s.yml">s.yml</a> and installs the Cobbler provisioning tool and loads this information into Cobbler. This sets each node to <a href="netboot-enabled: true">netboot-enabled: true</a> in Cobbler. Each node will be automatically marked as <a href="netboot-enabled: false">netboot-enabled: false</a> when it completes its operating system install successfully. Even if the node tries to PXE boot subsequently, Cobbler will not serve it. This is deliberate so that you can't reimage a live node by accident.

The <u>cobbler-deploy.yml</u> playbook prompts for a password - this is the password that will be encrypted and stored in Cobbler, which is associated with the user running the command on the lifecycle manager, that you will use to log in to the nodes via their consoles after install. The username is the same as the user set up in the initial dialogue when installing the lifecycle manager from the iso, and is the same user that is running the cobbler-deploy play.

1. Run the following playbook which confirms that there is iLo connectivity for each of your nodes so that they are accessible to be re-imaged in a later step:

cd ~/helion/hos/ansible

```
ansible-playbook -i hosts/localhost bm-power-status.yml
```

2. Run the following playbook to deploy Cobbler:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost cobbler-deploy.yml
```

#### Part Two: Image the Nodes

This phase of the install process goes through a number of distinct steps:

- 1. Powers down the nodes to be installed
- 2. Sets the nodes hardware boot order so that the first option is a network boot.
- 3. Powers on the nodes. (The nodes will then boot from the network and be installed using infrastructure set up in the previous phase)
- 4. Waits for the nodes to power themselves down (this indicates a success install). This can take some time.
- 5. Sets the boot order to hard disk and powers on the nodes.
- 6. Waits for the nodes to be ssh-able and verifies that they have the signature expected.

The reimage command is:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost bm-reimage.yml [-e nodelist=node1,node2,node3]
```

If a nodelist is not specified then the set of nodes in cobbler with <a href="netboot-enabled">netboot-enabled</a>: True is selected. The playbook pauses at the start to give you a chance to review the set of nodes that it is targeting and to confirm that it's correct.

You can use the command below which will list all of your nodes with the <a href="netboot-enabled">netboot-enabled</a>: True flag set:

```
sudo cobbler system find --netboot-enabled=1
```

For any troubleshooting information regarding these steps, see Section 18.2, "Issues while Provisioning your Baremetal Nodes".

## 10.6 Running the Configuration Processor

Once you have your configuration files setup you need to run the configuration processor to complete your configuration.

When you run the configuration processor you will be prompted for two passwords. Enter the first password to make the configuration processor encrypt its sensitive data, which consists of the random inter-service passwords that it generates and the ansible <a href="mailto:group\_vars">group\_vars</a> and <a href="mailto:host\_vars">host\_vars</a> that it produces for subsequent deploy runs. You will need this password for subsequent ansible deploy and configuration processor runs. If you wish to change an encryption password that you have already used when running the configuration processor then enter the new password at the second prompt, otherwise just press <a href="mailto:line">Enter</a> to bypass this.

Run the configuration processor with this command:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

For automated installation (e.g. CI) you can specify the required passwords on the ansible command line. For example, the command below will disable encryption by the configuration processor

```
ansible-playbook -i hosts/localhost config-processor-run.yml -e encrypt="" -e rekey=""
```

If you receive an error during this step then there is probably an issue with one or more of your configuration files. We recommend that you verify that all of the information in each of your configuration files is correct for your environment and then commit those changes to git using the instructions in the previous section before re-running the configuration processor again.

For any troubleshooting information regarding these steps, see Section 18.3, "Issues while Updating Configuration Files".

# 10.7 Configuring TLS



This section is optional, but recommended, for a HPE Helion OpenStack installation.

After you run the configuration processor the first time, the IP addresses for your environment will be generated and populated in the \_-/helion/my\_cloud/info/address\_info.yml file. It's at this point that you will want to take into consideration whether or not you want to configure TLS and setup a SSL certificate for your environment. Please read Chapter 25, Configuring Transport Layer Security (TLS) before proceeding for how to achieve this.

# 10.8 Deploying the Cloud

1. Use the playbook below to create a deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

2. [OPTIONAL] - Run the wipe\_disks.yml playbook to ensure all of your partitions on your nodes are completely wiped before continuing with the installation. If you are using fresh machines this step may not be necessary.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts wipe_disks.yml
```

If you have used an encryption password when running the configuration processor use the command below and enter the encryption password when prompted:

```
ansible-playbook -i hosts/verb_hosts wipe_disks.yml --ask-vault-pass
```

3. Run the site.yml playbook below:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts site.yml
```

If you have used an encryption password when running the configuration processor use the command below and enter the encryption password when prompted:

```
ansible-playbook -i hosts/verb_hosts site.yml --ask-vault-pass
```

99 Deploying the Cloud



The step above runs <u>osconfig</u> to configure the cloud and <u>hlm-deploy</u> to deploy the cloud. Therefore, this step may run for a while, perhaps 45 minutes or more, depending on the number of nodes in your environment.

4. Verify that the network is working correctly. Ping each IP in the <a href="//etc/hosts">/etc/hosts</a> file from one of the controller nodes.

For any troubleshooting information regarding these steps, see Section 18.4, "Issues while Deploying the Cloud".

# 10.9 Configuring a Block Storage Backend (Optional)

HPE Helion OpenStack supports VSA, 3PAR, and Ceph as block storage backend options. You can utilize one or more of these as setting up multiple block storage backends and multiple volume types is supported.

Regardless of whether you have a single or multiple block storage backends defined in your <a href="cinder.conf.j2">cinder.conf.j2</a> file then you can create one or more volume types using the specific attributes associated with the backend. You can find details on how to do that for each of the supported backend types here:

- Section 17.2, "Configuring for VSA Block Storage Backend"
- Section 17.1, "Ceph Overview"
- Section 17.3, "Configuring for 3PAR Block Storage Backend"

# 10.10 Post-Installation Verification and Administration

We recommend verifying the installation using the instructions in *Chapter 22, Cloud Verification*. There are also a list of other common post-installation administrative tasks listed in the *Chapter 28, Other Common Post-Installation Tasks* list.

# 11 Installation for Helion Entry Scale ESX (with OVS-vApp, KVM with VSA Model)

This document describes the procedure for the deployment of an ESX cloud using input model and preparing and deploying ESX Computes and OVSvAPPs.

It contains the following topics:

- •
- •
- •

## Important

Before you start your ESX cloud deployment make sure to you read the following instructions carefully.

## 11.1 Important Notes

- If you are looking for information about when to use the GUI installer and when to use the CLI, see the *Installation Overview*.
- Review the Book "Planning", Chapter 2 "Hardware and Software Support Matrix" that we have listed.
- Review the release notes to make yourself aware of any known issues and limitations.
- The installation process can occur in different phases. For example, you can install the control plane only and then add Compute nodes afterwards if you would like.
- If you run into issues during installation, we have put together a list of *Chapter 18, Troubleshooting the Installation* you can reference.
- Make sure all disks on the system(s) are wiped before you begin the install. (For Swift, refer to )

101 Important Notes

- There is no requirement to have a dedicated network for OS-install and system deployment, this can be shared with the management network. More information can be found in *Book "Planning"*, *Chapter 10 "Example Configurations"*.
- You may see the terms deployer and lifecycle manager used interchangeably. These are referring to the same nodes in your environment.
- When running the Ansible playbook in this installation guide, if a runbook fails you will see in the error response to use the <a href="e--limit">--limit</a> switch when retrying a playbook. This should be avoided. You can simply re-run any playbook without this switch.
- DVR is not supported with ESX compute.
- When you attach a Cinder volume to the VM running on the ESXi host, the volume will
  not get detected automatically. Make sure to set the image metadata vmware\_adaptertype=lsiLogicsas for image before launching the instance. This will help to discover the
  volume change appropriately.

### 11.2 Before You Start

We have put together a *Chapter 2, Pre-Installation Checklist* that should help with the recommended pre-installation tasks.

## 11.3 Prerequisites

ESX/vCenter integration is not fully automatic, vCenter administrators are advised of the following responsibilities to ensure secure operation:

- 1. The VMware administrator is responsible for administration of the vCenter servers and the ESX nodes using the VMware administration tools. These responsibilities include:
  - a. Installing and configuring vCenter server
  - b. Installing and configuring ESX server and ESX cluster

**102** Before You Start

- c. Installing and configuring shared datastores
- d. Establishing network connectivity between the ESX network and the HPE Helion OpenStack® management network
- 2. The VMware administration staff is responsible for the review of vCenter logs. These logs are not automatically included in HPE Helion OpenStack® centralized logging.
- 3. The VMware administrator is responsible for administration of the vCenter servers and the ESX nodes using the VMware administration tools.
- 4. Logging levels for vCenter should be set appropriately to prevent logging of the password for the HPE Helion OpenStack® message queue.
- 5. The vCenter cluster and ESX Compute nodes must be appropriately backed up.
- 6. Backup procedures for vCenter should ensure that the file containing the HPE Helion OpenStack® configuration as part of Nova and Cinder volume services is backed up and the backups are protected appropriately.
- 7. Since the file containing the HPE Helion OpenStack® message queue password could appear in the swap area of a vCenter server, appropriate controls should be applied to the vCenter cluster to prevent discovery of the password via snooping of the swap area or memory dumps
- 8. It is recommended to have a common shared storage for all the ESXi hosts in a particular cluster.
- 9. Ensure that you have enabled HA (High Availability) and DRS (Distributed Resource Scheduler) settings in a cluster configuration before running the installation. DRS/HA is disabled only for OVSvApp. This is done so that it does not move to a different host. If you do not enable DRS/HA prior to installation then you will not be able to disable it only for OVSvApp. As a result DRS/HA can migrate OVSvApp to different host, which will create a network loop.



#### Note

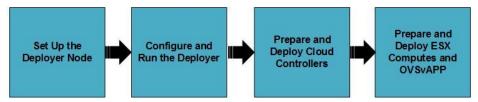
No two clusters should have the same name across datacenters in a given vCenter.

10. L3 HA VRRP is enabled by default.

**103** Prerequisites

# 11.4 Deploy ESX Cloud with OVSvApp

The following diagram depicts the procedure to configure and deploy ESX cloud.



## 11.5 Procedure to Deploy ESX Cloud with OVSvApp

### 11.5.1 Setting Up the Lifecycle Manager

#### **Installing the Lifecycle Manager**

The lifecycle manager will contain the installation scripts and configuration files to deploy your cloud. You can set up the lifecycle manager on a dedicated node or you do so on your first controller node. The default choice is to use the first controller node as the lifecycle manager.

- 1. Sign in and download the product and signature files at the link below:
  - a. Software Entitlement Portal (http://www.hpe.com/software/entitlements)
- 2. You can verify the download was complete via the signature verification process outlined here (https://h20392.www2.hpe.com/portal/swdepot/displayProductInfo.do?product-Number=HPLinuxCodeSigning) .
- 3. Boot your lifecycle manager from the ISO contained in the download.
- 4. Enter "install" to start installation.



#### Note

"install" is all lower case

- 5. Select the language. Note that only the English language selection is currently supported.
- 6. Select the location.

- 7. Select the keyboard layout.
- 8. Select the primary network interface, if prompted:
  - a. Assign IP address, subnet mask, and default gateway
- 9. Create new account:
  - a. Enter a username.
  - b. Enter a password.
  - c. Enter time zone.

Once the initial installation is finished, complete the lifecycle manager setup with these steps:

- 1. Ensure your lifecycle manager has a valid DNS nameserver specified in <a href="//etc/resolv.conf">/etc/resolv.conf</a>.
- 2. Set the environment variable LC\_ALL:

export LC\_ALL=C



Note

This can be added to ~stack/.bashrc or /etc/bash.bashrc.

At the end of this section you should have a node set up with Linux for HPE Helion on it.

#### Configure and Run the Lifecycle Manager



#### **Important**

It is critical that you don't run any of the commands below as the <u>root</u> user or use <u>sudo</u>, unless it is stated explicitly in the steps. Run then as the user you just created (or <u>stack</u> if you left the default of "stack").

 Log into your lifecycle manager as the user you created and mount the install media at / media/cdrom. It may be necessary to use wget or another file transfer method to transfer the install media to the lifecycle manager before completing this step. Here is the command to mount the media:

Example for HPE Helion OpenStack® 5.0:

```
sudo mount HelionOpenStack-5.0.iso /media/cdrom
```

2. Unpack the tarball that is in the <a href="media/cdrom/hos/">media/cdrom/hos/</a> directory: Example for HPE Helion OpenStack 5.0

```
tar xvf /media/cdrom/hos/hos-5.0.0-<timestamp>.tar
```

3. Run the hos-init.bash script which is included in the build: Example for HPE Helion OpenStack® 5.0

```
~/hos-5.0.0/hos-init.bash
```

You will be prompted to enter an optional SSH passphrase when running <a href="hos-init.bash">hos-init.bash</a>. This passphrase is used to protect the key used by Ansible when connecting to its client nodes. If you do not want to use a passphrase then just press <a href="Enter">Enter</a> at the prompt. For automated installation (e.g. CI) it is possible to disable SSH passphrase prompting by setting the <a href="https://hos-init.bash">HOS\_INIT\_AUTO</a> environment variable before running <a href="hos-init.bash">hos-init.bash</a>, like this:

```
export HOS_INIT_AUTO=y
```

If you have protected the SSH key with a passphrase then execute the following commands to avoid having to enter the passphrase on every attempt by Ansible to connect to its client nodes:

```
eval $(ssh-agent)
ssh-add ~/.ssh/id_rsa
```

At the end of this section you should have a local directory structure, as described below:

```
~/helion/
~/helion/examples/
clouds
~/helion/my_cloud/definition/
~/helion/my_cloud/config/
~/helion/hos/ansible directory
~/helion/hos/
Directory contains the config input files
Directory contains the config input files
Directory contains .j2 files which are symlinks to the
Directory contains files used by the installer
```



#### Warning

It is important that you do not add any extra files in the  $\frac{-/\text{helion}}{}$  directory on your lifecycle manager. This includes temporary save files. Doing so will cause errors during the installation.

# Important

Be aware that the files in the <u>~/helion/my\_cloud/config/</u> directory are symlinks to the <u>~/helion/hos/ansible/</u> directory. For example, <u>~/helion/my\_cloud/config/keepalived/defaults.yml</u> is a symbolic link to <u>~/helion/hos/ansible/roles/keepalived/defaults/main.yml</u>.

```
ls -al ~/helion/my_cloud/config/keepalived/defaults.yml
lrwxrwxrwx 1 stack stack 55 May 24 20:38 /home/stack/helion/my_cloud/config/
keepalived/defaults.yml -> ../../hos/ansible/roles/keepalived/defaults/main.yml
```

If you are using a tool like <u>sed</u> to make edits to files in this directory, you might break the symbolic link and create a new copy of the file. To maintain the link, you will need to force **sed** to follow the link:

```
sed -i --follow-symlinks \
's$keepalived_vrrp_offset: 0$keepalived_vrrp_offset: 2$' \
~/helion/my_cloud/config/keepalived/defaults.yml
```

Alternatively, directly edit the target of the link <u>~/helion/hos/ansible/roles/</u>keepalived/defaults/main.yml.

For any troubleshooting information regarding these steps, see Section 18.1, "Issues during Lifecy-cle-manager Setup".

## 11.5.2 Prepare and Deploy Cloud Controllers

- 1. Setup your configuration files, as follows:
  - a. Copy the example configuration files into the required setup directory and edit them as required:

```
cp -r ~/helion/examples/entry-scale-esx-kvm-vsa/* ~/helion/my_cloud/definition/
```

See a sample set of configuration files in the <u>~/helion/examples/entry-scale-esx-kvm-vsa</u> directory. The accompanying README.md file explains the contents of each of the configuration files.

b. Begin inputting your environment information into the configuration files in the \_~/ helion/my\_cloud/definition directory.



#### Note

If you want to use a dedicated deployer node in your ESX deployment, add **eon-client** service-component, to manage vCenter via EON operation from the deployer node, in the <u>control\_plane.yml</u> file as shown in the following example.

2. Commit your cloud deploy configuration to the *Chapter 3, Using Git for Configuration Management*, as follows:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```



#### Note

This step needs to be repeated any time you make changes to your configuration files before you move onto the following steps. See *Chapter 3, Using Git for Configuration Management* for more information.

### 11.5.3 Provisioning Your Baremetal Nodes

To provision the baremetal nodes in your cloud deployment you can either use the automated operating system installation process provided by HPE Helion OpenStack or you can use the 3rd party installation tooling of your choice. We will outline both methods below:

#### **Using 3rd Party Baremetal Installers**

If you do not wish to use the automated operating system installation tooling included with 5.0 then the requirements that have to be met using the installation tooling of your choice are:

- The operating system must be installed via the HPE Linux for HPE Helion OpenStack ISO
  provided on the Software Entitlement Portal (http://www.hpe.com/software/entitlements)
- Each node must have SSH keys in place that allows the same user from the lifecycle manager node who will be doing the deployment to SSH to each node without a password.
- Passwordless sudo needs to be enabled for the user.
- There should be a LVM logical volume as /root on each node.
- If the LVM volume group name for the volume group holding the "root" LVM logical volume is hlm-vg then it will align with the disk input models in the examples.
- Ensure that openssh-server, python, python-apt, and rsync are installed.

If you chose this method for installing your baremetal hardware, skip forward to the *Chapter 10, Installing Mid-scale and Entry-scale KVM* step.

If you would like to use the automated operating system installation tools provided by 5.0 then complete all of the steps below.

# Using the Automated Operating System Installation Provided by HPE Helion OpenStack Part One: Deploy Cobbler

This phase of the install process takes the baremetal information that was provided in <a href="s.yml">s.yml</a> and installs the Cobbler provisioning tool and loads this information into Cobbler. This sets each node to <a href="netboot-enabled: true">netboot-enabled: true</a> in Cobbler. Each node will be automatically marked as <a href="netboot-enabled: false">netboot-enabled: false</a> when it completes its operating system install successfully. Even if the node tries to PXE boot subsequently, Cobbler will not serve it. This is deliberate so that you can't reimage a live node by accident.

The <u>cobbler-deploy.yml</u> playbook prompts for a password - this is the password that will be encrypted and stored in Cobbler, which is associated with the user running the command on the lifecycle manager, that you will use to log in to the nodes via their consoles after install. The username is the same as the user set up in the initial dialogue when installing the lifecycle manager from the iso, and is the same user that is running the cobbler-deploy play.

1. Run the following playbook which confirms that there is iLo connectivity for each of your nodes so that they are accessible to be re-imaged in a later step:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost bm-power-status.yml
```

2. Run the following playbook to deploy Cobbler:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost cobbler-deploy.yml
```

#### Part Two: Image the Nodes

This phase of the install process goes through a number of distinct steps:

- 1. Powers down the nodes to be installed
- 2. Sets the nodes hardware boot order so that the first option is a network boot.
- 3. Powers on the nodes. (The nodes will then boot from the network and be installed using infrastructure set up in the previous phase)
- 4. Waits for the nodes to power themselves down (this indicates a success install). This can take some time.
- 5. Sets the boot order to hard disk and powers on the nodes.
- 6. Waits for the nodes to be ssh-able and verifies that they have the signature expected.

The reimage command is:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost bm-reimage.yml [-e nodelist=node1,node2,node3]
```

If a nodelist is not specified then the set of nodes in cobbler with <a href="netboot-enabled: True">netboot-enabled: True</a> is selected. The playbook pauses at the start to give you a chance to review the set of nodes that it is targeting and to confirm that it's correct.

You can use the command below which will list all of your nodes with the <a href="netboot-enabled">netboot-enabled</a>: True flag set:

```
sudo cobbler system find --netboot-enabled=1
```

For any troubleshooting information regarding these steps, see Section 18.2, "Issues while Provisioning your Baremetal Nodes".

### 11.5.4 Running the Configuration Processor

Once you have your configuration files setup you need to run the configuration processor to complete your configuration.

When you run the configuration processor you will be prompted for two passwords. Enter the first password to make the configuration processor encrypt its sensitive data, which consists of the random inter-service passwords that it generates and the ansible <a href="mailto:group\_vars">group\_vars</a> and <a href="mailto:host\_vars">host\_vars</a> that it produces for subsequent deploy runs. You will need this password for subsequent ansible deploy and configuration processor runs. If you wish to change an encryption password that you have already used when running the configuration processor then enter the new password at the second prompt, otherwise just press <a href="mailto:line">Enter</a> to bypass this.

Run the configuration processor with this command:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

For automated installation (e.g. CI) you can specify the required passwords on the ansible command line. For example, the command below will disable encryption by the configuration processor

```
ansible-playbook -i hosts/localhost config-processor-run.yml -e encrypt="" -e rekey=""
```

If you receive an error during this step then there is probably an issue with one or more of your configuration files. We recommend that you verify that all of the information in each of your configuration files is correct for your environment and then commit those changes to git using the instructions in the previous section before re-running the configuration processor again.

For any troubleshooting information regarding these steps, see Section 18.3, "Issues while Updating Configuration Files".

## 11.5.5 Deploying the Cloud

1. Use the playbook below to create a deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

2. [OPTIONAL] - Run the wipe\_disks.yml playbook to ensure all of your partitions on your nodes are completely wiped before continuing with the installation. If you are using fresh machines this step may not be necessary.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts wipe_disks.yml
```

If you have used an encryption password when running the configuration processor use the command below and enter the encryption password when prompted:

```
ansible-playbook -i hosts/verb_hosts wipe_disks.yml --ask-vault-pass
```

3. Run the site.yml playbook below:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts site.yml
```

If you have used an encryption password when running the configuration processor use the command below and enter the encryption password when prompted:

```
ansible-playbook -i hosts/verb_hosts site.yml --ask-vault-pass
```



#### Note

The step above runs <u>osconfig</u> to configure the cloud and <u>hlm-deploy</u> to deploy the cloud. Therefore, this step may run for a while, perhaps 45 minutes or more, depending on the number of nodes in your environment.

4. Verify that the network is working correctly. Ping each IP in the <a href="//etc/hosts">/etc/hosts</a> file from one of the controller nodes.

For any troubleshooting information regarding these steps, see Section 18.4, "Issues while Deploying the Cloud".

112 Deploying the Cloud

# 11.6 Prepare and Deploy ESX Computes and OVSvAPPs

The following sections describe the procedure to install and configure ESX compute and OVS-vAPPs on vCenter.

#### **Preparation for ESX Cloud Deployment**

This section describes the procedures to prepare and deploy the ESX computes and OVSvAPPs for deployment.

Login to the lifecycle manager.
 Source <u>service.osrc</u>.
 a.
 b.

#### Manage vCenters and Clusters

The following section describes the detailed procedure on managing the vCenters and clusters.

#### Register a vCenter Server



## Note

C.

- If ESX related roles are added in the input model after the controller cluster is up, you must run hlm-reconfigure.yml before adding the eon resource-manager.
- Make sure to provide single quotes ('') or backslash (\) for the special characters (&!; "'() | \>) that are used when setting a password for vCenter registration. For example: '2the!Moon' or 2the\!Moon. You can use either of the format to set the vCenter password.

vCenter provides centralized management of virtual host and virtual machines from a single console.

1. Add a vCenter using EON python client.

#### where:

- Resource Manager Name the identical name of the vCenter server.
- Resource Manager IP address the IP address of the vCenter server.
- Resource Manager Username the admin privilege username for the vCenter.
- Resource Manager Password the password for the above username.
- Resource Manager Port the vCenter server port. By default it is 443.
- Resource Manager Type specify vcenter.



### **Important**

Please do not change the vCenter Port unless you are certain it is required to do so.

#### Sample Output:

#### Show vCenter

1. Show vCenter using EON python client.

```
# eon resource-manager-show <RESOURCE_MANAGER_ID>
```

#### Sample Output:

```
eon resource-manager-show BC9DED4E-1639-481D-B190-2B54A2BF5674
+-----+
| Property
         | Value
+-----+
| Clusters | Cluster1, virtClust, Cluster2
| Datacenters | DC1, DC2
    | BC9DED4E-1639-481D-B190-2B54A2BF5674 |
| IPv4Address | 10.1.200.38
| Name | vc01
| Password | <SANITIZED>
| Port | 443
State
         | registered
| Type | vcenter
| Username | administrator@vsphere.local
```

#### Create and edit an activation template

This involves getting a sample network information template and modifying the details of the template. You will use the template to register the cloud network configuration for the vCenter.

1. Execute the following command to generate a sample activation template:

#### Sample Output:

```
eon get-activation-template --filename activationtemplate.json --type esxcluster

Saved the sample network file in /home/user/activationtemplate.json
```

2. Change to the /home/user/ directory:

```
cd /home/user/
```

3. Modify the template (json file) as per your environment. See the sample file in *Chapter 12, Sample* activationtemplate.json *File for ESX Compute*.

#### **Activate Clusters**

This involves using the activation template to register the cloud network configuration for the vCenter.

This process spawns one compute proxy VM per cluster and one OVSvApp VM per host and configures the networking as defined in the JSON template. This process also updates the input model with the service VM details, executes the ansible playbooks on the new nodes, and performs a post activation check on the service VMs.

1. Activate the cluster for the selected vCenter.

```
# eon resource-activate <RESOURCE_ID> --config-json /home/user/
<ACTIVATION_JSON_NAME>
```



#### Note

To retrieve the resource ID, execute eon resource-list.



#### Note

Activating the first cluster in a datacenter requires you to provide an activation template JSON. You can provide the same activation template or a new activation template to activate the subsequent clusters.



#### Note

Minimum disk space required for a cluster activation is (number of hosts in that cluster +1)\* 45 GB.

2. Execute the following command to view the status of the cluster:

```
| 1228fce5-df5d-445c-834e-ae633ac7e426 | Cluster2 | domain-c184 |
BC9DED4E-1639-481D-B190-2B54A2BF5674 | UNSET | UNSET | esxcluster | imported
|
| 469710f6-e9f2-48a4-aace-1f00cbd60487 | virtClust | domain-c943 |
BC9DED4E-1639-481D-B190-2B54A2BF5674 | UNSET | UNSET | esxcluster | activated
|
| a3003a32-6e3a-4d89-a072-ec64a4247fb0 | Cluster1 | domain-c21 |
BC9DED4E-1639-481D-B190-2B54A2BF5674 | UNSET | UNSET | esxcluster | imported
|
+------+
```

When the state is <u>activated</u>, the input model is updated with the service VM details, a git commit has been performed, the required playbooks and the post activation checks are completed successfully.



#### Note

In multi-region setups, monitoring services will be deployed/running in a shared control-plane, unlike the normal deployment, where all services will be in one single control-plane. During ESX cluster activation in the multi-region scenarios, <a href="mailto:gener-ate\_hosts\_file">gener-ate\_hosts\_file</a> monitoring playbooks will be skipped due to absence of monitoring services in the ESX control-plane. To overcome this, you must run the playbook manually without limit.

```
cd ~/scratch/ansible/next/hos/ansible/
ansible-playbook -i hosts/verb_hosts site.yml --tag "generate_hosts_file"
```



#### Note

Whenever an ESX compute is activated or deactivated, you must run the following playbook to reflect the changes in opsconsole.

```
cd scratch/ansible/next/hos/ansible/
ansible-playbook -i hosts/verb_hosts ops-console-reconfigure.yml
```

#### **Modify the Volume Configuration File**

Once the cluster is activated you must configure the volume.

Perform the following steps to modify the volume configuration files:

1. Change the directory. The cinder.conf.j2 is present in following directories:

```
cd /home/stack/helion/hos/ansible/roles/_CND-CMN/templates
```

OR

```
cd /home/stack/helion/my_cloud/config/cinder
```

It is recommended to modify the <a href="cinder.conf.j2">cinder.conf.j2</a> available in <a href="https://home/stack/he-lion/my cloud/config/cinder">/home/stack/he-lion/my cloud/config/cinder</a>

2. Modify the cinder.conf.j2 as follows:

```
# Configure the enabled backends
enabled backends=<unique-section-name>
# Start of section for VMDK block storage
# If you have configured VMDK Block storage for cinder you must
# uncomment this section, and replace all strings in angle brackets
# with the correct values for vCenter you have configured. You
# must also add the section name to the list of values in the
# 'enabled_backends' variable above. You must provide unique section
# each time you configure a new backend.
#[<unique-section-name>]
#vmware_api_retry_count = 10
#vmware tmp dir = /tmp
#vmware image transfer timeout secs = 7200
#vmware_task_poll_interval = 0.5
#vmware_max_objects_retrieval = 100
#vmware_volume_folder = cinder-volumes
#volume_driver = cinder.volume.drivers.vmware.vmdk.VMwareVcVmdkDriver
#vmware_host_ip = <ip_address_of_vcenter>
#vmware host username = <vcenter username>
#vmware host password = <password>
#volume_backend_name = <vmdk-backend-name>
# End of section for VMDK block storage
```

#### **Commit your Cloud Definition**

1. Add the cloud deployment definition to git:

```
cd /home/stack/helion/hos/ansible;
git add -A;
git commit -m 'Adding ESX Configurations or other commit message';
```

2. Prepare your environment for deployment:

```
ansible-playbook -i hosts/localhost config-processor-run.yml;
ansible-playbook -i hosts/localhost ready-deployment.yml;
cd /home/stack/scratch/ansible/next/hos/ansible;
```

#### Configuring VMDK block storage

Execute the following command to configure VMDK block storage:

```
ansible-playbook -i hosts/verb_hosts cinder-reconfigure.yml
```

If there are more backends (like VSA) defined, you must create and use specific volume type to ensure that volume is created in ESX, as shown below:

```
# cinder type-create "ESX VMDK Storage"
...
# cinder type-key "ESX VMDK Storage" set volume_backend_name=<name of VMDK backend
  selected during installation>
...
# cinder create --volume-type "ESX VMDK Storage" 1
...
```

## 11.6.1 Validate the block storage

You can validate that the VMDK block storage is added to the cloud successfully using the following command:

```
| cinder-scheduler | ha-volume-manager | nova | enabled | down |
2016-06-13T08:47:07.000000 |
| cinder-scheduler |
                  ha-volume-manager
                                  | nova | enabled |
2016-06-14T08:44:57.000000 |
| cinder-volume |
                  ha-volume-manager
                                  | nova | enabled |
2016-06-13T07:36:41.000000 |
             | ha-volume-manager@vmdk1 | nova | enabled |
| cinder-volume
2016-06-14T08:44:55.000000 |
| cinder-volume | ha-volume-manager@vsa1 | nova | enabled |
2016-06-14T08:44:50.000000 |
+----+
```

### 11.6.2 Validate the compute

You can validate that the ESX compute cluster is added to the cloud successfully using the following command:

```
# nova service-list
+----+
| Id | Binary
                    | Host
                                             Zone
                                                     | Status | State |
Updated at
                      | Disabled Reason |
+----+
| 3 | nova-conductor | esxhos-joh-core-m1-mgmt
                                             | internal | enabled | up
2016-03-30T12:57:47.000000 | -
                                             | internal | enabled | up
| 63 | nova-scheduler
                    | esxhos-joh-core-m1-mgmt
2016-03-30T12:57:43.000000 | -
                                             | internal | enabled | up
| 66 | nova-conductor | esxhos-joh-core-m2-mgmt
2016-03-30T12:57:48.000000 | -
| 111 | nova-conductor
                    | esxhos-joh-core-m3-mgmt
                                             | internal | enabled | up
2016-03-30T12:57:41.000000 | -
                                             | internal | enabled | up
| 129 | nova-scheduler
                    | esxhos-joh-core-m3-mgmt
2016-03-30T12:57:41.000000 | -
| 132 | nova-consoleauth | esxhos-joh-core-m1-mgmt
                                             | internal | enabled | up
2016-03-30T12:57:44.000000 | -
                                             | internal | enabled | up
| 135 | nova-scheduler
                    | esxhos-joh-core-m2-mgmt
2016-03-30T12:57:47.000000 | -
| 138 | nova-compute
                    | esxhos-joh-esx-comp0001-mgmt | nova
                                                    | enabled | up
2016-03-30T12:57:41.000000 | -
+----+
```

120 Validate the compute

You can validate the hypervisor list using the following command:

#### 11.6.3 Validate the neutron

You can validate that the networking of the ESX cluster using the following command:

```
# neutron agent-list
+-----
+-----
                           | agent_type
                                          | host
   | alive | admin state up | binary
                                         1
+-----
+-----+
| :-) | True
                     | neutron-metadata-agent
| 2b255256-9505-489e-93bf-0d37f7ff83e4 | DHCP agent
                                          | esxhos-joh-core-m3-mgmt
   | :-) | True | neutron-dhcp-agent
| esxhos-joh-esx-ovsvapp0002-
mgmt | :-) | True
                     | ovsvapp-agent
| 4f0be657-fe9e-4bab-bef8-d8e688b6573e | L3 agent
                                          | esxhos-joh-core-m3-mgmt
   | :-) | True
                    | neutron-vpn-agent
| 6db7d585-604e-4205-b29c-b9684bfb0cb2 | 0VSvApp Agent
                                          | esxhos-joh-esx-ovsvapp0001-
mgmt | :-) | True | ovsvapp-agent
| 700f1a18-d237-4605-bf54-145951bd12db | DHCP agent
                                          | esxhos-joh-core-m1-mgmt
   | :-) | True | neutron-dhcp-agent
| a723fb23-2c87-4f69-b272-643df870f0b6 | DHCP agent
                                          | esxhos-joh-core-m2-mgmt
                     | neutron-dhcp-agent
   | :-) | True
| bd512ed5-9f39-4ad6-a505-148f58cf2e64 | L3 agent
                                          | esxhos-joh-core-m1-mgmt
   | :-) | True
                     | neutron-vpn-agent
| e3d4b8be-077e-4503-bed2-031871f3829e | Open vSwitch agent | esxhos-joh-core-m2-mgmt
   | :-) | True
                     | neutron-openvswitch-agent |
| e5d75917-1d26-4443-99f4-45d8c574e3f0 | Open vSwitch agent | esxhos-joh-core-m1-mgmt
   | :-) | True
                     | neutron-openvswitch-agent |
| e6a40540-adad-4ee7-b194-db18ac7287bd | Metadata agent
                                         | esxhos-joh-core-m1-mgmt
   | :-) | True | neutron-metadata-agent |
```

121 Validate the neutron

e734a20e-159c-450b-9fb8-dc46824ef12e   Metadata agent	esxhos-joh-core-m2-mgmt
:-)   True   neutron-metadata-agent	
f019c268-a6bd-4f66-8a46-6762c720a5bd   L3 agent	esxhos-joh-core-m2-mgmt
:-)   True   neutron-vpn-agent	
fc83c8dd-ead4-4ad5-aed3-95ab41e567bc   Open vSwitch agent	t   esxhos-joh-core-m3-mgmt
:-)   True   neutron-openvswitch-agent	
+	· <del>-</del>
+	++
· ·	

122 Validate the neutron

# 12 Sample activationtemplate.json File for ESX Compute

The process for installing the Helion Entry Scale ESX, KVM with VSA Model requires creating a sample network information template and modifying the details of the template. You will use the template to register the cloud network configuration for the vCenter. The template is in the JSON format.

Execute the following command to get the activation template.

```
eon get-activation-template [--filename <Activate JSON>] --type <Type of the resource>
```

#### For example:

EON supports two network driver for ESX deployment. Based on your deployment of ESX cloud, the template (activation template. json) is automatically pre-configured specific to the driver

- 1. OVSvApp driver: For ESX with OVSvApp driver (default)
- 2. NoOp driver: For ESX with DCN driver

A sample of <u>activationtemplate.json</u> file for OVSvApp driver (default) and NoOp are shown as follows.

#### OVSvApp template

A Sample json file for OVSvApp driver where networking solution is provided by OVSvApp appliances (default).

```
{
  "input_model": {
    "server_group": "RACK1"
},
  "network_properties": {
    "switches": [
    {
        "type": "dvSwitch",
        "name": "MGMT-DVS",
        "physical_nics": "vmnic1",
        "mtu": "1500"
```

```
},
 "type": "dvSwitch",
 "name": "TRUNK-DVS",
 "physical_nics": "",
 "mtu": "1500"
}
],
"vm_config": [
 "nics": [
   {
    "device": "eth0",
   "type": "vmxnet3",
    "pci_id": "",
    "portGroup": "ESX-CONF"
   },
   "device": "eth1",
   "type": "vmxnet3",
    "pci_id": "",
   "portGroup": "MGMT"
  }
  "server_role": "ESX-COMPUTE-ROLE",
 "memory_in_mb": "4096",
 "cpu": "4"
 },
 {
  "nics": [
    "device": "eth0",
    "type": "vmxnet3",
    "pci_id": "",
    "portGroup": "ESX-CONF"
   },
   "device": "eth1",
    "type": "vmxnet3",
    "pci_id": "",
    "portGroup": "MGMT"
    "device": "eth2",
    "type": "vmxnet3",
    "pci_id": "",
    "portGroup": "GUEST"
```

```
},
   {
    "device": "eth3",
   "type": "vmxnet3",
   "pci_id": "",
    "portGroup": "TRUNK"
  }
 ],
 "server_role": "OVSVAPP-ROLE",
 "memory_in_mb": "4096",
 "cpu": "4"
}
],
"portGroups": [
 "nic_teaming": {
  "network_failover_detection": "1",
   "notify_switches": "yes",
  "load_balancing": "1",
   "active_nics": "vmnic1"
 },
 "vlan_type": "trunk",
 "vlan": "33",
 "name": "ESX-CONF",
 "switchName": "MGMT-DVS"
},
{
 "nic_teaming": {
  "network_failover_detection": "1",
   "notify_switches": "yes",
  "load_balancing": "1",
  "active_nics": "vmnic1"
 },
 "vlan_type": "none",
  "vlan": "",
 "name": "MGMT",
 "switchName": "MGMT-DVS"
},
 "vlan": "34",
 "name": "GUEST",
  "vlan_type": "trunk",
 "switchName": "MGMT-DVS",
  "nic_teaming": {
  "network_failover_detection": "1",
   "notify_switches": "yes",
   "load_balancing": "1",
```

```
"active_nics": "vmnic1"
    },
    "cloud_network_type": "vxlan"
   },
    "vlan_type": "trunk",
    "vlan": "1-4094",
   "name": "TRUNK",
    "switchName": "TRUNK-DVS"
  }
  ],
  "template_info": {
   "upload_to_cluster": false
 },
  "esx_conf_net": {
   "start_ip": "",
   "cidr": "10.20.18.0/23",
   "end_ip": "",
   "gateway": "10.20.18.1",
   "portGroup": "ESX-CONF"
 }
}
}
```

#### NoOp template

A sample json file for NoOp driver where networking solution is provided by DCN.

```
{
"input_model": {
 "server_group": "RACK1"
},
 "network_properties": {
 "switches": [
   "type": "dvSwitch",
   "name": "MGMT-DVS",
   "physical_nics": "vmnic1",
   "mtu": "1500"
  }
 ],
  "vm_config": [
   "nics": [
      "device": "eth0",
      "type": "vmxnet3",
```

```
"pci_id": "",
    "portGroup": "ESX-CONF"
   },
   "device": "eth1",
   "type": "vmxnet3",
   "pci_id": "",
   "portGroup": "MGMT"
 ],
  "server_role": "ESX-COMPUTE-ROLE",
 "memory_in_mb": "4096",
 "cpu": "4"
}
],
"portGroups": [
{
 "nic teaming": {
  "network_failover_detection": "1",
   "notify_switches": "yes",
   "load_balancing": "1",
  "active_nics": "vmnic1"
  "vlan_type": "trunk",
  "vlan": "33",
 "name": "ESX-CONF",
 "switchName": "MGMT-DVS"
 },
 {
 "nic_teaming": {
  "network_failover_detection": "1",
   "notify_switches": "yes",
  "load_balancing": "1",
  "active_nics": "vmnic1"
 },
 "vlan_type": "none",
 "vlan": "",
 "name": "MGMT",
 "switchName": "MGMT-DVS"
}
],
"template_info": {
"upload_to_cluster": false
"esx_conf_net": {
"start_ip": "",
 "cidr": "10.20.18.0/23",
```

```
"end_ip": "",
    "gateway": "10.20.18.1",
    "portGroup": "ESX-CONF"
    }
}
```

The above template consists of the following sections:

- Input model
  - Server group
- Network properties
  - Port Groups to be created or reused
  - Distributed Virtual Switches to be created or reused
  - Template information
  - Virtual machine configuration (RAM, CPU and network interfaces)
  - Ansible network

The main purpose of the activation template is to create an environment to host <a href="nova-com-pute-proxy">nova-compute-proxy</a> (in case of OVSvApp driver) or only <a href="nova-compute-proxy">nova-compute-proxy</a> (in case of NoOp driver). The environment here refers to the required virtual networking (Distributed Virtual Switch (DVS) and port groups), VMs (RAM, CPU, NIC and associated server role) going to host the Neutron and Nova services.

#### • Input model

It describes the configuration of each entities in the input model.

In the following example we have specified the value of server group as RACK1.

```
{
  "input_model": {
    "server_group": "RACK1"
},
```

#### Server group

This key takes values specified in <a href="server\_groups.yml">server\_groups.yml</a> file defined in the input model. Default server groups defined in the input model are "RACK1", "RACK2", and "RACK3" based on different network definitions. This implies that all the member hosts in the cluster are part of the same RACK.

#### Network properties

You must define and configure the following network properties.

#### Port Groups

EON creates the portGroups (PGs) if the specified PG is unavailable. The portGroups key contains a JSON list of one or more entries.

You can set the tenant network type as VLAN or VXLAN or both VLAN and VXLAN. The cloud network type must be set as per the portGroup.



#### **Important**

Do not modify the VLAN range for trunk PG. The value should remain as 1-4094

In the following example there are three portGroups with a same tenant network type with an active NIC. Also, if there is any failover in the switches then it is notified as the value is set to **yes** with the failover detection set to 1.

```
"portGroups": [
  {
    "nic teaming": {
    "network_failover_detection": "1",
     "notify_switches": "yes",
     "load_balancing": "1",
    "active nics": "vmnic1"
   },
    "vlan_type": "trunk",
    "vlan": "33",
   "name": "ESX-CONF",
   "switchName": "MGMT-DVS"
  },
    "nic teaming": {
     "network_failover_detection": "1",
     "notify_switches": "yes",
     "load_balancing": "1",
     "active_nics": "vmnic1"
```

```
},
 "vlan_type": "trunk",
 "vlan": "3307",
 "name": "MGMT",
 "switchName": "MGMT-DVS"
},
{
 "vlan": "3285",
"name": "GUEST",
 "vlan_type": "trunk",
 "switchName": "MGMT-DVS",
 "nic_teaming": {
  "network_failover_detection": "1",
  "notify_switches": "yes",
  "load_balancing": "1",
  "active_nics": "vmnic1"
},
 "cloud network type": "vxlan"
},
```

In the following example a portGroup **GUEST** has an active NIC and **both** tenant network types. Also, if there is any failover in the switches then it is notified as the value is set to **yes** with the failover detection set to 1. If you want only tenant vlan, you can define this.

In the following example a portGroup **GUEST** has an active NIC with vlan as a tenant network type. Also, if there is any failover in the switches then it is notified as the value is set to **yes** with the failover detection set to 1.

In the following example a portGroup **GUEST** has an active NIC with vxlan as a tenant network types. Also, if there is any failover in the switches then it is notified as the value is set to **yes** with the failover detection set to 1.

The following table provides the list of parameters that must be configured.

Parameter	Description
name (String)	Unique name for the portGroup to be created or used.

Parameter	Description
vlan_type (String)	The following are acceptable VLAN types:
	• none = Untagged port
	<ul> <li>vlan = VLAN port and only the VLAN ID mentioned in "VLAN" key is allowed</li> </ul>
	<ul> <li>trunk = Port allows multiple</li> <li>VLAN networks as mentioned by</li> <li>"VLAN" Key</li> </ul>
vlan	Uplinks for the switch (comma separated string). The vlan is configured on the NIC. For example, "vlan": "33" or "vlan": "0,33-35".  Example: Two uplinks "vmnic0, vm-
	nic1". Single uplink – 'vmnic0'. No uplinks – ".
switchName (String)	Name of the Distributed Virtual Switch (DVS) under which this port-group should be created.
load_balancing (Int)	Specify how to choose an uplink. Acceptable values are 1 to 5.
	Route based on the originating     virtual port
	2. Route based on IP hash
	3. Route based on source MAC hash
	4. Route based on physical NIC load
	5. Use explicit failover order

Parameter	Description
notify_switches (Boolean)	Select whether to notify switches in the case of failover.
network_failover_detection (Int)	Specify the method to use for failover detection.  • Link Status only  • Beacon Probing
active_nics	(Comma separated string) Specify the uplinks to be used.
cloud_network_type (Srting)	(Optional) Specify the tenant network type ("vlan" or "vxlan", or both ("vlan", "vxlan"). In case of VLAN, the security policy for the port group is modified to enable 'Promiscuous Mode' and 'Forged transmits'.

Configuration of PG carrying Data in the single NIC VXLAN model.

,	+   Net_conf.json Vxlan +	
tagged-vlan: true   vlanid: 33	•	 
tagged-vlan: false   vlanid:	"vlan": "",   "vlan_type":"none", +	
tagged-vlan: false   vlanid: 33     	<pre>If the physical uplink switch is tagged to particular vlan "vlan": "33", "vlan_type":"vlan",</pre>	

Configuration of PG carrying Data in the single NIC VLAN model.

```
+-----+
```

#### Distributed Virtual Switches

EON creates the virtual switch if the specified switch does not exists in the vCenter. User needs to provide the same kind of input required in creating one from the vSphere web-client.

When 2 networks (PXE & CLM) are using a single dvSwitch(DVS-MGMT) and their MTU setting are different, the MTU setting of the network which have higher value should be assigned to the dvSwitch(DVS-MGMT) in the ESX activation.json.

**OVSvApp**: A minimum of two switch entries are required. One for OVSvApp trunk DVS with no physical\_nics and other for management switch with physical\_nics. In the following example we have created two DVS with a name MGMT-DVS (with vmnic1 as the uplink) and TRUNK-DVS (with no uplink)

In the following example we have created two DVS with a name MGMT-DVS and TRUNK-DVS and used vmnic1 as the uplink for the switch.

**NoOp:** One management switch entry is required with physical\_nics In the following example we have created a switch with a name MGMT-DVS and used vmnicl as the uplink for the switch.

```
"switches": [
     {
        "type": "dvSwitch",
```

```
"name": "MGMT-DVS",
    "physical_nics": "vmnic1",
    "mtu": "1500"
    },
]
```

The following table provides the list of parameters that must be configured.

Parameter	Description
type	The category of switch. It can be Virtual Standard Switch (VSS) or Distributed Virtual Switch (DVS). The supported type of switch is DVS.
name (String)	Unique name for the switch.
physical_nics	Uplinks for the switch (comma separated string). Example: Two uplinks "vm-nic0, vmnic1". Single uplink - 'vmnic0'. No uplinks - " ".
MTU (Optional)	Maximum Transmission Unit (MTU) packet size. Provide "mtu" value within the range of 1500 to 9000.

#### Template information

This key takes a dedicated boolean value (true or false) for the parameter <u>upload\_to\_cluster</u>. If the value is set to true, every activation of a cluster will upload the OVA template to the cluster. By default, it is set to false. In this case, the OVA is uploaded to the first cluster which is being activated. The format of the template name is **hlm-shell-vm-**<**hlm\_version**>-<**dc-name**>, user given "template\_name" will not be considered.

```
"template_info": {"upload_to_cluster": false}
```

#### • Virtual machine configuration

EON creates one <u>nova-compute-proxy</u> virtual machine and 'N' number of OVSvApp virtual machines where N represents the number of ESXi hosts participating in the cluster. In other words, a OVSvApp VM will be deployed per host.

In the following example we have mentioned the server role as **ESX-Compute-Role** and **OVSVAPP-ROLE**. These roles are defined in the <a href="server\_role.yml">server\_role.yml</a> file of the input model. We require a template to create a virtual machine. 4 vCPU is configured along with 4096 MB memory. The name of the network interfaces are eth0, eth1, eth2, and eth3 and interface model is vmxnet3.

```
"vm_config": [
 {
  "nics": [
    "device": "eth0",
    "type": "vmxnet3",
    "pci_id": "",
    "portGroup": "ESX-CONF"
   },
    "device": "eth1",
    "type": "vmxnet3",
    "pci_id": "",
     "portGroup": "MGMT"
   }
  ],
   "server_role": "ESX-COMPUTE-ROLE",
  "memory_in_mb": "4096",
  "cpu": "4"
 },
 {
  "nics": [
    "device": "eth0",
    "type": "vmxnet3",
    "pci_id": "",
     "portGroup": "ESX-CONF"
   },
   {
     "device": "eth1",
    "type": "vmxnet3",
    "pci id": "",
     "portGroup": "MGMT"
   },
     "device": "eth2",
     "type": "vmxnet3",
     "pci_id": "",
     "portGroup": "GUEST"
   },
```

```
{
    "device": "eth3",
    "type": "vmxnet3",
    "pci_id": "",
    "portGroup": "TRUNK"
    }
],
    "server_role": "0VSVAPP-ROLE",
    "memory_in_mb": "4096",
    "cpu": "4"
}
],
```

EON creates only one nova-compute-proxy virtual machine for NoOp driver.

```
"vm_config": [
 {
   "nics": [
     "device": "eth0",
     "type": "vmxnet3",
     "pci_id": "",
    "portGroup": "ESX-CONF"
   },
     "device": "eth1",
     "type": "vmxnet3",
    "pci_id": "",
     "portGroup": "MGMT"
   }
  ],
   "server_role": "ESX-COMPUTE-ROLE",
  "memory_in_mb": "4096",
   "cpu": "4"
```

The following table provides the list of parameters that must be configured.

Parameter	Description
server_role (String)	Specify the role associated to this appliance in the cloud input model.
CPU (Int)	Number of vCPUs to be configured.

Parameter	Description
memory_in_mb (Int)	Amount of memory to be configured in MB.
Device (String)	Network interface name. Generally in the format eth0, eth1 ethX.
portGroup (String)	Name of the virtual network (port-group) to be attached to the interface.
type (String)	Type of interface model. Supported models are 'vmxnet3' and 'e1000'.
pci_id	Device ID of the ESXi host when passthrough for a Network Device on the host is enabled.
nics	List of network interfaces to be connected to the virtual machine.

#### Ansible network

In ESX cloud, you require a separate network for running ansible commands instead of sharing the management network. IPAM is managed by EON and assigns IP addresses to virtual machines from this network.

The following example provides a network configuration that is defined under  $\underline{\text{es-}}$  x\_conf\_net.

```
"esx_conf_net": {
    "start_ip": "",
    "cidr": "10.20.18.0/23",
    "end_ip": "",
    "gateway": "10.20.18.1",
    "portGroup": "ESX-CONF"
},
```

The following table provides the list of parameters that must be configured.

Parameter	Description
portGroup (String)	Name of the network interface for which
	IP address needs to be assigned.

Parameter	Description
CIDR (String)	Classless Inter-Domain Routing of ESX Configuration Network.
Start_ip (String)	First IPv4 allocation.
end_ip (String)	Last IPv4 allocation.
Gateway (String)	IPv4 address of the network gateway.

Return to installing .

## 13 Installing Baremetal (Ironic)

Bare Metal as a Service is enabled in this release for deployment of nova instances on bare metal nodes using flat networking. HPE DL and SL line of servers are supported.

# 13.1 Installation for HPE Helion Entry-scale Cloud with Ironic Flat Network

This page describes the installation step requirements for the HPE Helion Entry-scale Cloud with Ironic Flat Network.

#### 13.1.1 Before You Start

We have put together a *Chapter 2, Pre-Installation Checklist* that should help with the recommended pre-installation tasks.

## 13.1.2 Setting Up the Lifecycle Manager

#### **Installing the Lifecycle Manager**

The lifecycle manager will contain the installation scripts and configuration files to deploy your cloud. You can set up the lifecycle manager on a dedicated node or you do so on your first controller node. The default choice is to use the first controller node as the lifecycle manager.

- 1. Sign in and download the product and signature files at the link below:
  - a. Software Entitlement Portal (http://www.hpe.com/software/entitlements)
- 2. You can verify the download was complete via the signature verification process outlined here (https://h20392.www2.hpe.com/portal/swdepot/displayProductInfo.do?product-Number=HPLinuxCodeSigning) .
- 3. Boot your lifecycle manager from the ISO contained in the download.
- 4. Enter "install" to start installation.



#### "install" is all lower case

- 5. Select the language. Note that only the English language selection is currently supported.
- 6. Select the location.
- 7. Select the keyboard layout.
- 8. Select the primary network interface, if prompted:
  - a. Assign IP address, subnet mask, and default gateway
- 9. Create new account:
  - a. Enter a username.
  - b. Enter a password.
  - c. Enter time zone.

Once the initial installation is finished, complete the lifecycle manager setup with these steps:

- 1. Ensure your lifecycle manager has a valid DNS nameserver specified in <a href="//etc/resolv.conf">/etc/resolv.conf</a>.
- 2. Set the environment variable LC\_ALL:

export LC\_ALL=C



#### Note

This can be added to ~stack/.bashrc or /etc/bash.bashrc.

At the end of this section you should have a node set up with Linux for HPE Helion on it.

Configure and Run the Lifecycle Manager

## Important

It is critical that you don't run any of the commands below as the <u>root</u> user or use <u>sudo</u>, unless it is stated explicitly in the steps. Run then as the user you just created (or <u>stack</u> if you left the default of "stack").

1. Log into your lifecycle manager as the user you created and mount the install media at / media/cdrom. It may be necessary to use wget or another file transfer method to transfer the install media to the lifecycle manager before completing this step. Here is the command to mount the media:

Example for HPE Helion OpenStack® 5.0:

```
sudo mount HelionOpenStack-5.0.iso /media/cdrom
```

2. Unpack the tarball that is in the <a href="media/cdrom/hos/">/media/cdrom/hos/</a> directory: Example for HPE Helion OpenStack 5.0

```
tar xvf /media/cdrom/hos/hos-5.0.0-<timestamp>.tar
```

3. Run the hos-init.bash script which is included in the build: Example for HPE Helion OpenStack® 5.0

```
~/hos-5.0.0/hos-init.bash
```

You will be prompted to enter an optional SSH passphrase when running <a href="hos-init.bash">hos-init.bash</a>. This passphrase is used to protect the key used by Ansible when connecting to its client nodes. If you do not want to use a passphrase then just press <a href="Enter">Enter</a> at the prompt. For automated installation (e.g. CI) it is possible to disable SSH passphrase prompting by setting the <a href="https://hos-init.bash">HOS\_INIT\_AUTO</a> environment variable before running <a href="hos-init.bash">hos-init.bash</a>, like this:

```
export HOS_INIT_AUTO=y
```

If you have protected the SSH key with a passphrase then execute the following commands to avoid having to enter the passphrase on every attempt by Ansible to connect to its client nodes:

```
eval $(ssh-agent)
ssh-add ~/.ssh/id_rsa
```

At the end of this section you should have a local directory structure, as described below:

```
~/helion/
~/helion/examples/
clouds
~/helion/my_cloud/definition/
~/helion/my_cloud/config/
~/helion/hos/ansible directory
~/helion/hos/
Top level directory
Contains the config input files of the example clouds
Top level directory
Contains the config input files
Directory contains .j2 files which are symlinks to the
Top level directory
Contains the config input files
Top level directory
Top level dir
```



#### Warning

It is important that you do not add any extra files in the  $\frac{\text{-/helion}}{\text{-/helion}}$  directory on your lifecycle manager. This includes temporary save files. Doing so will cause errors during the installation.

#### **Important**

Be aware that the files in the <u>~/helion/my\_cloud/config/</u> directory are symlinks to the <u>~/helion/hos/ansible/</u> directory. For example, <u>~/helion/my\_cloud/config/keepalived/defaults.yml</u> is a symbolic link to <u>~/helion/hos/ansible/roles/keepalived/defaults/main.yml</u>.

```
ls -al ~/helion/my_cloud/config/keepalived/defaults.yml
lrwxrwxrwx 1 stack stack 55 May 24 20:38 /home/stack/helion/my_cloud/config/
keepalived/defaults.yml -> ../../hos/ansible/roles/keepalived/defaults/main.yml
```

If you are using a tool like <u>sed</u> to make edits to files in this directory, you might break the symbolic link and create a new copy of the file. To maintain the link, you will need to force **sed** to follow the link:

```
sed -i --follow-symlinks \
's$keepalived_vrrp_offset: 0$keepalived_vrrp_offset: 2$' \
~/helion/my_cloud/config/keepalived/defaults.yml
```

Alternatively, directly edit the target of the link <u>~/helion/hos/ansible/roles/</u>keepalived/defaults/main.yml.

For any troubleshooting information regarding these steps, see Section 18.1, "Issues during Lifecy-cle-manager Setup".

## 13.1.3 Configure Your Environment

Prior to deploying an operational environment with Ironic, operators need to be aware of the nature of TLS certificate authentication. As pre-built deployment agent ramdisks images are supplied, these ramdisk images will only authenticate known third-party TLS Certificate Authorities in the interest of end-to-end security. As such, uses of self-signed certificates and private certificate authorities will be unable to leverage ironic without modifying the supplied ramdisk images.

- 1. Setup your configuration files, as follows:
  - a. See the sample sets of configuration files in the <u>~/helion/examples/</u> directory. Each set will have an accompanying README.md file that explains the contents of each of the configuration files.
  - b. Copy the example configuration files into the required setup directory and edit them to contain the details of your environment:

```
cp -r \sim/helion/examples/entry-scale-ironic-flat-network/* \sim/helion/my_cloud/definition/
```

- 2. [OPTIONAL] You can use the <a href="hosencrypt.py">hosencrypt.py</a> script to encrypt your iLo passwords. This script uses OpenSSL.
  - a. Change to the Ansible directory:

```
cd ~/helion/hos/ansible
```

b. Put the encryption key into the following environment variable:

```
export HOS_USER_PASSWORD_ENCRYPT_KEY=<encryption key>
```

c. Run the python script below and follow the instructions. Enter a password that you want to encrypt.

```
./hosencrypt.py
```

- d. Take the string generated and place it in the <a href=""">"ilo-password"</a> field in your <a href="">~/he-lion/my\_cloud/definition/data/servers.yml">-/he-lion/my\_cloud/definition/data/servers.yml</a> file, remembering to enclose it in quotes.
- e. Repeat the above for each server.



Before you run any playbooks, remember that you need to export the encryption key in the following environment variable: <a href="mailto:export HOS\_USER\_PASS-WORD\_ENCRYPT\_KEY=<encryption">encryption</a> key>

3. Commit your configuration to the local git repo (*Chapter 3, Using Git for Configuration Management*), as follows:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```



#### **Important**

This step needs to be repeated any time you make changes to your configuration files before you move onto the following steps. See *Chapter 3, Using Git for Configuration Management* for more information.

## 13.1.4 Provisioning Your Baremetal Nodes

To provision the baremetal nodes in your cloud deployment you can either use the automated operating system installation process provided by HPE Helion OpenStack or you can use the 3rd party installation tooling of your choice. We will outline both methods below:

#### **Using 3rd Party Baremetal Installers**

If you do not wish to use the automated operating system installation tooling included with 5.0 then the requirements that have to be met using the installation tooling of your choice are:

- The operating system must be installed via the HPE Linux for HPE Helion OpenStack ISO provided on the Software Entitlement Portal (http://www.hpe.com/software/entitlements) .
- Each node must have SSH keys in place that allows the same user from the lifecycle manager node who will be doing the deployment to SSH to each node without a password.
- Passwordless sudo needs to be enabled for the user.
- There should be a LVM logical volume as /root on each node.

- If the LVM volume group name for the volume group holding the "root" LVM logical volume is hlm-vg then it will align with the disk input models in the examples.
- Ensure that openssh-server, python, python-apt, and rsync are installed.

If you chose this method for installing your baremetal hardware, skip forward to the Chapter 10, Installing Mid-scale and Entry-scale KVM step.

If you would like to use the automated operating system installation tools provided by 5.0 then complete all of the steps below.

## Using the Automated Operating System Installation Provided by HPE Helion OpenStack Part One: Deploy Cobbler

This phase of the install process takes the baremetal information that was provided in <a href="s.yml">s.yml</a> and installs the Cobbler provisioning tool and loads this information into Cobbler. This sets each node to <a href="netboot-enabled: true">netboot-enabled: true</a> in Cobbler. Each node will be automatically marked as <a href="netboot-enabled: false">netboot-enabled: false</a> when it completes its operating system install successfully. Even if the node tries to PXE boot subsequently, Cobbler will not serve it. This is deliberate so that you can't reimage a live node by accident.

The <u>cobbler-deploy.yml</u> playbook prompts for a password - this is the password that will be encrypted and stored in Cobbler, which is associated with the user running the command on the lifecycle manager, that you will use to log in to the nodes via their consoles after install. The username is the same as the user set up in the initial dialogue when installing the lifecycle manager from the iso, and is the same user that is running the cobbler-deploy play.

1. Run the following playbook which confirms that there is iLo connectivity for each of your nodes so that they are accessible to be re-imaged in a later step:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost bm-power-status.yml
```

2. Run the following playbook to deploy Cobbler:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost cobbler-deploy.yml
```

Part Two: Image the Nodes

This phase of the install process goes through a number of distinct steps:

- 1. Powers down the nodes to be installed
- 2. Sets the nodes hardware boot order so that the first option is a network boot.
- 3. Powers on the nodes. (The nodes will then boot from the network and be installed using infrastructure set up in the previous phase)
- 4. Waits for the nodes to power themselves down (this indicates a success install). This can take some time.
- 5. Sets the boot order to hard disk and powers on the nodes.
- 6. Waits for the nodes to be ssh-able and verifies that they have the signature expected.

The reimage command is:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost bm-reimage.yml [-e nodelist=node1,node2,node3]
```

If a nodelist is not specified then the set of nodes in cobbler with <a href="netboot-enabled">netboot-enabled</a>: <a href="True">True</a> is selected. The playbook pauses at the start to give you a chance to review the set of nodes that it is targeting and to confirm that it's correct.

You can use the command below which will list all of your nodes with the <a href="netboot-enabled">netboot-enabled</a>: True flag set:

```
sudo cobbler system find --netboot-enabled=1
```

For any troubleshooting information regarding these steps, see Section 18.2, "Issues while Provisioning your Baremetal Nodes".

## 13.1.5 Running the Configuration Processor

Once you have your configuration files setup you need to run the configuration processor to complete your configuration.

When you run the configuration processor you will be prompted for two passwords. Enter the first password to make the configuration processor encrypt its sensitive data, which consists of the random inter-service passwords that it generates and the ansible <a href="group\_vars">group\_vars</a> and host vars that it produces for subsequent deploy runs. You will need this password for sub-

sequent ansible deploy and configuration processor runs. If you wish to change an encryption password that you have already used when running the configuration processor then enter the new password at the second prompt, otherwise just press [Enter] to bypass this.

Run the configuration processor with this command:

```
cd ~/helion/hos/ansible ansible-playbook -i hosts/localhost config-processor-run.yml
```

For automated installation (e.g. CI) you can specify the required passwords on the ansible command line. For example, the command below will disable encryption by the configuration processor

```
ansible-playbook -i hosts/localhost config-processor-run.yml -e encrypt="" -e rekey=""
```

If you receive an error during this step then there is probably an issue with one or more of your configuration files. We recommend that you verify that all of the information in each of your configuration files is correct for your environment and then commit those changes to git using the instructions in the previous section before re-running the configuration processor again.

For any troubleshooting information regarding these steps, see Section 18.3, "Issues while Updating Configuration Files".

## 13.1.6 Deploying the Cloud

1. Use the playbook below to create a deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

2. [OPTIONAL] - Run the wipe\_disks.yml playbook to ensure all of your partitions on your nodes are completely wiped before continuing with the installation. If you are using fresh machines this step may not be necessary.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts wipe_disks.yml
```

If you have used an encryption password when running the configuration processor use the command below and enter the encryption password when prompted:

```
ansible-playbook -i hosts/verb_hosts wipe_disks.yml --ask-vault-pass
```

148 Deploying the Cloud

3. Run the site.yml playbook below:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts site.yml
```

If you have used an encryption password when running the configuration processor use the command below and enter the encryption password when prompted:

```
ansible-playbook -i hosts/verb_hosts site.yml --ask-vault-pass
```



#### Note

The step above runs <u>osconfig</u> to configure the cloud and <u>hlm-deploy</u> to deploy the cloud. Therefore, this step may run for a while, perhaps 45 minutes or more, depending on the number of nodes in your environment.

4. Verify that the network is working correctly. Ping each IP in the <a href="//etc/hosts">/etc/hosts</a> file from one of the controller nodes.

For any troubleshooting information regarding these steps, see Section 18.4, "Issues while Deploying the Cloud".

## 13.1.7 Ironic configuration

Run the ironic-cloud-configure.yml playbook below:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ironic-cloud-configure.yml
```

This step configures ironic flat network, uploads glance images and sets the ironic configuration.

To see the images uploaded to glance, run:

```
$ source ~/service.osrc
$ glance image-list
```

This will produce output like the following example, showing three images that have been added by Ironic:

149 Ironic configuration

```
| b759a1f0-3b33-4173-a6cb-be5706032124 | ir-deploy-kernel-H0S3.0 |
| ce5f4037-e368-46f2-941f-c01e9072676c | ir-deploy-ramdisk-H0S3.0 |
+-----+
```

To see the network created by Ironic, run:

```
$ neutron net-list
```

This returns details of the "flat-net" generated by the Ironic install:

## 13.1.8 Node Configuration

#### 13.1.8.1 DHCP

Once booted, nodes obtain network configuration via DHCP. If multiple interfaces are to be utilized, a user may wish to pre-build images with settings to execute DHCP on all interfaces. An easy utility to build custom images is the diskimage-builder (https://github.com/openstack/diskimage-builder) utility which has a dhcp-all-interfaces (https://github.com/openstack/diskimage-builder/tree/master/elements/dhcp-all-interfaces) element that can be utilized to initiate DHCP on all interfaces.

The diskimage-builder tool does not support building UEFI images. The following patchset can be used to build UEFI images: https://review.openstack.org/#/c/287784 ♣.

### 13.1.8.2 Configuration Drives



#### Caution

Configuration Drives are stored unencrypted and should not include any sensitive data.

150 Node Configuration

Users may wish to utilize Configuration Drives (http://docs.openstack.org/user-guide/cli\_config\_drive.html) to store metadata for initial boot setting customization. While Configuration Drives are extremely useful for initial machine configuration, as a general security practice, they should not include any sensitive data. Configuration Drives should only be trusted upon the initial boot of an instance. <a href="cloud-init">cloud-init</a> utilizes a lock file for this purpose. Custom instance images should not rely upon the integrity of a Configuration Drive beyond the initial boot of a host as an administrative user with-in a deployed instance can potentially modify a configuration drive once written to disk and released for use.

## 13.1.9 TLS Certificates with Ironic Python Agent (IPA) Images

As part of HPE Helion OpenStack 5.0, Ironic Python Agent, better known as IPA in the OpenStack community, images are supplied and loaded into glance. Two types of image exist. One is a traditional boot ramdisk which is used by the <u>agent\_ipmitool</u>, <u>pxe\_ipmitool</u>, and <u>pxe\_ilo</u> drivers. The other is an ISO image that is supplied as virtual media to the host when using the agent\_ilo driver.

As these images are built in advance, they are unaware of any private certificate authorities. Users attempting to utilize self-signed certificates or a private certificate authority will need to inject their signing certificate(s) into the image in order for IPA to be able to boot on a remote node, and ensure that the TLS endpoints being connected to in HPE Helion OpenStack can be trusted. This is not an issue with publicly signed certificates.

As two different types of images exist, below are instructions for disassembling the image ramdisk file or the ISO image. Once this has been done, you will need to re-upload the files to glance, and update any impacted node's <u>driver\_info</u>, for example, the <u>deploy\_ramdisk</u> and <u>ilo\_deploy\_iso</u> settings that were set when the node was first defined. Respectively, this can be done with the

```
ironic node-update <node> replace driver_info/deploy_ramdisk=<glance_id>
```

or

```
ironic node-update <node> replace driver_info/ilo_deploy_iso=<glance_id>
```

### 13.1.9.1 Adding a certificate into a ramdisk image

As root, from a folder where the ramdisk image is present, perform the following steps.

- 1. Download the ramdisk image to /tmp/ and name the file ironic-deploy.initramfs
- 2. Change your working directory to /tmp/

```
cd /tmp/
```

3. Create a temporary folder that will hold the temporarily extracted image contents.

```
mkdir new_deploy_ramdisk
```

4. Change your shell working directory to the temporary folder

```
cd /tmp/new_deploy_ramdisk
```

5. Extract the original deployment archive file.

```
zcat /tmp/ironic-deploy.initramfs | cpio --extract --make-directories
```

6. Append your CA certificate to the file located at <a href="https://example.cal/lib/python2.7/dist-packages/requests/cacert.pem">/tmp/new\_deploy\_ramdisk/usr/lo-cal/lib/python2.7/dist-packages/requests/cacert.pem</a>, example below:

```
cat your_ca_certificate.pem >> /tmp/new_deploy_ramdisk/usr/local/lib/python2.7/dist-
packages/requests/cacert.pem
```

your\_ca\_certificate.pem is /etc/ssl/certs/ca-certificates.crt which can be
obtained through

```
cat service.osrc | grep CACERT
```

7. Package a new ramdisk file. From with-in the <a href="http://tmp/new\_deploy\_ramdisk">/tmp/new\_deploy\_ramdisk</a> folder, execute the following command:

```
find . | cpio --create --format='newc' | gzip -c -9 > /tmp/updated-ironic-
deploy.initramfs
```

Once completed, the new image can be found at <a href="///tmp/updated-ironic-deploy.initramfs">/tmp/updated-ironic-deploy.initramfs</a>, and will need to be uploaded to Glance. New Glance IDs will need to be recorded for any instances requiring this new image, as noted in the parent paragraph.

#### 13.1.9.2 Adding a certificate into an ISO image

As root, proceed with the following steps to disassemble the ISO image, and insert a new Certificate Authority or Signing Certificate into the provided image.

- 1. Download the deployment ISO image from glance and place it in <a href="tmp/">/tmp/</a> as <a href="ironic-de-ploy.iso">ironic-de-ploy.iso</a>.
- 2. Make a temporary folder to mount the <u>ironic-deploy.iso</u> image to in order to extract its contents.

```
mkdir /mnt/deploy_iso
```

3. Mount the ironic-deploy.iso image file to the folder you just created.

```
mount -o loop /tmp/ironic-deploy.iso /mnt/deploy_iso/
```

4. Make a temporary location to stage the new image.

```
mkdir /tmp/new_deploy_iso
```

5. Copy the contents from the previous image into the new temporary staging location.

```
cp -ra /mnt/deploy_iso/. /tmp/new_deploy_iso/
```

6. Make a folder to house the ramdisk that will need to be extracted from the ISO image.

```
mkdir /tmp/new_deploy_ramdisk
```

7. Change your working directory to /tmp/new\_deploy\_ramdisk

```
cd /tmp/new_deploy_ramdisk
```

8. Unpack the ramdisk image.

```
zcat /tmp/new_deploy_iso/initrd | cpio --extract --make-directories
```

9. Append your CA certificate to the file located at <a href="https://example.cal/lib/python2.7/dist-packages/requests/cacert.pem">/tmp/new\_deploy\_ramdisk/usr/local/lib/python2.7/dist-packages/requests/cacert.pem</a>, example below:

```
cat your_ca_certificate.pem >> /tmp/new_deploy_ramdisk/usr/local/lib/python2.7/dist-
packages/requests/cacert.pem
```

10. Replace the pre-existing initial ramdisk in the extracted disk image.

```
find . | cpio --create --format='newc' | gzip -c -9 > /tmp/new_deploy_iso/initrd
```

11. Change directory back to the new deployment ISO image folder.

```
cd /tmp/new_deploy_iso
```

- 12. Ensure that the program <u>xorriso</u> is available, you may need to execute <u>apt-get install</u> xorriso to install it.
- 13. Create the new ISO deployment image utilizing the following command.

```
xorriso -as mkisofs -b isolinux/isolinux.bin -c boot.cat -V VMEDIA_BOOT_ISO -r
-J -no-emul-boot -boot-load-size 4 -boot-info-table -eltorito-alt-boot --efi-
boot boot/grub/efi.img -isohybrid-gpt-basdat -isohybrid-apm-hfsplus -o /tmp/
new_ironic_deploy.iso ./
```

Once completed, the new image can be found at <a href="tmp/new\_ironic\_deploy.iso">/tmp/new\_ironic\_deploy.iso</a>, and will need to be uploaded to Glance. New Glance IDs will need to be recorded for any instances requiring this new image, as noted in the parent paragraph

## 13.2 Ironic in Multiple Control Plane

HPE Helion OpenStack 5.0 introduces the concept of multiple control planes and multiple regions - see the Input Model documentation for the relevant and . This document covers the use of an Ironic region in a multiple control plane cloud model in HPE Helion OpenStack.

## 13.2.1 Networking for Baremetal in Multiple Control Plane

**IRONIC-FLAT-NET** is the network configuration for baremetal control plane.

You need to set the environment variable **OS\_REGION\_NAME** to the ironic region in baremetal control plane. This will setup the ironic flat networking in neutron.

```
export OS_REGION_NAME=<ironic_region>
```

To see details of the <u>IRONIC-FLAT-NETWORK</u> created during configuration, use the following command:

```
neutron net-list
```

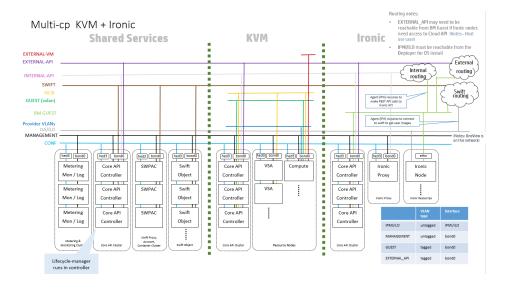


FIGURE 13.1: ARCHITECTURE OF MULTIPLE CONTROL PLANE WITH IRONIC

## 13.2.2 Handling Optional Swift Service

Swift is very resource-intensive and as a result, it is now optional in the HPE Helion OpenStack control plane. A number of services depend on Swift, and if it is not present, they must provide a fallback strategy. Fore example, Glance can use the filesystem in place of Swift for its backend store.

In Ironic, agent-based drivers require Swift - if it is not present, it is necessary to disable access to this Ironic feature in the control plane. The <a href="mailto:enable\_agent\_driver">enable\_agent\_driver</a> flag has been added to the ironic configuration data and can have values of <a href="mailto:true">true</a> or <a href="mailto:false">false</a>. Setting this flag to <a href="mailto:false">false</a> will disable swift configurations and the agent based drivers in the ironic control plane.

## 13.2.3 Instance Provisioning

In a multiple control plane cloud setup, changes for glance container name in the swift name-space of <u>ironic-conductor.conf</u> introduces a conflict with the one in <u>glance-api.conf</u>. Provisioning with agent based drivers requires the container name to be the same in ironic and glance. Hence, on instance provisioning with agent based drivers (swift enabled), the agent is not able to fetch the images from glance store and fails at that point.

You can resolve this issue using the following steps:

- 1. Copy the value of <a href="mailto:swift\_store\_container">swift\_store\_container</a> from the file <a href="mailto://opt/stack/service/glance-api/etc/glance-api.conf">/opt/stack/service/glance-api/etc/glance-api.conf</a>
- 2. Login to the lifecycle manager and use the value for <a href="mailto:swift\_container">swift\_container</a> in glance namespace of <a href="mailto:~/container">~/scratch/ansible/next/hos/ansible/roles/ironic-common/tem-plates/ironic-conductor.conf.j2</a>
- 3. Run the following playbook:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ironic-reconfigure.yml
```

## 13.3 Provisioning Baremetal Nodes with Flat Network Model

## Important

The steps outlined in **must** be performed.

A number of drivers are available to provision and manage bare-metal machines. The drivers are named based on the deployment mode and the power management interface. HPE Helion OpenStack has been tested with the following drivers:

- agent\_ilo
- agent\_ipmi
- pxe\_ilo
- pxe\_ipmi

For HPE servers, agent ipmitool and agent ilo drivers are supported.

Before you start, you should be aware that:

- 1. Node Cleaning is enabled for all the drivers in HPE Helion OpenStack 5.0.
- 2. Node parameter settings must have matching flavors in terms of cpus, local\_gb, and memory\_mb,boot\_mode and cpu\_arch.

- 3. It is advisable that nodes enrolled for ipmitool drivers are pre-validated in terms of bios settings, in terms of boot mode, prior to setting capabilities.
- 4. Network cabling and interface layout should also be pre-validated in any given particular boot mode or configuration that is registered.
- 5. The use of <u>agent</u> drivers is predicated upon Glance images being backed by a Swift image store, specifically the need for the temporary file access features. Use of Ceph or the filesystem as a Glance back-end image store means that the agent\_ drivers cannot be used.
- 6. Manual Cleaning (RAID) and Node inspection is supported by ilo drivers (<u>agent\_ilo</u> and pxe ilo)

## 13.3.1 Supplied Images

As part of the Helion Entry-scale Ironic Cloud installation, Ironic Python Agent (IPA) images are supplied and loaded into Glance. To see the images that have been loaded, execute the following commands on the deployer node:

```
$ source ~/service.osrc
glance image-list
```

This will produce output like the following example, showing three images that have been added by Ironic:

The <u>ir-deploy-ramdisk-HOS4.0</u> image is a traditional boot ramdisk used by the <u>agent\_ip-mitool</u>, <u>pxe\_ipmitool</u>, and <u>pxe\_ilo</u> drivers while <u>ir-deploy-iso-HOS4.0</u> is an ISO image that is supplied as virtual media to the host when using the agent\_ilo driver.

## 13.3.2 Provisioning a node

The information required to provision a node varies slightly depending on the driver used. In general the following details are required.

157 Supplied Images

- Network access information and credentials to connect to the management interface of the node.
- Sufficient properties to allow for nova flavor matching.
- A deployment image to perform the actual deployment of the guest operating system to the baremetal node.

A combination of the <u>ironic node-create</u> and <u>ironic node-update</u> commands are used for registering a node's characteristics with the Ironic service. In particular, <u>ironic node-update</u> <nodeid> <u>add</u> and <u>ironic node-update</u> <nodeid> <u>replace</u> can be used to modify the properties of a node after it has been created while <u>ironic node-update</u> <nodeid> <u>remove</u> will remove a property.

## 13.3.3 Creating a node using agent\_ilo

If you want to use a boot mode of BIOS as opposed to UEFI, then you need to ensure that the boot mode has been set correctly on the iLO:

While the iLO driver can automatically set a node to boot in UEFI mode via the <u>boot\_mode</u> defined capability, it cannot set BIOS boot mode once UEFI mode has been set.

Use the <u>ironic node-create</u> command to specify the <u>agent\_ilo</u> driver, network access and credential information for the iLO, properties of the node and the Glance ID of the supplied ISO IPA image. Note that memory size is specified in megabytes while disk size is gigabytes.

```
ironic node-create -d agent_ilo -i ilo_address=10.1.196.117 -i ilo_username=Administrator
  -i ilo_password=******* \
  -p cpus=2 -p cpu_arch=x86_64 -p memory_mb=64000 -p local_gb=99 \
  -i ilo_deploy_iso=b9499494-7db3-4448-b67f-233b86489c1f
```

This will generate output similar to the following:

Now update the node with boot\_mode and boot\_option properties:

```
ironic node-update ea7246fd-eld6-4637-9699-0b7c59c22e67 add properties/
capabilities="boot_mode:bios,boot_option:local"
```

The ironic node-update command returns details for all the node's characteristics.

```
+-----+
            | Value
| Property
+----+
| target_power_state | None
| extra
              | {}
             | None
| last error
| updated_at
          | None
| maintenance_reason | None
| clean_step
          | {}
| uuid
              | ea7246fd-e1d6-4637-9699-0b7c59c22e67
| console_enabled | False
| target_provision_state | None
| provision_updated_at | None
| maintenance
          | False
| inspection_started_at | None
| inspection_finished_at | None
```

```
| power_state
                          | None
| driver
                          | agent_ilo
                         | None
| reservation
                          | {u'memory_mb': 64000, u'cpu_arch': u'x86_64', u'local_gb': 99,
| properties
u'cpus': |
                          | 2, u'capabilities': u'boot_mode:bios,boot_option:local'}
| instance uuid
                          | None
| name
                          | None
                         | {u'ilo_address': u'10.1.196.117', u'ilo_password': u'******',
| driver_info
                          | u'ilo_deploy_iso': u'b9499494-7db3-4448-b67f-233b86489c1f',
                          | u'ilo_username': u'Administrator'}
| created_at
                         | 2016-03-11T10:17:10+00:00
| driver_internal_info
                         | {}
| chassis_uuid
| instance_info
                         | {}
```

## 13.3.4 Creating a node using agent\_ipmi

Use the <u>ironic node-create</u> command to specify the <u>agent\_ipmi</u> driver, network access and credential information for the iLO, properties of the node and the Glance IDs of the supplied kernel and ramdisk images. Note that memory size is specified in megabytes while disk size is gigabytes.

```
ironic node-create -d agent_ipmitool \
-i ipmi_address=10.1.196.117 -i ipmi_username=Administrator -i ipmi_password=***** \
-p cpus=2 -p memory_mb=64000 -p local_gb=99 -p cpu_arch=x86_64 \
-i deploy_kernel=8bee92b7-98ae-4242-b80e-1201a475361a \
-i deploy_ramdisk=0c889803-469d-4aad-8cf6-3501e39c532c
```

This will generate output similar to the following:

```
+-----
| Property | Value
+-----
| driver_info | {u'deploy_kernel': u'8bee92b7-98ae-4242-b80e-1201a475361a',
| u'ipmi_address': u'10.1.196.117', u'ipmi_username': u'Administrator',
| u'ipmi_password': u'******', u'deploy_ramdisk': u'0c889803-469d-4aad-
        | 8cf6-3501e39c532c'}
extra
        | {}
| driver | agent_ipmitool
| chassis_uuid |
| properties | {u'memory_mb': 64000, u'cpu_arch': u'x86_64', u'local_gb': 99, u'cpus':
        | 2}
| name
        | None
+-----
```

Now update the node with boot mode and boot option properties:

```
ironic node-update cf25b19d-098d-406c-b7al-e0c40cc84a12 add properties/
capabilities="boot_mode:bios,boot_option:local"
```

The ironic node-update command returns details for all the node's characteristics.

```
| extra
                         | {}
| last_error
                         | None
| updated_at
                         | None
| maintenance_reason
                         | None
| provision_state
                        | available
| clean_step
                        | {}
| uuid
                        | cf25b19d-098d-406c-b7a1-e0c40cc84a12
| console_enabled
                        | False
| target_provision_state | None
| provision_updated_at
                        None
| maintenance
                        | False
| inspection_started_at | None
| inspection_finished_at | None
| power_state
                        | None
| driver
                        | agent_ipmitool
| reservation
                        | None
| properties
                        | {u'memory_mb': 64000, u'cpu_arch': u'x86_64', u'local_gb': 99,
u'cpus': |
                        | 2, u'capabilities': u'boot_mode:bios,boot_option:local'}
                         | None
| instance_uuid
name
                         | None
| driver_info
                         | {u'ipmi_password': u'******', u'ipmi_address':
u'10.1.196.117',
                         | u'ipmi_username': u'Administrator', u'deploy_kernel':
                         | u'8bee92b7-98ae-4242-b80e-1201a475361a', u'deploy_ramdisk':
u'0c889803 |
```

For more information on node enrollment, see the OpenStack documentation at http://docs.open-stack.org/developer/ironic/deploy/install-guide.html#enrollment ▶.

#### 13.3.5 Create Flavor

Nova uses flavors when fulfilling requests for bare metal nodes. The Nova scheduler attempts to match the requested flavor against the properties of the created Ironic nodes. So an administrator needs to set up flavors that correspond to the available bare metal nodes using the <a href="movaer-create">novaer-create</a> command.

To see a list of all the available flavors, run nova flavor-list

163 Create Flavor

+			+	-+	++	
+	-+		+			
1			m1.tiny	512	1	0
1	.   1.0	True				
2			m1.small	2048	20	0
1	.   1.0	True				
3			ml.medium	4096	40	0
2	!   1.0	True		L 0102	1 00 1	0
	1.0	True	m1.large	8192	80	0
4	11.0	True	ı   m1.xlarge	16384	160	0 1
	3   1.0	True	l mirkturge	10504	100	· I
6	, , , , , , ,	1 1140	m1.baremetal	I 4096	80	0
	!   1.0	True	 	•		·
6450	le08d-2bc6-43f1	8a5f-2315a75b1348	8   bmtest	64000	99	0
2	!   1.0	True	I			
+			+	-+	++	
+	-+		+			

Now set the cpu architecture and boot mode and boot option capabilities:

```
nova flavor-key 645de08d-2bc6-43f1-8a5f-2315a75b1348 set cpu_arch=x86_64
nova flavor-key 645de08d-2bc6-43f1-8a5f-2315a75b1348 set capabilities:boot_option="local"
nova flavor-key 645de08d-2bc6-43f1-8a5f-2315a75b1348 set capabilities:boot_mode="bios"
```

For more information on flavor creation, see the OpenStack documentation at http://docs.open-stack.org/developer/ironic/deploy/install-guide.html#flavor-creation ₹.

## 13.3.6 Create network port

Register the MAC addresses of all connected physical network interfaces intended for use with the baremetal node.

```
ironic port-create -a 5c:b9:01:88:f0:a4 -n ea7246fd-e1d6-4637-9699-0b7c59c22e67
```

## 13.3.7 Create Glance Image

```
DIB_RELEASE=trusty disk-image-create -o ubuntu-trusty.qcow2 vm ubuntu
```

164 Create network port

Then load the image into glance:

```
glance image-create --name='ubuntu' --disk-format=qcow2 --container-format=bare --file ~/
ubuntu.qcow2
+-----+
        | Value
+-----+
| container_format | bare
| 17e4915a-ada0-4b95-bacf-ba67133f39a7 |
| id
       | 0
| min_disk
        | 0
| min_ram
| name
        | ubuntu
| virtual_size
        | None
| visibility
        | private
+-----
```

This image will be used subsequently to boot the baremetal node.

## 13.3.8 Generate Key Pair

Create a keypair that you will use when you login to the newly booted node.

```
nova keypair-add ironic_kp > ironic_kp.pem
```

#### 13.3.9 Determine the neuron network ID

165 Generate Key Pair

#### 13.3.10 Boot the node

Before booting, it is advisable to power down the node:

```
ironic node-set-power-state ea7246fd-e1d6-4637-9699-0b7c59c22e67 off
```

You can now boot the baremetal node with the information complied in the preceding steps, using the neutron network id, the whole disk image id, the matching flavor and the key name.

```
nova boot --nic net-id=c010267c-9424-45be-8c05-99d68531ca8c \
--image 17e4915a-ada0-4b95-bacf-ba67133f39a7 \
--flavor 645de08d-2bc6-43f1-8a5f-2315a75b1348 \
--key-name ironic_kp ubuntu
```

This command returns information about the state of the node that is booting.

```
| Value
| Property
+-----+
| OS-EXT-AZ:availability zone
| OS-EXT-SRV-ATTR:host
| OS-EXT-SRV-ATTR:hypervisor_hostname | -
| OS-EXT-SRV-ATTR:instance_name | instance-00000001
                                  | 0
| OS-EXT-STS:power_state
                                  | scheduling
| OS-EXT-STS:task_state
| OS-EXT-STS:vm state
                                  | building
| OS-SRV-USG:launched at
| OS-SRV-USG:terminated at
| accessIPv4
| accessIPv6
| adminPass
                                  | adpHw3KKTjHk
| config_drive
| created
                                  | 2016-03-11T11:00:28Z
| flavor
                                  | bmtest (645de08d-2bc6-43f1-8a5f-2315a75b1348)
| hostId
                                  | a90122ce-bba8-496f-92a0-8a7cb143007e
| id
                                  | ubuntu (17e4915a-ada0-4b95-bacf-ba67133f39a7) |
| image
| key_name
                                  | ironic_kp
```

166 Boot the node

metadata	1	{}	1
name	1	ubuntu	
os-extended-volumes:volumes_attached	1	[]	1
progress	1	0	1
security_groups	1	default	1
status	1	BUILD	1
tenant_id	1	d53bcaf15afb4cb5aea3adaedbaa60dd	1
updated	1	2016-03-11T11:00:28Z	1
user id	I	e580c645bfec4faeadef7dbd24aaf990	1

The boot procedure can typically take up to 10 minutes to complete. You can monitor the progress with the iLO console and by using the <a href="novalist">nova show <nova\_node\_id></a> and ironic node-show <ironic\_node\_id> commands.

During the boot procedure, a login prompt will appear for Linux for HPE Helion:

Just ignore this login and wait for the login screen of your target Operating System to appear:

If you now run the nova list command, it should show the node in the running state:

You can now log in to the booted node using the key you generated earlier. (You may be prompted to change the permissions on your private key files so that they are not accessible by others).

167 Boot the node

### 13.3.11 Create Glance Images using diskimage-builder (DIB)

The following sections show you how to create your own images using the OpenStack diskim-age-builder tool - for information on how to install the tool see http://docs.openstack.org/developer/diskimage-builder/user\_guide/installation.html ...

### 13.3.12 Creating BIOS images for RHEL

1. Install dib-utils:

```
pip install dib-utils
```

2. Download the qcow2 guest image from RHEL portal and define DIB\_LOCAL\_IMAGE:

```
export DIB_LOCAL_IMAGE=/home/stack/rhel-guest-image-7.2-20151102.0.x86_64.qcow2
```

3. Create a local yum repo and define DIB YUM REPO CONF to be the path to the repo:

```
export DIB_YUM_REPO_CONF = <<path_to_repo_file>>
```

4. Use disk-image-create to create the image:

```
cd diskimage-builder/bin
./disk-image-create -o rhel.qcow2 vm dhcp-all-interfaces rhel7
```

## 13.3.13 Creating BIOS images for Ubuntu

```
DIB_RELEASE=trusty ./disk-image-create -o ubuntu.qcow2 vm dhcp-all-interfaces ubuntu
```

## 13.3.14 Creating UEFI images

To create UEFI images, you will need to setup sudo to not prompt for passwords and configure web-proxy. Once you have sudo and web-proxy configured, follow these steps:

1. Clone the diskimage-builder git repository.

```
git clone https://git.gozer.hpcloud.net/openstack/diskimage-builder.git -b feature/
v2
```

2. Apply the patch to your cloned <u>diskimage-builder</u> repository (https://review.open-stack.org/#/c/391850/6 → and then clone the diskimage-builder utilities (dib-utils) repository.

```
cd diskimage-builder/
git fetch git://git.openstack.org/openstack/diskimage-builder refs/
changes/50/391850/6 && git cherry-pick FETCH_HEAD
cd ..
git clone https://git.gozer.hpcloud.net/openstack/dib-utils.git
```

3. Add execute permissions for disk-image-create.

```
chmod a+x /root/diskimage-builder/diskimage_builder/lib/disk-image-create
```

4. Export your path to include diskimage-builder utilities.

```
{\tt export\ PATH=/root/diskimage-builder/diskimage\_builder/lib:/root/dib-utils/bin:\$PATH}
```

5. Verify that the diskimage-build utilities are in your cloned path by locating <a href="dib-run-parts">dib-run-parts</a>.

```
which dib-run-parts
/root/dib-utils/bin/dib-run-parts
```

6. Verify that the diskimage-build is in your cloned path by locating disk-image-create.

```
which disk-image-create /root/diskimage-builder/lib/disk-image-create
```

7. Install diskimage-builder.

```
cd diskimage-builder/
python setup.py install
```

8. Export the diskimage-builder lib files.

```
export _LIB=/root/diskimage-builder/diskimage_builder/lib
```

9. Create the UEFI Image.

169 Creating UEFI images

• To create an Ubuntu UEFI image:

```
DIB_RELEASE=trusty DIB_UEFI_SUPPORT=true disk-image-create -o ubuntu-trusty-uefi.qcow2 vm ubuntu dhcp-all-interfaces
```

• To create a Fedora UEFI image:

```
\label{lem:disk-image-create} \begin{tabular}{ll} DIB\_UEFI\_SUPPORT=true \ disk-image-create -o \ fedora-uefi.qcow2 \ vm \ fedora \ dhcp-all-interfaces \end{tabular}
```

10. For OS support, refer to <u>supported\_distributions.rst</u> in <u>diskimage-builder</u> user documentation.

# 13.4 Provisioning Baremetal Nodes with Multi-Tenancy

1. Create a network and a subnet:

\$ neutron net-create guest-ne	et - 1	
Created a new network:		
+	+	+
Field	Value	
+	+	+
admin_state_up	True	
availability_zone_hints		
availability_zones		
created_at	2017-06-10T02:49:56Z	-
description		-
id	256d55a6-9430-4f49-8a4c-cc5192f5321e	
ipv4_address_scope		-
ipv6_address_scope		
mtu	1500	
name	guest-net-1	
project_id	57b792cdcdd74d16a08fc7a396ee05b6	
provider:network_type	vlan	1
provider:physical_network	physnet1	
provider:segmentation_id	1152	
revision_number	2	
router:external	False	1
shared	False	1
status	ACTIVE	
subnets		1

```
| tags
             | 57b792cdcdd74d16a08fc7a396ee05b6
| tenant_id
| updated at
                  | 2017-06-10T02:49:57Z
$ neutron subnet-create guest-net-1 200.0.0.0/24
Created a new subnet:
+-----
            | Value
| Field
| allocation pools | {"start": "200.0.0.2", "end": "200.0.0.254"} |
| cidr | 200.0.0.0/24
| created_at
            | 2017-06-10T02:53:08Z
| description
| dns_nameservers |
| enable_dhcp | True
| gateway_ip
             | 200.0.0.1
| host_routes
             | 53accf35-ae02-43ae-95d8-7b5efed18ae9
| id
| ip_version | 4
| ipv6_address_mode |
| ipv6_ra_mode
| name
| network_id
           | 256d55a6-9430-4f49-8a4c-cc5192f5321e
| revision_number | 2
| service_types |
```

#### 2. Review glance image list

#### 3. Create Ironic node

```
$ ironic --ironic-api-version 1.22 node-create -d agent_ipmitool \
```

```
-n test-node-1 -i ipmi_address=192.168.9.69 -i ipmi_username=ilo_user \
 -i ipmi_password=XXXXXXXX --network-interface neutron -p memory_mb=4096 \
 -p cpu_arch=x86_64 -p local_gb=80 -p cpus=2 \
 -p capabilities=boot mode:bios,boot option:local \
 -p root_device='{"name":"/dev/sda"}' \
 -i deploy kernel=db3d131f-2fb0-4189-bb8d-424ee0886e4c \
 -i deploy_ramdisk=304cae15-3fe5-4f1c-8478-c65da5092a2c
          | Value
| Property
 +-----
| chassis_uuid |
  | driver
            | agent_ipmitool
   | u'ipmi_address': u'192.168.9.69', u'ipmi_username': u'gozer',
                | u'ipmi_password': u'*****', u'deploy_ramdisk':
                | u'304cae15-3fe5-4f1c-8478-c65da5092a2c'}
               | {}
| extra
   Т
               | test-node-1
| name
| network_interface | neutron
properties | {u'cpu_arch': u'x86_64', u'root_device': {u'name': u'/dev/
sda'}, |
                | u'cpus': 2, u'capabilities':
u'boot_mode:bios,boot_option:local', |
                | u'memory_mb': 4096, u'local_gb': 80}
| resource_class | None
 | uuid
               | cb4dda0d-f3b0-48b9-ac90-ba77b8c66162
```

ipmi\_address, ipmi\_username and ipmi\_password are iLo access parameters for baremetal Ironic node. Adjust memory\_mb, cpus, local\_gb to your node size requirements. They also need to be reflected in flavor setting (see below). Use capabilities boot\_mode:bios for baremetal nodes operating in Legacy BIOS mode. For UEFI baremetal nodes, use boot\_mode:uefi lookup deploy\_kernel and deploy\_ramdisk in glance image list output above.



### **Important**

Since we are using Ironic API version 1.22, node is created initial state **enroll**. It needs to be explicitly moved to **available** state. This behavior changed in API version 1.11

#### 4. Create port

```
$ ironic --ironic-api-version 1.22 port-create --address f0:92:1c:05:6c:40 --
node cb4dda0d-f3b0-48b9-ac90-ba77b8c66162 -l switch id=e8:f7:24:bf:07:2e -l
switch_info=hp59srv1-a-11b -l port_id="Ten-GigabitEthernet 1/0/34" --pxe-enabled
+-----
| Property
                   | Value
         +-----
| address
                   | f0:92:1c:05:6c:40
                    | {}
| extra
| local_link_connection | {u'switch_info': u'hp59srv1-a-11b', u'port_id': u'Ten-
GigabitEthernet |
                    | 1/0/34', u'switch_id': u'e8:f7:24:bf:07:2e'}
                    | cb4dda0d-f3b0-48b9-ac90-ba77b8c66162
| node_uuid
| pxe enabled
                    | True
| uuid
                    | a49491f3-5595-413b-b4a7-bb6f9abec212
```

- for --address, use MAC of 1st NIC of ironic baremetal node, which will be used for PXE boot
- for --node, use ironic node uuid (see above)
- for -l switch\_id, use switch management interface MAC address. It can be retrieved by pinging switch management IP and looking up MAC address in 'arp -l -n' command output.
- for -l switch\_info, use switch\_id from data/ironic/ironic\_config.yml file. If you have several switch config definitions, use the right switch your baremetal node is connected to.
- for -l port\_id, use port ID on the switch
- 5. Move ironic node to manage and then available state

```
$ ironic node-set-provision-state test-node-1 manage
$ ironic node-set-provision-state test-node-1 provide
```

6. Once node is successfully moved to available state, it's resources should be included into Nova hypervisor statistics

```
$ nova hypervisor-stats
+----+
| Property | Value |
+----+
| count
            | 1
current_workload | 0
| disk available least | 80
| free_disk_gb | 80
| memory_mb_used
| running_vms
             | 0
            | 0
| vcpus
             | 2
| vcpus_used
             | 0
+----+
```

7. Prepare a keypair, which will be used for logging into the node

```
$ nova keypair-add ironic_kp > ironic_kp.pem
```

8. Obtain user image and upload it to glance. Please refer to OpenStack documentation on user image creation: https://docs.openstack.org/project-install-guide/baremetal/draft/configure-glance-images.html ...



#### Note

Deployed images are already populated by HPE Helion OpenStack installer.

```
$ glance image-create --name='Ubuntu Trusty 14.04' --disk-format=qcow2 --container-
format=bare --file ~/ubuntu-trusty.qcow2
+-----+
| Property | Value
+----+
| container_format | bare
| id
         | 9fdd54a3-ccf5-459c-a084-e50071d0aa39 |
+-----+
$ glance image-list
+----+
+-----+
| 73205fb7-64f1-4243-8d0f-a0dd2e493c9d | cirros-0.3.4-x86_64
| 83eecf9c-d675-4bf9-a5d5-9cf1fe9ee9c2 | ir-deploy-iso-HOS5.0
| db3d131f-2fb0-4189-bb8d-424ee0886e4c | ir-deploy-kernel-HOS5.0 |
| 304cae15-3fe5-4f1c-8478-c65da5092a2c | ir-deploy-ramdisk-HOS5.0 |
| 9fdd54a3-ccf5-459c-a084-e50071d0aa39 | Ubuntu Trusty 14.04
+----+
```

9. Create a baremetal flavor and set flavor keys specifying requested node size, architecture and boot mode. A flavor can be re-used for several nodes having the same size, architecture and boot mode

```
$ nova flavor-create m1.ironic auto 4096 80 2
+----+
                      | Name | Memory_MB | Disk | Ephemeral |
| ID
Swap | VCPUs | RXTX_Factor | Is_Public |
+-----
+----+
| ab69881b-0d5a-497d-93ae-6e28694c87bf | ml.ironic | 4096 | 80 | 0
  | 2 | 1.0
              | True
                     +-----
+----+
$ nova flavor-key ab69881b-0d5a-497d-93ae-6e28694c87bf set cpu arch=x86 64
$ nova flavor-key ab69881b-0d5a-497d-93ae-6e28694c87bf set
capabilities:boot_option="local"
$ nova flavor-key ab69881b-0d5a-497d-93ae-6e28694c87bf set
capabilities:boot_mode="bios"
```

Parameters must match parameters of ironic node above. Use <a href="capabilities:boot\_mode="mode="uefi" capabilities:boot\_mode="uefi" capabilities:boot\_mode

10. Boot the node

```
| OS-EXT-STS:power_state
                                       | 0
| OS-EXT-STS:task_state
                                      | scheduling
| OS-EXT-STS:vm_state
                                      | building
| OS-SRV-USG:launched_at
| OS-SRV-USG:terminated_at
| accessIPv4
| accessIPv6
                                       | XXXXXXXXXXX
| adminPass
| config_drive
                                       | 2017-06-14T21:25:18Z
| created
| flavor
                                       | ml.ironic
(ab69881b-0d5a-497d-93ae-6e28694c87bf)
| hostId
| id
                                       | f1a8c63e-da7b-4d9a-8648-b1baa6929682
                                       | Ubuntu Trusty 14.04 (9fdd54a3-ccf5-459c-
| image
a084-e50071d0aa39) |
| key_name
                                       | ironic_kp
| metadata
                                       | {}
| name
                                       | test-node-1
| os-extended-volumes:volumes_attached | []
                                       | 0
| progress
| security_groups
                                       | default
                                       | BUILD
| status
                                       | 57b792cdcdd74d16a08fc7a396ee05b6
| tenant_id
                                       | 2017-06-14T21:25:17Z
| updated
```

- for --image, use ID of user image created at previous step
- for --nic net-id, use id of tenant network created at the beginning



### Note

During the node provisioning, the following is happening in the background:

Neutron connects to switch management interfaces and assigns provisioning VLAN to baremetal node port on the switch. Ironic powers up the node using iLo interface. Node is booting IPA image via PXE. IPA image is writing provided user image onto specified root device (/dev/sda) and powers node down. Neutron connects to switch management interfaces and assigns tenant VLAN to baremetal node port on the switch. A VLAN ID is selected from provided range. Ironic powers up the node using iLo interface. Node is booting user image from disk.

11. Once provisioned, node will join the private tenant network. Access to private network from other networks is defined by switch configuration.

# 13.5 View Ironic System Details

# 13.5.1 View details about the server using **nova show <nova-node-id>**

```
| OS-EXT-SRV-ATTR:host
                                | helion-cp1-ir-compute0001-mgmt
| instance-0000000a
| OS-EXT-SRV-ATTR:instance_name
| OS-EXT-STS:power_state | 1
| OS-EXT-STS:task_state
| OS-EXT-STS:vm_state | active
| OS-SRV-USG:launched_at
                              | 2016-03-11T12:26:25.000000
| OS-SRV-USG:terminated_at | -
| accessIPv4
| accessIPv6
| config_drive
                                | 2016-03-11T12:17:54Z
| created
| flat-net network
                       | 192.3.15.14
                                | bmtest (645de08d-2bc6-43f1-8a5f-2315a75b1348)
| flavor
| hostId
ecafa4f40eb5f72f7298de0ef51eb7769e3bad47cbc01aa0a076114f |
                                | a90122ce-bba8-496f-92a0-8a7cb143007e
                                | ubuntu (17e4915a-ada0-4b95-bacf-ba67133f39a7)
| image
                                | ironic_kp
| key_name
| metadata
                                | {}
| name
                                | ubuntu
| os-extended-volumes:volumes_attached | []
                                | 0
| progress
                                | default
| security_groups
```

# 13.5.2 View detailed information about a node using **ironic** node-show <ironic-node-id>

```
ironic node-show ea7246fd-e1d6-4637-9699-0b7c59c22e67
| Property
             | Value
 -----+
| {}
| extra
| last_error | None
         | 2016-03-11T12:26:25+00:00
| updated_at
| maintenance_reason | None
| clean_step
        | {}
            | ea7246fd-e1d6-4637-9699-0b7c59c22e67
| uuid
| console_enabled | False
| target_provision_state | None
```

```
| maintenance
                         | False
| inspection_started_at
                        | None
| inspection_finished_at | None
| power_state
                         | power on
| driver
                         | agent_ilo
| reservation
                         | None
| properties
                         | {u'memory_mb': 64000, u'cpu_arch': u'x86_64', u'local_gb': 99,
u'cpus': |
                         | 2, u'capabilities': u'boot_mode:bios,boot_option:local'}
| instance_uuid
                         | a90122ce-bba8-496f-92a0-8a7cb143007e
                         | None
| name
                         | {u'ilo_address': u'10.1.196.117', u'ilo_password': u'******',
| driver_info
                         | u'ilo deploy iso': u'b9499494-7db3-4448-b67f-233b86489c1f',
                         | u'ilo_username': u'Administrator'}
| created_at
                         | 2016-03-11T10:17:10+00:00
                         | {u'agent_url': u'http://192.3.15.14:9999',
| driver_internal_info
u'is_whole_disk_image': True, |
                         | u'agent_last_heartbeat': 1457699159}
| chassis_uuid
                         | {u'root_gb': u'99', u'display_name': u'ubuntu',
| instance info
u'image_source': u
                         | '17e4915a-ada0-4b95-bacf-ba67133f39a7', u'capabilities':
u'{"boot_mode": |
                         | "bios", "boot_option": "local"}', u'memory_mb': u'64000',
u'vcpus':
                | u'2', u'image_url': u'http://192.168.12.2:8080/v1/
AUTH_ba121db7732f4ac3a |
                         | 50cc4999a10d58d/glance/17e4915a-ada0-4b95-bacf-ba67133f39a7?
temp_url_sig |
=ada691726337805981ac002c0fbfc905eb9783ea&temp_url_expires=1457699878', |
```

# 13.5.3 View detailed information about a port using **ironic port-show <ironic-port-id>**

# 13.5.4 View detailed information about a hypervisor using **nova hypervisor-list** and **nova hypervisor-show**

cpu_info	
current_workload	
disk_available_least	0
free_disk_gb	0
free_ram_mb	0
host_ip	192.168.12.6
hypervisor_hostname	ea7246fd-e1d6-4637-9699-0b7c59c22e67
hypervisor_type	ironic
hypervisor_version	1
id	541
local_gb	99
local_gb_used	99
memory_mb	64000
memory_mb_used	64000
running_vms	1
service_disabled_reason	None
service_host	helion-cp1-ir-compute0001-mgmt
service_id	25
state	up
status	enabled
vcpus	2
vcpus_used	2

# 13.5.5 View a list of all running services using **nova service- list**

nova service-list				
	+	+	+	-
Id   Binary Updated_at	Host   Disabled Reason	·	Status   State	
	++	+	+	-
1   nova-conduct 2016-03-11T12:44:5	or   helion-cp1-c1-m1-mgmt 1.000000   -	internal	enabled   up	I
7   nova-conduct 2016-03-11T12:44:5	or   helion-cp1-c1-m2-mgmt 3.000000   -	internal	enabled   up	I
10   nova-conduct 2016-03-11T12:44:5	or   helion-cp1-c1-m3-mgmt 2.000000   -	internal	enabled   up	I
13   nova-schedul 2016-03-11T12:44:5	er   helion-cp1-c1-m1-mgmt 3.000000   -	internal	enabled   up	I
16   nova-schedul 2016-03-11T12:44:5	er   helion-cp1-c1-m3-mgmt 7.000000   -	internal	enabled   up	I

# 13.6 Troubleshooting Ironic Installation

Sometimes the <u>nova boot</u> command does not succeed and when you do a <u>nova list</u>, you will see output like the following:

You should execute the <u>nova show <nova-node-id></u> and <u>ironic node-show <ironic-node-id></u> id> commands to get more information about the error.

# 13.6.1 No valid host was found. There are not enough hosts available.

This error is typically seen when performing the <u>nova boot</u> where there is a mismatch between the properties set on the node and the flavor used. For example, the output from a <u>nova show</u> command may look like this:

```
nova show ee08f82e-8920-4360-be51-a3f995624e5f
+-----+
```

```
| Property
                             | Value
| OS-EXT-AZ:availability_zone |
  | OS-EXT-SRV-ATTR:host | -
| OS-EXT-SRV-ATTR:hypervisor_hostname | -
| OS-EXT-SRV-ATTR:instance_name | instance-00000001
| OS-EXT-STS:power_state | 0
| OS-EXT-STS:task_state | -
                  | error
| OS-EXT-STS:vm_state
| OS-SRV-USG:launched_at | -
| OS-SRV-USG:terminated_at | -
| accessIPv4
```

```
| accessIPv6
| config_drive
| created
                                     | 2016-03-11T11:00:28Z
| fault
                                     | {"message": "No valid host was found. There
are not enough hosts available.", "code": 500, "details": " File \"/opt/stack/venv/
nova-20160308T002421Z/lib/python2.7/site-packages/nova/conductor/manager.py\", line 739,
in build_instances |
                                         request_spec, filter_properties)
         | File \"/opt/stack/venv/nova-20160308T002421Z/
lib/python2.7/site-packages/nova/scheduler/utils.py\", line 343, in wrapped
return func(*args, **kwargs)
                                     | File \"/opt/stack/venv/nova-20160308T002421Z/
lib/python2.7/site-packages/nova/scheduler/client/__init__.py\", line 52, in
select_destinations
                                     context, request_spec, filter_properties)
| File \"/opt/stack/venv/nova-20160308T002421Z/
lib/python2.7/site-packages/nova/scheduler/client/__init__.py\", line 37, in __run_method
                                     return getattr(self.instance, __name)(*args,
**kwargs)
                                     File \"/opt/stack/venv/nova-20160308T002421Z/
lib/python2.7/site-packages/nova/scheduler/client/query.py\", line 34, in
```

```
select_destinations
1
                                      context, request_spec, filter_properties)
         1
                                          File \"/opt/stack/venv/nova-20160308T002421Z/
lib/python2.7/site-packages/nova/scheduler/rpcapi.py\", line 120, in select_destinations
                                        request spec=request spec,
filter_properties=filter_properties)
                                          File \"/opt/stack/venv/nova-20160308T002421Z/
lib/python2.7/site-packages/oslo_messaging/rpc/client.py\", line 158, in call
                                          retry=self.retry)
Τ
         File \"/opt/stack/venv/nova-20160308T002421Z/
lib/python2.7/site-packages/oslo_messaging/transport.py\", line 90, in _send
timeout=timeout, retry=retry)
         I
                                      File \"/opt/stack/venv/nova-20160308T002421Z/
lib/python2.7/site-packages/oslo_messaging/_drivers/amqpdriver.py\", line 462, in send
                                            retry=retry)
Ι
                                      File \"/opt/stack/venv/nova-20160308T002421Z/
lib/python2.7/site-packages/oslo_messaging/_drivers/amqpdriver.py\", line 453, in _send
                                            raise result
1
```

```
| ", "created": "2016-03-11T11:00:29Z"}
| flavor
                                       | bmtest (645de08d-2bc6-43f1-8a5f-2315a75b1348)
| hostId
| id
                                       | ee08f82e-8920-4360-be51-a3f995624e5f
| image
                                       | ubuntu (17e4915a-ada0-4b95-bacf-ba67133f39a7)
| key_name
                                       | ironic_kp
| metadata
                                       | {}
| name
                                       | ubuntu
| os-extended-volumes:volumes_attached | []
| status
                                       | ERROR
| tenant_id
                                       | d53bcaf15afb4cb5aea3adaedbaa60dd
| updated
                                       | 2016-03-11T11:00:28Z
```

You can find more information about the error by inspecting the log file at <a href="https://www.no-va-scheduler.log">/war/log/nova/no-va-scheduler.log</a> or alternatively by viewing the error location in the source files listed in the stack-trace (in bold above).

To find the mismatch, compare the properties of the ironic node:

```
| Property
              | Value
   +-----+
| extra
         | {}
   | last_error | None
         | None
| updated_at
| maintenance_reason | None
| clean_step | {}
| uuid
              | ea7246fd-e1d6-4637-9699-0b7c59c22e67
| console_enabled
              | False
| target_provision_state | None
| provision_updated_at | None
| maintenance | False
```

```
| inspection_started_at | None
| inspection_finished_at | None
| power_state | None
| driver
             | agent_ilo
| reservation
                    | None
             | {u'memory_mb': 64000, u'local_gb': 99, u'cpus': 2,
| properties
u'capabilities': |
                     | u'boot_mode:bios,boot_option:local'}
| instance_uuid
                    | None
| name
                      | None
                     | {u'ilo_address': u'10.1.196.117', u'ilo_password': u'******',
| driver_info
                      | u'ilo_deploy_iso': u'b9499494-7db3-4448-b67f-233b86489c1f',
                      | u'ilo_username': u'Administrator'}
            | 2016-03-11T10:17:10+00:00
| created_at
| driver_internal_info | {}
| chassis_uuid
| instance_info | {}
```

#### with the flavor characteristics:

```
| OS-FLV-DISABLED:disabled | False
| OS-FLV-EXT-DATA:ephemeral | 0
| disk
                             | 99
                             | {"capabilities:boot_option": "local", "cpu_arch":
| extra_specs
 "x86_64", "capabilities:boot_mode": "bios"} |
                             | 645de08d-2bc6-43f1-8a5f-2315a75b1348
| name
                              | bmtest
| os-flavor-access:is_public | True
                             | 64000
| ram
| rxtx_factor
                             | 1.0
| swap
| vcpus
                             | 2
```

In this instance, the problem is caused by the absence of the "cpu\_arch": "x86\_64" property on the ironic node. This can be resolved by updating the ironic node, adding the missing property:

```
ironic node-update ea7246fd-e1d6-4637-9699-0b7c59c22e67 add properties/cpu_arch=x86_64
```

and then re-running the nova boot command.

# 13.6.2 Deployment to a node fails and in "ironic node-list" command, the power\_state column for the node is shown as "None"

**Possible cause:** The IPMI commands to the node take longer to change the power state of the server.

**Resolution:** Check if the node power state can be changed using the following command

```
ironic node-set-power-state $NODEUUID on
```

If the above command succeeds and the power\_state column is updated correctly, then the following steps are required to increase the power sync interval time.

191 node is shown as "None"

On the first controller, reconfigure Ironic to increase the power sync interval time. In the example below, it is set to 120 seconds. This value may have to be tuned based on the setup.

1. Go to the <u>~/helion/my\_cloud/config/ironic/</u> directory and edit <u>ironic-conductor.conf.j2</u> to set the sync\_power\_state\_interval value:

```
[conductor]
sync_power_state_interval = 120
```

2. Save the file and then run the following playbooks:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
ansible-playbook -i hosts/localhost ready-deployment.yml
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ironic-reconfigure.yml
```

### 13.6.3 Error Downloading Image

If you encounter the error below during the deployment:

```
"u'message': u'Error downloading image: Download of image id 77700b53-9e15-406c-b2d5-13e7d9b96551 failed: Image download failed for all URLs.', u'code': 500, u'type': u'ImageDownloadError', u'details': u'Download of image id 77700b53-9e15-406c-b2d5-13e7d9b96551 failed: Image download failed for all URLs.'"
```

you should visit the Single Sign-On Settings in the Security page of iLO and change the Single Sign-On Trust Mode setting from the default of "Trust None (SSO disabled)" to "Trust by Certificate".

# 13.6.4 Using node-inspection can cause temporary claim of IP addresses

**Possible cause:** Running <u>node-inspection</u> on a node discovers all the NIC ports including the NICs that don't have any connectivity. This causes a temporary consumption of the network IPs and increased usage of the allocated quota. As a result, other nodes are deprived of IP addresses and deployments can fail.

**Resolution:**You can add node properties manually added instead of using the inspection tool.

Note: Upgrade  $\underline{\text{ipmitool}}$  to a version >=1.8.15 or it may not return detailed information about the NIC interface for node-inspection.

### 13.6.5 Node permanently stuck in deploying state

#### Possible causes:

- Ironic conductor service associated with the node could go down.
- There might be a properties mismatch. MAC address registered for the node could be incorrect.

**Resolution:** To recover from this condition, set the provision state of the node to <u>Error</u> and maintenance to True. Delete the node and re-register again.

# 13.6.6 The NICs in the baremetal node should come first in boot order

**Possible causes:** By default, the boot order of baremetal node is set as NIC1, HDD and NIC2. If NIC1 fails, the nodes starts booting from HDD and the provisioning fails.

**Resolution:** Set boot order so that all the NICs appear before the hard disk of the baremetal as NIC1, NIC2..., HDD.

# 13.6.7 Increase in the number of nodes can cause power commands to fail

**Possible causes:**Ironic periodically performs a power state sync with all the baremetal nodes. When the number of nodes increase, ironic does not get sufficient time to perform power operations.

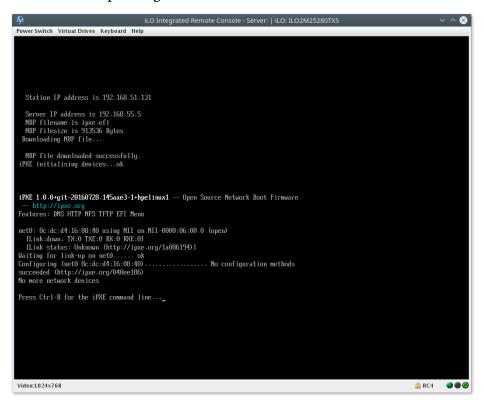
**Resolution:** The following procedure gives a way to increase sync power state interval:

- Edit the file \_~/helion/my\_cloud/config/ironic/ironic-conductor.conf.j2 and navigate to the section for [conductor]
- Increase the <u>sync\_power\_state\_interval</u>. For example, for 100 nodes, set <u>sync\_power\_state\_interval</u> = 90 and save the file.
- Execute the following set of commands to reconfigure Ironic:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
ansible-playbook -i hosts/localhost ready-deployment.yml
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ironic-reconfigure.yml
```

#### 13.6.8 DHCP succeeds with PXE but times out with iPXE

If you see DHCP error "No configuration methods succeeded" in iPXE, right after successful DHCP performed by embedded NIC firmware (please refer to screenshot below), there may be an issue with Spanning Tree Protocol on the switch.



To avoid this error, Rapid Spanning Tree Protocol needs to be enabled on the switch. If this is not an option due to conservative loop detection strategies, use the steps outlined below to install the iPXE binary with increased DHCP timeouts.

1. Clone iPXE source code

```
$ git clone git://git.ipxe.org/ipxe.git
$ cd ipxe/src
```

2. Modify lines 22-25 in file <u>config/dhcp.h</u>, which declare reduced DHCP timeouts (1-10 secs). Comment out lines with reduced timeouts and uncomment normal PXE timeouts (4-32)

```
//#define DHCP_DISC_START_TIMEOUT_SEC 1
//#define DHCP_DISC_END_TIMEOUT_SEC 10
#define DHCP_DISC_START_TIMEOUT_SEC 4 /* as per PXE spec */
#define DHCP_DISC_END_TIMEOUT_SEC 32 /* as per PXE spec */
```

3. Make undionly.kpxe (BIOS) and ipxe.efi (UEFI) images

```
$ make bin/undionly.kpxe
$ make bin-x86_64-efi/ipxe.efi
```

4. Copy iPXE images to lifecycle manager

```
$ scp bin/undionly.kpxe bin-x86_64-efi/ipxe.efi stack@10.0.0.4: stack@10.0.0.4's password: undionly.kpxe 100% 66KB 65.6KB/s 00:00 100% 918KB 918.2KB/s 00:00
```

5. From deployer, distribute image files onto all 3 controllers

```
stack@helion-cp1-c1-m1-mgmt:~$ cd ~/scratch/ansible/next/hos/ansible/
stack@helion-cp1-c1-m1-mgmt:~/scratch/ansible/next/hos/ansible$ ansible -i hosts/
verb_hosts IRN-CND -m copy -b -a 'src=/home/stack/ipxe.efi dest=/tftpboot'
...
stack@helion-cp1-c1-m1-mgmt:~/scratch/ansible/next/hos/ansible$ ansible -i hosts/
verb_hosts IRN-CND -m copy -b -a 'src=/home/stack/undionly.kpxe dest=/tftpboot'
...
```

There is no need to restart services. With next PXE boot attempt, iPXE binary with the increased timeout will be downloaded to the target node via TFTP.

# 13.7 Node Cleaning

Cleaning is the process by which data is removed after a previous tenant has used the node. Cleaning requires use of ironic's agent\_drivers. It is extremely important to note that if the pxe\_drivers are utilized, no node cleaning operations will occur, and a previous tenant's data could be found on the node. The same risk of a previous tenant's data possibly can occur if cleaning is explicitly disabled as part of the installation.

By default, cleaning attempts to utilize ATA secure erase to wipe the contents of the disk. If secure erase is unavailable, the cleaning functionality built into the Ironic Python Agent falls back to an operation referred to as "shred" where random data is written over the contents of the disk, and then followed up by writing "0"s across the disk. This can be a time-consuming process.

An additional feature of cleaning is the ability to update firmware or potentially assert new hardware configuration, however, this is an advanced feature that must be built into the Ironic Python Agent image. Due to the complex nature of such operations, and the fact that no one size fits all, this requires a custom Ironic Python Agent image to be constructed with an appropriate hardware manager. For more information on hardware managers, see http://docs.openstack.org/developer/ironic-python-agent/#hardware-managers ?

### 13.7.1 Setup

Cleaning is enabled by default in ironic when installed via the Lifecycle Manager. You can verify this by examining the ironic-conductor.conf file. Look for:

[conductor]
clean\_nodes=true

196 Node Cleaning

### 13.7.2 In use

When enabled, cleaning will be run automatically when nodes go from active to available state or from manageable to available. To monitor what step of cleaning the node is in, run <u>ironic</u> node-show:

```
stack@helion-cp1-c1-m1-mgmt:~$ ironic node-show 4e6d4273-2535-4830-a826-7f67e71783ed
+-----+
| Property
                      | Value
| target_power_state
                    | None
| extra
                     | {}
| last_error
                 | None
| updated_at
                    | 2016-04-15T09:33:16+00:00
| maintenance_reason
                      | None
| provision_state
                      | cleaning
| clean_step
                      | {}
| uuid
                      | 4e6d4273-2535-4830-a826-7f67e71783ed
| console_enabled
                      | False
| target_provision_state | available
| provision_updated_at | 2016-04-15T09:33:16+00:00
| maintenance
                      | False
| inspection_started_at | None
| inspection_finished_at | None
                      | power off
| power_state
| driver
                      | agent_ilo
```

197 In use

```
| reservation
                         | helion-cp1-c1-m1-mgmt
                         | {u'memory_mb': 4096, u'cpu_arch': u'amd64', u'local_gb': 80,
| properties
u'cpus': 2, |
                         | u'capabilities': u'boot_mode:uefi,boot_option:local'}
                         | None
| instance_uuid
name
                         None
                         | {u'ilo deploy iso': u'249bf095-e741-441d-bc28-0f44a9b8cd80',
| driver info
                         | u'ipmi_username': u'Administrator', u'deploy_kernel':
                         | u'3a78c0a9-3d8d-4764-9300-3e9c00e167a1', u'ilo_address':
                         | u'10.1.196.113', u'ipmi_address': u'10.1.196.113',
u'deploy ramdisk':
                         | u'd02c811c-e521-4926-9f26-0c88bbd2ee6d', u'ipmi_password':
                         | u'ilo_password': u'******', u'ilo_username': u'Administrator'}
                         | 2016-04-14T08:30:08+00:00
| created at
| driver_internal_info
                         | {u'clean_steps': None, u'hardware_manager_version':
                         | {u'generic_hardware_manager': u'1.0'}, u'is_whole_disk_image':
True,
                         | u'agent_erase_devices_iterations': 1, u'agent_url':
                         | u'http://192.168.246.245:9999', u'agent_last_heartbeat':
1460633166}
| chassis_uuid
| instance_info
                         | {}
```

The status will be in the <u>driver\_internal\_info</u> field. You will also be able to see the clean steps list there.

198 In use

### 13.7.3 Troubleshooting

If an error occurs during the cleaning process, the node will enter the clean failed state so that it is not deployed. The node remains powered on for debugging purposes. The node can be moved to the manageable state to attempt a fix using the following command:

```
ironic node-set-provision-state <node id> manage
```

Once you have identified and fixed the issue, you can return the node to available state by executing the following commands:

```
ironic node-set-maintenance <node id> false
ironic node-set-provision-state <node id> provide
```

This will retry the cleaning steps and set the node to available state upon their successful completion.

### 13.7.4 Disabling Node Cleaning

To disable node cleaning, you need to edit <a href="helion/my\_cloud/definition/data/iron-ic/ironic\_config.yml">helion/my\_cloud/definition/data/iron-ic/ironic\_config.yml</a> and set enable\_node\_cleaning to false.

Commit your changes:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "disable node cleaning"
```

Deploy these changes by re-running the configuration processor and reconfigure the ironic installation:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
ansible-playbook -i hosts/localhost ready-deployment.yml
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ironic-reconfigure.yml
```

199 Troubleshooting

## 13.8 Ironic and OneView

### 13.8.1 Enabling Ironic OneView driver in HPE Helion OpenStack

Edit the file \_/helion/my\_cloud/definition/data/ironic/ironicconfig.yml and set the value

```
enable_oneview: true
```

This will enable the OneView driver for Ironic in HPE Helion OpenStack.

## 13.8.2 Adding OneView Appliance Credentials

```
manage_url: https://<Onview appliance URL>
oneview_username: "<Appliance username>"
oneview_encrypted_password: "<Encrypted password>"
oneview_allow_insecure_connections: <true/false>
tls_cacert_file: <CA certificate for connection>
```

### 13.8.3 Encrypting the OneView Password

Encryption can be applied using <a href="hosencrypt.py">hosencrypt.py</a> or using <a href="openssl">openssl</a>. On the lifecyclemanager node, export the key used for encryption as environment variable:

```
export HOS_USER_PASSWORD_ENCRYPT_KEY="<key for encryption>"
```

And then execute the following commands:

```
cd ~/helion/hos/ansible python hosencrypt.py
```

Enter password to be encrypted when prompted. The script uses the key that was exported in the <a href="https://example.com/HOS\_USER\_PASSWORD\_ENCRYPT\_KEY">HOS\_USER\_PASSWORD\_ENCRYPT\_KEY</a> to do the encryption.

For more information, see the documentation for .

200 Ironic and OneView

### 13.8.4 Decrypting the OneView Password

Before running the <u>site.yml</u> playbook, export the key used for encryption as environment variable:

```
export HOS_USER_PASSWORD_ENCRYPT_KEY="<key for encryption>"
```

The decryption of the password is then automatically handled in ironic-ansible playbooks.

## 13.8.5 Registering Baremetal Node for OneView Driver

```
ironic node-create -d agent_pxe_oneview
```

Update node driver-info:

ironic node-update \$NODE\_UUID add driver\_info/server\_hardware\_uri=\$SH\_URI

### 13.8.6 Updating Node Properties

```
ironic node-update $NODE_UUID add \
  properties/capabilities=server_hardware_type_uri:$SHT_URI,enclosure_group_uri:
$EG_URI,server_profile_template_uri=$SPT_URI
```

## 13.8.7 Creating Port for Driver

```
ironic port-create -n $NODE_UUID -a $MAC_ADDRESS
```

### 13.8.8 Creating a Node

#### Create Node:

```
ironic node-create -n ovbay7 -d agent_pxe_oneview
```

#### Update driver info:

```
ironic node-update $ID add driver_info/server_hardware_uri="/rest/server-
hardware/30373237-3132-4753-4835-32325652464B" driver_info/deploy_kernel="$KERNELDISK"
    driver_info/deploy_ramdisk="$RAMDISK"
```

### Update node properties:

```
ironic node-update $ID add properties/local_gb=10 ironic node-update $ID add properties/cpus=24 properties/memory_mb=262144 properties/cpu_arch=x86_64
```

ironic node-update \$ID add properties/capabilities=server\_hardware\_type\_uri:'/rest/
server-hardware-types/B3141E03-2F96-4A55-9B77-86CB484FF69E',enclosure\_group\_uri:'/rest/
enclosure-groups/80efe99d-8f2d-4942-8add-9691bc8b79fa',server\_profile\_template\_uri:'/
rest/server-profile-templates/faafc3c0-6c81-47ca-a407-f67d11262da5'

### 13.8.9 Getting Data using REST API

GET login session auth id:

```
curl -k https://<oneview_manger_url>rest/login-sessions -H "content-type:application/
json" -X POST -d '{"userName":"<username>", "password":"<password>"}'
```

GET complete node details in json format:

```
curl -k "https://<manager_url>//rest/server-
hardware/30373237-3132-4753-4835-32325652464B" -H "content-type:application/json" -H
"Auth:<auth_session_id>"| python -m json.tool
```

### 13.8.10 Ironic OneView CLI

<u>ironic-oneview-cli</u> is already installed in <u>ironicclient</u> venv with a symbolic link to it. To generate an rc file for the OneView CLI, follow these steps:

1. Run the following commands:

```
source ~/service.osrc
glance image-list
```

2. Note the <u>deploy-kernel</u> and <u>deploy-ramdisk</u> UUID and then run the following command to generate the rc file:

```
ironic-oneview genrc
```

You will be prompted to enter:

- OneView Manager URL
- Username
- deploy-kernel
- deploy-ramdisk
- allow\_insecure\_connection
- cacert file

The <u>ironic-onefview.rc</u> file will be generated in the current directory, by default. It is possible to specify a different loaction.

3. Source the generated file:

```
source ironic-onefview.rc
```

You can now use the CLI for node and flavor creation as follows:

```
ironic-oneview node-create
ironic-oneview flavor-create
```

# 13.9 RAID Configuration for Ironic

1. Node Creation:

Check the raid capabilities of the driver:

```
ironic --ironic-api-version 1.15 driver-raid-logical-disk-properties pxe_ilo
```

This will generate output similar to the following:



```
controller | Controller to use for this logical disk. If not
specified, the driver will choose a suitable RAID controller |
                        on the bare metal node. Optional.
                        | The type of disk preferred. Valid values are 'hdd' and
| disk_type
'ssd'. If this is not specified, disk type will not
                        | be a selection criterion for choosing backing physical
disks. Optional.
| interface_type
                   | The interface type of disk. Valid values are 'sata',
'scsi' and 'sas'. If this is not specified, interface
                         | type will not be a selection criterion for choosing
backing physical disks. Optional.
| is_root_volume
                        | Specifies whether this disk is a root volume. By
default, this is False. Optional.
| number_of_physical_disks | Number of physical disks to use for this logical disk.
By default, the driver uses the minimum number of
                                                   | disks required for that RAID level. Optional.
| physical disks
                        | The physical disks to use for this logical disk. If not
specified, the driver will choose suitable physical |
                        | disks to use. Optional.
| raid level
                        | RAID level for the logical disk. Valid values are '0',
'1', '2', '5', '6', '1+0', '5+0' and '6+0'. Required. |
| share_physical_disks
                        | Specifies whether other logical disks can share
physical disks with this logical disk.
                        | By default, this is False. Optional.
                        | Size in GiB (Integer) for the logical disk. Use 'MAX'
| size gb
                                              as size_gb if this logical disk is supposed to
                         | use the rest of the space available. Required.
                        | Name of the volume to be created. If this is not
| volume name
specified, it will be auto-generated. Optional.
```

#### Node State will be Available

```
ironic node-create -d pxe_ilo -i ilo_address=<ip_address> \
    -i ilo_username=<username> -i ilo_password=<password> \
    -i ilo_deploy_iso=<iso_id> -i deploy_kernel=<kernel_id> \
    -i deploy_ramdisk=<ramdisk_id> -p cpus=2 -p memory_mb=4096 \
    -p local_gb=80 -p cpu_arch=amd64 \
    -p capabilities="boot_option:local,boot_mode:bios"
```

```
ironic port-create -a <port> -n <node-uuid>
```

2. Apply the target raid configuration on the node:

See the OpenStack documentation for RAID configuration at http://docs.openstack.org/developer/ironic/deploy/raid.html  $\mathbb{Z}$ .

Set the target RAID configuration by editing the file <a href="raid\_conf.json">raid\_conf.json</a> and setting the appropriate values, for example:

```
{ "logical_disks": [ {"size_gb": 5, "raid_level": "0", "is_root_volume": true} ] }
```

and then run the following command:

```
ironic --ironic-api-version 1.15 node-set-target-raid-config <node-uuid>
  raid_conf.json
```

The output produced should be similar to the following:

```
| Property
                         | Value
                | chassis_uuid
                         | {}
| clean_step
| console_enabled
                         | False
| created_at
                         | 2016-06-14T14:58:07+00:00
| driver
                         | pxe_ilo
| driver_info
                         | {u'ilo_deploy_iso': u'd43e589a-07db-4fce-
a06e-98e2f38340b4',
                         | u'deploy_kernel': u'915c5c74-1ceb-4f78-
bdb4-8547a90ac9c0',
                         | u'ilo_address': u'10.1.196.116', u'deploy_ramdisk':
                         | u'154e7024-bf18-4ad2-95b0-726c09ce417a', u'ilo_password':
u'*****',
                         | u'ilo_username': u'Administrator'}
```

```
| driver_internal_info
                        | {u'agent_cached_clean_steps_refreshed': u'2016-06-15
07:16:08.264091',
                         | u'agent_cached_clean_steps': {u'raid': [{u'interface':
u'raid',
                         | u'priority': 0, u'step': u'delete_configuration'},
{u'interface':
                         | u'raid', u'priority': 0, u'step':
u'create_configuration'}], u'deploy': |
                         | [{u'priority': 10, u'interface': u'deploy',
u'reboot requested': False,
                         | u'abortable': True, u'step': u'erase devices'}]},
u'clean_steps': None,
                         | u'hardware_manager_version':
{u'generic_hardware_manager': u'3'},
                         | u'agent_erase_devices_iterations': 1, u'agent_url':
                         | u'http://192.168.245.143:9999', u'agent_last_heartbeat':
1465974974}
| extra
                         | {}
| inspection_finished_at | None
| inspection_started_at | None
| instance_info
                         | {u'deploy_key': u'XXN20N0V9ER429MECETJMUG5YHTK0Q0Z'}
| instance_uuid
                         | None
| last_error
                         None
| maintenance
                         | False
| maintenance_reason
                         None
| name
                         None
| power_state
                         | power off
| properties
                         | {u'cpu_arch': u'amd64', u'root_device': {u'wwn':
u'0x600508b1001ce286'},
                         | u'cpus': 2, u'capabilities':
                         | u'boot_mode:bios,raid_level:6,boot_option:local',
u'memory_mb': 4096,
                         | u'local_gb': 4}
```

```
| provision_state
                         | available
| provision_updated_at
                         | 2016-06-15T07:16:27+00:00
| reservation
                         | padawan-ironic-cp1-c1-m2-mgmt
| target_power_state
                         | power off
| target_provision_state | None
| target_raid_config
                         | {u'logical_disks': [{u'size_gb': 5, u'raid_level': u'6',
                         | u'is_root_volume': True}]}
                         | 2016-06-15T07:44:22+00:00
| updated_at
| uuid
                         | 22ab9f85-71a1-4748-8d6b-f6411558127e
```

Now set the state of the node to manageable:

```
ironic --ironic-api-version latest node-set-provision-state <node-uuid> manage
```

#### 3. Manual cleaning steps:

Manual cleaning is enabled by default in production - the following are the steps to enable cleaning if the manual cleaning has been disabled.

- Provide cleaning\_network\_uuid in ironic-conductor.conf
- Edit ~/helion/my\_cloud/definition/data/ironic/ironic\_config.yml file and set enable\_node\_cleaning to true
- Then run the following set of commands:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "enabling node cleaning"
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
ansible-playbook -i hosts/localhost ready-deployment.yml
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ironic-reconfigure.yml
```

After performing these steps, the state of the node will become **Cleaning**.

#### Run the following command:

```
ironic --ironic-api-version latest node-set-provision-state 2123254e-8b31-4259-
b632-fb232f5aa6fd clean --clean-steps '[{ "interface": "raid", "step":
   "delete_configuration"},{ "interface": "raid", "step": "create_configuration"}]'
```

#### Node-information after a Manual cleaning:

```
+-----+
| Property
                      | Value
             | chassis_uuid
| clean_step
                     | {}
| console_enabled
                      | False
| created_at
                      | 2016-06-14T14:58:07+00:00
| driver
                      | pxe_ilo
| driver_info
                      | {u'ilo_deploy_iso': u'd43e589a-07db-4fce-
a06e-98e2f38340b4',
                      | u'deploy_kernel': u'915c5c74-1ceb-4f78-
bdb4-8547a90ac9c0',
                      | u'ilo_address': u'10.1.196.116', u'deploy_ramdisk':
                      u'154e7024-bf18-4ad2-95b0-726c09ce417a', u'ilo_password':
                      | u'ilo_username': u'Administrator'}
| driver_internal_info
                      | {u'agent_cached_clean_steps_refreshed': u'2016-06-15
07:16:08.264091',
                      | u'agent_cached_clean_steps': {u'raid': [{u'interface':
u'raid',
                      | u'priority': 0, u'step': u'delete_configuration'},
{u'interface':
                      | u'raid', u'priority': 0, u'step':
u'create_configuration'}], u'deploy': |
                      | [{u'priority': 10, u'interface': u'deploy',
u'reboot_requested': False, |
```

```
| u'abortable': True, u'step': u'erase_devices'}]},
u'clean steps': None,
                         | u'hardware_manager_version':
{u'generic_hardware_manager': u'3'},
                         | u'agent_erase_devices_iterations': 1, u'agent_url':
                         | u'http://192.168.245.143:9999', u'agent_last_heartbeat':
1465974974}
| extra
                         | {}
| inspection_finished_at | None
| inspection_started_at | None
                         | {u'deploy_key': u'XXN20N0V9ER429MECETJMUG5YHTK0Q0Z'}
| instance_info
| instance_uuid
                         | None
| last_error
                         | None
| maintenance
                         | False
| maintenance_reason
                         None
| name
                         | None
                         | power off
| power_state
                         | {u'cpu_arch': u'amd64', u'root_device': {u'wwn':
| properties
u'0x600508b1001ce286'},
                         | u'cpus': 2, u'capabilities':
                         | u'boot_mode:bios,raid_level:6,boot_option:local',
u'memory_mb': 4096,
                         | u'local_gb': 4}
                         | manageable
| provision_state
| provision_updated_at
                         | 2016-06-15T07:16:27+00:00
| raid_config
                         | {u'last_updated': u'2016-06-15 07:16:14.584014',
u'physical_disks':
                         | [{u'status': u'ready', u'size_gb': 1024,
u'interface_type': u'sata',
                         | u'firmware': u'HPGC', u'controller': u'Smart Array P440ar
in Slot 0
```

```
| (Embedded)', u'model': u'ATA MM1000GBKAL',
u'disk_type': u'hdd',
                        | u'id': u'1I:3:3'}, {u'status': u'ready', u'size_gb':
1024,
                        | u'interface_type': u'sata', u'firmware': u'HPGC',
u'controller': u'Smart |
                        | Array P440ar in Slot 0 (Embedded)', u'model': u'ATA
MM1000GBKAL', |
                        | u'disk_type': u'hdd', u'id': u'1I:3:1'}, {u'status':
u'active',
                        | u'size gb': 1024, u'interface type': u'sata',
u'firmware': u'HPGC',
                        | u'controller': u'Smart Array P440ar in Slot 0
(Embedded)', u'model':
                                    MM1000GBKAL', u'disk type': u'hdd', u'id':
                        | u'ATA
u'1I:3:2'},
                        | {u'status': u'active', u'size_gb': 1024,
u'interface type': u'sata',
                        | u'firmware': u'HPGC', u'controller': u'Smart Array P440ar
in Slot 0
                        | (Embedded)', u'model': u'ATA MM1000GBKAL',
u'disk_type': u'hdd',
                        | u'id': u'2I:3:6'}, {u'status': u'active', u'size_gb':
1024,
                        | u'interface_type': u'sata', u'firmware': u'HPGC',
u'controller': u'Smart |
                        | Array P440ar in Slot 0 (Embedded)', u'model': u'ATA
MM1000GBKAL', |
                        | u'disk_type': u'hdd', u'id': u'2I:3:5'}, {u'status':
u'active',
                        | u'size_gb': 1024, u'interface_type': u'sata',
u'firmware': u'HPGC',
                        | u'controller': u'Smart Array P440ar in Slot 0
(Embedded)', u'model':
                                    MM1000GBKAL', u'disk_type': u'hdd', u'id':
                        | u'ATA
u'1I:3:4'}],
                        | u'logical_disks': [{u'size_gb': 4, u'physical_disks':
[u'1I:3:2',
                        | u'2I:3:6', u'2I:3:5', u'1I:3:4'], u'raid level': u'6',
                        | u'is_root_volume': True, u'root_device_hint': {u'wwn':
                        | u'0x600508b1001ce286'}, u'controller': u'Smart Array
P440ar in Slot 0
                        | (Embedded)', u'volume_name':
u'015E795CPDNLH0BRH8N406AAB7'}]}
```

After the manual cleaning, run the following command to change the state of a node to available:

```
ironic --ironic-api-version latest node-set-provision-state <node-uuid> provide
```

# 13.10 Audit Support for Ironic

### 13.10.1 API Audit Logging

Audit middleware supports delivery of CADF audit events via Oslo messaging notifier capability. Based on <a href="mailto:notification\_driver">notification\_driver</a> configuration, audit events can be routed to messaging infrastructure (<a href="mailto:notification\_driver">notification\_driver</a> = messagingv2) or can be routed to a log file (<a href="mailto:notification\_driver">notification\_driver</a> = log).

Audit middleware creates two events per REST API interaction. The first event has information extracted from request data and the second one contains information on the request outcome (response).

211 Audit Support for Ironic

### 13.10.2 Enabling API Audit Logging

You can enable audit logging for Ironic by changing the configuration in the input model. Edit the file \_/helion/my\_cloud/definition/cloudConfig.yml and in the audit-settings section, change the \_default \_value to \_enabled. The ironic-ansible playbooks will now enable audit support for Ironic.

API audit events will be logged in the corresponding audit directory, for example, <a href="//var/audit/ironic/ironic-api-audit.log">/var/audit/ironic/ironic-api-audit.log</a>. An audit event will be logged in the log file for every request and response for an API call.

### 13.10.3 Sample Audit Event

The following output is an example of an audit event for an ironic node-list command:

```
"event_type": "audit.http.request",
"timestamp": "2016-06-15 06:04:30.904397",
"payload":{
   "typeURI": "http://schemas.dmtf.org/cloud/audit/1.0/event",
   "eventTime": "2016-06-15T06:04:30.903071+0000",
   "target":{
      "id":"ironic",
      "typeURI": "unknown",
      "addresses":[
         {
            "url": "http://{ironic_admin_host}:6385",
            "name": "admin"
         },
        {
            "url": "http://{ironic_internal_host}:6385",
            "name":"private"
        },
            "url": "http://{ironic_public_host}:6385",
            "name": "public"
        }
      ],
      "name":"ironic"
   "observer":{
      "id":"target"
   "tags":[
```

```
"correlation_id?value=685flabb-620e-5d5d-b74a-b4135fb32373"
      ],
      "eventType": "activity",
      "initiator":{
         "typeURI": "service/security/account/user",
         "name": "admin",
         "credential":{
            "token":"***",
            "identity_status":"Confirmed"
         },
         "host":{
            "agent": "python-ironicclient",
            "address":"10.1.200.129"
         },
         "project_id":"d8f52dd7d9e1475dbbf3ba47a4a83313",
         "id": "8c1a948bad3948929aa5d5b50627a174"
      },
      "action": "read",
      "outcome": "pending",
      "id": "061b7aa7-5879-5225-a331-c002cf23cb6c",
      "requestPath":"/v1/nodes/?associated=True"
   },
   "priority":"INFO",
   "publisher_id":"ironic-api",
   "message_id":"2f61ebaa-2d3e-4023-afba-f9fca6f21fc2"
}
```

213 Sample Audit Event

# 14 Installation for HPE Helion Entry-scale Cloud with Swift Only

This page describes the installation step requirements for the HPE Helion Entry-scale Cloud with Swift Only model.

## 14.1 Important Notes

- If you are looking for information about when to use the GUI installer and when to use the CLI, see the *Installation Overview*.
- Review the Book "Planning", Chapter 2 "Hardware and Software Support Matrix" that we have listed.
- Review the release notes to make yourself aware of any known issues and limitations.
- The installation process can occur in different phases. For example, you can install the control plane only and then add Compute nodes afterwards if you would like.
- If you run into issues during installation, we have put together a list of *Chapter 18, Troubleshooting the Installation* you can reference.
- Make sure all disks on the system(s) are wiped before you begin the install. (For Swift, refer to )
- There is no requirement to have a dedicated network for OS-install and system deployment, this can be shared with the management network. More information can be found in *Book "Planning"*, *Chapter 10 "Example Configurations"*.
- You may see the terms deployer and lifecycle manager used interchangeably. These are referring to the same nodes in your environment.
- When running the Ansible playbook in this installation guide, if a runbook fails you will see in the error response to use the <a href="e--limit">--limit</a> switch when retrying a playbook. This should be avoided. You can simply re-run any playbook without this switch.

214 Important Notes

- DVR is not supported with ESX compute.
- When you attach a Cinder volume to the VM running on the ESXi host, the volume will
  not get detected automatically. Make sure to set the image metadata vmware\_adaptertype = lsiLogicsas for image before launching the instance. This will help to discover the
  volume change appropriately.

### 14.2 Before You Start

We have put together a *Chapter 2, Pre-Installation Checklist* that should help with the recommended pre-installation tasks.

# 14.3 Setting Up the Lifecycle Manager

#### **Installing the Lifecycle Manager**

The lifecycle manager will contain the installation scripts and configuration files to deploy your cloud. You can set up the lifecycle manager on a dedicated node or you do so on your first controller node. The default choice is to use the first controller node as the lifecycle manager.

- 1. Sign in and download the product and signature files at the link below:
  - a. Software Entitlement Portal (http://www.hpe.com/software/entitlements)
- 2. You can verify the download was complete via the signature verification process outlined here (https://h20392.www2.hpe.com/portal/swdepot/displayProductInfo.do?product-Number=HPLinuxCodeSigning) .
- 3. Boot your lifecycle manager from the ISO contained in the download.
- 4. Enter "install" to start installation.



#### Note

"install" is all lower case

- 5. Select the language. Note that only the English language selection is currently supported.
- 6. Select the location.

215 Before You Start

- 7. Select the keyboard layout.
- 8. Select the primary network interface, if prompted:
  - a. Assign IP address, subnet mask, and default gateway
- 9. Create new account:
  - a. Enter a username.
  - b. Enter a password.
  - c. Enter time zone.

Once the initial installation is finished, complete the lifecycle manager setup with these steps:

- 1. Ensure your lifecycle manager has a valid DNS nameserver specified in <a href="//etc/resolv.conf">/etc/resolv.conf</a>.
- 2. Set the environment variable LC\_ALL:

export LC\_ALL=C



Note

This can be added to ~stack/.bashrc or /etc/bash.bashrc.

At the end of this section you should have a node set up with Linux for HPE Helion on it.

#### Configure and Run the Lifecycle Manager



### **Important**

It is critical that you don't run any of the commands below as the <u>root</u> user or use <u>sudo</u>, unless it is stated explicitly in the steps. Run then as the user you just created (or <u>stack</u> if you left the default of "stack").

 Log into your lifecycle manager as the user you created and mount the install media at / media/cdrom. It may be necessary to use wget or another file transfer method to transfer the install media to the lifecycle manager before completing this step. Here is the command to mount the media:

Example for HPE Helion OpenStack® 5.0:

```
sudo mount HelionOpenStack-5.0.iso /media/cdrom
```

2. Unpack the tarball that is in the <a href="media/cdrom/hos/">/media/cdrom/hos/</a> directory: Example for HPE Helion OpenStack 5.0

```
tar xvf /media/cdrom/hos/hos-5.0.0-<timestamp>.tar
```

3. Run the hos-init.bash script which is included in the build: Example for HPE Helion OpenStack® 5.0

```
~/hos-5.0.0/hos-init.bash
```

You will be prompted to enter an optional SSH passphrase when running <a href="hos-init.bash">hos-init.bash</a>. This passphrase is used to protect the key used by Ansible when connecting to its client nodes. If you do not want to use a passphrase then just press <a href="Enter">Enter</a> at the prompt. For automated installation (e.g. CI) it is possible to disable SSH passphrase prompting by setting the <a href="https://hos-init.bash">HOS\_INIT\_AUTO</a> environment variable before running <a href="hos-init.bash">hos-init.bash</a>, like this:

```
export HOS_INIT_AUTO=y
```

If you have protected the SSH key with a passphrase then execute the following commands to avoid having to enter the passphrase on every attempt by Ansible to connect to its client nodes:

```
eval $(ssh-agent)
ssh-add ~/.ssh/id_rsa
```

At the end of this section you should have a local directory structure, as described below:

```
~/helion/
~/helion/examples/
clouds
~/helion/my_cloud/definition/
~/helion/my_cloud/config/
~/helion/hos/ansible directory
Top level directory
Directory contains the config input files
Directory contains the config input files
Directory contains .j2 files which are symlinks to the
Allouder to the config input files
Directory contains .j2 files which are symlinks to the
Allouder to the config input files
Directory contains .j2 files which are symlinks to the
Allouder to the config input files
Directory contains .j2 files which are symlinks to the
Allouder to the example config input files of the examp
```



### Warning

It is important that you do not add any extra files in the  $\frac{\text{-/helion}}{\text{-/helion}}$  directory on your lifecycle manager. This includes temporary save files. Doing so will cause errors during the installation.

# •

### **Important**

Be aware that the files in the \_-/helion/my\_cloud/config/ directory are symlinks to the \_-/helion/hos/ansible/ directory. For example, \_-/helion/my\_cloud/config/keepalived/defaults.yml is a symbolic link to \_-/helion/hos/ansible/roles/keepalived/defaults/main.yml.

```
ls -al ~/helion/my_cloud/config/keepalived/defaults.yml
lrwxrwxrwx 1 stack stack 55 May 24 20:38 /home/stack/helion/my_cloud/config/
keepalived/defaults.yml -> ../../hos/ansible/roles/keepalived/defaults/main.yml
```

If you are using a tool like <u>sed</u> to make edits to files in this directory, you might break the symbolic link and create a new copy of the file. To maintain the link, you will need to force **sed** to follow the link:

```
sed -i --follow-symlinks \
's$keepalived_vrrp_offset: 0$keepalived_vrrp_offset: 2$' \
~/helion/my_cloud/config/keepalived/defaults.yml
```

Alternatively, directly edit the target of the link <u>~/helion/hos/ansible/roles/</u>keepalived/defaults/main.yml.

For any troubleshooting information regarding these steps, see Section 18.1, "Issues during Lifecy-cle-manager Setup".

# 14.4 Configure Your Environment

This part of the install is going to depend on the specific cloud configuration you are going to use.

- 1. Setup your configuration files, as follows:
  - a. See the sample sets of configuration files in the <u>~/helion/examples/</u> directory. Each set will have an accompanying README.md file that explains the contents of each of the configuration files.
  - b. Copy the example configuration files into the required setup directory and edit them to contain the details of your environment:

```
cp -r ~/helion/examples/entry-scale-swift/* ~/helion/my_cloud/definition/
```

c. Begin inputting your environment information into the configuration files in the \_~/ helion/my\_cloud/definition directory.

Full details of how to do this can be found here: .

In many cases, the example models provide most of the data you need to create a valid input model. However, there are two important aspects you must plan and configure before starting a deploy as follows:

- Check the disk model used by your nodes. Specifically, check that all disk drives are correctly named and used as described in .
- Select an appropriate partition power for your rings. Detailed information about this is provided at .

Optionally, you can use the <a href="hosencrypt.py">hosencrypt.py</a> script to encrypt your iLo passwords. This script uses OpenSSL.

a. Change to the Ansible directory:

```
cd ~/helion/hos/ansible
```

b. Put the encryption key into the following environment variable:

```
export HOS_USER_PASSWORD_ENCRYPT_KEY=<encryption key>
```

c. Run the python script below and follow the instructions. Enter a password that you want to encrypt.

```
hosencrypt.py
```

- d. Take the string generated and place it in the <a href=""">"ilo\_password"</a> field in your <a href="">~/he-lion/my\_cloud/definition/data/servers.yml"</a> file, remembering to enclose it in quotes.
- e. Repeat the above for each server.



#### Note

Before you run any playbooks, remember that you need to export the encryption key in the following environment variable: <a href="mailto:export HOS\_USER\_PASSWORD\_ENCRYP-T\_KEY=<encryption">export HOS\_USER\_PASSWORD\_ENCRYP-T\_KEY=<encryption</a> key>

Commit your configuration to the local git repo (*Chapter 3, Using Git for Configuration Management*), as follows:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```



### **Important**

This step needs to be repeated any time you make changes to your configuration files before you move onto the following steps. See *Chapter 3, Using Git for Configuration Management* for more information.

# 14.5 Running the Configuration Processor

Once you have your configuration files setup you need to run the configuration processor to complete your configuration.

When you run the configuration processor you will be prompted for two passwords. Enter the first password to make the configuration processor encrypt its sensitive data, which consists of the random inter-service passwords that it generates and the ansible <a href="mailto:group\_vars">group\_vars</a> and <a href="mailto:host\_vars">host\_vars</a> that it produces for subsequent deploy runs. You will need this password for subsequent ansible deploy and configuration processor runs. If you wish to change an encryption password that you have already used when running the configuration processor then enter the new password at the second prompt, otherwise just press <a href="mailto:Enter">Enter</a> to bypass this.

Run the configuration processor with this command:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

For automated installation (e.g. CI) you can specify the required passwords on the ansible command line. For example, the command below will disable encryption by the configuration processor

```
ansible-playbook -i hosts/localhost config-processor-run.yml -e encrypt="" -e rekey=""
```

If you receive an error during this step then there is probably an issue with one or more of your configuration files. We recommend that you verify that all of the information in each of your configuration files is correct for your environment and then commit those changes to git using the instructions in the previous section before re-running the configuration processor again.

For any troubleshooting information regarding these steps, see Section 18.3, "Issues while Updating Configuration Files".

# 14.6 Provisioning Your Baremetal Nodes

To provision the baremetal nodes in your cloud deployment you can either use the automated operating system installation process provided by HPE Helion OpenStack or you can use the 3rd party installation tooling of your choice. We will outline both methods below:

#### **Using 3rd Party Baremetal Installers**

If you do not wish to use the automated operating system installation tooling included with 5.0 then the requirements that have to be met using the installation tooling of your choice are:

- The operating system must be installed via the HPE Linux for HPE Helion OpenStack ISO
  provided on the Software Entitlement Portal (http://www.hpe.com/software/entitlements) .
- Each node must have SSH keys in place that allows the same user from the lifecycle manager node who will be doing the deployment to SSH to each node without a password.
- Passwordless sudo needs to be enabled for the user.
- There should be a LVM logical volume as /root on each node.
- If the LVM volume group name for the volume group holding the "root" LVM logical volume is hlm-vg then it will align with the disk input models in the examples.
- Ensure that openssh-server, python, python-apt, and rsync are installed.

If you chose this method for installing your baremetal hardware, skip forward to the *Chapter 10, Installing Mid-scale and Entry-scale KVM* step.

If you would like to use the automated operating system installation tools provided by 5.0 then complete all of the steps below.

# Using the Automated Operating System Installation Provided by HPE Helion OpenStack Part One: Deploy Cobbler

This phase of the install process takes the baremetal information that was provided in <a href="s.yml">s.yml</a> and installs the Cobbler provisioning tool and loads this information into Cobbler. This sets each node to <a href="netboot-enabled: true">netboot-enabled: true</a> in Cobbler. Each node will be automatically marked as <a href="netboot-enabled: false">netboot-enabled: false</a> when it completes its operating system install successfully. Even if the node tries to PXE boot subsequently, Cobbler will not serve it. This is deliberate so that you can't reimage a live node by accident.

The <u>cobbler-deploy.yml</u> playbook prompts for a password - this is the password that will be encrypted and stored in Cobbler, which is associated with the user running the command on the lifecycle manager, that you will use to log in to the nodes via their consoles after install. The username is the same as the user set up in the initial dialogue when installing the lifecycle manager from the iso, and is the same user that is running the cobbler-deploy play.

1. Run the following playbook which confirms that there is iLo connectivity for each of your nodes so that they are accessible to be re-imaged in a later step:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost bm-power-status.yml
```

2. Run the following playbook to deploy Cobbler:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost cobbler-deploy.yml
```

#### Part Two: Image the Nodes

This phase of the install process goes through a number of distinct steps:

- 1. Powers down the nodes to be installed
- 2. Sets the nodes hardware boot order so that the first option is a network boot.
- 3. Powers on the nodes. (The nodes will then boot from the network and be installed using infrastructure set up in the previous phase)

- 4. Waits for the nodes to power themselves down (this indicates a success install). This can take some time.
- 5. Sets the boot order to hard disk and powers on the nodes.
- 6. Waits for the nodes to be ssh-able and verifies that they have the signature expected.

The reimage command is:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost bm-reimage.yml [-e nodelist=node1,node2,node3]
```

If a nodelist is not specified then the set of nodes in cobbler with <a href="netboot-enabled: True">netboot-enabled: True</a> is selected. The playbook pauses at the start to give you a chance to review the set of nodes that it is targeting and to confirm that it's correct.

You can use the command below which will list all of your nodes with the <a href="netboot-enabled">netboot-enabled</a>: True flag set:

```
sudo cobbler system find --netboot-enabled=1
```

For any troubleshooting information regarding these steps, see Section 18.2, "Issues while Provisioning your Baremetal Nodes".

# 14.7 Deploying the Cloud

1. Use the playbook below to create a deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

2. [OPTIONAL] - Run the wipe\_disks.yml playbook to ensure all of your partitions on your nodes are completely wiped before continuing with the installation. If you are using fresh machines this step may not be necessary.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts wipe_disks.yml
```

If you have used an encryption password when running the configuration processor use the command below and enter the encryption password when prompted:

```
ansible-playbook -i hosts/verb_hosts wipe_disks.yml --ask-vault-pass
```

223 Deploying the Cloud

3. Run the site.yml playbook below:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts site.yml
```

If you have used an encryption password when running the configuration processor use the command below and enter the encryption password when prompted:

```
ansible-playbook -i hosts/verb_hosts site.yml --ask-vault-pass
```



### Note

The step above runs <u>osconfig</u> to configure the cloud and <u>hlm-deploy</u> to deploy the cloud. Therefore, this step may run for a while, perhaps 45 minutes or more, depending on the number of nodes in your environment.

4. Verify that the network is working correctly. Ping each IP in the <a href="//etc/hosts">/etc/hosts</a> file from one of the controller nodes.

For any troubleshooting information regarding these steps, see Section 18.4, "Issues while Deploying the Cloud".

### 14.8 Post-Installation Verification and Administration

We recommend verifying the installation using the instructions in *Chapter 22, Cloud Verification*. There are also a list of other common post-installation administrative tasks listed in the *Chapter 28, Other Common Post-Installation Tasks* list.

# 15 Installing SLES Compute

# 15.1 SLES Compute Node Installation Overview

HPE Helion OpenStack 5.0 supports SLES compute nodes, specifically SLES 12 SP2. HPE does not ship a SLES ISO with HPE Helion OpenStack so you will need to download a copy of the SLES iso (SLE-12-SP2-Server-DVD-x86\_64-GM-DVD1.iso) and the SLES SDK ISO (SLE-12-SP2-SDK-DVD-x86\_64-GM-DVD1.iso) from SUSE. You can use the following link to download the ISO. To do so, you will need to sign in or create a SUSE customer service account before downloading: https://www.suse.com/products/server/download/ ...

There are two approaches for deploying SLES compute nodes in HPE Helion OpenStack:

- Using the lifecycle manager to automatically deploy SLES Compute nodes.
- Provisioning SLES nodes yourself, either manually or using a 3rd party tool, and then provide the relevant information to the lifecycle manager.

These two approaches can be used whether you are installing a cloud for the first time or adding a compute node to an existing cloud. Regardless of your approach, you should be certain to register your SLES compute nodes in order to get product updates as they come available. For more information, see *Book "Planning"*, *Chapter 1 "Registering SUSE Linux"*.

# 15.2 SLES Support

SUSE Linux Enterprise Server (SLES) Host OS KVM and/or supported SLES guests have been tested and qualified by HPE to run on HPE Helion OpenStack. Refer to the published compatibility matrix for your version of HPE Helion OpenStack at the following URL http://docs.hpcloud.com/#home.html . HPE is one of SUSE's largest OEMs with follow the sun global support coverage.

• One Number to Call

HPE customers who have purchased both HPE Helion OpenStack and SLES subscriptions with support from HPE will have one number to call for troubleshooting, fault isolation and support from HPE technical support specialists in HPE Helion OpenStack and SUSE technologies. If the problem is isolated to SLES software itself the issue will be replicated on a SUSE certified platform and escalated to SUSE for resolution.

#### A Dual Support Model

HPE will troubleshoot and fault isolate an issue at the Helion software level. If HPE Helion OpenStack software is excluded as the cause of the problem, then customers who did not purchase SLES support from HPE will be directed to the vendor from whom they purchased SLES for continued support.

# 15.3 Using the Lifecycle Manager to Deploy SLES Compute Nodes

The method used for deploying SLES compute nodes using Cobbler on the lifecycle manager uses legacy BIOS.



#### Note

UEFI and Secure boot are currently not supported on SLES compute.

### 15.3.1 Deploying legacy BIOS SLES compute nodes

The installation process for SLES nodes is almost identical to that of HPE Linux nodes as described in the topic for . The key differences are:

• The standard SLES ISO (SLE-12-SP2-Server-DVD-x86\_64-GM-DVD1.iso) must be accessible via /home/stack/sles12.iso. Rename the ISO or create a symbolic link:

```
mv SLE-12-SP2-Server-DVD-x86 64-GM-DVD1.iso /home/stack/sles12.iso
```

The contents of the SLES SDK ISO (SLE-12-SP2-SDK-DVD-x86\_64-GM-DVD1.iso) must be
mounted or copied to /opt/hlm\_packager/hlm/sles12/zypper/SDK/. If you choose to
mount the ISO, we recommend creating an /etc/fstab entry to ensure the ISO is mounted
after a reboot.

- You must identify the node(s) on which you want to install SLES, by adding the key/ value pair <u>distro-id</u>: <u>sles12sp2-x86\_64</u> to server details in <u>servers.yml</u>. You will also need to update <u>net\_interfaces.yml</u>, <u>server\_roles.yml</u>, <u>disk\_compute.yml</u> and <u>control\_plane.yml</u>. For more information on configuration of the Input Model for SLES, see.
- HPE Helion OpenStack playbooks currently do not take care of the SDK, therefore it needs to be added manually. The following command needs to be run on every SLES compute node:

```
deployer_ip=192.168.10.254
zypper addrepo --no-gpgcheck --refresh http://$deployer_ip:79/hlm/sles12/zypper/SDK
SLES-SDK
```

There is no need to add OS repo as in case of *Section 15.4, "Provisioning SLES Yourself"*, because SLES already comes with OS repo populated after a Cobbler-managed install.

# 15.4 Provisioning SLES Yourself

#### 15.4.1 Introduction

This article outlines the steps needed to manually provision a SLES node so that it can be added to a new or existing HPE Helion OpenStack 5.0 cloud.

### 15.4.2 Configure Lifecycle Manager to Enable SLES

- Take note of the lifecycle manager's IP address. It will be used below during .
- Mount or copy the contents of SLE-12-SP2-Server-DVD-x86\_64-GM-DVD1.iso to /opt/ hlm\_packager/hlm/sles12/zypper/0S/
- Mount or copy the contents of SLE-12-SP2-SDK-DVD-x86\_64-GM-DVD1.iso to /opt/ hlm packager/hlm/sles12/zypper/SDK/



If you choose to mount an ISO, we recommend creating an /etc/fstab entry to ensure the ISO is mounted after a reboot.

#### 15.4.3 Install SLES 12 SP2

Install SLES 12 SP2 using the standard iso (SLE-12-SP2-Server-DVD-x86\_64-GM-DVD1.iso)

### 15.4.4 Assign a static IP

1. Use the ip addr command to find out what network devices are on your system:

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
link/loopback 00:00:00:00:00 brd 00:00:00:00:00
inet 127.0.0.1/8 scope host lo
valid_lft forever preferred_lft forever
inet6 ::1/128 scope host
valid_lft forever preferred_lft forever

2: eno1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP qlen 1000
link/ether f0:92:1c:05:89:70 brd ff:ff:ff:ff
inet 10.13.111.178/26 brd 10.13.111.191 scope global eno1
valid_lft forever preferred_lft forever
inet6 fe80::f292:lcff:fe05:8970/64 scope link
valid_lft forever preferred_lft forever

3: eno2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP qlen 1000
link/ether f0:92:1c:05:89:74 brd ff:ff:ff:ff:ff:ff
```

2. Identify the one that matches the MAC address of your server and edit the corresponding config file in /etc/sysconfig/network-scripts.

```
vi /etc/sysconfig/network-scripts/ifcfg-enol
```

3. Edit the <u>IPADDR</u> and <u>NETMASK</u> values to match your environment. Note that the <u>IPADDR</u> is used in the corresponding stanza in <u>servers.yml</u>. You may also need to set <u>BOOTPROTO</u> to none.

```
TYPE=Ethernet

BOOTPROTO=none

DEFROUTE=yes
```

228 Install SLES 12 SP2

```
PEERROUTES=yes
PEERROUTES=yes
IPV4_FAILURE_FATAL=no
IPV6INIT=yes
IPV6_AUTOCONF=yes
IPV6_DEFROUTE=yes
IPV6_PEERDNS=yes
IPV6_PEERROUTES=yes
IPV6_FAILURE_FATAL=no
NAME=eno1
UUID=36060f7a-12da-469b-alda-ee730a3b1d7c
DEVICE=eno1
ONBOOT=yes
NETMASK=255.255.255.192
IPADDR=10.13.111.14
```

4. [OPTIONAL] Reboot your SLES node and ensure that it can be accessed from the lifecycle manager.

### 15.4.5 Add stack user and home directory

```
useradd -m stack
passwd stack
```

### 15.4.6 Allow user stack to sudo without password

Setting up sudo on SLES is covered in the SUSE documentation, Configuring sudo (https://www.suse.com/documentation/sles-12/book\_sle\_admin/data/sec\_sudo\_conf.html) ₽.

The recommendation is to create user specific sudo config files under /etc/sudoers.d, therefore creating an /etc/sudoers.d/stack config file with the following content will allow sudo commands without the requirement of a password.

```
stack ALL=(ALL) NOPASSWD:ALL
```

# 15.4.7 Add zypper repository

Using the ISO-based repositories created above, add the zypper repositories.

Follow these steps. Update the value of deployer\_ip as necessary.

```
deployer_ip=192.168.10.254
zypper addrepo --no-gpgcheck --refresh http://$deployer_ip:79/hlm/sles12/zypper/OS SLES-
OS
zypper addrepo --no-gpgcheck --refresh http://$deployer_ip:79/hlm/sles12/zypper/SDK SLES-
SDK
```

To verify that the repositories have been added, run:

```
zypper repos --detail
```

You can find more information on Using Zypper in the SLES 12 Admin guide: https://www.suse.com/documentation/sles-12/book\_sle\_admin/data/sec\_zypper.html

### 15.4.8 Add Required Packages

As documented in the section of , you will need to add some extra packages that are required. Ensure that openssh-server, python, and rsync are installed.

### 15.4.9 Set up passwordless SSH access

Once you have started your installation using the lifecycle manager, or if you are adding a SLES node to an existing cloud, you need to copy the lifecycle manager public key to the SLES node. One way of doing this is to copy the <a href="https://home/stack/.ssh/authorized\_keys">/home/stack/.ssh/authorized\_keys</a> from another node in the cloud to the same location on the SLES node. If you are installing a new cloud, this file will be available on the nodes after running the bm-reimage.yml playbook.



Ensure that there is global read access to the file /home/stack/.ssh/authorized\_keys.

Now test passwordless ssh from the deployer and check your ability to remotely execute sudo commands:

```
ssh stack@<<ip of sles node>> "sudo tail -5 /var/log/messages"
```

230 Add Required Packages

# 15.5 Using SLES as a Ceph Client



### Warning

SUSE Linux Enterprise Server 12 SP2 comes with Ceph packages from the Ceph Jewel release. HPE Helion OpenStack 5.0 Ceph server is based on the Ceph Hammer release, therefore it is not advised that client related Ceph management commands from a SLES compute node be performed. Doing so could trigger an update of the Ceph config which would not be compatible with other non-SLES Ceph clients in the same HPE Helion OpenStack 5.0 system.

# 16 HLM-Hypervisor instructions

### 16.1 Introduction

You can read an overview of the design on the Helion lifecycle manager creation of VMs (https://wiki.hpcloud.net/display/core/HLM+creation+of+VMs) page. This page describes the steps needed to set up and use a HLM-Hypervisor (previously known as a VM-factory) and related functionality of HPE Helion OpenStack 5.0. This comprises model changes and running a couple of new playbooks to bring up the VMs.

# 16.2 Model changes

### 16.2.1 passthrough-network-groups

With PB3 the specification of <a href="mailto:passthrough-network-groups">passthrough-network-groups</a> is now supported within the Interfaces section of the input model.

A <u>passthrough-network</u> group is a network group that can be accessed by a VM hosted by the hypervisor system regardless of whether the network group is required on the hypervisor system itself.

#### network-groups:

List of network groups that this server may require access to.

Access to the network is only configured if the network-group is required by at least one service hosted on this server.

#### forced-network-groups:

List of network-groups that this server requires access to.

Access to the network is configured regardless of whether the network-group is required by any service hosted on this server.

Network groups should not appear in both <a href="network-groups">network-groups</a> and <a href="forced-net-work-groups">forced-net-work-groups</a>.

#### passthrough-network-groups:

List of network-groups that need to be passed through this server to hosted VMs.

232 Introduction

No access is configured for this server if the network group is listed <u>only</u> in <u>passthrough</u>-network-groups.

The <a href="mailto:passthrough-network-groups">passthrough-network-groups</a> may include networks from either <a href="mailto:network-groups">network-groups</a> but may also include network groups not used directly by this server. If a <a href="mailto:passthrough-network-groups">passthrough-network-groups</a> is listed in either of <a href="mailto:network-groups">network-groups</a> or <a href="mailto:forced-network-groups">forced-network-groups</a> the access to the network group will be such that both the server and hosted VMs are network peers.

### 16.2.1.1 Example:

The following hypervisor interface-model example specifies:

- bond0
  - PXE available to hypervisor and Hosted-VMs
  - CLM available to hypervisor and Hosted-VMs
  - CAN available to Hosted-VMs only
- bond1
  - VxLAN-VLAN1-TUL, EXT and VLAN2-TUL available to Hosted-VMs only

```
interface-models:
  - name: HLM-HYPERVISOR-INTERFACES
    network-interfaces:
      - name: BOND0
        device:
          name: bond0
        bond-data:
          options:
            mode: active-backup
           miimon: 200
            primary: hed5
          provider: linux
          devices:
            - name: hed5
        network-groups:
          - PXE
          - CLM
        passthrough-network-groups:
          - PXE
```

```
- CLM
    - CAN
- name: BOND1
 device:
   name: bond1
 bond-data:
   options:
     mode: active-backup
     miimon: 200
     primary: hed1
   provider: linux
   devices:
     - name: hed1
 passthrough-network-groups:
   - V×LAN-VLAN1-TUL
    - EXT
    - VLAN2-TUL
```

### 16.2.2 hlm-hypervisor

With PB3 we now support the specification of <a href="https://hlm-hypervisor">hlm-hypervisor</a> boolean within the Servers section of the input model. This setting is used to identify nodes that can host VMs (non-NOVA). This is related to the specification of <a href="https://hypervisor-id">hypervisor-id</a> that is required to identify servers to be instantiated as hosted-VMs.

- Any server referenced as a target of a <a href="https://hypervisor.id">hypervisor.id</a> must have <a href="https://hupervisor.id">hlm-hypervisor</a> set to True.
- A server with <a href="https://htm-hypervisor">https://htm-hypervisor</a> of True setting may or may not have Helion lifecycle manager Hosted-VMs associated with it.

### 16.2.2.1 Example

In the following example from servers.yml

- three nodes (<u>controller1</u>, <u>controller2</u> and <u>controller3</u>) are declared as Hosted-VMs on hypervisor1 (hypervisor-id: hypervisor1)

```
- id: hypervisor1
```

```
ip-addr: 10.225.107.74
 role: HLM-HYPERVISOR-ROLE
 server-group: RACK2
 hlm-hypervisor: True
 nic-mapping: HP-BL460-6PORT
 mac-addr: 9c:7e:96:22:9e:d8
 ilo-ip: 10.1.18.42
 ilo-password: whatever
 ilo-user: whatever
- id: controller1
 ip-addr: 10.225.107.75
 role: CONTROLLER-ROLE
 hypervisor-id: hypervisor1
 nic-mapping: VIRTUAL-1-PORT
- id: controller2
 ip-addr: 10.225.107.76
 role: CONTROLLER-ROLE
 hypervisor-id: hypervisor1
 nic-mapping: VIRTUAL-1-PORT
- id: controller3
 ip-addr: 10.225.107.77
 role: CONTROLLER-ROLE
 hypervisor-id: hypervisor1
 nic-mapping: VIRTUAL-1-PORT
- id: compute1
 ip-addr: 10.225.107.78
 role: COMPUTE-ROLE
 server-group: RACK1
 nic-mapping: HP-BL460-6PORT
 mac-addr: 6c:9e:96:22:7e:d8
 ilo-ip: 10.1.18.44
 ilo-password: whatever
 ilo-user: whatever
```

The basic steps are as follows, with examples below. In all of these examples, it is assumed that hed1 is the physical interface to which the VMs need to be connected.

1. Create a new network interfaces group for the Hypervisor nodes, appropriate to the hardware in use. It is called HLM-HYPERVISOR-INTERFACES in this example. You may also want to create nic mappings for these nodes.

```
@@ -94,3 +94,12 @@ network-groups:
```

```
- MANAGEMENT

+ - name: HLM-HYPERVISOR-INTERFACES
+ network-interfaces:
+ - name: hed1
+ device:
+ name: hed1
+ network-groups:
+ EXTERNAL-VM
+ GUEST
+ MANAGEMENT
```

2. Create a new disk model for the Hypervisor nodes. Here's an example for a Hypervisor node with 4 disks. In this specific example, we've included storage to be set-aside in the HPE Helion OpenStack input model for gluster.

```
product:
  version: 2
disk-models:
- name: HLM-HYPERVISOR-DISKS
  volume-groups:
    - name: hlm-vg
      physical-volumes:
       - /dev/sda_root
       - /dev/sdb
      logical-volumes:
        - name: root
          size: 35%
          fstype: ext4
          mount: /
        - name: log
          size: 50%
          mount: /var/log
          fstype: ext4
          mkfs-opts: -0 large_file
        - name: crash
          size: 10%
          mount: /var/crash
          fstype: ext4
          mkfs-opts: -0 large_file
```

```
# optional VG dedicated to VM images to keep VM IOPS off the OS disk
  # The consumer stanza allows user defined location of where to store the gcows
  - name: hlm-images
    physical-volumes:
      - /dev/sdc
      - /dev/sdd
    logical-volumes:
      - name: hlm-images
        size: 95%
        mount: /var/lib/hlm-images
        fstype: ext4
        mkfs-opts: -0 large_file
        consumer:
    name: hlm-hypervisor
    usage: hlm-hypervisor-images
# Optionally define some storage in the input model to be used
# by gluster
device-groups:
  - name: gluster
    devices:
      - name: /dev/sde
      - name: /dev/sdf
    consumer:
      name: gluster
      suppress-warnings: True
```

Note the <u>suppress-warnings</u>: <u>True</u> flag for "gluster". This indicates to the configuration processor that is should not issue warnings when it cannot find a "gluster" service definition.

3. Create a new role that uses this interface and disk model.

4. Define a new type of node in the "clusters" section of control\_plane.yml so that nodes with this role will get their own CP group.

```
@@ -135,3 +135,10 @@
```

```
- ntp-client
- vsa

+ - name: hypervisor
+ resource-prefix: vmf
+ server-role: HLM-HYPERVISOR-ROLE
+ allocation-policy: strict
+ min-count: 0
+ service-components:
- lifecycle-manager-target
+ ntp-server
```

If the lifecycle manager is also a hypervisor then you should add the component <u>lifecy-cle-manager</u> instead of <u>lifecycle-manager-target</u>. It is required that the hypervisor nodes have the ntp-server component.

- 5. Assign the HLM-HYPERVISOR-ROLE to the appropriate nodes in servers.yml
  - a. If the lifecycle manager node is also going to be a Hypervisor, then give it the HLM-HYPERVISOR-ROLE in servers.yml
  - b. Having a baremetal node that is both a Hypervisor and a Controller is not supported.
  - c. Also add definitions for all of the control plane VMs that you intend to create later. Note that you don't need iLO information or mac-addresses for VMs.
  - d. You will need to modify the IP addresses etc in this example to match your machine.

```
@ -80,7 +80,7 @
    - id: hypervisor1
       ip-addr: 10.225.107.74
       role: HLM-HYPERVISOR-ROLE
+
+
       hlm-hypervisor: True
       server-group: RACK2
+
       nic-mapping: HP-BL460-6PORT
       mac-addr: 9c:7e:96:22:9e:d8
       ilo-ip: 10.1.18.42
       ilo-password: whatever
       ilo-user: whatever
+
     - id: controller1
+
       ip-addr: 10.225.107.75
       role: CONTROLLER-ROLE
       hypervisor-id: hypervisor1
       nic-mapping: VIRTUAL-1-PORT
```

```
- id: controller2
+
       ip-addr: 10.225.107.76
       role: CONTROLLER-ROLE
       hypervisor-id: hypervisor1
       nic-mapping: VIRTUAL-1-PORT
     - id: controller3
       ip-addr: 10.225.107.77
       role: CONTROLLER-ROLE
       hypervisor-id: hypervisor1
       nic-mapping: VIRTUAL-1-PORT
     - id: compute1
+
+
       ip-addr: 10.225.107.78
       role: COMPUTE-ROLE
       server-group: RACK1
       nic-mapping: HP-BL460-6PORT
       mac-addr: 6c:9e:96:22:7e:d8
       ilo-ip: 10.1.18.44
       ilo-password: whatever
+
       ilo-user: whatever
```

6. Add a definition of this new nic-mapping to nic\_mappings.yml:

```
name: VIRTUAL-1-PORT
physical-ports:
bus-address: '0000:01:01.0'
logical-name: hed1
type: simple-port
```

See the section NIC Mapping for Helion lifecycle manager Virtual-Controllers (#nic\_map-ping\_for\_hlm\_virtual\_controllers) below for more information and examples.

- 7. You will also need to make model changes for the vms. The vm role above is CONTROLLER-ROLE. You need to add information to the model that specifies:
  - a. The number of vcpus this is done by adding a <a href="cpu-model">cpu-model</a> to the vm role with a <a href="vm-size">vm-size</a> stanza
  - b. The amount of RAM to be used this is done by adding a <u>memory-model</u> to the vm role with a vm-size stanza
  - c. The sizes of the virtual disks this is done by adding a <u>vm-size</u> stanza to the disk-model used with the vm-role.

239 hlm-hypervisor

The CONTROLLER-ROLE definition will look like:

```
    name: CONTROLLER-ROLE
    interface-model: CONTROLLER-INTERFACES
    disk-model: CONTROLLER-DISKS
    cpu-model: CONTROLLER-CPU
    memory-model: CONTROLLER-MEMORY
```

The cpu-model will look like the following - place it in a yaml file and set <u>vcpus</u> to the value that you want to use:

```
product:
version: 2

cpu-models:
- name: CONTROLLER-CPU
vm-size:
vcpus: 4
```

The memory-model will look like the following - place it in a yaml file and set <u>ram</u> to the value that you want to use (valid values are in "KMG"):

```
product:
version: 2
memory-models:
- name: CONTROLLER-MEMORY
vm-size:
ram: 16G
```

The disk-model will have a stanza like the following added to it:

240 hlm-hypervisor

<u>disks</u> is a list of all the physical-volumes in volume-groups and devices in device-groups used in the disk-model, one entry for each with the name in the <u>name</u>: field and the size in the size: field. Valid values for size are KMGT

- 8. Note hat the network that is associated with ansible traffic needs be on the network interface in the vm which is on the "lowest" pci address. In the nic-mapping description below this would be hed1.
- 9. You do need to have a vda root described in your input model,

## 16.3 Bootstrap Instructions

Start with one bare-metal node and install it directly from the HPE Helion OpenStack 5.0 iso e.g. through virtual media on its iLO. Follow the usual install instructions e.g. but make sure you use a model with the changes described above. Follow all the usual steps, including cobbler-deploy, bm-reimage and config-processor-run. Stop after running <a href="ready-deployment.yml">ready-deployment.yml</a> though, and move to the instructions in the next section, to bring up the virtual control plane nodes.

## 16.4 Customizing for VCP

The <u>ready-deployment.yml</u> playbook will have created <u>~/scratch/ansible/next/hos/ansible</u> so change to that directory and perform the rest of your work from that directory.

```
cd ~/scratch/ansible/next/hos/ansible
```

There are a number of installation choices at this point:

- 1. You are installing a VCP system which is fully described and managed by input model.
- 2. You are installing a VCP system where additional steps need to be performed after provisioning the VCP VMs before running site.yml
- 3. You are installing a VCP system where additional steps need to be performed between configuring the hypervisor nodes, deploying the hypervisors, provisioning the VCP VMs, or running site.yml

241 Bootstrap Instructions

## 16.4.1 Installing a fully input model managed Virtual Control Plane based HPE Helion OpenStack Cloud

The plays to setup the hlm-hypervisor have been encapsulated into a single play that can be run before running site.yml to bring up the full cloud.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts hlm-hypervisor-setup.yml
ansible-playbook -i hosts/verb_hosts site.yml
```

This will configure and deploy the Helion lifecycle manager Hypervisor nodes, provision the VCP VMs, and fully deploy the HPE Helion OpenStack Cloud.

## 16.4.2 Provisioning just the Virtual Control Plane VMs

If your HPE Helion OpenStack Cloud installation requires that you perform additional actions between the provisioning of the VCP VMs and the running of the site.yml playbook then you can run the hlm-hypervisor-setup.yml playbook to provision the VCP VMs only, as follows:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts hlm-hypervisor-setup.yml
```

This will configure and deploy the Helion lifecycle manager Hypervisor nodes, and then provision the VCP VMs. Once you have performed the additional actions you can complete the installation by running the site.yml playbook.

## 16.4.3 Manually running each phase of the Virtual Control Plane provisioning

If you need explicit control over when the different phases of the configuration and deployment of the HLM Hypervisor nodes and the provisioning of the VCP VMs happen then you can run through the following steps, performing any required additional actions before moving on to the next step.

## 16.4.3.1 Configure OS and Networking on HLM Hypervisor nodes

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts osconfig-run.yml --limit=hlm-hypervisors
```

This will configure the OS environment and any relevant networking infrastructure on the HLM Hypervisor nodes that will be needed to support the associated VCP VMs.

### 16.4.3.2 Deploy and Configure the HLM Hypervisor nodes

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts hlm-hypervisor-deploy.yml
```

This will install and configure any additional software that needs to be in place before you can provision the VCP VMs.

Provision the VCP VMs

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts hlm-hypervisor-vms-deploy.yml
```

This will provision and start the VCP VMs, but will not perform any additional setup operations on them.

When you have completed any additional actuons you can then run <u>site.yml</u> to complete the installation of the HPE Helion OpenStack Cloud.

## 16.5 NIC Mapping for HLM Virtual-Controllers

For a bare-metal system the set of <u>nic-mappings</u> specified within <u>nic\_mappings.yml</u> file define the mapping of PCI bus-addresses to device name.

For a HLM Virtual-Controller the set of <a href="nic-mappings">nic-mappings</a> specified within <a href="nic\_mappings.yml">nic\_mappings.yml</a> file defines the mapping of PCI bus-addresses to device name, but is also used to construct the HLM Virtual-Controller VM with network interfaces at those PCI bus-addresses.

When defining the network-interface model for a HLM Virtual-Controller it will become apparent how many network interfaces are required for each controller and therefore how many NICs need to be defined in the <a href="nic-mappings">nic-mappings</a> for this VM.

It is recommended that you define a separate <a href="nic-mappings">nic-mappings</a> for each number of NICs that you reference. For example the following defines three mappings, one with a single NIC, a second with two NICs and a third with eight NICs:

```
...
nic-mappings:
...
```

```
- name: HOS-VIRTUAL-ONE-PORT
 physical-ports:
    - logical-name: hed1
      bus-address: '0000:01:01.0'
      type: simple-port
- name: HOS-VIRTUAL-TWO-PORT
 physical-ports:
    - logical-name: hed1
      type: simple-port
      bus-address: "0000:01:01.0"
    - logical-name: hed2
      type: simple-port
      bus-address: "0000:01:02.0"
- name: HOS-VIRTUAL-EIGHT-PORT
 physical-ports:
    - bus-address: '0000:01:01.0'
      logical-name: hed1
      type: simple-port
    - bus-address: '0000:01:02.0'
      logical-name: hed2
      type: simple-port
    - bus-address: '0000:01:03.0'
      logical-name: hed3
      type: simple-port
    - bus-address: '0000:01:04.0'
      logical-name: hed4
      type: simple-port
    - bus-address: '0000:01:05.0'
      logical-name: hed5
      type: simple-port
    - bus-address: '0000:01:06.0'
      logical-name: hed6
      type: simple-port
    - bus-address: '0000:01:07.0'
      logical-name: hed7
      type: simple-port
    - bus-address: '0000:01:08.0'
      logical-name: hed8
      type: simple-port
```

## 16.5.1 Recommended Mappings

It is recommended that the nic-mappings for HLM Virtual-Controllers use the PCI-bus addresses of the form "0000:01:XX.0", where XX is the value 01..1f (1..31 in hexadecimal).

This style isolates the virtual NIC interfaces onto PCI bus  $\underline{01}$  and should prevent bus-address clashes with other PCI devices on the system.

## 16.6 Notification of actions to be taken post upgrade

As part of an upgrade play actions that need to be taken - that cannot be automatically done since sharing the host with other vms will be placed in /var/run/hos.

The actions that are currently notified are

```
/var/run/hos/reload_apparmor

/var/run/hos/reload_libvirtd

/var/run/hos/restart_libvirtd
```

## 16.7 Reboot of a HLM-Hypervisor node

Rebooting a HLM-Hypervisor node follows the standard procedure in the HPE Helion OpenStack documentation .

The sequence here is to target the HPE Helion OpenStack vms first then the host itself - then shutdown the vms before reboot. Note that there is an ansible host group in verb hosts called <hlm-hypervisor node name>-vms which targets the vms on the node with name of hlm-hypervisor-node-name. In this case liam-vcp-cp1-vmf-m1-mgmt-vms

```
ansible-playbook -i hosts/verb_hosts --limit liam-vcp-cpl-vmf-ml-mgmt-vms hlm-stop.yml ansible-playbook -i hosts/verb_hosts --limit liam-vcp-cpl-vmf-ml-mgmt hlm-stop.yml ansible-playbook -i hosts/verb_hosts --limit liam-vcp-cpl-vmf-ml-mgmt hlm-hypervisor-vms-stop.yml # ie. shutdown the vms.
```

#### reboot node

```
ansible-playbook -i hosts/verb_hosts --limit liam-vcp-cp1-vmf-m1-mgmt hlm-hypervisor-vms-
start.yml (the vms should now be on autoboot)
ansible-playbook -i hosts/verb_hosts --limit liam-vcp-cp1-vmf-m1-mgmt-vms hlm-start.yml
ansible-playbook -i hosts/verb_hosts --limit liam-vcp-cp1-vmf-m1-mgmt hlm-start.yml
```

# 16.8 Staged install on HLM-Hypervisor individual play books.

This is a call out of the set of various playbooks to allow for a staged deployment of the hlm-hypervisor to allow for NFV vms to be installed as well at an appropriate time.

Once the node has been installed and the input model created from the NFV seed VM, the next stage is to configure this node and deploy the other nodes in the control plane.

The install of the other BM nodes can follow standard practice of running the cobbler deploy and the bm-reimage playbooks - then you need to run the osconfig playbook and hlm-hypervisor on these nodes.

This brings up the networking on the nodes and deploys a minimal set of packages to allow for virtual machine deployment. This is done by

```
ansible-playbook -i hosts/verb_hosts --limit hlm-hypervisors osconfig-run.yml ansible-playbook -i hosts/verb_hosts --limit hlm-hypervisors hlm-hypervisors-deploy.yml
```

At this stage you now have the nodes up with the networking configured and the basic set of packages installed to allow for subsequent deploy of the NFV vms.

Once this is done we then would need to deploy the HPE Helion OpenStack vms - this can be done with

```
ansible-playbook - i hosts/verb_hosts --limit hlm-hypervisors hlm-hypervisor-vms-
deploy.yml
```

Note that these 3 plays have been encapsulated into a single play that can be run with

```
ansible-playbook -i hosts/verb_hosts hlm-hypervisor-setup.yml
```

At this point you now have all the nodes to be able to deploy the cloud - this can now be achieved by running

```
ansible-playbook -i hosts/verb_hosts site.yml
```

## 16.9 Virtual Controller Replacement

There are several use cases here, all related to repair or replacement of VCP infrastructure. Essentially in each case you start by recovering the hypervisor (if necessary) and then create new VMs to replace the ones that were lost. Once they complete <a href="https://doi.org/nlm-hypervisor-vms-deploy.yml">https://doi.org/nlm-hypervisor-vms-deploy.yml</a>, it is the equivalent of having just completed <a href="maje.yml">bm-reimage.yml</a> for a physical node replacement, and can proceed with the existing instructions for controller recovery.

- 1. Repair or replacement of a single bad VM. One that has been damaged or lost, not just a simple reboot.
- 2. Of a single hypervisor, not causing e.g. RabbitMQ to lose quorum or similar effects, because all the VMs on this host are part of 3-node clusters and the other nodes are still up.
- 3. Disaster recovery loss of all hypervisors (and therefore control plane VMs)

#### Use case 1 - single bad VM.

```
ansible-playbook -i hosts/verb_hosts --limit hlm-hypervisors osconfig-run.yml ansible-playbook -i hosts/verb_hosts --limit hlm-hypervisors hlm-hypervisors-deploy.yml ansible-playbook - i hosts/verb_hosts --limit hlm-hypervisors hlm-hypervisor-vms-deploy.yml
```

## Use case 2 - single lost hypervisor but services are still running on the rest of the control plane

This starts with recovering the hypervisor, which is very similar to recovering a dead physical controller.

- 1. You'll need to update servers.yml and cobbler if the node's mac address or ilo details have changed (because of physical repairs/replacements).
- 2. Reimage just this node and then run the hypervisor plays on it

```
ansible-playbook -i hosts/localhost bm-reimage.yml --limit=hypervisor1
ansible-playbook -i hosts/verb_hosts --limit hypervisor1 osconfig-run.yml
ansible-playbook -i hosts/verb_hosts --limit hypervisor1 hlm-hypervisors-deploy.yml
ansible-playbook - i hosts/verb_hosts --limit hypervisor1 hlm-hypervisor-vms-
deploy.yml
```

At this stage, the VMs are up with a bare operating system, and you should follow the physical controller recovery instructions from here.

#### Use case 3 - lost all hypervisors and thus the complete control plane

Currently this use case is not correctly documented, even for a completely physical system - see DOCS-2921 (https://jira.hpcloud.net/browse/DOCS-2921) ▶. However once there is a working procedure for physical it could be adapted using similar techniques to the above.

# 16.10 Ansible host-specific variables for network-group access

From a 3rd-party view-point there are two important scenarios for accessing the network:

- 3rd-Party service on a HLM managed Server.
- 3rd-Party VM on a HLM-Hypervisor Server.

Each requires different sub-sets of data to establish the network configuration and potentially the end-point to which they need to attach. The following table indicates the set of data that is made available, and the scenarios that may find this data useful. The network-group specifications are provided under an attribute: <a href="https://doi.org/network\_groups:">host.my\_network\_groups:</a> containing a list of dicts, one per network-group.

	Consumed by 3rd-Party scenario				
Index	Attribute	Value	3rd-Party VM	3rd-Party Service	Comment
<net- work-group name&gt;</net- 			+	+	name of the net- work-group
	net- work-name	<net- work-name&gt;</net- 	+	+	name of the network
	passthroug device	h <bridge- name&gt;</bridge- 	+		open- vswitch-brid to attach to for un- filtered ac- cess to

			Consumed scenario	by 3rd-Party	
Index	Attribute	Value	3rd-Party VM	3rd-Party Service	Comment
					the spec- ified net- work
	tagged- vlan	true/ false	+		is this net- work run- ning tagged
	vlanid	<value></value>	+		which vlanid: is it tagged with
	device	<dev- name&gt;</dev- 		+	Device used for lo- cal-services
	address	<ip-ad- dress&gt;</ip-ad- 		+	Address on- ly used for lo- cal-services
	cidr	<cidr></cidr>	+	+	CIDR in use by this net- work
	gate- way-ip	<ip-ad- dress&gt;</ip-ad- 	+	+	gateway-ip in use on this net- work < op- tional >
	<u>mtu</u>	<mtu></mtu>	+	+	the MTU specified or inherited

			Consumed by 3rd-Party scenario		
Index	Attribute	Value	3rd-Party VM	3rd-Party Service	Comment
					on this net- work

### For Example:

```
host:
    my_network_groups:
        BLS:
            address: 10.244.47.102
            cidr: 10.244.47.0/24
            device: br-vlan3537
            gateway-ip: 10.244.47.1
            network-name: BLS-NET
            passthrough-device: br-bond0
            tagged-vlan: true
            vlanid: 3537
        CAN:
            cidr: 10.244.45.0/24
            device: br-bond0
            gateway-ip: 10.244.45.1
            network-name: CAN-NET
            passthrough-device: br-bond0
            tagged-vlan: true
            vlanid: 3535
        CLM:
            address: 10.244.44.102
            cidr: 10.244.44.0/24
            device: vlan3534
            gateway-ip: 10.244.44.1
            network-name: CLM-NET
            passthrough-device: br-bond0
            tagged-vlan: true
            vlanid: 3534
        PXE:
            address: 10.244.43.102
            cidr: 10.244.43.0/24
            device: br-bond0
            gateway-ip: 10.244.43.1
            network-name: PXE-NET
```

```
passthrough-device: br-bond0
    tagged-vlan: false
    vlanid: 3533
VLAN1-TUL:
    device: br-bond1
    network-name: VLAN1-TUL-NET
    passthrough-device: br-bond1
    tagged-vlan: false
VLAN2-TUL:
    device: br-bond1
    network-name: VLAN2-TUL-NET
    passthrough-device: br-bond1
    tagged-vlan: false
VxLAN-TUL:
   cidr: 10.244.48.0/24
    device: br-bond1
    gateway-ip: 10.244.48.1
    network-name: VxLAN-TUL-NET
    passthrough-device: br-bond1
    tagged-vlan: false
    vlanid: 3538
```

#### Which decodes to mean:

- BLS: available locally (address: 10.244.47.102) and passthrough to hosted-VMs (passthrough-device: br-bond0, vlanid: 3537, cidr:, gateway: )
- CAN: not-available locally, but passthrough to Hosted-VMs (passthrough-device: br-bond0, vlanid: 3535, cidr:, gateway: )
- CLM: available locally (address: 10.244.44.102) and passthrough to hosted-VMs (passthrough-device: br-bond0, vlanid: 3534, cidr:, gateway: )
- PXE: available locally (address: 10.244.43.1102) and passthrough to hosted-VMs (passthrough-device: br-bond0, vlanid: 3533, cidr:, gateway:)
- VLAN1-TUL: not available locally, but passthrough to Hosted-VMs (<a href="mailto:passthrough-de-vice">passthrough-de-vice</a>: br-bond1, tagged-vlan: false)
- VLAN2-TUL: not available locally, but passthrough to Hosted-VMs (passthrough-device: br-bond1, tagged-vlan: false)
- VxLAN-TUL: not available locally, but passthrough to Hosted-VMs (passthrough-de-vice: br-bond1, tagged-vlan: false, cidr:, gateway:)

A 3rd-Party application can provide an ansible-playbook that processes the specified set of ansible-vars and extracts the information required to support the configuration of the 3rd-party application.

## 16.10.1 Example 3rd-Party Ansible for Network Configuration Access

The following is an example ansible playbook to find the name of the <u>device</u> associated with the <u>MANAGEMENT</u> network-group. It can be easily customised to select any known network and any attribute of that network shown in the table above.

```
# (c) Copyright 2016 Hewlett Packard Enterprise Development LP
# Licensed under the Apache License, Version 2.0 (the "License"); you may
# not use this file except in compliance with the License. You may obtain
# a copy of the License at
# http://www.apache.org/licenses/LICENSE-2.0
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS, WITHOUT
# WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the
# License for the specific language governing permissions and limitations
# under the License.
# Example playbook to interrogate and recover network specific
# configuration data for Third-Party applications
- name: third-party | network-details | display all available network groups
 debug:
   var: host.my network groups
- name: third-party | network-details | select a network group
 set fact:
   tp_network_group_name: 'MANAGEMENT'
- name: third-party | network-details | display which network group we're interrogating
 debug:
   var: tp_network_group_name
- name: third-party | network-details | get the details of the selected network group
 set fact:
   tp network group data: "{{ host.my network groups | item(tp network group name) }}"
- name: third-party | network-details | display the details of the selected network group
 debug:
   var: tp_network_group_data
```

```
    name: third-party | network-details | get the 'device' for the selected network group set_fact:
        tp_network_group_device: "{{ tp_network_group_data[0] | item('device') }}"
    name: third-party | network-details | display the 'device' for the selected network group debug:
        var: tp_network_group_device
```

## 16.11 Ansible host-specific variables for gluster

If disks have been set aside for gluster in the hypervisor input model as outlined earlier, then a list of these disks can be found in host.my\_device\_groups.gluster.devices, for example:

```
host:

my_device_groups:

gluster:

consumer:

name: gluster

suppress_warnings: true

devices:

name: /dev/sde

name: /dev/sdf

name: gluster
```

## 16.12 "Empty" HLM Hypervisor Node

In order to create an empty hlm-hypervisor node then specifying <a href="https://hlm-hypervisor">hlm-hypervisor</a>: True is necessary.

```
- id: hypervisor1
ip-addr: 10.225.107.74
role: HLM-HYPERVISOR-ROLE
server-group: RACK2
hlm-hypervisor: True
nic-mapping: HP-BL460-6PORT
mac-addr: 9c:7e:96:22:9e:d8
ilo-ip: 10.1.18.42
ilo-password: whatever
ilo-user: whatever
```

In addition, if these nodes are being used in scenarios where there are no HPE Helion OpenStack servers to consume log data etc., you will need to ensure that in the input model the only common-service-components would be the lifecycle-manager-target.

## 16.13 Monasca Monitoring of HLM Hypervisors

If your build includes a top-level play <a href="https://example.com/htm-hyperviros-monitoring-deploy.yml">https://example.com/htm-hyperviros-monitoring-deploy.yml</a> then you have the option of enabling Monasca monitoring for the HLM Hypervisor nodes by running that playbook as follows:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts hlm-hypervisor-monitoring-deploy.yml
```

#### 16.13.1 Monitored Metrics

Currently two types of metrics are gathered:

- 1. Overall summary state for a HLM Hypervisor node hlm-hypervisor.vcp\_vms
  - a. Reports 0 indicating healthy so long as all monitored VMs, and associated networks, are running/active.
  - b. If any of the VMs or associated networks are not running/active, reports a value of 1.
  - c. Associated additional dimensions:
    - i. service: hlm-hypervisor
    - ii. component: vcp
- 2. Summary state for each VCP VM running on a HLM Hypervisor node <a href="https://hunning.ncm/hlm-hypervisor.vcp\_vm.</a> <a href="https://hunning.ncm/hlm-hypervisor.vcp\_vm.</a> <a href="https://hunning.ncm/hlm-hypervisor.vcp\_vm.</a> <a href="https://hunning.ncm/hlm-hypervisor.vcp\_vm.</a> <a href="https://hunning.ncm/hypervisor.vcp\_vm.</a> <a href="https://hunning.ncm/hyperviso

- a. Reports 0 indicating health so long as a VM, and its associated networks, are running/active.
- b. If the VM or any of its networks are not running/active, reports a value of 1.
- c. Associated additional dimensions:

```
i. service: hlm-hypervisor
```

ii. component: vcp\_vm

iii. domain: < vm\_name >

## 16.13.2 Configured Alarms

Alarms are configured to track these metrics; an alarm will be triggered if the metric starts reporting a state of 1.

If an alarm is triggered the metric's measurement value\_meta will include a detail section outlining the detected issues.

## 16.13.3 HLM Hypervisor Monitoring Configuration

The monitoring mechanism is driven by per-VM JSON configuration files in /etc/hlm-hyper-visor/vms with the following format:

```
{
    "hhv_vm_config": {
        "name": "<hostname>",
        "domain": "<vm name>",
        "networks": [
            "<vnet_1>",
            "<vnet 2>"
        ],
        "dimensions": {
             "service": "hlm-hypervisor",
            "component": "vcp vm",
         "domain": "<vm_name>"
        },
        "check_type": "vm"
    }
}
```

255 Configured Alarms

The HLM Hypervisor Monasca detection plugin <a href="https://hhv.py">hhv.py</a>, exposes two classes, <a href="https://hummary.nc.nlm.nc.">HLMHypervisorVMs</a>, which respectively generate the appropriate Monasca configuration to support both types of metric.

The HLM Hypervisor Monasca check plugin, <a href="https://hww.py">hhv.py</a>, processes the configuration generated by the detection plugin and updates the relevant metrics.

## 17 Integrations

Once you have completed your cloud installation, these are some of the common integrations you may want to perform.

## 17.1 Ceph Overview

### 17.1.1 Overview

HPE Helion OpenStack® 5.0 supports the Hammer version of Ceph cluster. Ceph cluster is a distributed object storage solution that can scale horizontally up to multiple petabytes and is based on the Reliable Autonomic Distributed Object Store (RADOS) object-based storage system. Ceph cluster stores all data as objects. It also provides components that support other storage protocols. For example: RBD ( RADOS Block Devices) supports block storage protocol, RADOS Gateway support S3/Swift API protocol for object storage access and so on. For example, RBD (RADOS Block Devices) supports the block storage protocol, RADOS Gateway supports the S3/Swift API protocol for object storage access, and so on. The following are Ceph's key features:

- 1. Uses of any commodity hardware.
- 2. Scales horizontally up to petabytes (there is no theoretical limit).
- 3. Self healing, self managing.
- 4. No single point of failure.
- 5. Supports various client protocol for block device access, object storage access using the S3/Swift API and more.

Ceph deployment is flexible. Although it supports a variety of storage protocols based on deployed components, you can add extra components (except mandatory ones) *only* if you need to support specific storage protocols. For example, if you are not going to support the S3/Swift protocol, then you might choose not to deploy the RADOS Gateway. For easy readability, the

257 Ceph Overview

various aspects of cluster deployment, such as hardware configuration, service configuration parameters, and deployment architecture, are categorized into separate sections. Each major section is further categorized into following segments:

- Core Ceph
- RADOS Gateway

Each sub-section provides recommendations vs. supported configurations. We recommend the production configuration, although all supported configurations are tested and validated. Use alternative supported configurations only if you have a strong need for them after evaluating their pros and cons.

#### Core Ceph

Core Ceph has two primary components, Object Storage Daemon (OSD) and Monitor, which are mandatory for cluster function. The following table briefly describes these components.

Components	Description
OSD	A Ceph OSD Daemon (OSD) stores data, handles data replication, recovery, backfilling, and rebalancing, and provides monitoring information to Ceph Monitors by checking other Ceph daemons for an activity.
Monitor	Ceph Monitor maintains maps of the cluster state including the monitor map, the OSD map, the placement group (PG) map, and the CRUSH map. It also maintains a history (called an "epoch") of each state change in the Ceph Monitors, Ceph OSD Daemons, and PGs.

#### **RADOS Gateway**

The RADOS Gateway service is an object storage interface that enables user to perform HTTP-based CRUD operations on an object. It supports both the OpenStack Swift and Amazon S3 REST APIs. It has the following two types of users, unlike the rest of OpenStack services, which rely completely on Keystone for user management (see more details at ).

- 1. Keystone
- 2. RADOS Gateway user (managed by Ceph itself and does not require Keystone)

258 Overview

RADOS Gateway is an optional component. Deploy it only if you need to access objects storage functionality using Swift or S3 API. Features include the following:

- 1. RADOS Gateway is configured to run in a simple (non-federated or single region) mode.
- 2. The HAProxy on the HPE Helion OpenStack® controller node acts as a load balancer for RADOS Gateway servers (in least connection mode, the load balancer selects the server with the least number of connections) for RADOS Gateway servers.
- 3. Provides OpenStack Keystone integration (users having a configured set of roles can access Swift APIs served by radosgw).
- 4. The default HPE Helion OpenStack® configuration installs RADOS Gateway on standalone nodes.
- 5. Access to Amazon S3 APIs is limited to RADOS Gateway users.
- 6. The RADOS Gateway external and internal endpoints are SSL/TLS- enabled, including the public end points represented by HAProxy.

## 17.1.2 Deployment Architecture

Consider the following issues when deploying Ceph:

- Ceph networking
- Placement of service components (such as OSD, monitor, RADOS Gateway) across nodes.
   For example, RADOS Gateway and monitor can be deployed on standalone nodes or together.

## 17.1.3 Ceph Networking

Ceph clients transmits traffic directly to OSD daemons for storage operations instead of the client routing requests to a specific gateways. OSD daemons perform data replication and participate in recovery activities. In general, a storage pool is configured with a replica count of 3, causing daemons to transact three sets of client data over the cluster network. As a result, 4 MB of write

traffic results in a total of 12 MB of data movement in the Ceph clusters (4 MB \* 3 replicas = 12MB). Also, Ceph clusters routinely share data among themselves. It is important to segregate Ceph data traffic into three segments, as follows:

- Management traffic, which includes monitoring and logging.
- Client traffic (often termed as data traffic) includes client requests sent to OSD daemons.
- Cluster traffic (often termed as replication traffic), which includes replication and recovery data traffic among OSD daemons.

For a high-performance clusters, a proper network configuration is very important. Use multiple networks for different data traffics. For a cluster of a reasonable size (a few terabytes), we recommend having a cluster with at least two networks, such as a single network for management and client data traffic (front-side) and a cluster (back-side) network. For large Ceph clusters, we recommend segregating all three traffic configurations. Segregation enables helps make for a secure connection because a cluster network is not required to be connected to the Internet directly. It allows OSD daemons to keep communicating without intervention so that placement groups can be brought to active and clean states relatively easily, whenever required. Apart from network separation (using VLANs), be sure to consider NICs used for Ceph servers too. We strongly recommend using bonded NICs to prevent single points of failure. Note that the network (and hence VLAN) separation is different from NIC separation even though they are linked to each other. In a multi-network model, emphasize VLAN separation, which ideally should be complemented by NIC separation, using the following priority order:

- 1. Separation of VLAN
- 2. Separation of NIC for VLANs.

The number of NICs you have will depend on your number of bonded interfaces, as follows:

- 1. Three bonded interfaces with six NICs: first for management, second for OSD client, and third for OSD internal (preferred setup).
- 2. Two bonded interfaces with four NICs: management VLANs hooked to the first bonded interface, and OSD networks hooked to the second bonded interface (most-used setup).
- 3. Only one bonded interface: two NICs for all VLANs (for a few NICs only).

260 Ceph Networking

HPE Helion OpenStack® Ceph software offers significant flexibility when defining and deploying OpenStack-based clouds, letting you implement a wide variety of different Ceph configurations. This allows you to design, model, and deploy cloud based on your requirements. You have the following VLAN choices based on traffic types:

- 1. Single: A single VLAN is used for all traffic, meant primarily for a small cluster.
- 2. Double deployment: One VLAN is used for cloud management and client traffic and another VLAN is used for Ceph internal traffic.
- 3. Triple deployment: Separate VLANs are used for management, client, and internal traffic. As mentioned previously, you can link all VLANs to the same bonded interface, to a separate bonded interface, or a combination of interfaces. For a large cluster, you can use a separate bonded interface which in turn necessitates having at least six NICs for OSD nodes and four NICs for monitor nodes.

## 17.1.4 Placement of service component

Although you can place service components in multiple ways, this section focuses on recommended deployments only, covering Core Ceph (OSD and monitors) and the RADOS Gateway. For details on alternative supported deployment architecture, please refer to .

- Core Ceph (i.e. OSD and monitors)
- Rados Gateway

#### • Core Ceph

We recommend the following deployment composition is recommended to avoid single points of failure for the Ceph cluster deployment.

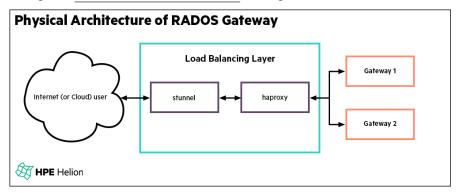
- 1. Three monitor nodes, to retain the odd number criteria of the monitor quorum.
- 2. At least three OSD nodes, to ensure that the object is replicated on three separate physical nodes, if the replica count of the pool is set to three (the recommended storage pool configuration is to set the replica count to three).

As mentioned in , we recommend using separate VLANS to separate Ceph traffic. The cloud management network is used for logging and/or monitoring, the OSD client network is used for Ceph client traffic, and the OSD internal network is used for internal Ceph traffic, such as replication.

#### RADOS Gateway

We recommend deploying at least two instances of RADOS Gateway on a standalone node front- ended by HAProxy.

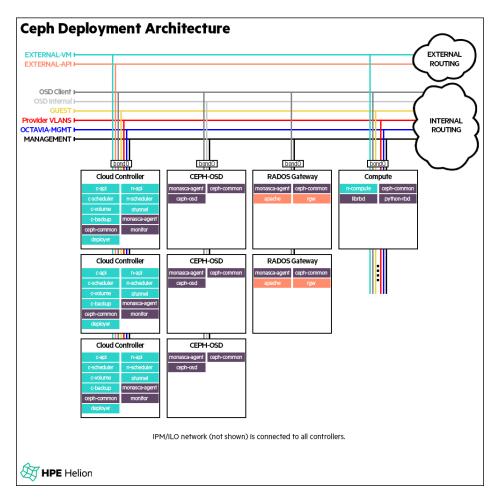
The following diagram illustrates the physical architecture of the RADOS Gateway and using the entry-scale-kvm-ceph configuration.



The default example model (<a href="entry-scale-kvm-ceph">entry-scale-kvm-ceph</a>) is the recommended mechanism to deploy the Ceph cluster with the RADOS Gateway.

## 17.1.5 Ceph Deployment Architecture

The following diagram illustrates the Ceph deployment architecture.



The preceding diagram illustrates the Ceph deployment scenario of <a href="entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> input model. Monitors are deployed on three controller nodes. Three standalone nodes are used for OSD and two standalone nodes are used for the RADOS Gateway, front-ended by HAProxy running on controller nodes. Management, client, and internal data traffic is separated using independent VLANs.

## 17.1.6 Alternative supported architecture

HPE Helion OpenStack® 5.0 also supports an alternate deployment architecture, for which the following points are relevant.

- Core networking
- Placement of service components
  - Core Ceph

- Monitor on standalone node
- RadosGateway
  - Co-hosted on cluster nodes hosting monitor service components
  - Co-hosted on controller nodes

## 17.1.7 Core networking

Ceph clients transmit traffic directly to OSD daemons for storage operations instead of the client routing requests to a specific gateway. For more information, refer to .

## 17.1.8 Placement of service components

#### Core Ceph

The architecture choices which is supported in HPE Helion OpenStack® 5.0 is as follows.

 Monitor on standalone node
 You can deploy the Ceph monitor service on a one or more dedicated clusters or resource nodes. Ensure that you modify your environment after installing the lifecycle

manager. For more details, refer to .

#### RADOS Gateway

RADOS Gateway service can be co-hosted with other HPE Helion OpenStack® services as follows.

- Alternate RADOS Gateway deployment architecture choice
  - Co-hosted on cluster nodes hosting monitor service components
  - Co-hosted on controller nodes

The default <u>entry-scale-kvm-ceph</u> input model deploys the <u>radosgw</u> services on two dedicated cluster nodes. However, you can also install the RADOS Gateway on a cluster node hosting the Ceph monitor service or on controller nodes. Refer to or for more details.

264 Core networking



Because the RADOS Gateway service shares server resources with multiple services, these alternate configurations will result in sub-optimal performance, as compared to the default configuration.

### 17.1.9 Hardware recommendations

For hardware recommendations, refer to .

## 17.1.10 Ceph Deployment and Configurations

HPE Helion OpenStack® 5.0 Ceph deployment leverages the cloud lifecycle operations supported by Helion lifecycle management. It provides the simplified lifecycle management of critical cluster operations such as service check, upgrade, and reconfiguring service components. This section assumes that you understand cloud input models and highlights only important aspects of cloud input models pertaining to Ceph. We focus on deployment aspects of the <a href="mailto:entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> input model, which is the most widely used configuration. You can see for the deployment of Ceph with various supported options. To ensure the proper deployment and verification of Ceph, it is important to read the topics and perform the steps in order. This section provides insight on how to alter the <a href="mailto:entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> input model to deploy Ceph with various supported options. We recommend that you deploy your supported choice only after evaluating all pros and cons.

- 1. a. Define an OSD Disk Model for an OSD Disk
  - b. Customize Your Service Configuration

2.

3.

### 17.1.10.1 Predeployment

Before you start deploying the HPE Helion OpenStack® cloud with Ceph, you must understand the following aspects of Ceph clusters,

#### Define an OSD Disk Model for an OSD Disk

This section focus on expressing the storage requirements of an OSD (object-storage daemon) node. OSD nodes have system, data, and journal disks.

System disks are used for OSD components, logging, and other tasks. The configuration of data and journal disks in important for Ceph deployment. A sample disk model file for the entry-scale-kvm-ceph cloud is as follows.

```
product:
 version: 2
disk-models:
- name: OSD-DISKS
 # Disk model to be used for Ceph OSD nodes
 # /dev/sda_root is used as a volume group for /, /var/log and /var/crash
 # sda_root is a templated value to align with whatever partition is really used
  # This value is checked in os config and replaced by the partition actually used
  # on sda e.g. sda1 or sda5
  volume-groups:
    - name: hlm-vg
      physical-volumes:
        - /dev/sda_root
      logical-volumes:
      # The policy is not to consume 100% of the space of each volume group.
      # 5% should be left free for snapshots and to allow for some flexibility.
        - name: root
          size: 30%
         fstype: ext4
         mount: /
        - name: log
          size: 45%
         mount: /var/log
          fstype: ext4
         mkfs-opts: -0 large_file
        - name: crash
          size: 20%
         mount: /var/crash
          fstype: ext4
          mkfs-opts: -0 large file
      consumer:
         name: os
  # Disks to be used by Ceph
```

```
# Additional disks can be added if available
device-groups:
  - name: ceph-osd-data-and-journal
    devices:
      - name: /dev/sdc
    consumer:
      name: ceph
       attrs:
         usage: data
         journal_disk: /dev/sdd
  - name: ceph-osd-data-and-shared-journal-set-1
    devices:
      - name: /dev/sde
    consumer:
      name: ceph
       attrs:
         usage: data
         journal_disk: /dev/sdg
  - name: ceph-osd-data-and-shared-journal-set-2
    devices:
      - name: /dev/sdf
    consumer:
       name: ceph
       attrs:
         usage: data
         journal_disk: /dev/sdg
```

The disk model has the following parameters:

Value	Description
device-groups	The name of the device group. There can be several device groups, which allows different sets of disks to be used for different purposes.
name	An arbitrary name for the device group.  The name must be unique.

Value	Description
devices	A list of devices allocated to the device group. A <a href="mailto:name">name</a> field containing <a href="mailto://dev/sdc">/dev/sdc</a> , <a a="" dev="" href="mailto://dev/sde" sde"="" sde<=""> and <a a="" dev="" href="mailto://dev/sde" sde"="" sde<=""> and <a a="" dev="" href="mailto://dev/sde" sde"="" sde<=""> indicates that the device group is used by Ceph.</a></a></a>
consumer	The service that uses the device group. A <a href="mailto:name">name</a> field containing <b>ceph</b> indicates that the device group is used by Ceph.
attrs	Attributes associated with the consumer.
usage	Devices for a particular service can have several uses. In the preceding sample, the usage field contains data, which indicates that the device is used for data storage.
journal_disk [OPTIONAL]	The disk to be used for storing journal data. When running multiple Ceph OSDs on a single node, a journal disk can be shared between OSDs of the node.  If you do not specify this value, Ceph stores the journal on the OSD's data disk (in a separate partition).

The preceding sample file represents the following:

- The first disk is used for OS and system purposes.
- There are three OSD data disks (sdc, sde, and sdf) and two journal disks (sdd and sdg). This configuration shows that we can share journal disks for multiple OSDs. It is recommended to use an OSD journal disk for four OSD data disks. See for more details.
- The drive type is not mentioned for the journal or data disks. You can consume any drive type but we **recommend** using an SSD (solid-state drive) for the journal disk.

Although the preceding model illustrates mixed use of the journal disk, we strongly advise that you keep journal data separate from OSD data, which means that your disk model **should not** have journal disks shared on the same data disks. For more information, see *Usage of Journal Disk*.

#### **Usage of Journal Disk**

HPE Helion OpenStack® 5.0 recommends storing the Ceph OSD journal on an SSD and the OSD object data on a separate hard disk drive. SSD drives are costly, so it saves money to use multiple partitions in a single SSD drive for multiple OSD journals. We recommend not more than four or five OSD journals on each SSD disk as a reasonable balance between cost and optimal performance. If you have too many OSD journals on a single SSD, and the journal disk crashes, you might lose your data on those disks. Also, too many journals in a single SSD can negatively affect performance.

Using an OSD journal as a partition on the data disk itself is supported. However, you might see a significant decline in Ceph performance because each client request to store an object is first written to the journal disk before sending an acknowledgment to the client. The Ceph OSD journal size defaults to 5120 MB (5 GB) in HPE Helion OpenStack® 5.0. This value can be changed, but it does not apply to any existing journal partitions. It will affect new OSDs created after the journal size is changed (whether the journal is on the same disk or a separate disk than the data disk). To change the journal size, edit the <a href="mailto:osd\_journal\_size">osd\_journal\_size</a> parameter in the <a href="mailto:~/helion/my\_cloud/config/ceph/settings.yml">~/helion/my\_cloud/config/ceph/settings.yml</a> file. To summarize:

- 1. Use SSD for the journal disk.
- 2. The ratio of OSD data disks to the journal disk is recommended to 4:1.
- 3. The default journal partition size is 5 GB, which you can change. Actual journal size depends upon your disk drive rpm and expected throughput. The formula is:

```
OSD_JOURNAL_SIZE = 2 * (EXPECTED_THROUGHPUT * FILESTORE_MAX_SYNC_INTERVAL)
```

- 4. The journal size for previously configured OSD disks does not change even if you change the <a href="mailto:osd\_journal\_size">osd\_journal\_size</a> parameter in the <a href="mailto:~/helion/my\_cloud/config/ceph/settings.yml">~/helion/my\_cloud/config/ceph/settings.yml</a> file. If you want to resize the journal partition of previously configured OSD disks, you should flush journal data, remove the OSD from the cluster, and then add it again.
- Customize Your Service Configuration

You must customize the parameters in the following files:

• Customize the parameters in the <u>~/helion/my\_cloud/config/ceph/set-tings.yml</u> file.

#### **Core Service Parameters**

Parameter	Description	Default Value	Recommendation
ceph_cluster	The name of the Ceph clusters. The default value is Ceph.	Ceph	Customize to suit your requirements.
ceph_release	The name of the Ceph release.	hammer	Do not change the default value.
osd_pool_de- fault_size	The number of replicas for objects in the pool.	3	Do not lower the default value. The value can be increased to the maximum number of OSD nodes in the environment (increasing it beyond this limit will cause the cluster to never reach an active + clean state).

Parameter	Description	Default Value	Recommendation
osd_pool_de- fault_pg_num	The default number of placement groups for a pool. This value changes based on the number of OSDs available.	128	The value can be changed based on the number of OSD servers/nodes in the deployment.  Refer to the Ceph PG calculator at http://ceph.com/pgcalc/ to customize it.
fstype	Storage filesystem type for OSDs.	xfs	Only the xfs file system is certified.
zap_data_disk	Zap partition table and contents of the disk.	True	Not recommended to change the default value.
persist_mountpoint	Place to persist OSD data disk mount point.	fstab	Not recommended to change the default value (as it ensures that the OSD data disks are mounted automatically on a system reboot).
osd_settle_time	The time in seconds to wait for after starting/restarting Ceph OSD services.	10 seconds	Increase this value only if the number of OSD servers is more than three or the servers have a slow network.

Parameter	Description	Default Value	Recommendation
osd_journal_size	The size of the journal in megabytes.	5120	You can increase this value to achieve optimal use of the journal disk (if it is shared between multiple OSDs).
data_disk_poll_at- tempts	The maximum number of attempts before attempting to activate an OSD for a new disk (default value 5).	5	Increase this value only if the OSD data disk drives are under-performing or slower than expected. Because this parameter and data_disk_poll_linterval (following) have a combined effect, we recommend that you consider both while tweaking either of them.
data_disk_poll_in- terval	The time interval in seconds to wait between data_disk_pol-l_attempts.	12	You can customize this value to suit your requirements. However, because this parameter and data_disk_pol-l_attempts (preceding) have a combined effect, we recommend

Parameter	Description	Default Value	Recommendation
			that you consider both while tweak- ing either of them.
osd_max_open_files	Maximum number of file descriptors for OSD.	32768	Do not change the default value.
mon_default_dir	Directory to store monitor data.	/var/lib/ceph/ mon/ <ceph_cluster></ceph_cluster>	Do not change the default value.
mon_max_open_files	Maximum number of file descriptors for monitor.	16384	Do not change the default value.
root_bucket	Ceph CRUSH map root bucket name.	default	Not recommended to change the default value. Changing this value affects CRUSH map and data placement. If you change the default value ensure to create a new rule set with a new root bucket and map the Ceph storage pools to use a new rule set. It is strongly recommended not to change this value post day-zero deployment. Chang-

Parameter	Description	Default Value	Recommendation
			ing the value af-
			ter deployment re-
			sults in a newly
			added OSD nodes
			to go to a newer
			root_bucket.It
			affects placement
			of the placement
			groups (data) of
			the storage pools.
			If you are up-
			grading cluster(s)
			from HPE Helion
			OpenStack® 3.0
			to HPE Helion
			OpenStack® 5.0,
			you are strong-
			ly advice not to
			change the de-
			fault value of
			<u>root_bucket</u> .

## **RADOS Gateway (RGW) Parameters**

Parameter	Description	Default Value	Recommendation
radosgw_user	The name of the Ceph client user for <u>radosgw</u> .	gateway	Customize to suit your requirements.
radosgw_admin_e- mail	The email address of the server administrator.	admin@hpe.com	Update the email address of the server administrator.

Parameter	Description	Default Value	Recommendation
rgw_keystone_ser- vice_type	<b>DEPRECATED</b> To configure RADO	S Gateway before dep	oloyment refer to .
rgw_keystone_ac- cepted_roles	Only users having either of the roles listed here will be able to access the Swift APIs of radosgw.	admin,_member_	Do not change the default value.



#### Note

The default service password for the RADOS Gateway service can be modified by following the steps documented at .

## Configuring the RADOS Gateway service type in the Keystone catalog Configuring the service type before deployment

1. You can configure the RADOS Gateway service type in Keystone catalog by replacing the <a href="mailto:ceph-object-store">ceph-object-store</a> with desired value in <a href="mailto:~/helion/hos/ser-vices/ceph/rgw.yml">~/helion/hos/ser-vices/ceph/rgw.yml</a> file on the life cycle manager node.

```
advertises-to-services:
    - service-name: KEY-API
    entries:
    - service-name: ceph-rgw
        service-type: ceph-object-store
        service-description: "Ceph Object Storage Service"
        url-suffix: "/swift/v1"
```

2. After you have modified the **service-type**, commit the change to the local git repository.

```
cd ~/helion
git checkout site
git add ~/helion/hos/services/ceph/rgw.yml
git commit -m "Updating the RADOS Gateway service type"
```

3. Rerun the configuration processor.

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

4. Rerun the deployment area preparation playbooks.

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

5. Run reconfiguration playbook in deployment area.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts keystone-reconfigure.yml
```

#### Configuring the RADOS Gateway service type after deployment

To update the RADOS Gateway service type in a deployed or a running cloud, you must delete the <u>ceph-rgw</u> service from the Keystone catalog and perform the same steps as mentioned in the preceding section ().

1. To delete the <u>ceph-rgw</u> service, you must know the service-id. Execute the following command from a controller node.

```
source ~/keystone.osrc
openstack service list |grep ceph-rgw | awk '{print $2}'
openstack service delete <service-id>
```

#### Ceph client parameters

Value	Description	Default Value	Recommendation
pg_active_de- lay_time	The delay time for Ceph PGs to come into active state.	10	You can increase this value if the number of OSD servers/nodes in the deployment is more than three. Because this parameter and pg_active_retries (following) have a

Value	Description	Default Value	Recommendation
			combined effect,
			we recommend
			that you consider
			both while tweak-
			ing either of them.
pg_active_retries	The number of	5	You can customize
	retries for Ceph		this value to suit
	placement groups		your requirements.
	to come into ac-		However, because
	tive state with		this parameter
	a duration of		and pg_active_de-
	pg_active_de-		lay_time (preced-
	lay_time seconds		ing) have a com-
	between entries.		bined effect, we
			recommend that
			you consider both
			while tweaking ei-
			ther of them.

The following table provides parameter descriptions. You can edit each parameter in the main.yml file.

Value	Description
fsid	A unique identifier, File System ID, for the Ceph cluster that you should gener-
	ate prior to deploying a cluster (use the

Value	Description
	uuidgen command to generate a new
	FSID). When set, this value cannot be
	changed.

## 17.1.10.2 Deploying Ceph

To deploy a new HPE Helion OpenStack® Ceph cloud using the default <u>entry-scale-kvm-ceph model</u>, follow these steps .



### Note

- In a multi-region cloud, Ceph can be deployed only as a shared service. In other words, Ceph services (Monitor, OSD, and RADOS Gateway servers) should be deployed in the shared control plane (such that those services will run in the same control plane as the Keystone service) ONLY.
- The non-shared control plane(s) need to ensure that the ceph-monitor service is specified in the service\_components section of uses and imports sections of the control plane definition.
- Ensure that all the regions that need access to Ceph have the shared control plane included.

#### **Edit Your Ceph Environment Input Files**

Perform the following steps:

- 1. Log in to the lifecycle manager.
- 2. Copy the example configuration files into the required setup directory and edit them to contain the details of your environment:

```
cp -r ~/helion/examples/entry-scale-kvm-ceph/* ~/helion/my_cloud/definition/
```

Enter your environment information into the configuration files in the  $\sim$ /helion/my\_cloud/definition directory.

You can find details of how to do this at.

3. Edit the \_~/helion/my\_cloud/definition/data/servers.yml file and enter details. If you are using alternative RADOS Gateway deployments, see before editing the \_server-s.yml file.

```
# Ceph OSD Nodes
    - id: osd1
      ip-addr: 192.168.10.9
      role: OSD-ROLE
      server-group: RACK1
      nic-mapping: MY-2PORT-SERVER
      mac-addr: "8b:f6:9e:ca:3b:78"
      ilo-ip: 192.168.9.9
      ilo-password: password
      ilo-user: admin
    - id: osd2
      ip-addr: 192.168.10.10
      role: OSD-ROLE
      server-group: RACK2
      nic-mapping: MY-2PORT-SERVER
      mac-addr: "8b:f6:9e:ca:3b:79"
      ilo-ip: 192.168.9.10
      ilo-password: password
      ilo-user: admin
    - id: osd3
      ip-addr: 192.168.10.11
      role: OSD-ROLE
      server-group: RACK3
      nic-mapping: MY-2PORT-SERVER
      mac-addr: "8b:f6:9e:ca:3b:7a"
      ilo-ip: 192.168.9.11
      ilo-password: password
      ilo-user: admin
# Ceph RGW Nodes
   - id: rgw1
     ip-addr: 192.168.10.12
     role: RGW-ROLE
     server-group: RACK1
     nic-mapping: MY-2PORT-SERVER
     mac-addr: "8b:f6:9e:ca:3b:62"
     ilo-ip: 192.168.9.12
     ilo-password: password
```

ilo-user: admin

- id: rgw2
 ip-addr: 192.168.10.13
 role: RGW-ROLE
 server-group: RACK2
 nic-mapping: MY-2PORT-SERVER
 mac-addr: "8b:f6:9e:ca:3b:63"
 ilo-ip: 192.168.9.13
 ilo-password: password
 ilo-user: admin

The preceding sample file contains three OSD nodes and two RADOS Gateway nodes.

4. Edit the \_~/helion/my\_cloud/definition/data/disks\_osd.yml file to align the disk model to fit the server specification in your environment. For details on disk models, refer to .

Ceph service configuration parameters can be modified as described in the preceding section.

5. Commit your configuration to the local git repo (*Chapter 3, Using Git for Configuration Management*):

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```

6. After you set up your configuration files, continue with the installation procedure from .



#### Note

For any troubleshooting information regarding the OSD node failure, see .

### 17.1.10.3 Verifying Ceph Cluster Status

If you have deployed RADOS Gateway with core Ceph, then you need to ensure that all service components including RADOS Gateway are functioning as expected.

#### **Verify Core Ceph**

Perform the following steps to check the status of the Ceph cluster:

- 1. Log in to the monitor node.
- 2. Execute the following command and make sure that the result is HEALTH\_OK or HEALTH\_WARN:

```
$ ceph health
```

Optionally, you can also set up the lifecycle manager as a Ceph client node (refer to ) and execute the preceding command from the lifecycle manager.

#### **Verify RADOS Gateway**

To make sure that a Keystone user can access RADOS Gateway using Swift, perform the following steps:

- 1. Log in to a controller node.
- 2. Source the service.osrc file:

```
source ~/service.osrc
```

3. Execute the following command to generate a list of the containers associated with the user:

```
swift --os-service-type ceph-object-store list
```

If the containers are listed, this indicates that RADOS Gateway is accessible.

## 17.1.11 Alternative Supported Choices

This section provides insight on how to alter the <a href="entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> input model to deploy Ceph with various supported options. We recommend that you deploy your supported choice only after evaluating all pros and cons. For technical details, please consult with the technical support team. The choices available can impact the performance and scaling of clusters. Choices are illustrated for reference purposes and you can combine one or more of them as needed. The content is categorized as follows:

1. Core Ceph

- •
- •
- •
- 2. RADOS Gateway
  - •
  - •
  - •
- 3.

### 17.1.11.1 Core Ceph

#### **Installing the Monitor Service on Standalone Nodes**

The following section provides the procedure for installing the monitor service on standalone nodes instead of installing on controller nodes, as mentioned in entry-scale-kvm-ceph.



### **Important**

If you want to install the monitor service as a dedicated resource node, you must decide before deploying Ceph. HPE Helion OpenStack 5.0 does not support deployment transition. After Ceph is deployed, you cannot migrate the monitor service from controller nodes to dedicated resource nodes.

#### **Prerequisite**

Perform the following steps to install the Ceph monitor on a dedicated node. Note that the Ceph requires at least three monitoring servers to form a cluster in case of a node failure.

- 1. Log in to the lifecycle manager.
- 2. Copy the <a href="entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> input model to the <a href="entry-scale-kvm-ceph">~/helion/my\_cloud/definition</a> directory before you begin the editing process:

```
cp -r ~/helion/examples/entry-scale-kvm-ceph/* ~/helion/my_cloud/definition/
```

3. Edit the <u>control\_plane.yml</u> to create a new cluster, such as with <u>ceph-mon</u>, as shown here.

```
clusters:
  - name: cluster1
    cluster-prefix: cl
    server-role: CONTROLLER-ROLE
    member-count: 3
    allocation-policy: strict
    service-components:
      - lifecycle-manager
      - ntp-server
  - name: ceph-mon
    cluster-prefix: ceph-mon
    server-role: CEP-MON-ROLE
    min-count: 3
    allocation-policy: strict
    service-components:
      - ntp-client
      - ceph-monitor
  - name: rgw
    cluster-prefix: rgw
    server-role: RGW-ROLE
```

4. Edit the \_/helion/my\_cloud/definition/data/servers.yml file to define the Ceph monitor node (monitor services). The following example shows three nodes for monitor services. We recommend using an odd number of monitor nodes.

```
# Ceph Monitor Nodes
- id: ceph-mon1
ip-addr: 10.13.111.141
server-group: RACK1
role: CEP-MON-ROLE
nic-mapping: MY-4PORT-SERVER
mac-addr: "f0:92:1c:05:69:10"
ilo-ip: 10.12.8.217
ilo-password: password
ilo-user: admin

- id: ceph-mon2
ip-addr: 10.13.111.142
server-group: RACK2
```

```
role: CEP-MON-ROLE
  nic-mapping: MY-4PORT-SERVER
  mac-addr: "83:92:1c:55:69:b0"
  ilo-ip: 10.12.8.218
  ilo-password: password
  ilo-user: admin
- id: ceph-mon3
  ip-addr: 10.13.111.143
  server-group: RACK3
  role: CEP-MON-ROLE
  nic-mapping: MY-4PORT-SERVER
 mac-addr: "d9:92:1c:25:69:e0"
  ilo-ip: 10.12.8.219
 ilo-password: password
  ilo-user: admin
# Ceph RGW Nodes
- id: rgw1
  . . .
```

5. Edit the \_~/helion/my\_cloud/definition/data/net\_interfaces.yml file to define a new network interface set for your Ceph monitors, as shown here.

```
## This defines the interface used for management
## traffic such as logging, monitoring, etc.
- name: CEP-MON-INTERFACES
  network-interfaces:
    - name: BOND0
      device:
          name: bond0
      bond-data:
          options:
              mode: active-backup
              miimon: 200
              primary: hed1
          provider: linux
          devices:
            - name: hed1
            - name: hed2
      network-groups:
        - MANAGEMENT
- name: RGW-INTERFACES
  network-interfaces:
```

6. Edit \_~/helion/my\_cloud/definition/data/disks\_ceph\_monitor.yml to define the disk model for monitor nodes.

```
disk-models:
- name: CEP-MON-DISKS
 # Disk model to be used for Ceph monitor nodes
 # /dev/sda_root is used as a volume group for /, /var/log and /var/crash
 # sda_root is a templated value to align with whatever partition is really used
 # This value is checked in os config and replaced by the partition actually used
 # on sda e.g. sda1 or sda5
 volume-groups:
    - name: hlm-vg
     physical-volumes:
        - /dev/sda_root
     logical-volumes:
     # The policy is not to consume 100% of the space of each volume group.
     # 5% should be left free for snapshots and to allow for some flexibility.
        - name: root
          size: 30%
          fstype: ext4
         mount: /
        - name: log
          size: 45%
          mount: /var/log
          fstype: ext4
          mkfs-opts: -0 large_file
        - name: crash
          size: 20%
          mount: /var/crash
          fstype: ext4
          mkfs-opts: -0 large_file
      consumer:
         name: os
```

7. Edit the ~/helion/my\_cloud/definition/data/server\_roles.yml file to define a new server role for your Ceph monitors:

```
- name: CEP-MON-ROLE
interface-model: CEP-MON-INTERFACES
disk-model: CEP-MON-DISKS
```

8. Commit your configuration:

```
cd ~/helion/hos/ansible
```

```
git add -A
git commit -m "<commit message>"
```

9. Run the following playbook to add your nodes into Cobbler:

```
cd ~/helion/hos/ansible/
ansible-playbook -i hosts/localhost cobbler-deploy.yml
```

10. To reimage all the nodes using PXE, run the following playbook:

```
cd ~/helion/hos/ansible/
ansible-playbook -i hosts/localhost bm-reimage.yml
```

11. Run the configuration processor:

```
cd ~/helion/hos/ansible/
ansible-playbook -i hosts/localhost config-processor-run.yml
```

12. Update your deployment directory with this playbook:

```
cd ~/helion/hos/ansible/
ansible-playbook -i hosts/localhost ready-deployment.yml
```

13. Deploy these changes:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts site.yml
```

#### Using a Single VLAN for All Ceph Traffic (Management, Client, and Internal OSD)

You can use a single VLAN to transmit all Ceph traffic. This configuration is recommended for a small cluster deployment.

Perform the following steps to configure Entry-scale-kvm-ceph-single-network.

- 1. Log in to the lifecycle manager.
- 2. Copy the <a href="entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> input model to the <a href="entry-scale-kvm-ceph">~/helion/my\_cloud/definition</a> directory before you begin the editing process:

```
cp -r ~/helion/examples/entry-scale-kvm-ceph/* ~/helion/my_cloud/definition/
```

- 3. Validate that NIC interfaces are correctly specified in <a href="nic\_mapping.yml">nic\_mapping.yml</a> for servers that are used in the cloud.
- 4. Ensure that you have at least two NICs for Ceph nodes to create a bonded interface for it.

5. Validate that your servers are mapped to a correct NIC interface specification in <a href="mailto:server-symbol">server-symbol</a>. The following is an example of a server node used for OSD deployment:

```
# Ceph OSD Nodes
- id: osd1
ip-addr: 192.168.10.9
role: OSD-ROLE
server-group: RACK1
nic-mapping: MY-2PORT-SERVER
mac-addr: "8b:f6:9e:ca:3b:78"
ilo-ip: 192.168.9.9
ilo-password: password
ilo-user: admin
```

- 6. Delete the OSD-INTERNAL and OSD-CLIENT network groups from <a href="network\_groups.yml">network\_groups.yml</a>. This is necessary because only the management network is used for Ceph traffic, thus OSD-INTERNAL and OSD-CLIENT network groups are not required.
- 7. Define net\_interfaces.yml to use only management network groups, as shown here.

```
- name: CONTROLLER-INTERFACES
   network-interfaces:
      - name: BOND0
        device:
            name: bond0
        bond-data:
            options:
                mode: active-backup
                miimon: 200
                primary: hed3
            provider: linux
            devices:
              - name: hed3
              - name: hed4
        network-groups:
          - EXTERNAL-API
          - EXTERNAL-VM
          - GUEST
          - MANAGEMENT
  - name: COMPUTE-INTERFACES
    network-interfaces:
      - name: BOND0
        device:
            name: bond0
        bond-data:
```

```
options:
              mode: active-backup
              miimon: 200
              primary: hed3
          provider: linux
          devices:
            - name: hed3
            - name: hed4
      network-groups:
        - EXTERNAL-VM
        - GUEST
        - MANAGEMENT
- name: OSD-INTERFACES
 network-interfaces:
    - name: BOND0
      device:
          name: bond0
      bond-data:
          options:
              mode: active-backup
              miimon: 200
              primary: hed3
          provider: linux
          devices:
            - name: hed3
            - name: hed4
      network-groups:
        - MANAGEMENT
- name: RGW-INTERFACES
 network-interfaces:
    - name: BOND0
      device:
          name: bond0
      bond-data:
          options:
              mode: active-backup
              miimon: 200
              primary: hed3
          provider: linux
          devices:
            - name: hed3
            - name: hed4
      network-groups:
        - MANAGEMENT
```

- 8. Delete VLAN information for OSD-INTERNAL-NET and OSD-CLIENT-NET from <a href="network-s.yml">network-s.yml</a>. Only Management VLANs are used.
- 9. After you set up your configuration files, perform in .

#### Using Two VLANs: For Management and Client Traffic and for Internal OSD Traffic

You can use dual VLANs to transmit Ceph traffic. In this configuration one VLAN transmits management and client traffic and the other VLAN transmits internal OSD traffic. A separate bonded interface for two VLANs is used with four NICs. This configuration provides two aspects:

- Use of two networks, such as VLANs.
- Use of separate bonded interfaces for each VLAN (different from what is provided in entry-scale-kvm-ceph).

The use of separate NICs segregates traffic at the interface level and requires your server to have at least four NICs. But using a separate bonded interface for each VLAN is not mandatory, and thus you can use a single bonded interface (or server with only two NICs) for Ceph deployment. Perform the following steps to configure Entry-scale-kvm-ceph-dual-network.

- 1. Log in to the lifecycle manager.
- 2. Copy the <a href="entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> input model to the <a href="entry-scale-kvm-ceph">~/helion/my\_cloud/definition</a> directory before you begin the editing process:

```
cp -r ~/helion/examples/entry-scale-kvm-ceph/* ~/helion/my_cloud/definition/
```

3. Validate that NIC interfaces are correctly specified in <a href="nic\_mapping.yml">nic\_mapping.yml</a> for servers that are used in the cloud. For Ceph OSD nodes, four port servers are required. You can use **HP-DL360-4PORT** as it is defined in <a href="nic\_mapping.yml">nic\_mapping.yml</a> of <a href="entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> or define a new NIC mapping (as shown here) for new sets of servers having four port servers.

```
type: simple-port
bus-address: "0000:09:00.0"

- logical-name: hed4
  type: simple-port
  bus-address: "0000:0a:00.0"
```

4. Modify OSD nodes to use four port servers, as shown here. Change the NIC mapping attribute from **HP-DL360-4PORT** to use any other name defined in nic\_mapping.yml.

```
# Ceph OSD Nodes
- id: osd1
ip-addr: 192.168.10.9
role: OSD-ROLE
server-group: RACK1
nic-mapping: HP-DL360-4PORT
mac-addr: "8b:f6:9e:ca:3b:78"
ilo-ip: 192.168.9.9
ilo-password: password
ilo-user: admin
```

- 5. Delete OSD-CLIENT from <a href="network\_groups.yml">network\_groups.yml</a>. Note that no dedicated network group exists for client traffic. Only the management network group is used for client traffic.
- 6. Edit net\_interfaces.yml with a bonded NIC as shown here.

```
- name: CONTROLLER-INTERFACES
    network-interfaces:
       - name: BOND0
        device:
             name: bond0
        bond-data:
            options:
                 mode: active-backup
                 miimon: 200
                 primary: hed1
             provider: linux
             devices:
               - name: hed1
               - name: hed2
        network-groups:
           - EXTERNAL-API
           - EXTERNAL-VM
           - GUEST
           - MANAGEMENT
```

```
- name: COMPUTE-INTERFACES
  network-interfaces:
    - name: BOND0
      device:
          name: bond0
      bond-data:
          options:
              mode: active-backup
              miimon: 200
              primary: hed1
          provider: linux
          devices:
            - name: hed1
            - name: hed2
      network-groups:
        - EXTERNAL-VM
        - GUEST
        - MANAGEMENT
- name: OSD-INTERFACES
  network-interfaces:
    - name: BOND0
      device:
          name: bond0
      bond-data:
          options:
              mode: active-backup
              miimon: 200
              primary: hed1
          provider: linux
          devices:
            - name: hed1
            - name: hed2
      network-groups:
        - MANAGEMENT
    - name: BOND1
      device:
          name: bond1
      bond-data:
          options:
              mode: active-backup
              miimon: 200
              primary: hed3
          provider: linux
          devices:
            - name: hed3
            - name: hed4
```

```
network-groups:
        - OSD-INTERNAL

    name: RGW-INTERFACES

network-interfaces:
   - name: BOND0
     device:
         name: bond0
     bond-data:
          options:
              mode: active-backup
              miimon: 200
              primary: hed1
          provider: linux
          devices:
            - name: hed1
            - name: hed2
     network-groups:
        - MANAGEMENT
```

- 7. Delete OSD-CLIENT from server\_groups.yml.
- 8. Delete VLAN information for OSD-CLIENT-NET from <a href="networks.yml">networks.yml</a>. Only management VLANs are used for client traffic.
- 9. After you set up your configuration files, perform in .

## 17.1.11.2 RADOS Gateway

#### Installing RADOS Gateway on Dedicated Cluster Nodes that Host the Ceph Monitor Service

You can configure RADOS Gateway to install on one or more dedicated cluster nodes hosting the Ceph monitor service as follows:

- 1. Remove the sections for servers in the \_~/helion/my\_cloud/definition/data/servers.yml file that have the \_role: RGW-ROLE attribute.
- 2. Edit the \_-/helion/my\_cloud/definition/data/control\_plane.yml file and add the following lines to the \_service-components section for the cluster nodes that have the server-role: MON-ROLE attribute.

```
- ceph-radosgw
- apache2
```

3. Edit the \_~/helion/my\_cloud/definition/data/net\_interfaces.yml file to remove the RGW-INTERFACES section. This section defines RADOS Gateway network interfaces, which are not required in this configuration:

```
- name: RGW-INTERFACES
 network-interfaces:
    - name: BOND0
      device:
         name: bond0
      bond-data:
         options:
            mode: active-backup
            miimon: 200
            primary: hed3
         provider: linux
         devices:
            - name: hed3
            - name: hed4
      network-groups:
        - MANAGEMENT
        - OSD-CLIENT
```

4. After you set up your configuration files, perform in .

#### **Installing RADOS Gateway on Controller Nodes**

You can configure RADOS Gateway to install on controller nodes. To do this, perform the following steps:

- 1. Remove the sections for servers in the \_~/helion/my\_cloud/definition/data/server-s.yml file that have the role: RGW-ROLE attribute.
- 2. Edit the \_~/helion/my\_cloud/definition/data/net\_interfaces.yml file to remove the RGW-INTERFACES section. This section defines RADOS Gateway network interfaces, which are not required in this configuration:

primary: hed3 provider: linux

devices:

- name: hed3 - name: hed4 network-groups: - MANAGEMENT - OSD-CLIENT

3. Edit the ~/helion/my cloud/definition/data/control plane.yml file and add the following line to service-components for the cluster with the server-role: CON-TROLLER-ROLE attribute.

```
- ceph-radosgw
```

4. After you set up your configuration files, perform in .

#### **Installing More than Two RADOS Gateway Servers**

To deploy more than two RADOS Gateway servers, you need to add a section to the ~/ helion/my cloud/definition/data/servers.yml file for each additional RADOS Gateway node.



#### Note

Installing additional RADOS Gateway servers is possible only if RADOS Gateway is installed on dedicated cluster nodes or on dedicated cluster nodes that host the Ceph monitor service. Additional RADOS Gateway servers cannot be added if RADOS Gateway is installed on a controller node.

#### Ceph Deployment with Virtual Control Plane 17.1.11.3

HPE Helion OpenStack® 5.0 supports the deployment of control plane elements on virtual machines which can be co-located on one baremetal machine or spread across three baremetal machines. The baremetal machine in this context is termed as VM factory host(s). The following aspects must be considered while deploying Ceph with the virtual control plane.

- 1. Deploy OSD and monitor node on a single VM factor host It is applicable to X1 cloud model. The number of VM factory host is one. Therefore, Ceph cluster will not have HA support as there will be only one instance of the monitor component of Ceph.
- 2. Deploy OSD and monitor on set of three VM factor hosts It is applicable to S1 cloud model.
- 3. Deploy monitor on set of three VM factor hosts but OSD nodes are deployed independently as a resource nodes It is applicable to M1 cloud model.

The following aspects must be considered while deploying Ceph with virtual control plane:

- 1. Scale-out of cluster Adding a new set of monitor or OSD nodes is not validated by the engineering team. Although, technically it is feasible but not recommended because it can have a significant performance impact. However, one can increase a cluster capacity by adding more disks to the VM factory host and configuring them as OSD (a scale-in path to increase capacity) nodes.
- 2. Deployment of RADOS Gateway on VM factory host is not formally supported.
- 3. Performance of Ceph components (OSD and monitor nodes) are sensitive to compute resources i.e. memory, core CPU, and so on. It is strongly recommended to plan and allocate minimum amount of resources for Ceph components to avoid resource contention because same set of machines will be running the control plane elements and the Ceph components. Starvation of resources might causes impact on the performance and the stability of the Ceph clusters health. Consider the following resource aspect for planning:
  - CPU
  - RAM
  - Disk space for monitoring logs

The following section focus on the change of the input model for X1 and S1 model ONLY. The change in the input model for M1 model is similar except that OSD node is deployed as the resource nodes. For other aspects of cluster management like upgrade, adding new set of disks, stopping and starting services and so on, you can follow the similar approach that is used for the deployment of Ceph cluster using HPE Helion OpenStack® .

#### Steps to deploy Ceph on VM factory host(s)

- 1. Log in to the lifecycle manager.
- 2. Go to ~/helion/my cloud/definition/.



We assume that you have already copied the cloud model representing control element deployment on VM factory host(s).

3. Edit your <u>control\_plane.yml</u> file to add <u>ceph-osd</u> and <u>ceph-monitor</u> components to the vmfactory nodes. For example:

```
- name: vmfactory
    resource-prefix: vmf
    server-role: HLM-HYPERVISOR-ROLE
    min-count:
    allocation-policy: strict
    service-components:
        - ntp-server
        - ceph-osd
        - ceph-monitor
        - lifecycle-manager
        - tempest
        - openstack-client
```

4. Edit <u>disks\_vmfactory.yml</u> file of VM factory hosts to define data and journal disks for OSD. For example, the following disk model illustrates the usage of <u>/dev/sdd</u>, <u>/dev/sde</u>, and <u>/dev/sdf</u> as data disks and <u>/dev/sdg</u> as journal disks for OSD. Disks allocated to OSD must not be used for any other purpose.

```
product:
version: 2

disk-models:
- name: HLM-HYPERVISOR-DISKS
volume-groups:
- name: hlm-vg
physical-volumes:
- /dev/sda_root
logical-volumes:
# The policy is not to consume 100% of the space of each volume group.
# 5% should be left free for snapshots and to allow for some flexibility.
- name: root
size: 35%
fstype: ext4
mount: /
```

```
- name: log
         size: 50%
         mount: /var/log
         fstype: ext4
         mkfs-opts: -0 large_file
       - name: crash
         size: 10%
         mount: /var/crash
         fstype: ext4
         mkfs-opts: -0 large_file
   - name: vg-images
     # this VG is dedicated to libvirt images to keep VM IOPS off the OS disk
     physical-volumes:
       - /dev/sdb
       - /dev/sdc
     logical-volumes:
       - name: images
         size: 95%
         mount: /var/lib/libvirt/images
         fstype: ext4
         mkfs-opts: -0 large_file
device-groups:
   - name: ceph-osd-disks
     devices:
    - name: /dev/sdd
    - name: /dev/sde
    - name: /dev/sdf
     consumer:
        name: ceph
        attrs:
          usage: data
          journal_disk: /dev/sdg
```

5. Commit your configuration to the local git repo.

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```

6. After setting up the configuration files, continue with the installation procedure mentioned at .



- a. Running <u>ceph-stop.yml</u> and <u>ceph-start.yml</u> playbooks on the VM factory host stops all the services (OSD and monitor) on the node. There is no playbook that can stop only one service.
- b. If you are deploying Ceph on a single VM factory host (i.e. X1 model), make the following changes before deployment.
  - Edit ~/helion/my\_cloud/config/ceph/settings.yml file to add the following content.

```
extra:

global:

osd_crush_chooseleaf_type: 0
```

 Commit your configuration to the local git repo and continue with the installation procedure mentioned at .

## 17.1.12 Usage of Ceph Storage

Ceph is a versatile storage technology that facilitates the consumption of storage by using multiple protocols. Ceph is closely integrated with HPE Helion OpenStack® services to support various storage use cases such as the use of storage pools for cinder-volume and glance data store. HPE Helion OpenStack® further simplifies administrator tasks by providing automated playbooks for various storage use cases.

- Creating storage pools
- Deploying Ceph client components on client nodes
- Configuring OpenStack services with storage pools

The following section guides you through the following common scenarios that you might come across while consuming Ceph for various storage use cases.

1.

2.

a.

b.

C.

d.

3.

Although you can also use Ceph storage pools as ephemeral storage backends for Nova, this practice is not formally supported.

## 17.1.12.1 Set Up Your Deployer as a Ceph Client Node

By default, your deployer node is not configured as a Ceph client node. As a result, you cannot start your Ceph operation here. It is usually more helpful to set up your deployer node as an admin client node to perform various Ceph operations. This will help you manage the Ceph cluster without logging in to the monitor node. You can turn your deployer node into a client node by executing the following playbook.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ceph-setup-deployer-as-client.yml
```

After you complete the preceding playbook, your deployer node will be configured with an admin key ring and thus act as an admin node.

## 17.1.12.2 Using Core Ceph OpenStack Services

The use of Cinder for glance, cinder-volume, and cinder-backup requires pool creation and setting up ceph-clients on respective OpenStack service nodes. To simplify the workflow, the playbook <a href="mailto:ceph-client-prepare.yml">ceph-client-prepare.yml</a> is provided. This playbook reads pool configurations following the specification provided in <a href="mailto:ceph\_user\_model.yml">ceph\_user\_model.yml</a> and sets up an OpenStack client node. Perform the following steps:

- 1. Define the user model.
- Run hosts/verb\_hosts ceph-client-prepare.yml.

The default user model provided with HPE Helion OpenStack® is mentioned below. The default configuration contains pool configuration for cinder-volume, glance, and cinder-backup. You can edit the file if you do not intend to use Ceph for all services or want to change pool attributes based on your cluster configuration. For example, if you have large number of disks then you can increase the Placement Group (PG) number for a given pool.

```
- - -
product:
   version: 2
ceph_user_models:
    - user:
        name: cinder
        type: openstack
        secret_id: 457eb676-33da-42ec-9a8c-9293d545c337
      pools:
        - name: volumes
          attrs:
            creation_policy: eager
            type: replicated
            replica_size: 3
            permission: rwx
            pg: 100
          usage:
            consumer: cinder-volume
        - name: vms
          attrs:
            creation_policy: eager
            type: replicated
            replica size: 3
            permission: rwx
            pg: 100
          usage:
            consumer: nova
        - name: images
          attrs:
            creation policy: lazy
            permission: rx
    - user:
        name: glance
        type: openstack
      pools:
        - name: images
          attrs:
            creation_policy: eager
```

300 Usage of Ceph Storage

```
type: replicated
        replica_size: 3
       permission: rwx
        pg: 128
     usage:
        consumer: glance-datastore
- user:
   name: cinder-backup
   type: openstack
 pools:
    - name: backups
     attrs:
       creation_policy: eager
       type: replicated
        replica_size: 3
       permission: rwx
       pg: 128
     usage:
        consumer: cinder-backup
```

The following table provides the descriptions of the preceding parameters.

Parameter	Value	Description	Recommendation
user.name	A user-defined string.	Defines the name of the user who can access a set of pools. A user is created with same name in the Ceph system.	Retain the default names for some of the well-known OpenStack users as described in the default user model. For example, cinder user for volume access who needs to access volumes, vms, and the images pool, as defined in the default model.
user.type	OpenStack   user.	Indicates whether the user is specific for OpenStack ser-	Use <b>openstack</b> for pools used by cinder-volume, cin-

301

Parameter	Value	Description	Recommendation
		vices. This parameter does not have semantic implications.	der-backup, glance, and Nova services. For other services, you can choose user.
user.secret_id	UUID instance.	A unique UUID.	Generate a value for your cluster before setting up the Cinder client. The <u>lib-virt</u> process needs this value to access the cluster while attaching a block device from Cinder.
pools.name	A user-defined string.	A user-defined name.	The name has to be unique in the Ceph namespace.
pools.attrs.cre- ation_policy	eager   lazy	If creation_policy is "eager," the playbooks will create the user. If the creation_policy is set to "lazy," the pool will be created externally (not by Ceph ansible playbooks) out of band.	Retain default values for pools meant to be consumed by OpenStack services.
pools.attrs.type	Replicated.	Defines the type of pool. Ceph supports erasure-coded and replicated pools.	Retain the value as replicated if you do want to use

302 Usage of Ceph Storage

Parameter	Value	Description	Recommendation
			ceph-client-pre-pare.yml for pool creation.  This does not mean that you cannot create an erasure-code pool with your cluster deployed using HPE Helion OpenStack®. It just means that erasure-code pools are not officially supported.
pools.attrs.repli- ca_size	1n	Replica count.	Use 3 as the replica count because it is used in most common scenarios providing the right level of protection and is the minimum setting that HPE Helion OpenStack supports.
pools.attrs.permission	rwx	Read, write, and execute the permission a user will have on given pool.	Retain the values for OpenStack user as mentioned in the default user model. Altering these values might affect your client setup configuration.

303 Usage of Ceph Storage

Parameter	Value	Description	Recommendation
pools.attrs.pg	Placement group number.	Number of placement groups.	The default value is 128. You can change the value of this parameter if your disk size is larger.
pools.usage.con- sumer	Predefined choices.	Defines pool consumers, and indicates which services are expected to consume a given pool with which access. It has semantic meaning for the following predefined choices:  a. nova  b. glance-datastore  c. cinder-volume  d. cinder-back-up	Do not pick the value from glancedatastore, nova, cinder-volume, or cinder-backup for user-defined pools. For user-defined pools, you can skip this parameter.

## 17.1.12.3 Prerequisites

To use Ceph as a volume backend, the nodes running these services should have the Ceph client installed on them. Use the <a href="ceph-client-prepare.yml">ceph-client-prepare.yml</a> playbook to deploy the Ceph client on these nodes.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ceph-client-prepare.yml
```

This playbook also creates Ceph users and Ceph pools on the resource nodes.



### Note

The following steps install packages and configure the existing client nodes (such as the Cinder, Glance, and Nova Compute nodes) required to use the Ceph cluster. For any new client nodes added later on that need to be configured to use the Ceph cluster, just execute the preceding playbook with the addition of the --limit <new-client-node> switch.

## 17.1.12.4 Use Ceph Storage Pools as Glance Data Stores

To enable Ceph as a Glance data store, perform the following steps:

- 1. Log in to the lifecycle manager.
- 2. Make the following changes in the \_~/helion/my\_cloud/config/glance/glance-api.conf.j2 file:
  - a. Copy the following content and paste it into the <u>glance-api.conf.j2</u> file under [glance\_store]:

```
[glance_store]
default_store = rbd
stores = rbd
rbd_store_pool = images
rbd_store_user = glance
rbd_store_ceph_conf = /etc/ceph/ceph.conf
rbd_store_chunk_size = 8
```

b. In the same file, comment out the following references to Swift:

```
stores = {{ glance_stores }}
default_store = {{ glance_default_store }}
```

- 3. **IMPORTANT:** If you have preexisting images in your Glance repo and want to use Ceph exclusively as a backend, use the Glance CLI or other tool to back up your images before configuring Ceph as your Glance backend:
  - a. Snapshot or delete all Nova instances using those images.
  - b. Download the images locally that you want to save.
  - c. Delete all the images from Glance.

After you finish the Ceph configuration you will need to add those images again.

**4.** Commit your configuration to the local git repo (*Chapter 3, Using Git for Configuration Management*), as follows:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```

5. Run the configuration processor:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

6. Update your deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

7. Run the Glance reconfigure playbook to configure Ceph as a Glance backend:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts glance-reconfigure.yml
```

After you execute the preceding command successfully, the glance is configured to use Ceph as a data store. You can use glance operations to upload images, which will be stored in the Ceph cluster.

## 17.1.12.5 Use Ceph Storage Pools as Cinder Volume Backends

To enable Ceph as a Cinder volume backend, follow these steps:

- 1. Log in to the lifecycle manager.
- 2. Make the following changes to the <u>~/helion/my\_cloud/config/cinder/cinder.con-f.j2</u> file:
  - a. Add your Ceph backend to the enabled\_backends section:

```
# Configure the enabled backends
enabled_backends=ceph1
```

306

# Important

If you are using multiple backend types, you can use a comma-delimited list here. For example, if you are going to use both VSA and Ceph backends, specify enabled backends=vsa-1, ceph1.

b. [optional] If you want your volumes to use a default volume type, enter the name of the volume type in the <a href="[DEFAULT]">[DEFAULT]</a> section with the following syntax. Make note of this value because you will need it when you create your volume type later.



## **Important**

If you do not specify a default type, then your volumes will default to a nonredundant RAID configuration. We recommend that you create a volume type that meets your environments needs and specify it here.

```
[DEFAULT]
# Set the default volume type
default_volume_type = <your new volume type>
```

c. Uncomment the <u>ceph</u> section and supply the values to match your cluster information. If you have more than one cluster, you will need to add another similar section with respective values. The following example adds only one cluster.

```
[ceph1]
rbd_secret_uuid = <secret-uuid>
rbd_user = <ceph-cinder-user>
rbd_pool = <ceph-cinder-volume-pool>
rbd_ceph_conf = <ceph-config-file>
rbd_cluster_name = <cluster_name>
volume_driver = cinder.volume.drivers.rbd.RBDDriver
volume_backend_name = <ceph-backend-name>
```

This example has the following values:

Value	Description
rbd_secret_uuid	The secret_id value from the ~/he-lion/my_cloud/config/ceph/user_model.yml file, as highlighted here:
	- user:     name: cinder     type: openstack     secret_id: <secret be="" here="" id="" will=""> pools:     - name: volumes  Important     You should generate and use     your own secret ID. You can use     any UUID generation tool.</secret>
rbd_user	The username value from the/ helion/my_cloud/config/ceph/ user_model.yml file, as highlighted here:  - user:
	<pre>name: cinder   type: openstack   secret_id: <secret be="" here="" id="" will=""> pools:</secret></pre>

Value	Description
rbd_pool	The pool name value from the/ helion/my_cloud/config/ceph/ user_model.yml file, as highlighted here:
	<pre>- user:     name: cinder     type: openstack     secret_id:     457eb676-33da-42ec-9a8c-9293d545c337 pools:     - name: volumes</pre>
rbd_ceph_conf	The Ceph configuration file location, usually /etc/ceph/ceph.conf.
rbd_cluster_name	Name of the Ceph cluster.
volume_driver	The Cinder volume driver. Leave this as the default value specified for Ceph.
volume_backend_name	The name given to the Ceph backend.

- 3. To enable attaching Ceph volumes to Nova provisioned instances, make the following changes to the ~/helion/my\_cloud/config/nova/kvm-hypervisor.conf.j2 file:
  - a. Uncomment the Ceph backend lines and edit them as follows:

```
[libvirt]
rbd_user = <ceph-user>
rbd_secret_uuid = <secret-uuid>
```

309 Usage of Ceph Storage

The values are as follows:

Value	Description
rbd_user	The username value from the ~/
	helion/my_cloud/config/ceph/
	user_model.yml file, as highlighted
	here:
	- user: name: <b>cinder</b>
	type: openstack
	secret_id:
	457eb676-33da-42ec-9a8c-9293d545c337
rbd_secret_uuid	The secret_id value from the _~/he-
	lion/my_cloud/config/ceph/
	user_model.yml file, as highlighted
	here:
	- user:
	name: cinder
	type: openstack
	secret_id: <b>457eb676-33da-42ec-9a8c-9293d5</b>



## Note

To attach a volume provisioned out of a newly added Ceph backend to an existing OpenStack instance, the instance must be rebooted after the new backend has been added.

## Important

Do not use <u>backend\_host</u> variable in <u>cinder.conf</u> file. If <u>backend\_host</u> is set, it will override the [DEFAULT]/host value which HPE Helion OpenStack 5.0 is dependent on.

**4.** Commit your configuration to the local git repo (*Chapter 3, Using Git for Configuration Management*), as follows:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```

5. Run the configuration processor:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

6. Update your deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

7. Run the Cinder reconfigure playbook:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts cinder-reconfigure.yml
```

8. If Nova has been configured to attach Ceph backend volumes, run the Nova reconfigure playbook:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts nova-reconfigure.yml
```

After you execute the preceding command successfully, the cinder-volume is configured to use Ceph as a data store. You can use volume lifecycle operations as described at *FIXME: broken external xref* 

### 17.1.12.6 Use Ceph Storage Pools as Cinder Backup Devices

To enable Cinder backup devices for Ceph, make the following changes to the <u>~/he-lion/my\_cloud/config/cinder/cinder.conf.j2</u> file:

1. Uncomment the ceph backup section and supply the values:

```
[DEFAULT]
backup_driver = cinder.backup.drivers.ceph
```

```
backup_ceph_conf = <ceph-config-file>
backup_ceph_user = <ceph-backup-user>
backup_ceph_pool = <ceph-backup-pool>
```

The values are as follows:

Value	Description
backup_driver	The Cinder volume driver. Leave this as the default value specified for Ceph.
backup_ceph_conf	The Ceph configuration file location, usually: <a href="mailto://etc/ceph/ceph.conf">/etc/ceph/ceph.conf</a> .
backup_ceph_user	The username value from the _~/he-lion/my_cloud/config/ceph/user_model.yml file, as highlighted here:  - user:     name: cinder-backup     type: openstack pools:     - name: backups
backup_ceph_pool	The pool name value from the ~/he-lion/my_cloud/config/ceph/user_model.yml file, as highlighted here:  pools:     name: backups

2. Commit your configuration to the local git repo (*Chapter 3, Using Git for Configuration Management*), as follows:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```

3. Run the configuration processor:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

4. Update your deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

5. Run the Cinder reconfigure playbook:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts cinder-reconfigure.yml
```

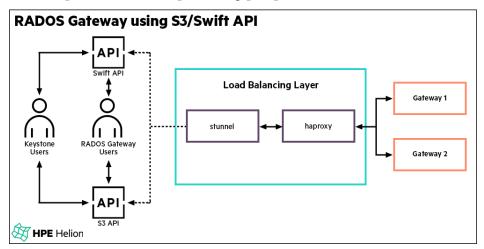
After you execute the preceding command successfully, the cinder-backup service is configured to use Ceph as a data store. You can use volume lifecycle operations as described in .

### 17.1.12.7 Use RADOS Gateway to Access Objects Using S3/Swift API

RADOS Gateway provides object storage functionality through S3 and Swift API. It has a built-in authentication mechanism and it can use Keystone as an authentication backend, unlike the OpenStack Services, where Keystone is the only authentication backend. Users primarily see the following combinations of object accessibility:

- 1. Keystone users using:
  - a. Swift API
  - b. S3 API
- 2. RADOS Gateway users using:
  - a. Swift API
  - b. S3 API

Here's a pictorial view of preceding perspective:



The first view uses Keystone as the authentication backend while the second one uses RADOS Gateway (internal authentication backend). Default deployment of RADOS Gateway in HPE Helion OpenStack® 5.0 enables 1.a, 2.a, and 2.b only.

Enabling choice 1.b for Keystone users accessing S3 API requires Keystone to be the default authentication backend. In this case all user authentication goes first to Keystone, with failover to RADOS Gateway even for non-Keystone users (such as in case 2.b), but this situation causes a serious performance drop for RADOS Gateway users.

Lab tests performed on a 10 TB Ceph cluster have shown that enabling S3 API access for Keystone users degrades the performance of S3 API access for RADOS Gateway users roughly by 50%. We ran the rest-bench tool for 10 seconds with a concurrency of 16 and an object size of 4096 (4K) bytes on Ceph deployed using the <a href="mailto:entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> input model. Our lab finding was that the system performed an average of 1350 (object size 4K) write operations when Keystone was enabled as the default authentication backend, in contrast to 2500 (object size 4K) operations in the default configuration, which was a degradation of roughly 49%. Note that these numbers are for illustration purposes to provide an insight on the degree of degradation. However, storing all credentials in Keystone provides the advantage of reducing the maintenance burden. It is not required to create or manage credentials in a database for S3 authentication. You can use standard administrative tools such as Horizon instead. Aspect 1.b does not apply if you do not intend to enable S3 API access for Keystone users.

#### 17.1.12.7.1 Steps to Access S3 API for RADOS Gateway Users

#### Perform these steps:

1. Create a user by executing the following command:

```
sudo radosgw-admin user create --uid="{username}" --display-name="{Display Name}" --
email="{email}"
```

#### Example:

```
sudo radosgw-admin user create --uid=rgwuser --display-name="RadosGatewayUser" -- email=rgwuser@test.com
```

2. Perform the object operation. You can use S3 clients. Example: python-boto.

#### 17.1.12.7.2 Steps to Access Swift API for RADOS Gateway Users

#### Perform these steps:

1. Create a user by executing the following command:

```
sudo radosgw-admin user create --uid="{username}" --display-name="{Display Name}" --
email="{email}"
```

#### Example:

```
sudo radosgw-admin user create --uid=rgwuser --display-name="RadosGatewayUser" --
email=rgwuser@test.com
```

2. Create a subuser by executing the following command:

```
sudo radosgw-admin subuser create --uid={username} --subuser={username}:
{subusername} --access=full --key-type=swift --gen-secret
```

#### Example:

```
sudo radosgw-admin subuser create --uid="rgwuser" --subuser="rgwuser:rgwswiftuser"
    --access=full --key-type=swift --gen-secret
```

3. Perform the object operation using Swift CLI or any compatible client.

#### 17.1.12.7.3 Steps to Access Swift API for Keystone Users

When RADOS Gateway is deployed per the default <a href="entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> model, it exists alongside OpenStack Swift (this is because the Keystone service type for RADOS Gateway defaults to <a href="ceph-object-store">ceph-object-store</a>). In this (coexisting) mode, OpenStack Horizon communicates only with OpenStack Swift for all the object store operations. However, the OpenStack Swift command line interface can be tuned to communicate with both OpenStack Swift and RADOS Gateway as follows:

• To interact with Swift (default):

```
export OS_SERVICE_TYPE=object-store
```

To interact with RADOS Gateway:

```
export OS_SERVICE_TYPE=ceph-object-store
```

You can use Swift CLI to perform object operations on Ceph.

#### Example:

Log in to the monitor node to list the objects in Ceph using the credentials available in the  $\frac{\sim}{}$  service.osrc file.

```
source ~/service.osrc
```

Execute the following command to list objects in the Ceph object store:

```
export OS_SERVICE_TYPE=ceph-object-store
swift list
```

Execute the following command to list objects in the Swift object store:

```
export OS_SERVICE_TYPE=object-store
swift list
```

#### 17.1.12.7.4 Steps to Access S3 API for Keystone Users

By default, a Keystone user can access the RADOS Gateway functionality using the Swift API. To configure a Keystone user for S3 API to access the RADOS Gateway, perform the following steps:

- 1. Login to lifecycle manager.
- 2. Edit the \_\_~/helion/my\_cloud/config/ceph/settings.yml file and set the \_rg-w\_s3\_auth\_use\_keystone value to **true**.

```
rgw_s3_auth_use_keystone: true
```

3. Commit your configuration to local repo.

```
cd ~/helion/hos/ansible
git add -A
git commit -m <"commit message">
```

4. Run ready-deployment playbook.

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

5. Run the ceph-reconfigure playbook.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ceph-reconfigure.yml
```

#### Create the EC2 credentials

a. Login to a controller node and source the admin user credentials.

```
source ~/service.osrc
```

b. Execute the following commands to generate the EC2 credentials for the OpenStack user.

```
openstack ec2 credentials create --project <project-name> --user <user-name>
```

For example, to generate the EC2 credentials for user demo for the project demo.

```
openstack ec2 credentials create --project demo --user demo
```

## Important

Please contact the Professional Services team for details on how to perform the preceding steps.

## 17.1.13 Managing Ceph Clusters After Deployment

After you successfully deploy your Ceph cluster and its is functioning as expected, you can perform various maintenance operations for the Ceph cluster when required. For example, you can change the configuration of Ceph service components, or scale up storage nodes. The following list provides details about various maintenance options that you might encounter for cluster management. Text links take you to the instructions for each step.

1. Modify service configuration.
2. Scale out cluster.
a
b
c. See .
3. Scale down cluster.
a
b
c
4. Repair scenarios.
a
b
c
d
e
5
Note that the preceding maintenance procedure is based on the <a href="entry-scale-kvm-ceph">entry-scale-kvm-ceph</a> input model. If you have altered your deployment, such as by deploying the monitor service on standalone nodes, you might have to adjust the procedure accordingly.

## 17.2 Configuring for VSA Block Storage Backend

This page describes how to configure your VSA backend for the Helion Entry-scale with KVM cloud model. This procedure can be performed during or after installation.

## 17.2.1 Prerequisites

- Your HPE Helion OpenStack cloud must be fully deployed using the KVM cloud model with VSA added. The Entry-Scale KVM with VSA cloud model is an example of such a model. For more details on the installation refer to *Chapter 10, Installing Mid-scale and Entry-scale KVM*.
- It's important that all of your systems have the correct date/time because the HPE Store-Virtual VSA license has a start date. If the start date is later than the system time then the installation will fail.

#### 17.2.2 Notes

- The license for the HPE StoreVirtual VSA license is bundled with HPE Helion OpenStack and comes with a free trial which allows a maximum limit of 50 TB per node. Hence the total amount of the configured storage on an individual VSA node should not exceed 50 TB. To extend the 50 TB per node limit, you can add nodes. A VSA cluster can support up to 16 nodes, which means configured storage on a VSA cluster can be as much as 800 TB. You can deploy VSA with Adaptive Optimization (AO) or without. The deployment process for each of these options is similar, you just need to make a change in the disk input model. For more detailed information, refer to VSA with or without Adaptive Optimization (AO) in Book "Planning", Chapter 11 "Modifying the Entry-scale KVM with VSA Model for Your Environment".
- A single VSA node can have a maximum of seven raw disks (excluding the operating system disks) attached to it, which is defined in the disk input model for your VSA nodes. It is expected that no more than seven disks are specified (including Adaptive Optimization disks) per VSA node. For example, if you want to deploy VSA with two disks for Adaptive Optimization then your disk input model should not specify more than five raw disks for data and two raw disks for Adaptive Optimization. Exceeding the disk limit causes VSA deployment failure.

- You can create more than one VSA cluster of same or different type by specifying the configuration in cloud model. For more details, refer to .
- Usernames and passwords designated during VSA configuration or repair must be alphabetic only.

Concerning using multiple backends: If you are using multiple backend options, ensure that you specify each of the backends you are using when configuring your <a href="cin-der.conf.j2">cin-der.conf.j2</a> file using a comma delimited list. An example would be <a href="enabled\_backend-backend-s=vsa-1,3par\_iSCSI,ceph1">enabled\_backend-backen

VSA driver has updated name: In the OpenStack Mitaka release, the VSA driver used for HPE Helion integration had its name updated from <a href="https://hplefthand\_">hplefthand\_</a> to <a href="https://hplefthand\_">hplefthand\_</a>. To prevent issues when upgrading from previous HPE Helion OpenStack releases, we left the names as-is in the release and provided a mapping so that the integration would continue to work. This will produce a warning in the <a href="cinder-volume.log">cinder-volume.log</a> file advising you of the deprecated name. The warning will look similar to this:

```
Option "hplefthand_username" from group "<your section>" is deprecated. Use option "hpelefthand_username" from group "<your section>" should be ignored.
```

These are just warnings and can be ignored.

## 17.2.3 Configure HPE StoreVirtual VSA

The process for configuring HPE StoreVirtual VSA for HPE Helion OpenStack involves two major steps:

- Creating a cluster
- Configuring VSA as the block storage backend

You can perform these tasks in an automated process using the provided Ansible playbooks or manually using the HPE StoreVirtual Centralized Management Console (CMC) GUI.

## N ■ N ■ N ■ N

### Note

Using Ansible playbooks to create clusters is strongly recommended. Cluster creation using HPE StoreVirtual Centralized Management Console (CMC) GUI is deprecated in HPE Helion OpenStack 5.0. Clusters created using CMC manually cannot be reconfigured using the Ansible playbooks. If you used CMC to create the clusters, you must continue to use CMC to administer those clusters.

## 17.2.4 Using Ansible

Perform the following steps to create the cluster using the provided Ansible playbooks.

- 1. Log in to the lifecycle manager.
- 2. Run the following playbook, which will both create your management group and your cluster:

```
cd ~/scratch/ansible/next/hos/ansible/
ansible-playbook -i hosts/verb_hosts vsalm-configure-cluster.yml
```

3. When prompted, enter an administrative user name and password you will use to administer the CMC utility.

Administrative user naming requirements:

- 3 to 30 characters
- Must begin with a letter
- Allowed characters: 0-9, a-z, A-Z, hyphen (-), underline ()
- Disallowed characters: Multibyte character set

Administrative user password requirements:

- 5 to 40 characters
- Must begin with a letter
- Many ASCII and extended ASCII characters, Multibyte character set
- Disallowed characters: dot (.), colon (:), forward slash (/), comma (,), backward slash (\), semi-colon (;), single-quote ('), space ( ), equals (=)

321 Using Ansible

## Important

You will need to remember these values to access the cluster using the CMC utility or to re-run this playbook later if further management of your VSA system is needed. Do not change the credentials for one or more clusters using CMC as it will make the cluster automation playbook ignore those clusters for reconfiguration.

4. In the output of the Ansible playbook you will see different steps where you can locate the values for your VSA cluster name, the virtual IP address, and the IP addresses of each of the VSA hosts in the cluster. You will use these values later when configuring your backend. The following is an example, with the important information bolded.

This step displays the VSA cluster name and the virtual IP address of the cluster:

```
TASK: [vsalm | _create_cluster | Display the status of cluster creation performed in
  above step] ***
ok: [helion-cp1-vsa0001-mgmt] => {
    "msg": "Created cluster - Name=cluster-vsa, IP=10.13.111.142"
}
```

- 5. You can view the added node to the cluster in CMC GUI as follows:
  - a. Launch the CMC utility. For more information, see Section 17.2.5, "Using the CMC Utility".
  - b. Find the new VSA system.
  - c. View the added node in the cluster.

## 17.2.5 Using the CMC Utility

In HPE Helion OpenStack, cluster creation using HPE StoreVirtual Centralized Management Console (CMC) GUI is deprecated. The clusters created by CMC manually cannot be configured using Ansible playbooks.

Perform the following steps to create the cluster using CMC.

#### Launching the CMC utility GUI

The HPE StoreVirtual Centralized Management Console (CMC) GUI is an interface where you can configure VSA.

The CMC utility requires a GUI to access it. You can use either of the following methods to launch the CMC GUI.

- Section 17.2.5.1, "RDP/VNC Connection Method"
- Section 17.2.5.2, "Install (any) X Display Tool"

#### 17.2.5.1 RDP/VNC Connection Method

Set up VNC connect on the Controller Node

The following steps will allow you to setup a VNC connect to your controller node so you can view the GUI.

- 1. Log in to your first controller node.
- 2. Run the following command to install the package that is required to launch CMC:

```
sudo apt-get install -y xrdp
```

3. Start <u>vnc4server</u> using the instructions below. You will be prompted for a password (minimum 6 characters). Enter a password and proceed. A sample output is shown below:

```
stack@helion-cpl-cl-m1-mgmt:~$ vnc4server

You will require a password to access your desktops.

Password:
Verify:
xauth: file /home/stack/.Xauthority does not exist

New 'helion-cpl-cl-m1-mgmt:3 (stack)' desktop is helion-cpl-cl-m1-mgmt:3

Creating default startup script /home/stack/.vnc/xstartup
Starting applications specified in /home/stack/.vnc/xstartup
Log file is /home/stack/.vnc/helion-cpl-cl-m1-mgmt:3.log
```



#### Note

If you use  $\underline{\mathbf{xrdp}}$  directly to connect to the first controller node without using the VNC server, then a remote session is created whenever you login. To avoid this, a dedicated VNC server instance is launched and connected to that instance by  $\underline{\mathbf{xrdp}}$ . This helps to maintain the session.

4. Run **netstat -anp | grep vnc** to determine the public port that VNC is using. In the example below, the port is 5903:

```
stack@helion-cp1-c1-m1-mgmt:~$ netstat -anp | grep vnc

(Not all processes could be identified, non-owned process info
will not be shown, you would have to be root to see it all.)

tcp 0 0 0.0.0.6003 0.0.0.0:* LISTEN

1413/Xvnc4

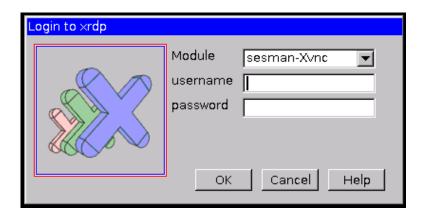
tcp6 0 0 :::5903 :::* LISTEN
```



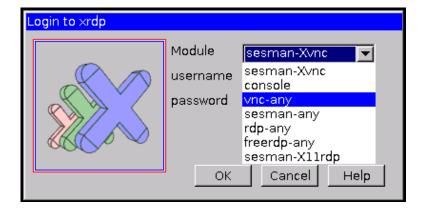
#### Note

If you reboot the controller node, repeat the instructions starting from *Step 3*.

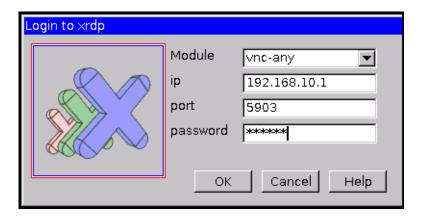
- 5. Connect to your controller node through any remote desktop or VNC client. We will show the **xrdp** method first and the VNC method is below it:
  - a. Connecting through remote desktop client
    - i. Login to your remote desktop. You will be prompted with xrdp login screen.



ii. Click the Module drop-down list and select vnc-any.

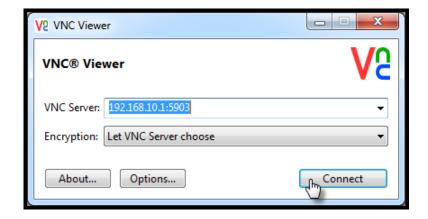


iii. Enter the IP address, port and password in the respective fields.



iv. Click Ok.

- b. Connecting through a VNC client, such as VNC Viewer (https://www.re-alvnc.com/download/viewer/) ▶:
  - Enter the IP address and port and click **Connect**. You will be prompted for your password once the connection is established.



A terminal emulator will be displayed where you can enter the CMC launch command. Note that the CMC launch with this method will have the following limitations: by default, CMC-xterm disables all the title bars and borders. This is an expected behavior.

### 17.2.5.2 Install (any) X Display Tool

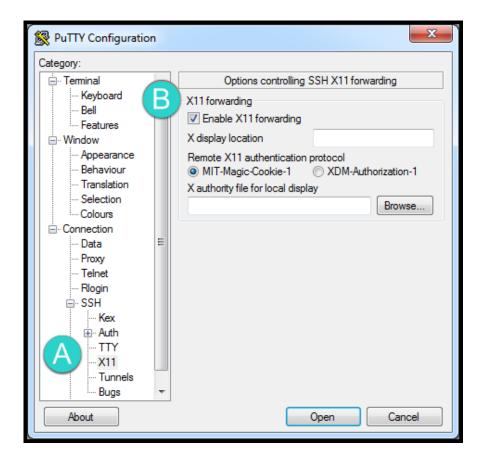


#### Note

You can use SSH to an X server but the performance may be poor.

You must configure an X display tool to launch CMC. You can the choose any X display tool. In this section we are using the tool Xming as an example for launching CMC. The following example provides the steps to installing Xming and launching CMC in a Windows environment. Another alternative (not shown in the documentation) is MobaXterm (http://mobaxterm.mobatek.net/) .

- 1. Download and install Xming on a Windows machine that can access the lifecycle manager. You can download Xming from http://sourceforge.net/projects/xming/ ▶.
- 2. Select *Enable X11 forwarding* checkbox on the PuTTY session for lifecycle manager. You can do this in PuTTY by:
  - a. Navigate to the *Connection* > *SSH* > *X11* option in PuTTY.
  - b. Make sure *Enable X11 forwarding* is activated.



3. SSH to first control plane node.

```
ssh -X
```

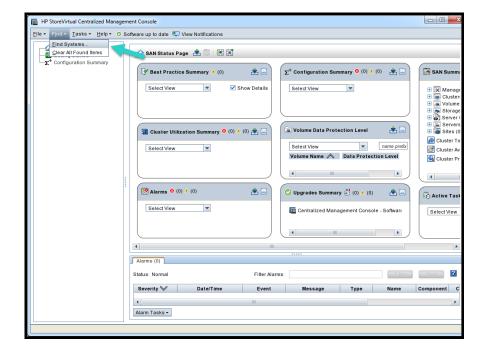
and enter the CMC command (as mentioned below) to launch CMC.

1. Run the following command from your first controller node which will open the CMC GUI on your local machine:

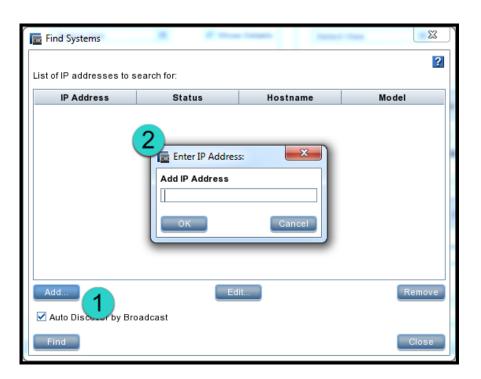
```
/opt/HP/StoreVirtual/UI/jre/bin/java -jar /opt/HP/StoreVirtual/UI/UI.jar
```

By default, the CMC GUI is configured to discover the VSA nodes in the subnet in which it is installed. This discovery functionality of VSA nodes using the CMC controller node is not supported in HPE Helion OpenStack. Instead, you must manually add each VSA node, as shown below.

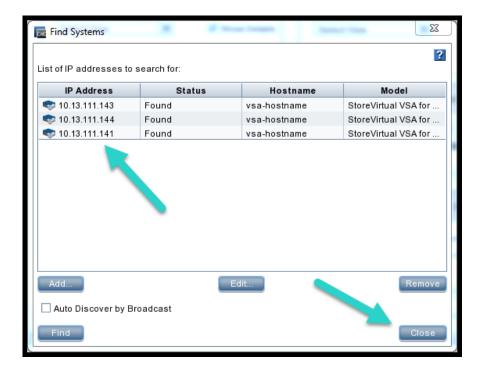
2. In the CMC GUI, click the Find menu and then select the Find Systems options.



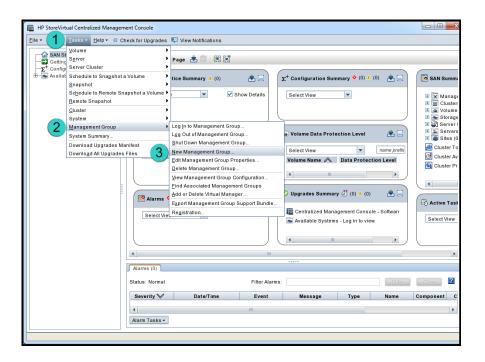
3. Click the *Add* button which will open the *Enter IP Address* dialogue box where you can enter the IP address of your VSA nodes which you noted earlier from your ~/scratch/ansible/next/my\_cloud/stage/info/net\_info.yml file.



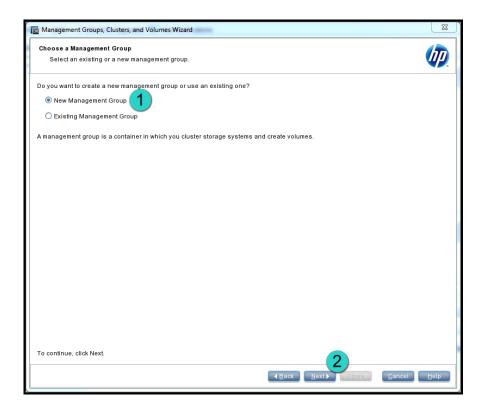
4. Once you have all of your VSA nodes entered, click the **Close** button.



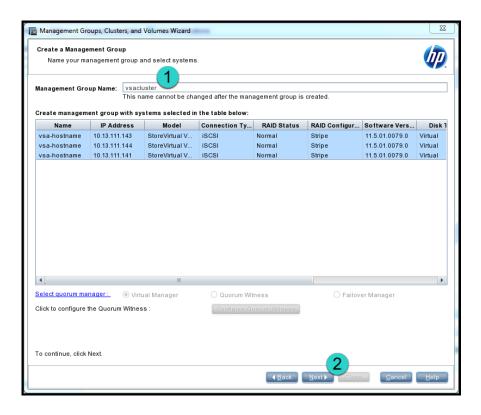
5. Next, click the *Tasks* menu and then navigate to the *Management Group* submenu and select the *New Management Group* option.



6. In the Management Group wizard, click *Next* and then select *New Management Group* and then *Next* again to continue.



7. Enter a name in the New Management Group Name field and then click Next.

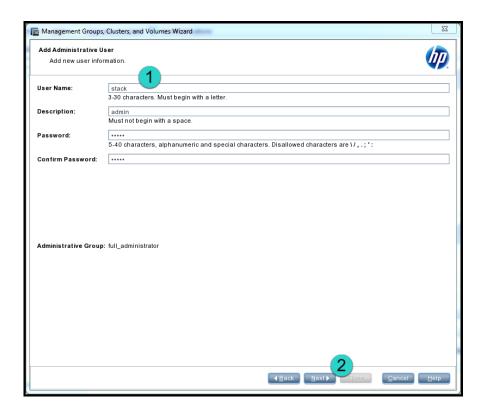


8. On the *Add Administrative User*, enter a username and password you will use to administer the CMC utility.

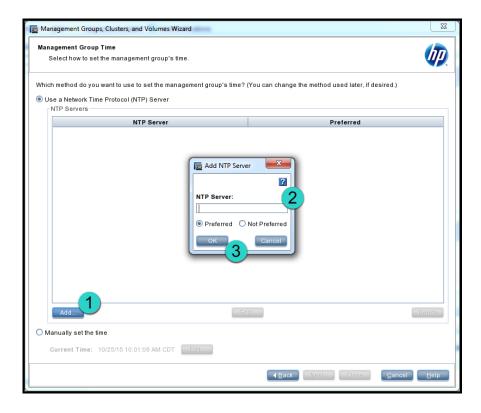


## **Important**

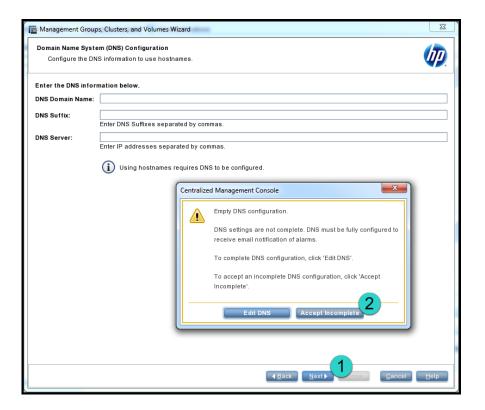
You will need to remember these values as you will input them into your <u>cin-der.conf.j2</u> file later.



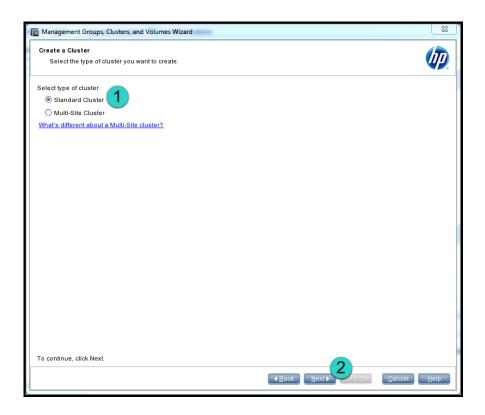
- 9. Click Nextto display the Management Group Time page.
- 10. Add your NTP server information and click Next



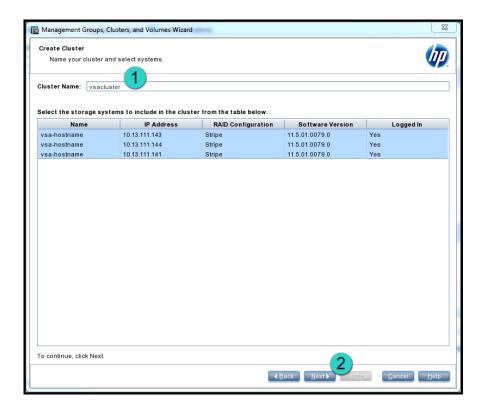
11. Skip the DNS and SMTP sections. To do so, click *Next* and a popup will display where you can choose the *Accept Incomplete* option. Repeat this to skip SMTP section as well.



12. On the *Create a Cluster* options, select *Standard Cluster* from the displayed options and click *Next*.



13. In the Cluster Name field, enter a name for the cluster and click Next.

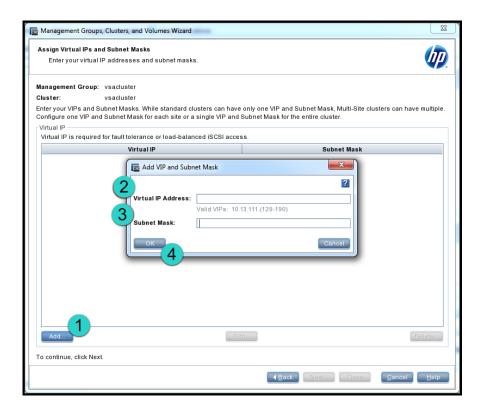


14. On the *Assign Virtual IPs and Subnet Masks* page, click *Add* and enter the virtual IP address and subnet mask of the cluster in the respective boxes and click *OK*.

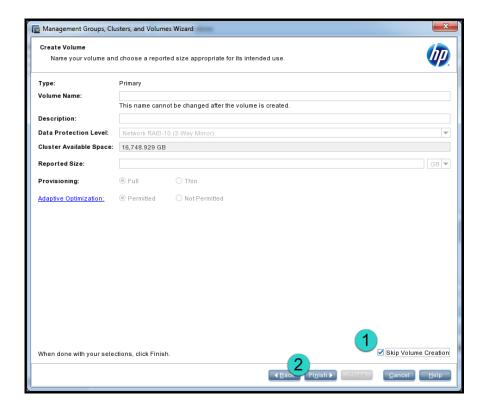


#### Note

The virtual IP address will be found as the <a href="cluster\_ip">cluster\_ip</a> value in your <a href="~/scratch/ansible/next/my\_cloud/stage/info/net\_info.yml">msible/next/my\_cloud/stage/info/net\_info.yml</a> file and your subnet mask will be the subnet address from the network your VSA nodes are attached to, usually your MANAGEMENT network.



- 15. The CMC utility will verify the virtual IP address information and then you can click the *Next* button.
- 16. Select the checkbox for *Skip Volume Creation* and click the *Finish* button which will display your VSA management cluster.



Cluster creation using CMC takes approximately 10 minutes or longer.



### Tip

You may get a pop-up notice telling you that your hostnames are not unique. This can be ignored by clicking the OK button.

17. If this process is successful you will see a summary page at the end which outlines what you have completed.



## **Important**

You can create more than one VSA cluster of same or different type by specifying the configuration in cloud model. For more details, refer to Modifying Cloud Model to Create Multiple Clusters

## 17.2.6 Configure VSA as the Backend

You will use the information you input to the CMC utility to configure your Cinder backend to use your VSA environment.

To update your Cinder configuration to add VSA storage you must modify the ~/he-lion/my\_cloud/config/cinder/cinder.conf.j2 file on your lifecycle manager as follows:

- 1. Log in to the lifecycle manager.
- 2. Make the following changes to the <u>~/helion/my\_cloud/config/cinder/cinder.con-f.j2</u> file:
  - a. Add your VSA backend to the enabled\_backends section:

# Configure the enabled backends
enabled\_backends=vsa-1



### **Important**

If you are using multiple backend types, you can use a comma delimited list here. For example, if you are going to use both VSA and Ceph backends, you would specify something like this: enabled\_backends=vsa-1,ceph1.

b. [OPTIONAL] If you want your volumes to use a default volume type, then enter the name of the volume type in the <a href="[DEFAULT]">[DEFAULT]</a> section with the syntax below. Remember this value when you create your volume type in the next section.



#### **Important**

If you do not specify a default type then your volumes will default to a non-redundant RAID configuration. It is recommended that you create a volume type and specify it here that meets your environments needs.

[DEFAULT]
# Set the default volume type
default\_volume\_type = <your new volume type>

c. Uncomment the <u>StoreVirtual (VSA) cluster</u> section and fill the values as per your cluster information. If you have more than one cluster, you will need to add another similar section with its respective values. In the following example only one cluster is added.

```
[vsa-1]
hplefthand_password: <vsa-cluster-password>
hplefthand_clustername: <vsa-cluster-name>
hplefthand_api_url: https://<vsa-cluster-vip>:8081/lhos
hplefthand_username: <vsa-cluster-username>
hplefthand_iscsi_chap_enabled: false
volume_backend_name: <vsa-backend-name>
volume_driver:
    cinder.volume.drivers.san.hp.hp_lefthand_iscsi.HPLeftHandISCSIDriver
hplefthand_debug: false
```

#### where:

Value	Description
hplefthand_password	Password entered during cluster creation in the CMC utility. If you have chosen to encrypt this password, enter the value in this format:  hplefthand_password: {{ ' <encrypted vsa-cluster-password="">'   hos_user_password_decrypt }}</encrypted>
	See for more details.
hplefthand_clustername	Name of the VSA cluster provided while creating a cluster in the CMC utility.
hplefthand_api_url	Virtual IP address of your VSA cluster, found in your <pre>~/scratch/an-sible/next/my_cloud/stage/in-fo/net_info.yml</pre> file.

Value	Description
hplefthand_username	Username given during cluster creation in the CMC utility.
hplefthand_iscsi_chap_enabled	If you set this option as <u>true</u> then the hosts will not be able to access the storage without the generated secrets. And if you set this option as <u>false</u> then no CHAP authentication is required for the ISCSI connection.
volume_backend_name	Name given to the VSA backend. You will specify this value later in the Associate the Volume Type to a Backend steps.
volume_driver	Cinder volume driver. Leave this as the default value for VSA.
hplefthand_debug	If you set this option as true then the Cinder driver for the VSA will generate logging in debug mode; these logging entries can be found in <a href="cinder-vol-ume.log">cinder-vol-ume.log</a> .

[OPTIONAL] HPE Helion OpenStack 5.0 supports VSA deployment for KVM hypervisor only but it can be used as pre-deployed (or out of the band deployed) Lefthand storage boxes or VSA appliances (running on ESX/hyper-v/KVM hypervisor). It also supports Cinder configuration of physical Lefthand storage device and VSA appliances. Depending upon your setup, you will have to edit the below section if your storage array is running LeftHand OS lower than version 11:

```
[<unique-section-name>]
volume_driver=cinder.volume.drivers.san.hp.hp_lefthand_iscsi.HPLeftHandISCSIDriver
volume_backend_name=lefthand-cliq
san_ip=<san-ip>
san_login=<san_username>
If adding a password here, then the password can be encrypted using the
```

mechanism specified in the documentation. If the password has been encrypted add the value and the hos\_user\_password\_decrypt filter like so:

san\_password= {{ '<encrypted san\_password>' | hos\_user\_password\_decrypt }}

Note that the encrypted value has to be enclosed in quotes

If you choose not to encrypt the password then the unencrypted password must be set as follows:

san\_password=<san\_password>
san\_ssh\_port=16022

san\_clustername=<vsa-cluster-name>

volume\_backend\_name=<vsa-backend-name>



## **Important**

Similar to your <a href="https://hydro.com/hplefthand\_password">hplefthand\_password</a> in the previous example, encryption for your <a href="mailto:san\_password">san\_password</a> is supported. If you chose to use encryption you would use the syntax below to express that:

```
san_password= {{ '<encrypted san_password>' |
hos_user_password_decrypt }}
```

See for more details.

# 0

#### Important

Do not use <u>backend\_host</u> variable in <u>cinder.conf</u> file. If <u>backend\_host</u> is set, it will override the [DEFAULT]/host value which HPE Helion OpenStack 5.0 is dependent on.

3. Commit your configuration to a local repository (Chapter 3, Using Git for Configuration Management):

```
cd ~/helion/hos/ansible
git add -A
git commit -m "configured VSA backend"
```



#### Note

Before you run any playbooks, remember that you need to export the encryption key in the following environment variable:

export HOS\_USER\_PASSWORD\_ENCRYPT\_KEY=ENCRYPTION\_KEY

For more information, see Chapter 10, Installing Mid-scale and Entry-scale KVM.

4. Run the configuration processor:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

5. Run the following command to create a deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

6. Run the Cinder Reconfigure Playbook:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts cinder-reconfigure.yml
```

### 17.2.7 Post-Installation Tasks

After you have configured VSA as your Block Storage backend, here are some tasks you will want to complete:

•

• Section 23.1, "Verifying Your Block Storage Backend"

## 17.3 Configuring for 3PAR Block Storage Backend

This page describes how to configure your 3PAR backend for the HPE Helion Entry-scale with KVM cloud model.

341 Post-Installation Tasks

## 17.3.1 Prerequisites

- You must have the license for the following software before you start your 3PAR backend configuration for the Helion Entry-scale with KVM cloud model:
  - Thin Provisioning
  - Virtual Copy
  - System Reporter
  - Dynamic Optimization
  - Priority Optimization
- Your HPE Helion Entry-scale KVM Cloud should be up and running. Installation steps can be found in Chapter 10, Installing Mid-scale and Entry-scale KVM.
- Your 3PAR Storage Array should be available in the cloud management network or routed to the cloud management network and the 3PAR FC and iSCSI ports configured.
- The 3PAR management IP and iSCSI port IPs must have connectivity from the controller and compute nodes.
- Please refer to the system requirements for 3PAR in the OpenStack configuration guide, which can be found here: 3PAR System Requirements (http://docs.openstack.org/liberty/config-reference/content/hp-3par-sys-reqs.html) ▶.

#### 17.3.2 Notes

**Encrypted 3Par Volume**: Attaching an encrypted 3Par volume is possible after installation by setting volume\_use\_multipath = true under the libvirt stanza in the nova/kvm-hypervisor.conf.j2 file and reconfigure nova.

Concerning using multiple backends: If you are using multiple backend options, ensure that you specify each of the backends you are using when configuring your <a href="cin-der.conf.j2">cin-der.conf.j2</a> file using a comma delimited list. An example would be <a href="enabled\_backend-backend-s=vsa-1,3par\_iSCSI,ceph1">enabled\_backend-backend-backend-backend-backend-backend-backend-backend</a> and is included in the steps below. You will also want to create multiple volume types so you can specify which backend you want to utilize when creating volumes. These instructions are included below as well. In addition to our documentation, you can also read the OpenStack documentation at Configure multiple storage backends (http://docs.openstack.org/admin-guide-cloud/blockstorage\_multi\_backend.html) as well.

342 Prerequisites

**Concerning iSCSI and Fiber Channel:** You should not configure cinder backends so that multipath volumes are exported over both iSCSI and Fiber Channel from a 3PAR backend to the same Nova compute server.

**3PAR driver has updated name:** In the OpenStack Mitaka release, the 3PAR driver used for HPE Helion integration had its name updated from HP3PARFCDriver and HP3PARISCSIDriver to HPE3PARFCDriver and HPE3PARISCSIDriver. To prevent issues when upgrading from previous HPE Helion OpenStack releases, we left the names as-is in the release and provided a mapping so that the integration would continue to work. This will produce a warning in the cinder-volume.log file advising you of the deprecated name. The warning will look similar to this:

```
Option "hp3par_api_url" from group "<your section>" is deprecated. Use option "hpe3par_api_url" from group "<your section>"'
```

These are just warnings and can be ignored.

## 17.3.3 Multipath Support

If you want to enable multipath for Cinder volumes carved from 3PAR FC/iSCSI storage system, please go through the ~/helion/hos/ansible/roles/multipath/README.md file on the lifecycle manager. The README.md file contains detailed procedures for configuring multipath for 3PAR FC/iSCSI Cinder volumes.

We have also included an additional set of steps needed if you are using 3PAR FC/iSCSI multipath which is included below.

If you are using 3PAR FC/iSCSI multipath, an additional configuration is required:



#### Note

If you are planning on attaching an encrypted 3Par volume after istallation, ensure that you <a href="volume\_use\_multipath">volume\_use\_multipath</a> = true under the libvirt section in the <a href="nova/kvm-hyper-visor.conf.j2">nova/kvm-hyper-visor.conf.j2</a> file before configuring Cinder.

343 Multipath Support

- 1. Log in to the lifecycle manager.
- 2. Edit the \_~/helion/my\_cloud/config/nova/kvm-hypervisor.conf.j2 file add this line under the [libvirt] section:

Example:

```
[libvirt]
...
iscsi_use_multipath=true
```

3. Edit the <u>~/helion/my\_cloud/config/cinder/cinder.conf.j2</u> file add this line under the [DEFAULT] section:

Example:

```
[DEFAULT]
...
use_multipath_for_image_xfer=true
```

4. Commit your configuration to the local git repo (*Chapter 3, Using Git for Configuration Management*), as follows:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```

5. Run the configuration processor:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

6. Use the playbook below to create a deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

7. Run the Nova reconfigure playbook:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts nova-reconfigure.yml
```

## 17.3.4 Configure 3PAR FC as a Cinder Backend

You must modify the cinder.conf.j2 to configure the FC details.

Perform the following steps to configure 3PAR FC as Cinder backend:

- 1. Log in to lifecycle manager.
- 2. Make the following changes to the <u>~/helion/my\_cloud/config/cinder/cinder.con-f.j2</u> file:
  - a. Add your 3PAR backend to the enabled\_backends section:

```
# Configure the enabled backends
enabled_backends=3par_FC
```



#### **Important**

If you are using multiple backend types, you can use a comma delimited list here. For example, if you are going to use both VSA and 3par backends, you would specify something like this: enabled\_backends=vsa-1,3par\_FC.

b. [OPTIONAL] If you want your volumes to use a default volume type, then enter the name of the volume type in the [DEFAULT] section with the syntax below. You will want to remember this value when you create your volume type in the next section.



### **Important**

If you do not specify a default type then your volumes will default to a non-redundant RAID configuration. It is recommended that you create a volume type and specify it here that meets your environments needs.

```
[DEFAULT]
# Set the default volume type
default_volume_type = <your new volume type>
```

c. Uncomment the <u>StoreServ (3par) iscsi cluster</u> section and fill the values per your cluster information. Here is an example:

```
[3par_FC]
san_ip: <3par-san-ipaddr>
san_login: <3par-san-username>
san_password: <3par-san-password>
```

```
hp3par_username: <3par-username>
hp3par_password: <hp3par_password>
hp3par_api_url: https://<3par-san-ipaddr>:8080/api/v1
hp3par_cpg: <3par-cpg-name-1>[,<3par-cpg-name-2>, ...]
volume_backend_name: <3par-backend-name>
volume_driver: cinder.volume.drivers.san.hp.hp_3par_fc.HP3PARFCDriver
```

### **Important**

Do not use <u>backend\_host</u> variable in <u>cinder.conf</u> file. If <u>backend\_host</u> is set, it will override the [DEFAULT]/host value which HPE Helion OpenStack 5.0 is dependent on.

3. Commit your configuration to the local git repo (*Chapter 3, Using Git for Configuration Management*), as follows:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```

4. Run the configuration processor:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

5. Update your deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

6. Run the following playbook to complete the configuration:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts cinder-reconfigure.yml
```

## 17.3.5 Configure 3PAR iSCSI as Cinder backend

You must modify the cinder.conf.j2 to configure the iSCSI details.

Perform the following steps to configure 3PAR iSCSI as Cinder backend:

- 1. Log in to lifecycle manager.
- 2. Make the following changes to the <u>~/helion/my\_cloud/config/cinder/cinder.con-f.j2</u> file:
  - a. Add your 3PAR backend to the enabled backends section:

```
# Configure the enabled backends
enabled_backends=3par_iSCSI
```

b. Uncomment the <u>StoreServ (3par) iscsi cluster</u> section and fill the values per your cluster information. Here is an example:

```
[3par_iSCSI]

san_ip: <3par-san-ipaddr>
san_login: <3par-san-username>
san_password: <3par-san-password>
hp3par_username: <3par-username>
hp3par_password: <hp3par_password>
hp3par_api_url: https://<3par-san-ipaddr>:8080/api/v1
hp3par_cpg: <3par-cpg-name-1>[,<3par-cpg-name-2>, ...]
volume_backend_name: <3par-backend-name>
volume_driver: cinder.volume.drivers.san.hp.hp_3par_iscsi.HP3PARISCSIDriver
hp3par_iscsi_ips: <3par-ip-address-1>[,<3par-ip-address-2>,<3par-ip-address-3>, ...]
hp3par_iscsi_chap_enabled=true
```

## •

#### **Important**

Do not use <u>backend\_host</u> variable in <u>cinder.conf</u> file. If <u>backend\_host</u> is set, it will override the [DEFAULT]/host value which HPE Helion OpenStack 5.0 is dependent on.

3. Commit your configuration your local git repository:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "<commit message>"
```

4. Run the configuration processor:

```
cd ~/helion/hos/ansible
```

```
ansible-playbook -i hosts/localhost config-processor-run.yml
```

When you run the configuration processor you will be prompted for two passwords. Enter the first password to make the configuration processor encrypt its sensitive data, which consists of the random inter-service passwords that it generates and the Ansible group\_vars and host\_vars that it produces for subsequent deploy runs. You will need this key for subsequent Ansible deploy runs and subsequent configuration processor runs. If you wish to change an encryption password that you have already used when running the configuration processor then enter the new password at the second prompt, otherwise press <code>Enter</code>. For CI purposes you can specify the required passwords on the ansible command line. For example, the command below will disable encryption by the configuration processor

```
ansible-playbook -i hosts/localhost config-processor-run.yml -e encrypt="" -e
rekey=""
```

If you receive an error during either of these steps then there is an issue with one or more of your configuration files. We recommend that you verify that all of the information in each of your configuration files is correct for your environment and then commit those changes to git using the instructions above.

5. Run the following command to create a deployment directory.

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

6. Run the following command to complete the configuration:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts cinder-reconfigure.yml
```

#### 17.3.6 Post-Installation Tasks

After you have configured 3PAR as your Block Storage backend, here are some tasks you will want to complete:

•

Section 23.1, "Verifying Your Block Storage Backend"

348 Post-Installation Tasks

### 17.4 Ironic OneView Integration

HPE Helion OpenStack 5.0 supports integration of Ironic (Baremetal) service with HPE OneView using agent\_pxe\_oneview driver. Please refer to OpenStack Documentation (https://docs.openstack.org/developer/ironic/drivers/oneview.html) for more information.

#### 17.4.1 Prerequisites

- 1. Installed HPE Helion OpenStack 5.0 with entry-scale-ironic-flat-network or entry-scale-ironic-multi-tenancy model.
- 2. OneView 3.0 instance is running and connected to management network.
- 3. OneView configuration is set into <u>definition/data/ironic/ironic\_config.yml</u> (and <u>ironic-reconfigure.yml</u> playbook ran if needed). This should enable *agent\_pxe\_oneview* driver in ironic conductor.
- 4. Managed node(s) should support PXE booting in legacy BIOS mode.
- 5. Managed node(s) should have PXE boot NIC listed first. That is, embedded 1Gb NIC must be disabled (otherwise it always goes first).

#### 17.4.2 Integrating with OneView

1. On the lifecycle manager, open the file \_~/helion/my\_cloud/definition/data/iron-ic/ironic\_config.yml

```
~$ cd ~/helion
~/helion$ vi my_cloud/definition/data/ironic/ironic_config.yml
```

- 2. Modify the settings listed below:
  - a. enable oneview: should be set to "true" for OneView integration
  - b. <u>oneview\_manager\_url</u>: HTTPS endpoint of Oneview management interface, for example: https://10.0.0.10/
  - c. oneview username: OneView username, for example: Administrator

- d. <a href="mailto:oneview\_encrypted\_password">oneview\_encrypted\_password</a>: OneView password in encrypted or clear text form. Encrypted form is distinguished by presence of '@hos@' at the beginning of the string. Encrypted form can be created by running hosencrypt.py program. This program is shipped as part of HPE Helion OpenStack and can be found in <a href="mailto:~/helion/hos/ansible">~/helion/hos/ansible</a> directory on lifecycle manager.
- e. <u>oneview\_allow\_insecure\_connections</u>: should be set to "true" if OneView is using self-generated certificate.
- 3. Once you have saved your changes and exited the editor, add files, commit changes to local git repository, and run <u>config-processor-run.yml</u> and <u>ready-deployment.yml</u> playbooks, as described in *Chapter 3, Using Git for Configuration Management*.

```
~/helion$ git add my_cloud/definition/data/ironic/ironic_config.yml
~/helion$ cd hos/ansible
~/helion/hos/ansible$ ansible-playbook -i hosts/localhost config-processor-run.yml
...
~/helion/hos/ansible$ ansible-playbook -i hosts/localhost ready-deployment.yml
```

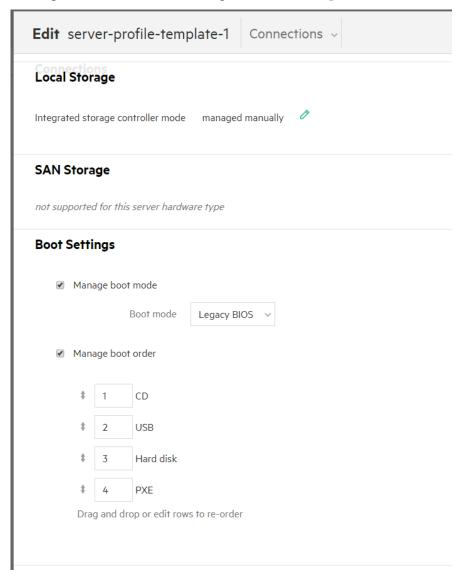
4. Run ironic-reconfigure.yml playbook.

350

#### 17.4.3 Registering Node in OneView

In the OneView web interface:

- 1. Navigate to Menu -> **Server Hardware**. Add new **Server Hardware** item, using managed node iLO IP and credentials. If this is the first node of this type being added, corresponding **Server Hardware Type** will be created automatically.
- 2. Navigate to Menu -> Server Profile Template. Add Server Profile Template. Use Server Hardware Type corresponding to node being registered. In BIOS Settings section, set Manage Boot Mode and Manage Boot Order options must be turned on:



3. Verify that node is powered off. Power the node off if needed.

#### 17.4.4 Provisioning Ironic Node

- 1. Login to the lifecycle manager and source respective credentials file (e.g. service.osrc for admin account).
- 2. Review glance images with glance image-list

Ironic deploy images ("ir-deploy-\*") were created automatically by HOS installer. We will be using agent\_pxe\_oneview Ironic driver which needs ir-deploy-kernel and ir-deploy-ramdisk images. Managed node will be deployed using Ubuntu Trusty 14.04 BIOS image.

3. Create node using agent pxe oneview driver.

```
$ ironic --ironic-api-version 1.22 node-create -d agent_pxe_oneview --name test-
node-1 --network-interface neutron -p memory_mb=131072 -p cpu_arch=x86_64 -p
local gb=80 -p cpus=2 \
   -p 'capabilities=boot_mode:bios,boot_option:local,server_hardware_type_uri:/
rest/server-hardware-types/E5366BF8-7CBF-48DF-
A752-8670CF780BB2, server_profile_template_uri:/rest/server-profile-
templates/00614918-77f8-4146-a8b8-9fc276cd6ab2' \
   -i 'server_hardware_uri=/rest/server-
hardware/32353537-3835-584D-5135-313930373046' \
   -i dynamic_allocation=True \
   -i deploy_kernel=633d379d-e076-47e6-b56d-582b5b977683 \
   -i deploy_ramdisk=d5828785-edf2-49fa-8de2-3ddb7f3270d5
+-----+
| Property | Value
  +----+
```

```
| chassis_uuid
| driver
                    | agent_pxe_oneview
| driver_info
                   | {u'server_hardware_uri': u'/rest/server-
                    | hardware/32353537-3835-584D-5135-313930373046',
u'dynamic_allocation':
                    | u'True', u'deploy_ramdisk': u'd5828785-
edf2-49fa-8de2-3ddb7f3270d5',
                    | u'deploy kernel': u'633d379d-e076-47e6-b56d-582b5b977683'}
                    | {}
| extra
| name
                    | test-node-1
| network_interface | neutron
                    | {u'memory_mb': 131072, u'cpu_arch': u'x86_64', u'local_gb':
| properties
80, u'cpus': |
                    | 2, u'capabilities':
                    | u'boot_mode:bios,boot_option:local,server_hardware_type_uri:/
rest
                    /server-hardware-types/E5366BF8-7CBF-
                    | 48DF-A752-8670CF780BB2,server_profile_template_uri:/rest/
server-profile- |
                    | templates/00614918-77f8-4146-a8b8-9fc276cd6ab2'}
| resource_class
                    | None
                    | c202309c-97e2-4c90-8ae3-d4c95afdaf06
| uuid
```

## Note

- For deployments created via Ironic/OneView integration, <a href="memory\_mb">memory\_mb</a> property must reflect physical amount of RAM installed in managed node. That is, for a server with 128 Gb of RAM it works out to 132\*1024=13072.
- Boot mode in capabilities property must reflect boot mode used by the server,
   i.e. 'bios' for Legacy BIOS and 'uefi' for UEFI.
- Values for <a href="server\_hardware\_type\_uri">server\_profile\_template\_uri</a> and <a href="server\_hardware\_uri">server\_hardware\_uri</a> can be grabbed from browser URL field while navigating to respective objects in OneView UI. URI corresponds to the part of URL which starts form the token '/rest'. That is, URL https://oneview.mycorp.net/#/profile-templates/show/overview/r/rest/server-profile-templates/12345678-90ab-cdef-0123-012345678901 corresponds to URI /rest/server-profile-templates/12345678-90ab-cdef-0123-012345678901.
- Grab IDs of deploy\_kernel and deploy\_ramdisk from glance image-list output above.

#### 4. Create port.

```
$ ironic --ironic-api-version 1.22 port-create \
  --address aa:bb:cc:dd:ee:ff \
  --node c202309c-97e2-4c90-8ae3-d4c95afdaf06 \
  -l switch_id=ff:ee:dd:cc:bb:aa \
  -l switch_info=MY_SWITCH \
  -l port id="Ten-GigabitEthernet 1/0/1" \
  --pxe-enabled true
+----
+-----
| Property
                | Value
   1
+-----
+-----+
| address
               | 8c:dc:d4:b5:7d:1c
   | extra
                | {}
```



#### Warning

Ironic Multi-Tenancy networking model is used in this example. Therefore, ironic port-create command contains information about the physical switch. OneView integration can also be performed using the Ironic Flat Networking model. For more information, see .

5. Move node to manageable provisioning state. The connectivity between Ironic and OneView will be verified, Server Hardware Template settings validated, and Server Hardware power status retrieved from OneView and set into the Ironic node.

```
$ ironic node-set-provision-state test-node-1 manage
```

6. Verify that node power status is populated.

```
| driver
                         | agent_pxe_oneview
| driver_info
                         | {u'server_hardware_uri': u'/rest/server-
                         | hardware/32353537-3835-584D-5135-313930373046',
u'dynamic_allocation':
                         | u'True', u'deploy_ramdisk': u'd5828785-
edf2-49fa-8de2-3ddb7f3270d5',
                         | u'deploy_kernel': u'633d379d-e076-47e6-
b56d-582b5b977683'}
| driver_internal_info
                         | {}
| extra
                         | {}
| inspection_finished_at | None
| inspection_started_at | None
| instance_info
                         | {}
| instance_uuid
                         | None
| last_error
                         None
| maintenance
                         | False
| maintenance_reason
                         None
                          | test-node-1
| name
| network_interface
                         | power off
| power_state
                         | {u'memory_mb': 131072, u'cpu_arch': u'x86_64',
| properties
u'local_gb': 80, u'cpus': |
                         | 2, u'capabilities':
u'boot_mode:bios,boot_option:local,server_hardware_type_uri:/rest
                         | /server-hardware-types/E5366BF8-7CBF-
                         | 48DF-A752-8670CF780BB2,server_profile_template_uri:/rest/
server-profile- |
                         | templates/00614918-77f8-4146-a8b8-9fc276cd6ab2'}
```

7. Move node to available provisioning state. The Ironic node will be reported to Nova as available.

```
$ ironic node-set-provision-state test-node-1 provide
```

8. Verify that node resources were added to Nova hypervisor stats.

+----+

#### 9. Create Nova flavor.



#### Note

All parameters (specifically, amount of RAM and boot mode) must correspond to ironic node parameters.

10. Create Nova keypair if needed.

```
$ nova keypair-add ironic_kp --pub-key ~/.ssh/id_rsa.pub
```

11. Boot Nova instance.

```
| OS-EXT-SRV-ATTR:hypervisor_hostname | -
| OS-EXT-SRV-ATTR:instance_name
| OS-EXT-STS:power_state
                                      | 0
| OS-EXT-STS:task_state
                                     | scheduling
| OS-EXT-STS:vm_state
                                       | building
| OS-SRV-USG:launched_at
| OS-SRV-USG:terminated_at
| accessIPv4
| accessIPv6
| adminPass
                                       | pE3m7wRACvYy
| config_drive
                                       | 2017-06-30T21:08:42Z
| created
                                       | ml.ironic (33c81884-b8aa-46b7-
| flavor
ab3f-076f3b72f8d8)
| hostId
| id
                                       | b47c9f2a-e88e-411a-abcd-6172aea45397
                                       | Ubuntu Trusty 14.04 BIOS
| image
(d6b5d971-42fd-492b-b33e-6e5f2d88e942) |
                                       | ironic_kp
| key_name
| metadata
                                       | {}
                                       | test-node-1
| name
| os-extended-volumes:volumes_attached | []
                                       | 0
| progress
| security_groups
                                       | default
                                       | BUILD
| status
```

tenant_id	c8573f7026d24093b40c769ca238fddc
updated	2017-06-30T21:08:42Z
user_id	2eae99221545466d8f175eeb566cc1b4
+	
+	+

During nova instance boot, the following operations will be performed by Ironic via OneView REST API.

- In OneView, new Server Profile is generated for specified Server Hardware, using specified Server Profile Template. Boot order in Server Profile is set to list PXE as the first boot source.
- Managed node is powered on and boots IPA image from PXE.
- IPA image writes user image onto disk and reports success back to Ironic.
- Ironic modifies Server Profile in OneView to list 'Disk' as default boot option.
- Ironic reboots the node (via OneView REST API call).

## 18 Troubleshooting the Installation

We have gathered some of the common issues that occur during installation and organized them by when they occur during the installation. These sections will coincide with the steps labeled in the installation instructions.

- Section 18.1, "Issues during Lifecycle-manager Setup"
- Section 18.2, "Issues while Provisioning your Baremetal Nodes"
- Section 18.3, "Issues while Updating Configuration Files"
- Section 18.4, "Issues while Deploying the Cloud"
- Chapter 19, Troubleshooting the Block Storage Backend Configuration

## 18.1 Issues during Lifecycle-manager Setup

## Issue: Running the hos-init.bash script when configuring your lifecycle manager does not complete

Part of what the ~/hos-3.0.0/hos-init.bash script does is install git and so if your DNS nameserver(s) is not specified in your /etc/resolv.conf file, is not valid, or is not functioning properly on your lifecycle manager then it won't be able to complete.

To resolve this issue, double check your nameserver in your <a href="//etc/resolv.conf">/etc/resolv.conf</a> file and then re-run the script.

## 18.2 Issues while Provisioning your Baremetal Nodes

#### Issue: The bm-power-status.yml playbook returns an error.

If the output of the bm-power-status.yml playbook gives you an error like the one below for one of your servers then the most likely cause is an issue with the information located in the  $\sim$ / helion/my\_cloud/definition/servers.yml file. It could be that the information is incorrect or the iLO username you are using does not have administrator privileges. Verify all of the information for that server and recommit the changes to git and re-run the playbook if the information was incorrect. If it's a rights issue, ensure the iLO user has administrative privileges.

#### Error:

```
failed: [vsa1] => {"cmd": "ipmitool -I lanplus -E -N 5 -R 12 -U iLOAdmin -H 10.12.11.185
power status", "failed": true, "rc": 1}
stderr: Error: Unable to establish IPMI v2 / RMCP+ session
```

This error below, which may come with running the <u>bm-power-status.yml</u> playbook, may be resolved by placing the <u>ilo-password</u> values in your <u>servers.yml</u> file in double quotes.

Error:

```
failed: [controller2] => {"failed": true}
msg: ipmi: 'int' object has no attribute 'startswith'
```

#### Cobbler Deployment Fails due to CIDR Error

The error will look similar to this:

The most likely cause of this is that your netmask entry in the ~/helion/my\_cloud/definition/data/servers.yml file is incorrect.

To resolve this issue ensure that you have the correct <u>subnet</u> and <u>netmask</u> entries in the <u>~/</u> helion/my\_cloud/definition/data/servers.yml for your Management network.

You can also verify these values exist in the <a href="//etc/network/interfaces.d/ethX">/etc/network/interfaces.d/ethX</a> file on your lifecycle manager. These should match the entries in the servers.yml file.

Once you make your corrections to your configuration files, commit the changes with these steps:

1. Ensuring that you stay within the ~/helion directory, commit the changes you just made:

```
cd ~/helion
git commit -a -m "commit message"
```

2. Redeploy Cobbler:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost cobbler-deploy.yml
```

Issue: Configuration changes needed after Cobbler deploy

If you've made a mistake or wish to change your <u>~/helion/my\_cloud/definition/da-ta/servers.yml</u> configuration file after you've already run the <u>cobbler-deploy.yml</u> playbook, follow these steps to ensure Cobbler gets updated with the new server information:

- 1. Ensure your servers.yml file is updated
- 2. Ensuring that you stay within the ~/helion directory, commit the changes you just made:

```
cd ~/helion
git commit -a -m "commit message"
```

3. Determine which nodes you have entered into Cobbler by using:

```
sudo cobbler system list
```

4. Remove the nodes that had the old information from Cobbler using:

```
sudo cobbler system remove --name <nodename>
```

5. Re-run the cobbler-deploy.yml playbook to update the new node definitions to Cobbler:

```
cd ~/helion/hos/ansible ansible-playbook -i hosts/localhost cobbler-deploy.yml
```



#### Note

If you've also already run the <u>bm-reimage.yml</u> playbook then read the section below for how to ensure your nodes get re-imaged when re-running this playbook.

#### Issue: bm-reimage.yml playbook doesn't find any nodes to image

You may receive this error when running the bm-reimage.yml playbook:

This behavior occurs when you don't specify the <u>-e nodelist</u> switch to your command and all of your nodes are marked as <u>netboot-enabled: false</u> in Cobbler. By default, without the <u>-e nodelist</u> switch, the <u>bm-reimage.yml</u> playbook will only reimage the nodes marked as <u>netboot-enabled: True</u>, which you can verify which nodes you have that are marked as such with this command:

```
sudo cobbler system find --netboot-enabled=1
```

If this is on a fresh install and none of your nodes have been imaged with the HPE Helion OpenStack ISO, then you should be able to remove all of your nodes from Cobbler and re-run the cobbler-deploy.yml playbook. You can do so with these commands:

1. Get a list of your nodes in Cobbler:

```
sudo cobbler system list
```

2. Remove each of them with this command:

```
sudo cobbler system remove --name SYSTEM_NAME
```

3. Confirm they are all removed in Cobbler:

```
sudo cobbler system list
```

4. Re-run cobbler-deploy.yml:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost cobbler-deploy.yml
```

5. Confirm they are all now set to netboot-enabled: True:

```
sudo cobbler system find --netboot-enabled=1
```

#### Issue: The --limit switch does not do anything

If the <u>cobbler-deploy.yml</u> playbook fails, it makes a mention that to retry you should use the --limit switch, as seen below:

```
to retry, use: --limit @/home/headmin/cobbler-deploy.retry
```

This is a standard Ansible message but it is not needed in this context. Please use our instructions described above to remove your systems, if necessary, from Cobbler before re-running the playbook as the <u>--limit</u> is not needed. Note that using the <u>--limit</u> switch does not cause any harm and won't prevent the playbook from completing, it's just not a necessary step.

#### **Issue: Dealing With Nodes that Fail to Install**

The <u>bm-reimage.yml</u> playbook will take every node as far as it can through the baremetal install process. Nodes that fail will not prevent the others from continuing to completion. At the end of the run you will get a list of the nodes (if any) that failed to install.

If you run <u>bm-reimage.yml</u> a second time, by default it will target only the failed nodes the second time round (because the others are already marked as "completed"). Alternatively you can target specific nodes for reimage using -e nodelist as described in the section below.

The places where you are most likely to see a node fail is timeout in the "wait for shutdown" step, which means that the node did not successfully install an operating system (e.g. it could be stuck in POST) or a timeout in "wait for ssh" at the end of the baremetal install. This means that the node did not come back up after being powered on. To fix the issues you'll need to connect to the nodes' consoles and investigate.

#### How to re-image an existing node

Once Linux for HPE Helion has been successfully installed on a node, that node will be marked as "installed" and subsequent runs of <a href="mage.yml">bm-reimage.yml</a> (and related playbooks) will not target it. This is deliberate so that you can't reimage the node by accident. If you do need to reimage existing nodes you will need to use the <a href="mage-e-e-nodelist">-e nodelist</a> option to target them specifically. For example:

ansible-playbook -i hosts/localhost bm-reimage.yml -e nodelist=cpn-0044,cpn-0045



Note

You can target all nodes with -e nodelist=all

This will power cycle the specified nodes and reinstall the operating system on them, using the existing settings stored in Cobbler.

If you want to change settings for a node in the configuration files, see Section 18.2, "Issues while Provisioning your Baremetal Nodes".

#### Issue: Wait for SSH phase in bm-reimage.yml hangs

This issue has been observed during deployment to systems configured with QLogic based BCM578XX network adapters utilizing the bnx2x driver and is currently under investigation. The symptom manifests following a cold boot during deployment at the "wait for SSH" phase in bm-reimage.yml which will result in a hang and eventually timeout, causing the baremetal

install to fail. The presence of this particular issue can be further confirmed by connecting to the remote console of the server via iLO or by checking the server's dmesg output for the presence of bnx2 panic dump messages, similar to the following:

This workaround is completed by rebooting the server.

#### Subsequent runs of bm-reimage.yml sometimes fail to trigger network install

If you've run the <u>site.yml</u> playbook during an install and have a reason to run the <u>bm-reimage.yml</u> playbook afterwards, it may fail to trigger the network install over PXE due to the name of the interface changing during this process.

To resolve this issue, you can run the <u>cobbler-deploy.yml</u> playbook to refresh these settings and then proceed to run bm-reimage.yml.

#### Issue: Soft lockup at imaging

When imaging your nodes, if you see a kernel error of the type "BUG: soft lockup - CPU#10 stuck...." you should reset the node and make sure it is imaged properly next time. Note that depending on when you get the error, you may have to rerun <a href="cobbler-deploy.yml">cobbler-deploy.yml</a>. If the bmreimage playbook says it failed to image the node, then cobbler knows this has occurred and you can reset the node. If not, you can follow the instructions above for .

#### Blank Screen Seen When Monitoring the Imaging Step

If you are watching the <u>os-install</u> process on a node via the console output, there can be a pause between 2-3 minutes in length where nothing gets reported to the console screen. This is just after the grub menu has been displayed on a UEFI-based system.

This is normal and nothing to be concerned with. The imaging process will continue on after this pause.

## Unable to SSH to a Compute Node after the bm-reimage Step Even Though SSH Keys Are In Place

If after running the <a href="mage.yml">bm-reimage.yml</a> playbook you get a reported failure when attempting to SSH to one or more Compute nodes stating that permission was denied and you've confirmed that the SSH keys are correctly coped from the lifecycle manager to the Compute node then you can follow the steps below to resolve this issue.

The suspected root cause of this issue is that the bm-reimage playbook is finding a different MAC address when attempting to SSH to the Compute node than was specified in the <a href="mailto:servers.yml">servers.yml</a> file. This could be because you entered an incorrect MAC address or you may have specified the same IP address to two different Compute nodes.

To resolve this issue you can follow these steps:

1. Remove the Compute node that failed from Cobbler using this command:

```
sudo cobbler system remove --name <node name>
```



#### Note

Use sudo cobbler system list to get a list of your nodes in Cobbler.

- 2. Correct the IP and MAC address information in your <u>~/helion/my\_cloud/definition/data/servers.yml</u> file.
- 3. Commit your changes to git:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```

4. Re-run the bm-reimage.yml playbook and confirm the error is not received again.

#### Server Crashes During Imaging

If your server suffers a kernel panic just after it gets PXE booted, a possible cause could be that you have an incorrect bus address for the server specified in the nic mappings.yml file.

To resolve this issue you will need to obtain the proper bus address. You can do this by installing HP Linux for OpenStack on the node and once you receive a terminal prompt you can use the command below to obtain the proper bus address:

```
sudo lspci -D | grep -i eth
```

Once you have the proper bus address, follow these steps to finish the reimage:

1. Remove the node that failed from Cobbler using this command:

```
sudo cobbler system remove --name <node name>
```



Use sudo cobbler system list to get a list of your nodes in Cobbler.

- 2. Correct the bus address in your ~/helion/my\_cloud/definition/data/nic\_map-pings.yml file.
- 3. Commit your changes to git:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "My config or other commit message"
```

4. Re-run the bm-reimage.yml playbook and confirm the error is not received again.

## 18.3 Issues while Updating Configuration Files

#### Configuration Processor Fails Due to Wrong yml Format

If you receive the error below when running the configuration processor then you may have a formatting error:

```
TASK: [fail msg="Configuration processor run failed, see log output above for details"]
```

First you should check the ansible log in the location below for more details on which yml file in your input model has the error:

```
~/.ansible/ansible.log
```

Check the configuration file to locate and fix the error, keeping in mind the following tips below. Check your files to ensure that they don't contain the following:

- Non-ascii characters
- Unneeded spaces

Once you have fixed the formatting error in your files, commit the changes with these steps:

1. Commit your changes to git:

```
cd ~/helion/hos/ansible
```

```
git add -A
git commit -m "My config or other commit message"
```

2. Re-run the configuration processor playbook and confirm the error is not received again.

#### Configuration processor fails with provider network OCTAVIA-MGMT-NET error

If you receive the error below when running the configuration processor then you have not correctly configured your VLAN settings for Octavia.

To resolve the issue, ensure that your settings in \_-/helion/my\_cloud/definition/data/neutron/neutron\_config.yml are correct for the VLAN setup for Octavia.

#### Changes Made to your Configuration Files

If you have made corrections to your configuration files and need to re-run the Configuration Processor, the only thing you need to do is commit your changes to your local git:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "commit message"
```

You can then re-run the configuration processor:

```
cd ~/helion/hos/ansible ansible-playbook -i hosts/localhost config-processor-run.yml
```

#### Configuration Processor Fails Because Encryption Key Does Not Meet Requirements

If you choose to set an encryption password when running the configuration processor, you may receive the following error if the chosen password does not meet the complexity requirements:

If you receive the above error, simply run the configuration processor again and select a password that meets the complexity requirements detailed in the error message:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

## 18.4 Issues while Deploying the Cloud

#### Issue: If the site.yml playbook fails, you can query the log for the reason

Ansible is good about outputting the errors into the command line output, however if you'd like to view the full log for any reason the location is:

```
~/.ansible/ansible.log
```

This log is updated real time as you run ansible playbooks.



Tip

Use grep to parse through the log. Usage: grep <text> ~/.ansible/ansible.log

#### Issue: How to Wipe the Disks of your Machines

If you have re-run the <u>site.yml</u> playbook, you may need to wipe the disks of your nodes

You would generally run the playbook below after re-running the <u>bm-reimage.yml</u> playbook but before you re-run the site.yml playbook.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts wipe_disks.yml
```

The playbook will show you the disks to be wiped in the output and allow you to confirm that you want to complete this action or abort it if you do not want to proceed. You can optionally use the --limit <NODE\_NAME> switch on this playbook to restrict it to specific nodes.

If you receive an error stating that <u>osconfig</u> has already run on your nodes then you will need to remove the <u>/etc/hos/osconfig-ran</u> file on each of the nodes you want to wipe with this command:

sudo rm /etc/hos/osconfig-ran

That will clear this flag and allow the disk to be wiped.

#### Issue: Errors during create\_db Task

If you are doing a fresh installation on top of a previously used system and you have not completely wiped your disk prior to the installation, you may run into issues when the installation attempts to create new Vertica databases. The recommendation here is to wipe your disks and re-start the installation, however we have included a playbook for cleaning your Vertica databases if they need to be re-used.

To run this playbook, follow these steps:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts monasca-vertica-dbclean.yml
```

#### Issue: Freezer installation fails if an independent network is used for the External API.

Currently the Freezer installation fails if an independent network is used for the External\_API. If you intend to deploy the External API on an independent network, the following changes need to be made:

In roles/freezer-agent/defaults/main.yml add the following line:

```
backup_freezer_api_url: "{{ FRE_API | item('advertises.vips.private[0].url', default='
    ') }}"
```

In roles/freezer-agent/templates/backup.osrc.j2 add the following line:

```
export OS_FREEZER_URL={{ backup_freezer_api_url }}
```

#### Error Received if Root Logical Volume is Too Small

When running the <u>site.yml</u> playbook, you may receive the error below if your root logical-volume is too small:

```
2015-09-29 15:54:03,022 p=26345 u=stack | stderr: New size given (7128 extents) not larger than existing size (7629 extents)
```

The specific part of this error to parse out and resolve is:

```
stderr: New size given (7128 extents) not larger than existing size (7629 extents)
```

The error also references the root volume:

```
"name": "root", "size": "10%"
```

The problem is that the root logical-volume, as specified in the <a href="disks\_controller.yml">disks\_controller.yml</a> file, is set to 10% of the overall physical volume and this value is too small.

To resolve this issue you need to ensure that the percentage is set properly for the size of your logical-volume. The default values in the configuration files is based on a 500GB disk, so if your logical-volumes are smaller you may need to increase the percentage so there is enough room.

#### Multiple Keystone Failures Received during Site.yml

If you receive the Keystone error below during your site.yml run then follow these steps:

```
TASK: [OPS-MON | _keystone_conf | Create Ops Console service in Keystone] *****

failed: [helion-cp1-c1-m1-mgmt] => {"failed": true}

msg: An unexpected error prevented the server from fulfilling your request. (HTTP 500)

(Request-ID: req-23a09c72-5991-4685-b09f-df242028d742), failed

FATAL: all hosts have already failed -- aborting
```

The most likely cause of this error is that the virtual IP address is having issues and the Keystone API communication through the virtual IP address is not working properly. You will want to check the Keystone log on the controller where you will likely see authorization failure errors.

Verify that your virtual IP address is active and listening on the proper port on all of your controllers using this command:

```
netstat -tplan | grep 35357
```

Ensure that your lifecycle manager did not pick the wrong (unusable) IP address from the list of IP addresses assigned to your Management network.

The lifecycle manager will take the first available IP address after the <a href="mailto:gateway-ip">gateway-ip</a> defined in your <a href="mailto:relation/my\_cloud/definition/data/networks.yml">relation/my\_cloud/definition/data/networks.yml</a> file. This IP will be used as the virtual IP address for that particular network. If this IP address is used and reserved for another purpose outside of your HPE Helion OpenStack deployment then you will receive the error above.

To resolve this issue we recommend that you utilize the <u>start-address</u> and possibly the <u>end-address</u> (if needed) options in your <u>networks.yml</u> file to further define which IP addresses you want your cloud deployment to use. See for more details.

After you have made changes to your <a href="networks.yml">networks.yml</a> file, follow these steps to commit the changes:

1. Ensuring that you stay within the ~/helion directory, commit the changes you just made:

```
cd ~/helion
git commit -a -m "commit message"
```

2. Run the configuration processor:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

3. Update your deployment directory:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost ready-deployment.yml
```

4. Re-run the site.yml playbook:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts site.yml
```

#### VSA installer failed during deployment

VSA deployment might fails sporadically due to a defect in the HPE StoreVirtual VSA installation files for the KVM host. The following error messages is displayed and VSA installation will hang for more than 45-60 minutes:



#### Note

VSA deployment can take approximately usually 15-30 minutes, depending on how many disks are configured. Do not terminate the installation until 60 minutes have passed and/ or you see the error.

#### To correct this issue:

- 1. Log into the failed VSA node.
- 2. Change to the root user:

```
sudo -i
```

3. Destroy the VSA VM:

```
virsh destroy <VSA_VM_Name>
virsh undefine <VSA_VM_Name>
```

#### For example:

```
virsh destroy VSA-VM-vsa2
virsh undefine VSA-VM-vsa2
```

To get the name of the VSA VM, execute the following command:

```
virsh list --all

Id Name State

2 VSA-VM-vsa2 running
```

4. Destroy the VSA network:

```
virsh net-destroy vsa-network
virsh net-undefine vsa-network
```

5. Destroy the VSA storage-pool:

```
virsh pool-destroy vsa-storage-pool
virsh pool-undefine vsa-storage-pool
```

6. Remove the vsa-installer directory:

```
rm -rf /home/vsa-installer
```

7. Remove the VSA image:

rm -rf /mnt/state/vsa-kvm-storage

8. Reboot the VSA VM:

reboot

- 9. Log into your lifecycle manager.
- 10. Run the deployment playbooks from the resulting scratch directory:

cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb\_hosts site.yml

## 19 Troubleshooting the Block Storage Backend Configuration

We have gathered some of the common issues that occur during the Block Storage configuration steps and organized them by product.

#### Issue: You have forgotten the identifying information about your VSA cluster

If you have forgotten the cluster name or IP address after creating the cluster in the CMC utility then you can perform the following steps to retrieve them:

- 1. Log in to CMC console using the <a href="https://opt/HP/StoreVirtual/UI/jre/bin/java -jar /opt/HP/StoreVirtual/UI/JI command from the lifecycle manager/deployer node.">https://opt/HP/StoreVirtual/UI/JI command from the lifecycle manager/deployer node.</a>
- 2. Log in to the Management group.
- 3. On the left side of the CMC console screen, click **mgmt**. It expands and you will be able to see the cluster name.
- 4. Click the cluster name. The page on the right hand side will populate with the cluster information.
- 5. Click <u>iSCSI</u> on the right hand side of the page. The virtual IP address will be displayed in the tabular form. You can view your cluster ID here.

#### Issue: Error During VSA Deployment Exhibits When Running site.yml or hlm-deploy.yml

When you are running either <u>site.yml</u> or <u>hlm-deploy.yml</u> and you get either of the errors below during the VSA deployment task, there is an error in your locale settings.

#### Errors:

```
File "/opt/stack/venv/storevirtual-installer-20151022T080854Z/lib/python2.7/site-packages/storevirtual_installer/deployer.py", line 479, in _roll_back_installation self._vsa_network_destroy()
File "/opt/stack/venv/storevirtual-installer-20151022T080854Z/lib/python2.7/site-packages/storevirtual_installer/deployer.py", line 493, in _vsa_network_destroy raise vsa_excs.NetworkDestroyFailed(msg)
storevirtual_installer.vsa_excs.NetworkDestroyFailed: Failed to destroy and undefine the network

FATAL: all hosts have already failed -- aborting
```

#### OR

#### Fix:

Make sure that the locale variables on the lifecycle manager are valid and then run the playbook again. You can set this variable to set your locale:

```
export LC_ALL=C
```

## 20 Troubleshooting the ESX

This section contains troubleshooting tasks for your HPE Helion OpenStack® 5.0 for ESX.

# 20.1 Issue: If hlm-ux-services.service is not running, the EON resource-activate and resource-deactivate commands fails

If you perform any maintenance work or reboot the lifecycle manager/deployer node, make sure to restart the HLM UX service for standalone deployer node and shared HLM/controller node based on your environment.

For standalone deployer node, execute <a href="https://execute.com/htm-start.yml">https://execute.com/htm-start.yml</a> playbook to restart the HLM UX services on the deployer node after a reboot.

For shared deployer/controller node, execute <a href="https://htm.start.yml">https://htm.start.yml</a> playbook on all the controllers to restart HLM UX services.

#### For example:

cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb\_hosts hlm-start.yml --limit <host name of the HLM node or
HLM Node/Shared Controller>

## 20.2 Issue: ESX onboard Compute

- 1. If there is a failure in provisioning phase
  - a. The spawned applicances will be deleted.
  - b. Check for <a href="//var/log/eon/eon-conductor.log">/var/log/eon/eon-conductor.log</a> for specific errors to troubleshoot the error.
- 2. If there is a failures in activating phase
  - a. The spawned appliances will be left intact, and you can try executing activate command again.
    - b. Check the  $\sim$ /.ansible/ansible.log or  $\sim$ /var/log/hlm-ux-service/service.log in the hlm-deployer.

Issue: If hlm-ux-services.service is not running, the EON resource-activate and resource-deacti-

378 vate commands fails

In case of SSH failure, check the background connectivity and try to ping again.

3. Failures leading to inconsistent Database (DB) state, execute the following command:

```
eon resource-deactivate <id> --forced True
```

## 20.3 Issue: Migrate eon-conductor

Currently eon-conductor service is running in the first node of the control plane.

#### Procedure to migrate eon-conductor

Perform the following steps to migrate eon-conductor service to any other node in the control plane.

- 1. SSH to lifecycle manager.
- 2. Change the directory.

```
cd /home/user/scratch/ansible/next/hos/ansible
```

3. Search where the eon conductor is running currently. In the following example, eon-conductor is running in helion-cpl-cl-ml-mgmt.

```
(helion-cp1-c1-m1-mgmt)
stack@helion-cp1-c1-m1-mgmt:~$ ps aux | grep eon
        11236 5.5 0.0 142104 52088 ?
                                               Ss
                                                    18:53
                                                            0:01 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
        11281 23.2 0.1 401964 238080 ?
                                                    18:53
                                                            0:06 /opt/stack/venv/
                                               Ss
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-conductor/venv/bin/eon-
conductor --config-file=/opt/stack/service/eon-conductor-20160503T082438Z/etc/eon/
eon-conductor.conf
         11590 1.5 0.0 144420 49084 ?
                                               S
                                                    18:53
                                                            0:00 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
        11591 1.4 0.0 144416 49056 ?
                                               S
                                                    18:53
                                                            0:00 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
         11592 1.1 0.0 144420 48828 ?
                                                    18:53
                                                            0:00 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
stack
         12086 0.0 0.0 12736 2064 pts/9
                                                    18:54
                                                            0:00 grep eon
```

```
(helion-cp1-c1-m2-mgmt)
```

```
stack@helion-cp1-c1-m2-mgmt:~$ ps aux | grep eon
         7050 1.9 0.0 142100 52052 ?
                                              Ss
                                                   18:53
                                                           0:03 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
         7677 1.1 0.0 144552 49156 ?
                                              S
                                                   18:53
                                                           0:02 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
         7678 1.2 0.0 144544 49256 ?
                                                   18:53
                                              S
                                                           0:02 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
          7679 1.2 0.0 144548 49136 ?
                                              S
                                                   18:53
                                                           0:02 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
        13523 0.0 0.0 12732 2076 pts/0
                                                   18:56
                                              S+
                                                           0:00 grep eon
```

#### 4. Stop the eon service.

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts eon-stop.yml
```

#### Example:

```
stack@helion-cp1-c1-m1-mgmt:~/scratch/ansible/next/hos/ansible$ ansible-playbook -i
hosts/verb_hosts eon-stop.yml
changed: [helion-cp1-c1-m3-mgmt]
changed: [helion-cp1-c1-m1-mgmt]
changed: [helion-cp1-c1-m2-mgmt]
changed: [helion-cp1-c1-m1-mgmt]
skipping: no hosts matched
eon-common | stop | stop eon-conductor service ------ 8.39s
eon-common | stop | stop eon-api service ------ 0.42s
                              ----- 8.99s
helion-cp1-c1-m1-mgmt
             : ok=2
                   changed=2
                         unreachable=0
                                 failed=0
helion-cp1-c1-m2-mgmt : ok=1 changed=1 unreachable=0 failed=0
```

```
helion-cp1-c1-m3-mgmt : ok=1 changed=1 unreachable=0 failed=0
```

5. Edit the /home/user/scratch/ansible/next/hos/ansible/hosts/verb\_hosts. In the following example, we change the node name from <a href="https://hosts/hosts/verb\_hosts">hetohosts/verb\_hosts</a>. In the following example, we change the node name from <a href="https://hosts/hosts/verb\_hosts">hetohosts/hosts/verb\_hosts</a>. In the following example, we change the node name from <a href="https://hosts/

```
[EON-CND--first-member:children] (ORIGINAL)
helion-cp1-c1-m1-mgmt <- original node that eon-conductor is setup on.

[EON-CND--first-member:children] (MODIFIED)
helion-cp1-c1-m2-mgmt <- new node that eon-conductor is setup on.</pre>
```

6. Execute the following playbook in sequence.

```
ansible-playbook -i hosts/verb_hosts eon-deploy.yml
ansible-playbook -i hosts/verb_hosts eon-start.yml
```

#### Example:

```
changed: [helion-cp1-c1-m2-mgmt]
eon-api | start | start eon-api service ------ 0.26s
eon-conductor | start | start eon-conductor service ----- 0.23s
eon-api | start | activate the latest installed version ----- 0.07s
eon-api | start | restart eon-api service ------ 0.07s
eon-conductor | start | restart eon-conductor service ------ 0.02s
eon-conductor | start | activate the latest installed version ------ 0.02s
     Total: ------ 1.09s
helion-cp1-c1-m1-mgmt
                : ok=1
                       changed=0
                               unreachable=0
                                         failed=0
helion-cp1-c1-m2-mgmt
                               unreachable=0
                                         failed=0 <-
                : ok=2
                       changed=0
Service is moved and started
helion-cp1-c1-m3-mgmt
                : ok=1
                       changed=0
                               unreachable=0
                                         failed=0
```

7. Ensure that the conductor is running on the node that you have migrated. In the following example observe that the eon-conductor have migrated to node helion-cpl-cl-m2-mgmt.

```
(helion-cp1-c1-m1-mgmt)
stack@helion-cp1-c1-m1-mgmt:~$ ps aux | grep eon
stack 2624 0.0 0.0 12732 2132 pts/9 S+ 18:46 0:00 grep eon
```

```
27786 1.3 0.0 142112 52064 ?
                                              Ss
                                                   16:21 1:55 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
         27816 1.2 0.0 149712 55060 ?
                                                    16:21
                                                           1:50 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
         27817 1.2 0.0 150380 55724 ?
                                                   16:21
                                              ς
                                                           1:49 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
         27818 1.2 0.0 148516 53844 ?
                                              S
                                                   16:21
                                                           1:49 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
(helion-cp1-c1-m2-mgmt)
stack@helion-cp1-c1-m2-mgmt:~$ ps aux | grep eon
         1567 1.3 0.0 142108 52076 ?
                                                   16:21
                                              Ss
                                                           1:56 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
         1606 1.2 0.0 144816 49888 ?
                                                   16:21
                                                           1:51 /opt/stack/venv/
                                              S
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
         1607 1.2 0.0 145068 49888 ?
                                              S
                                                   16:21
                                                           1:50 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
          1608 1.2 0.0 145064 49888 ?
                                              S
                                                    16:21
                                                           1:49 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-api/venv/bin/eon-api --
config-file=/opt/stack/service/eon-api-20160503T082438Z/etc/eon/eon-api.conf
         22720 0.0 0.1 404872 241220 ?
                                              Ss
                                                   16:15
                                                           0:08 /opt/stack/venv/
eon-20160503T082438Z/bin/python2 /opt/stack/service/eon-conductor/venv/bin/eon-
conductor --config-file=/opt/stack/service/eon-conductor-20160503T082438Z/etc/eon/
eon-conductor.conf
        31965 0.0 0.0 12732 2308 pts/0
                                              S+
stack
                                                   18:50
                                                           0:00 grep eon
```

# 20.4 Issue: ESX Cluster shows UNKNOWN in OpsConsole

In the OpsConsole Alarms dashboard, if all the alarms for ESX cluster are showing UNKNOWN then restart the monasca-agent running in ESX compute proxy.

- 1. SSH to the respective compute proxy. You can find the hostname of the proxy from the dimensions list shown against the respective alarm.
- 2. Restart the monasca-agent service.

# 20.5 Issue: Unable to view the VM console in Horizon UI

By default the gdbserver firewall is disabled in ESXi host which results in a Handshake error when accessing the VM instance console in the Horizon UI.

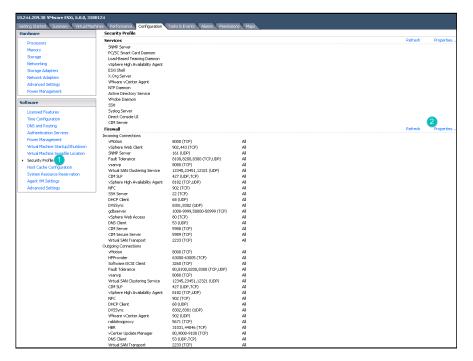


### Procedure to enable gdbserver

- 1. Login to vSphere Client.
- 2. Select the ESXi Host and click **Configuration** tab in the menu bar. You must perform the following actions on all the ESXi hosts in the compute clusters.

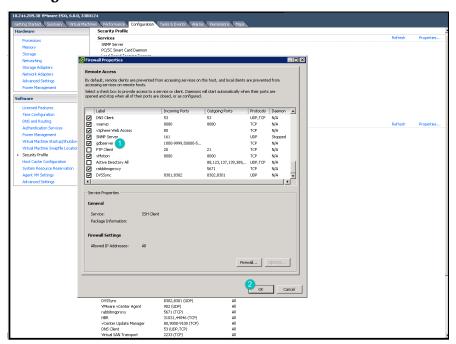


3. On the left hand side select **Security Profile** from the list of **Software**. Click **Properties** on the right hand side.



Firewall Properties box displays.

4. Select **gdbserver** checkbox and click **OK**.



## III Post-Installation

- 21 Overview 386
- 22 Cloud Verification 387
- 23 UI Verification 393
- 24 Installing OpenStack Clients 403
- 25 Configuring Transport Layer Security (TLS) 406
- 26 Configuring Availability Zones 422
- 27 Configuring Load Balancer as a Service 423
- 28 Other Common Post-Installation Tasks 448

## 21 Overview

Once you have completed your cloud deployment, these are some of the common post-installation tasks you may need to perform. Take a look at the descriptions below to determine which of these you need to do.

## 22 Cloud Verification

Once you have completed your cloud deployment, these are some of the common post-installation tasks you may need to perform to verify your cloud installation.

### 22.1 API Verification

HPE Helion OpenStack 5.0 provides a tool, Tempest, that you can use to verify that your cloud deployment completed successfully:

- Section 22.1.1, "Prerequisites"
- Section 22.1.2, "Tempest Integration Tests"
  - Section 22.1.3, "Running the Tests"
  - Section 22.1.4, "Viewing Test Results"
  - Section 22.1.5, "Customizing the Test Run"
- Section 23.1, "Verifying Your Block Storage Backend"
- Section 23.2, "Verify the Object Storage (Swift) Operations"

### 22.1.1 Prerequisites

The verification tests rely on you having an external network setup and a cloud image in your image (Glance) repository. Run the following playbook to configure your cloud:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts hlm-cloud-configure.yml
```



### Note

In HPE Helion OpenStack 5.0, the EXT\_NET\_CIDR setting for the external network is now specified in the input model - see .

387 API Verification

### 22.1.2 Tempest Integration Tests

Tempest is a set of integration tests for OpenStack API validation, scenarios, and other specific tests to be run against a live OpenStack cluster. In HPE Helion OpenStack 5.0, Tempest has been modeled as a service and this gives you the ability to locate Tempest anywhere in the cloud. It is recommended that you install Tempest on your lifecycle manager node - that is where it resides by default in a new installation.

A version of the upstream Tempest (http://docs.openstack.org/developer/tempest/) integration tests is pre-deployed on the HPE Helion OpenStack lifecycle manager node. For details on what Tempest is testing, you can check the contents of this file on your lifecycle manager:

```
/opt/stack/tempest/run_filters/ci.txt
```

You can use these embedded tests to verify if the deployed cloud is functional.

For more information on running Tempest tests, see Tempest - The OpenStack Integration Test Suite (https://git.openstack.org/cgit/openstack/tempest/tree/README.rst) ₽.



### **Important**

Running these tests requires access to the deployed cloud's identity admin credentials

Tempest creates and deletes test accounts and test resources for test purposes.

In certain cases Tempest might fail to clean-up some of test resources after a test is complete, for example in case of failed tests.

### 22.1.3 Running the Tests

To run the default set of Tempest tests:

- 1. Log in to the lifecycle manager.
- 2. Ensure you can access your cloud:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts cloud-client-setup.yml
source /etc/environment
```

3. Run the tests:

```
cd ~/scratch/ansible/next/hos/ansible
```

```
ansible-playbook -i hosts/verb_hosts tempest-run.yml
```

Optionally, you can Section 22.1.5, "Customizing the Test Run".

### 22.1.4 Viewing Test Results

Tempest is deployed under <a href="/>/opt/stack/tempest">/opt/stack/tempest</a>. Test results are written in a log file in the following directory:

```
/opt/stack/tempest/logs
```

A detailed log file is written to:

```
/opt/stack/tempest/tempest.log
```

The results are also stored in the testrepository database.

To access the results after the run:

- 1. Log in to the lifecycle manager.
- 2. Change to the tempest directory and list the test results:

```
cd /opt/stack/tempest
./venv/bin/testr last
```

### **Important**

If you encounter an error saying "local variable 'run\_subunit\_content' referenced before assignment", you may need to log in as the <u>tempest</u> user to run this command. This is due to a known issue reported at https://bugs.launchpad.net/testrepository/+bug/1348970 .

See Test Repository Users Manual (https://testrepository.readthedocs.org/en/latest/) for more details on how to manage the test result repository.

### 22.1.5 Customizing the Test Run

There are several ways available to customize which tests will be executed.

•

•

389 Viewing Test Results

### 22.1.6 Run Tests for Specific Services and Exclude Specific Features

Tempest allows you to test specific services and features using the <u>tempest.conf</u> configuration file.

A working configuration file with inline documentation is deployed under <a href="https://opt/stack/tem-pest/etc/">/opt/stack/tem-pest/etc/</a>.

To use this, follow these steps:

- 1. Log in to the lifecycle manager.
- 2. Edit the /opt/stack/tempest/configs/tempest\_region1.conf file.
- 3. To test specific service, edit the <a href="[service\_available" section">[service\_available</a>] section and clear the comment character # and set a line to true to test that service or false to not test that service.

```
cinder = true
neutron = false
```

4. To test specific features, edit any of the \*\_feature\_enabled sections to enable or disable tests on specific features of a service.

```
[volume-feature-enabled]
[compute-feature-enabled]
[identity-feature-enabled]
[image-feature-enabled]
[network-feature-enabled]
[object-storage-feature-enabled]
```

```
#Is the v2 identity API enabled (boolean value)
api_v2 = true
#Is the v3 identity API enabled (boolean value)
api_v3 = false
```

5. Then run tests normally

## 22.1.7 Run Tests Matching a Series of White and Blacklists

You can run tests against specific scenarios by editing or creating a run filter file. Run filter files are deployed under /opt/stack/tempest/run\_filters.

Use run filters to whitelist or blacklist specific tests or groups of tests:

- lines starting with # or empty are ignored
- lines starting with + are whitelisted
- lines starting with are blacklisted
- lines not matching any of the above conditions are blacklisted

If whitelist is empty, all available tests are fed to blacklist. If blacklist is empty, all tests from whitelist are returned.

Whitelist is applied first. The blacklist is executed against the set of tests returned by the whitelist.

To run whitelist and blacklist tests:

- 1. Log in to the lifecycle manager.
- 2. Make sure you can access the cloud:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts cloud-client-setup.yml
source /etc/environment
```

3. Run the tests:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts tempest-run.yml -e run_filter
<run_filter_name>
```

Note that the run\_filter\_name is the name of the run\_filter file except for the extension. For instance, to run using the filter from the file /opt/stack/tempest/run\_filters/ci.txt, use the following:

```
ansible-playbook -i hosts/verb_hosts tempest-run.yml -e run_filter=ci
```

Documentation on the format of white and black-lists is available at:

```
/opt/stack/tempest/tests2skip.py
```

#### Example:

The following entries run API tests, exclude tests that are less relevant for deployment validation, such as negative, admin, cli and thirdparty (EC2) tests:

```
+tempest\.api\.*
```

```
*[Aa]dmin.*
*[Nn]egative.*
- tempest\.cli.*
- tempest\.thirdparty\.*
```

## 23 UI Verification

Once you have completed your cloud deployment, these are some of the common post-installation tasks you may need to perform to verify your cloud installation.

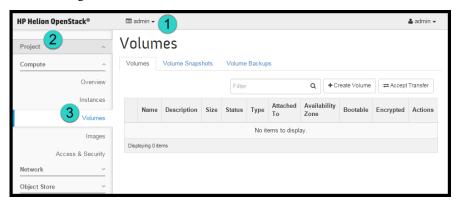
## 23.1 Verifying Your Block Storage Backend

The sections below will show you the steps to verify that your Block Storage backend was setup properly.

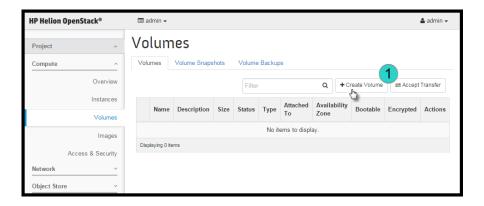
### 23.1.1 Create a Volume

Perform the following steps to create a volume using Horizon dashboard.

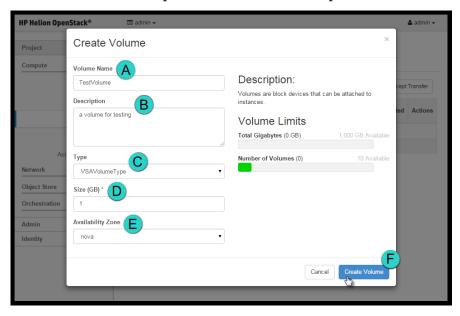
- 1. Log into the Horizon dashboard. See for details.
- 2. Under the **Project** menu in the navigation pane, click on **Volumes** under the **Compute** subheading.



3. On the **Volumes** tabs, click the **Create Volume** button to create a volume.



- 4. In the **Create Volume** options, enter the required details into the fields and then click the **Create Volume** button:
  - a. Volume Name This is the name you specify for your volume.
  - b. Description (optional) This is an optional description for the volume.
  - c. Type Select the volume type you have created for your volumes from the drop down.
  - d. Size (GB) Enter the size, in GB, you would like the volume to be.
  - e. Availability Zone You can either leave this at the default option of **Any Availability Zone** or select a specific zone from the drop down.



The dashboard will then show the volume you have just created.

394 Create a Volume

### 23.1.2 Attach Volume to an Instance

Perform the following steps to attach a volume to an instance:

- 1. Log into the Horizon dashboard. See for details.
- 2. Under the **Project** menu in the navigation pane, click the **Instances** under the **Compute** subheading.
- 3. In the **Action** column, choose the **Edit Attachments** in the menu drop down box next to the instance you want to attach the volume to.
- 4. In the Attach To Instance drop-down, select the volume that you want to attach.
- 5. Edit the **Device Name** if necessary.
- 6. Click **Attach Volume** to complete the action.
- 7. Verify that the volume you attached is displayed in the **Attached To** columns on the **Volumes** screen.

### 23.1.3 Detach Volume from Instance

Perform the following steps to detach the volume from instance:

- 1. Log into the Horizon dashboard. See for details.
- 2. Under the **Project** menu in the navigation pane, click the **Instances** under the **Compute** subheading.
- 3. Click the check box next to the name of the volume you want to detach.
- 4. In the **Action** column, choose the **Edit Attachments** in the menu drop down box next to the instance you want to attach the volume to.
- 5. Click **Detach Attachment**. A confirmation dialog box appears.
- 6. Click **Detach Attachment** to confirm the detachment of the volume from the associated instance.

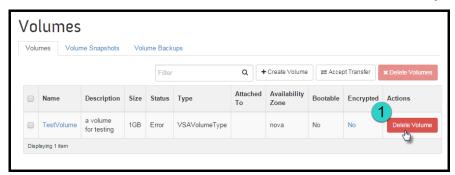
### 23.1.4 Delete Volume

Perform the following steps to delete a volume using Horizon dashboard:

- 1. Log into the Horizon dashboard. See for details.
- 2. Under the **Project** menu in the navigation pane, click on **Volumes** under the **Compute** subheading.



3. In the **Actions** column, click **Delete Volume** next to the volume you would like to delete.



4. To confirm and delete the volume, click Delete Volume again.



5. Verify that the volume was removed from the **Volumes** screen.

## 23.1.5 Verifying Your Object Storage (Swift)

Validate That All Servers Have Been Added to the Swift Rings

396 Delete Volume

1. Run the swift-compare-model-rings.yml playbook as follows:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts swift-compare-model-rings.yml
```

2. Search for output similar to the following. Specifically, look at the number of drives that are proposed to be added.

```
TASK: [swiftlm-ring-supervisor | validate-input-model | Print report] ********
ok: [helion-cp1-c1-m1-mgmt] => {
    "var": {
        "report.stdout_lines": [
            "Rings:",
            " ACCOUNT: ",
                 ring exists",
                 no device changes",
                 ring will be rebalanced",
               CONTAINER:",
                 ring exists",
                 no device changes",
                 ring will be rebalanced",
              OBJECT-0:",
                 ring exists",
                 no device changes",
                 ring will be rebalanced"
        ]
    }
}
```

- 3. If the text contains "no device changes" then the deploy was successful and no further action is needed.
- 4. If there are more drives need to be added, it indicates that the deploy failed on some nodes and that you restarted the deploy to include those nodes. However, the nodes are not in the Swift rings because enough time has not elapsed to allow the rings to be rebuilt. You have two options to continue:
  - a. Repeat the deploy. There are two steps involved as follows:
    - i. Delete the ring builder files as described in .
    - ii. Repeat the installation process starting by running the <u>site.yml</u> playbook as described in the section of the installation instructions.
  - b. Rebalance the rings several times until all drives are incorporated in the rings. This process may take several hours to complete (because you need to wait one hour between each rebalance). The steps are as follows:
    - i. Change the min-part-hours to 1 hour. See .
    - ii. Use the "First phase of ring rebalance" and "Final rebalance phase" as described in . The "Weight change phase of ring rebalance" does not apply because you have not set the weight-step attribute at this stage.
    - iii. Set the min-part-hours to the recommended 16 hours as described in .

If you receive errors during the validation, read the documentation.

## 23.2 Verify the Object Storage (Swift) Operations

Full details on how to verify the operations can be found in .

## 23.3 Uploading an Image for Use

To create a Compute instance, you need to obtain an image that you can use. The HPE Helion OpenStack lifecycle manager provides an Ansible playbook that will download a CirrOS Linux image, and then upload it as a public image to your image repository for use across your projects.

### 23.3.1 Running the Playbook

Use the following command to run this playbook:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts glance-cloud-configure.yml -e proxy=<PROXY>
```

The table below shows the optional switch that you can use as part of this playbook to specify environment-specific information:

Switch	Description
<pre>-e proxy="<pre>roxy_address:port&gt;"</pre></pre>	Optional. If your environment requires a proxy for the internet, use this switch to specify the proxy information.

### 23.3.2 How to Curate Your Own Images

OpenStack has created a guide to show you how to obtain, create, and modify images that will be compatible with your cloud:

OpenStack Virtual Machine Image Guide (http://docs.openstack.org/image-guide/content/) 

✓

## 23.3.3 Using the GlanceClient CLI to Create Images

You can use the GlanceClient on a machine accessible to your cloud or it's also installed automatically on your lifecycle manager.

The GlanceClient allows you to create, update, list, and delete images as well as manage your image member lists, which allows you to share access to images across multiple tenants. As with most of the OpenStack CLI tools, you can use the <a href="mailto:glance-help">glance help</a> command to get a full list of commands as well as their syntax.

If you would like to use the <u>--copy-from</u> option when creating an image, you will need to have your Administrator enable the http store in your environment using the instructions outlined at .

399 Running the Playbook

## 23.4 Creating an External Network

You must have an external network set up to allow your Compute instances to reach the internet. There are multiple methods you can use to create this external network and we provide two of them here. The HPE Helion OpenStack installer provides an Ansible playbook that will create this network for use across your projects. We also show you how to create this network via the command line tool from your lifecycle manager.

- Section 23.4.2, "Using the Ansible Playbook"
- Section 23.4.3, "Using the NeutronClient CLI"

### 23.4.1 Notes

If you have multiple regions in your cloud environment, it is important that when you set up your external networks in each region that you allocate unique IP ranges for each. If you have overlapping CIDR's then you risk having the same floating IP being allocated to two different virtual machines which will cause a conflict.

### 23.4.2 Using the Ansible Playbook

This playbook will query the Networking service for an existing external network, and then create a new one if you do not already have one. The resulting external network will have the name <a href="ext-net">ext-net</a> with a subnet matching the CIDR you specify in the command below.

If you need to specify more granularity, for example specifying an allocation pool for the subnet then you should utilize the Section 23.4.3, "Using the NeutronClient CLI".

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts neutron-cloud-configure.yml -e EXT_NET_CIDR=<CIDR>
```

The table below shows the optional switch that you can use as part of this playbook to specify environment-specific information:

Switch	Description
-e EXT_NET_CIDR= <cidr></cidr>	Optional. You can use this switch to specify the external network CIDR. If you choose not
	to use this switch, or use a wrong value, the
	VMs will not be accessible over the network.

Switch	Description
	This CIDR will be from the EXTERNAL VM network.  Note  If this option is not defined the default value is "172.31.0.0/16"

### 23.4.3 Using the NeutronClient CLI

For more granularity you can utilize the Neutron command line tool to create your external network.

- 1. Log in to the lifecycle manager.
- 2. Source the Admin credentials:

```
source ~/service.osrc
```

3. Create the external network and then the subnet using these commands below. Creating the network:

```
neutron net-create --router:external <external-network-name>
```

### Creating the subnet:

```
neutron subnet-create <external-network-name> <CIDR> --gateway <gateway> --
allocation-pool start=<IP_start>,end=<IP_end> [--disable-dhcp]
```

#### Where:

Value	Description
external-network-name	This is the name given to your external network. This is a unique value that you will choose. The value <a href="ext-net">ext-net</a> is usually used.

Value	Description
CIDR	You can use this switch to specify the external network CIDR. If you choose not to use this switch, or use a wrong value, the VMs will not be accessible over the network.  This CIDR will be from the EXTERNAL VM network.
gateway	Optional switch to specify the gateway IP for your subnet. If this is not included then it will choose the first available IP.
allocation-pool start end	Optional switch to specify a start and end IP address to use as the allocation pool for this subnet.
disable-dhcp	Optional switch if you want to disable DHCP on this subnet. If this is not specified then DHCP will be enabled.

## 23.4.4 Next Steps

Once the external network is created, users can create a Private Network to complete their networking setup. For instructions, see .

402 Next Steps

## 24 Installing OpenStack Clients

The way OpenStack clients are installed changed in HPE Helion OpenStack 4.0. If you have a standalone deployer, the OpenStack CLI and other clients will not be installed automatically on that node. If you require access to these clients, you will need to follow the procedure below to add the appropriate software.



### **Important**

In the <a href="entry-scale-esx-kvm-vsa">entry-scale-kvm-esx-vsa-mml</a> example configurations, <a href="eon-client">eon-client</a> is installed only on <a href="esx-compute">esx-compute</a> and <a href="esx-ovsvapp">esx-ovsvapp</a> resources. With the changes to the OpenStack client installation, <a href="eon-client">eon-client</a> won't get installed on either dedicated HLM node or Controller nodes. You need to install the <a href="eon-client">eon-client</a> on either the dedicated deployer node or controller nodes so that you can use it to add the vCenter and activate the ESX compute clusters.

1. [OPTIONAL] Connect to your standalone deployer and try to use the OpenStack CLI:

```
source ~/keystone.osrc
openstack project list

-bash: openstack: command not found
```

- 2. Edit the configuration file containing details of you Control Plane, typically  $\frac{\mbox{-/he-lion/my\_cloud/definition/data/control\_plane}}{\mbox{-/my\_cloud/definition/data/control\_plane}}$
- 3. Locate the stanza for the cluster where you want to install the client(s). For a standalone deployer, this will look like the following extract:

4. Choose the client(s) you wish to install from the following list of available clients:

```
- openstack-client
- ceilometer-client
- cinder-client
- designate-client
- glance-client
- heat-client
- ironic-client
- keystone-client
- neutron-client
- nova-client
- swift-client
- monasca-client
- barbican-client
```

5. Add the client(s) to the list of <u>service-components</u> - in this example, we add the <u>open-stack-client</u> to the standalone deployer:

```
clusters:
 - name: cluster0
    cluster-prefix: c0
    server-role: LIFECYCLE-MANAGER-ROLE
   member-count: 1
   allocation-policy: strict
    service-components:
      - ntp-server
      - lifecycle-manager
      - openstack-client
      - ceilometer-client
      - cinder-client
      - designate-client
      - glance-client
      - heat-client
      - ironic-client
      - keystone-client
      - neutron-client
      - nova-client
      - swift-client
      - monasca-client
      - barbican-client
```

6. Commit the configuration changes:

```
cd ~/helion/hos/ansible
git add -A
```

```
git commit -m "Add explicit client service deployment"
```

7. Run the configuration processor, followed by the ready-deployment playbook:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml -e encrypt="" -e
rekey=""
ansible-playbook -i hosts/localhost ready-deployment.yml
```

8. Add the software for the clients using the following command:

```
cd /home/stack/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts clients-upgrade.yml
```

9. Check that the software has been installed correctly. In this instance, connect to your standalone deployer and try to use the OpenStack CLI:

```
source ~/keystone.osrc openstack project list
```

You should now see a list of projects returned:

```
| Name |
```

## 25 Configuring Transport Layer Security (TLS)

TLS is enabled by default during the installation of HPE Helion OpenStack 5.0 and additional configuration options are available to secure your environment, as described below.

In HPE Helion OpenStack 5.0, you can provide your own certificate authority and certificates for internal and public virtual IP addresses (VIPs), and you should do so for any production cloud. The certificates automatically generated by HPE Helion OpenStack are useful for testing and setup, but you should always install your own for production use. Certificate installation is discussed below.

Please read the following if you're using the default <u>cert-name: my-public-cert</u> in your model.

The bundled test cert for public endpoints, located at <u>~/helion/my\_cloud/config/tls/certs/my-public-cert</u>, is now expired but was left in the product in case you changed the content with your valid cert. Please verify if the certificate is expired and generate your own by following the guidelines further down on this page or by using a generic instruction from the web.

You can verify the expiry by running this command:

```
openssl x509 -in ~/helion/my_cloud/config/tls/certs/my-public-cert -noout -enddate notAfter=Oct 8 09:01:58 2016 GMT
```

Before you begin, the following list of terms will be helpful when generating and installing certificates.

#### Helion generated public CA

A Helion OpenStack generated public CA ( $\underline{\text{helion\_frontend\_cacert.crt}}$ ) is available for you to use in /usr/local/share/ca-certificates.

#### Fully qualified domain name (FQDN) of the public VIP

The registered domain name. A FQDN is not mandatory. It is perfectly valid to have no FQDN and use IP addresses instead. Note that you can use FQDNs on public endpoints, and you may change them whenever the need arises.

#### Certificate authority (CA) certificate

Your certificates must be signed by a CA, such as your internal IT department or a public certificate authority. For this example we will use a self-signed certificate.

#### Server certificate

It is easy to confuse server certificates and CA certificates. Server certificates reside on the server and CA certificates reside on the client. A server certificate affirms that the server that sent it serves a set of IP addresses, domain names, and set of services. A CA certificate is used by the client to authenticate this claim.

#### SAN (subject-alt-name)

The set of IP addresses and domain names in a server certificate request: A template for a server certificate.

### Certificate signing request (CSR)

A blob of data generated from a certificate request and sent to a CA, which would then sign it, produce a server certificate, and send it back.

#### **External VIP**

External virtual IP address

#### Internal VIP

Internal virtual IP address

The major difference between an external VIP certificate and an internal VIP certificate is that the internal VIP has approximately 40 domain names in the SAN. This is because each service has a different domain name in HPE Helion OpenStack 5.0. So it is unlikely that you can create an internal server certificate before running the configuration processor. But after a configuration processor run, a certificate request would be created for each of your cert-names.

## 25.1 Configuring TLS in the input model

For this example certificate configuration, let's assume there's no FQDN for the external VIP and that you're going to use the default IP address provided by HPE Helion OpenStack 5.0. Let's also assume that for the internal VIP you will use the defaults as well. If you were to call your certificate authority "example-CA," the CA certificate would then be called "example-CA.crt" and the key would be called "example-CA.key." In the following examples, the external VIP certificate will be named "example-public-cert" and the internal VIP certificate will be named "example-internal-cert."



Any time you make a cert change when using your own CA:

- You should use a distinct name from those already existing in config/tls/cacerts. This also
  means that you should not reuse your CA names (and use unique and distinguishable names
  such as MyCompanyXYZ\_PrivateRootCA.crt). A new name is what indicates that a file is
  new or changed, so reusing a name means that the file is not considered changed even
  its contents have changed.
- You should not remove any existing CA files from config/tls/cacerts
- If you want to remove an existing CA you must
  - 1. First remove the file.
  - 2. Then run:

```
ansible -i hosts/verb_hosts FND-STN -a 'sudo keytool -delete -alias
debian:<filename to remove> \
-keystore /usr/lib/jvm/java-7-openjdk-amd64/jre/lib/security/cacerts -storepass
changeit'
```

## Important

Be sure to install your own certificate for all production clouds after installing and testing your cloud. If you ever want to test or troubleshoot later, you will be able to revert to the sample certificate to get back to a stable state for testing.



#### Note

Unless this is a new deployment, do not update both the certificate and the CA together. Add the CA first and then run a site deploy. Then update the certificate and run tls-reconfigure, FND-CLU-stop, FND-CLU-start and then hlm-reconfigure. If a playbook has failed, rerun it with -vv to get detailed error information. The configure, HAproxy restart, and reconfigure steps are included below. If this is a new deployment and you are adding your own certs/CA before running site.yml this caveat does not apply.

You can add your own certificate by following the instructions below. All changes must go into the following file:

```
~/helion/my_cloud/definition/data/network_groups.yml
```

Below are the entries for TLS for the internal and admin load balancers:

```
- provider: ip-cluster
        name: lb
        tls-components:
        - default
        components:
        # These services do not currently support TLS so they are not listed
        # under tls-components
        - nova-metadata
        roles:
        - internal
        - admin
        cert-file: helion-internal-cert
        # The helion-internal-cert is a reserved name and
        # this certificate will be autogenerated. You
        # can bring in your own certificate with a different name
        # cert-file: customer-provided-internal-cert
        # replace this with name of file in "config/tls/certs/"
```

The configuration processor will also create a request template for each named certificate under info/cert\_reqs/ This will be of the form:

```
info/cert_reqs/customer-provided-internal-cert
```

These request templates contain the subject <u>Alt-names</u> that the certificates need. You can add to this template before generating your certificate signing request.

You would then send the CSR to your CA to be signed, and once you receive the certificate, place it in config/tls/certs

When you bring in your own certificate, you may want to bring in the trust chains (or CA certificate) for this certificate. This is usually not required if the CA is a public signer that is typically bundled with the operating system. However, we suggest you include it anyway by copying the file into the directory config/cacerts/.

## 25.2 User-provided certificates and trust chains

HPE Helion OpenStack generates its own internal certificates but is designed to allow you to bring in your own certificates for the VIPs. Here is the general process.

- 1. You must have a server certificate and a CA certificate to go with it (unless the signer is a public CA and it's already bundled with most distributions).
- 2. You must decide the names of the server certificates and configure the <a href="network\_group-s.yml">network\_group-s.yml</a> file in the input model such that each load balancer provider has at least one cert-name associated with it.
- 3. Run the configuration processor. Note that you may or may not have the certificate file at this point. The configuration processor would create certificate request file artifacts under <a href="info/cert\_reqs/">info/cert\_reqs/</a> for each of the cert-name(s) in the <a href="network\_groups.yml">network\_groups.yml</a> file. While there's no special reason to use the request file created for an external endpoint VIP certificate, it is important to use the request files created for internal certificates since the canonical names for the internal VIP can be many and service specific and each of these need to be in the Subject Alt Names attribute of the certificate.
- 4. Create a certificate signing request for this request file and send it to your internal CA or a public CA to get it certified and issued with a certificate. You will now have a server certificate and possibly a trust chain or CA certificate.
- 5. Next, upload it to the lifecycle manager. Server certificates should be added to <u>config/tls/cacerts</u>. The file extension should be <u>.crt</u> for the CA certificate to be processed by HPE Helion OpenStack. Detailed steps are next.

# 25.3 Edit the input model to include your certificate files

Edit the load balancer configuration in <a href="helion/my\_cloud/definition/data/network\_group-s.yml">helion/my\_cloud/definition/data/network\_group-s.yml</a>:

```
load-balancers:
```

- provider: ip-cluster

name: lb

```
tls-components:
 - default
components:
 - vertica
 - nova-metadata
 roles:
- internal
 - admin
cert-file: example-internal-cert #<<<----- Certificate name for the internal VIP
- provider: ip-cluster
name: extlb
external-name: myhelion.test #<<----- Use just IP for the external VIP in this example
tls-components:
- default
 roles:
- public
cert-file: example-public-cert #<<----- Certificate name for the external VIP
```

Commit your changes to the local git repository and run the configuration processor:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "changed VIP certificates"
ansible-playbook -i hosts/localhost config-processor-run.yml
```

Verify that certificate requests have been generated by the configuration processor for every certificate file configured in the <a href="networks\_groups.yml">networks\_groups.yml</a> file. In this example, there are two files, as shown from the list command:

```
ls ~/helion/my_cloud/info/cert_reqs
example-internal-cert
example-public-cert
```

## 25.4 Generate a self-signed CA



### Note

In a production setting you will not perform this step. You will use your company's CA or a valid public CA.

This section demonstrates to how you can create your own self-signed CA and then use this CA to sign server certificates. This CA can be your organization's IT internal CA that is self-signed and whose CA certificates are deployed on your organization's machines. This way the server certificate becomes legitimate.



#### Note

Please use a unique CN for your example Certificate Authority and do not install multiple CA certificates with the same CN into your cloud.

Copy the commands below to the command line and execute. This will cause the two files, example-CA.key and example-CA.crt to be created:

```
export EXAMPLE_CA_KEY_FILE='example-CA.key'
export EXAMPLE_CA_CERT_FILE='example-CA.crt'
openssl req -x509 -batch -newkey rsa:2048 -nodes -out "${EXAMPLE_CA_CERT_FILE}" \
-keyout "${EXAMPLE_CA_KEY_FILE}" \
-subj "/C=UK/0=hp/CN=YourOwnUniqueCertAuthorityName" \
-days 365
```

You can tweak the subj and days settings above to meet your needs, or to test. For instance, if you want to test what happens when a CA expires, you can set 'days' to a very low value. Grab the configuration processor-generated request file from info/cert\_reqs/:

```
cat ~/helion/my_cloud/info/cert_reqs/example-internal-cert
```

Now, copy this file to your working directory and appended a .req extension to it.

```
cp ~/helion/my_cloud/info/cert_reqs/example-internal-cert example-internal-cert.req
```

#### **EXAMPLE 25.1: CERTIFICATE REQUEST FILE**

```
[req]
distinguished_name = req_distinguished_name
req_extensions = v3_req
prompt = no

[ req_distinguished_name ]
CN = "helion-vip"

[ v3_req ]
basicConstraints = CA:FALSE
keyUsage = nonRepudiation, digitalSignature, keyEncipherment
subjectAltName = @alt_names
```

```
[ alt_names ]
DNS.1 = "deployerincloud-ccp-c0-m1-mamt"
DNS.2 = "deployerincloud-ccp-vip-CEI-API-mgmt"
DNS.3 = "deployerincloud-ccp-vip-CND-API-mgmt"
DNS.4 = "deployerincloud-ccp-vip-DES-API-mgmt"
DNS.5 = "deployerincloud-ccp-vip-FND-MDB-mgmt"
DNS.6 = "deployerincloud-ccp-vip-FND-RMQ-mgmt"
DNS.7 = "deployerincloud-ccp-vip-FND-VDB-mgmt"
DNS.8 = "deployerincloud-ccp-vip-FRE-API-mgmt"
DNS.9 = "deployerincloud-ccp-vip-GLA-API-mgmt"
DNS.10 = "deployerincloud-ccp-vip-GLA-REG-mgmt"
DNS.11 = "deployerincloud-ccp-vip-HEA-ACF-mgmt"
DNS.12 = "deployerincloud-ccp-vip-HEA-ACW-mgmt"
DNS.13 = "deployerincloud-ccp-vip-HEA-API-mgmt"
DNS.14 = "deployerincloud-ccp-vip-HUX-SVC-mgmt"
DNS.15 = "deployerincloud-ccp-vip-HZN-WEB-mgmt"
DNS.16 = "deployerincloud-ccp-vip-KEY-API-mgmt"
DNS.17 = "deployerincloud-ccp-vip-KEYMGR-API-mgmt"
DNS.18 = "deployerincloud-ccp-vip-LOG-API-mgmt"
DNS.19 = "deployerincloud-ccp-vip-LOG-SVR-mgmt"
DNS.20 = "deployerincloud-ccp-vip-MON-API-mgmt"
DNS.21 = "deployerincloud-ccp-vip-NEU-SVR-mgmt"
DNS.22 = "deployerincloud-ccp-vip-NOV-API-mgmt"
DNS.23 = "deployerincloud-ccp-vip-NOV-MTD-mgmt"
DNS.24 = "deployerincloud-ccp-vip-OCT-API-mgmt"
DNS.25 = "deployerincloud-ccp-vip-OPS-WEB-mgmt"
DNS.26 = "deployerincloud-ccp-vip-SHP-API-mgmt"
DNS.27 = "deployerincloud-ccp-vip-SWF-PRX-mgmt"
DNS.28 = "deployerincloud-ccp-vip-admin-CEI-API-mgmt"
DNS.29 = "deployerincloud-ccp-vip-admin-CND-API-mgmt"
DNS.30 = "deployerincloud-ccp-vip-admin-DES-API-mgmt"
DNS.31 = "deployerincloud-ccp-vip-admin-FND-MDB-mgmt"
DNS.32 = "deployerincloud-ccp-vip-admin-FRE-API-mgmt"
DNS.33 = "deployerincloud-ccp-vip-admin-GLA-API-mgmt"
DNS.34 = "deployerincloud-ccp-vip-admin-HEA-ACF-mgmt"
DNS.35 = "deployerincloud-ccp-vip-admin-HEA-ACW-mgmt"
DNS.36 = "deployerincloud-ccp-vip-admin-HEA-API-mgmt"
DNS.37 = "deployerincloud-ccp-vip-admin-HUX-SVC-mgmt"
DNS.38 = "deployerincloud-ccp-vip-admin-HZN-WEB-mgmt"
DNS.39 = "deployerincloud-ccp-vip-admin-KEY-API-mgmt"
DNS.40 = "deployerincloud-ccp-vip-admin-KEYMGR-API-mgmt"
DNS.41 = "deployerincloud-ccp-vip-admin-MON-API-mgmt"
DNS.42 = "deployerincloud-ccp-vip-admin-NEU-SVR-mgmt"
DNS.43 = "deployerincloud-ccp-vip-admin-NOV-API-mgmt"
DNS.44 = "deployerincloud-ccp-vip-admin-OPS-WEB-mgmt"
DNS.45 = "deployerincloud-ccp-vip-admin-SHP-API-mgmt"
DNS.46 = "deployerincloud-ccp-vip-admin-SWF-PRX-mgmt"
```

```
DNS.47 = "192.168.245.5"

IP.1 = "192.168.245.5"

=======end of certificate request file.
```



### Note

In the case of a public VIP certificate, please add all the FQDNs you want it to support Currently Helion does not add the hostname for the external-name specified in <a href="mork\_groups.yml">met-work\_groups.yml</a> to the certificate request file . However, you can add it to the certificate request file manually. Here we assume that <a href="myhelion.test">myhelion.test</a> is your external-name. In that case you would add this line (to the full that is shown above):

```
DNS.48 = "myhelion.test"
```



### Note

Any attempt to use IP addresses rather than FQDNs in certificates must use subject alternate name entries that list both the IP address (needed for Google) and DNS with an IP (needed for a Python bug workaround). Failure to create the certificates in this manner will cause future installations of Go-based tools (such as Cloud Foundry, Stackato and other PaaS components) to fail.

## 25.5 Generate a certificate signing request

Now that you have a CA and a certificate request file, it's time to generate a CSR.

```
export EXAMPLE_SERVER_KEY_FILE='example-internal-cert.key'
export EXAMPLE_SERVER_CSR_FILE='example-internal-cert.csr'
export EXAMPLE_SERVER_REQ_FILE=example-internal-cert.req
openssl req -newkey rsa:2048 -nodes -keyout "$EXAMPLE_SERVER_KEY_FILE" -out
"$EXAMPLE_SERVER_CSR_FILE" -extensions v3_req -config "$EXAMPLE_SERVER_REQ_FILE"
```

Note that in production you would usually send the generated <a href="example-internal-cert.csr">example-internal-cert.csr</a> file to your IT department. But in this example you are your own CA, so sign and generate a server certificate.

### 25.6 Generate a server certificate



### Note

In a production setting you will not perform this step. You will send the CSR created in the previous section to your company CA or a to a valid public CA and have them sign and send you back the certificate.

This section demonstrates how you would use your own self-signed CA that your created earlier to sign and generate a server certificate. A server certificate is essentially a signed public key, the signer being a CA and trusted by a client. When you install this the signing CA's certificate (called CA certificate or trust chain) on the client machine, you are telling the client to trust this CA, and thereby implicitly trusting any server certificates that are signed by this CA, thus creating a trust anchor.

#### CA configuration file

When the CA signs the certificate, it uses a configuration file that tells it to verify the CSR. Note that in a production scenario the CA takes care of this for you.

Create a file called openssl.cnf and add the following contents to it.

```
# Copyright 2010 United States Government as represented by the
# Administrator of the National Aeronautics and Space Administration.
# All Rights Reserved.
#...
# OpenSSL configuration file.
# Establish working directory.
dir = .
[ ca ]
default_ca = CA_default
[ CA_default ]
serial = $dir/serial
database = $dir/index.txt
new certs dir = $dir/
certificate = $dir/cacert.pem
private_key = $dir/cakey.pem
unique_subject = no
```

```
default_crl_days = 365
default_days = 365
default md = md5
preserve = no
email_in_dn = no
nameopt = default ca
certopt = default_ca
policy = policy_match
copy_extensions = copy
# NOTE(dprince): stateOrProvinceName must be 'supplied' or 'optional' to
# work around a stateOrProvince printable string UTF8 mismatch on
# RHEL 6 and Fedora 14 (using openssl-1.0.0-4.el6.x86_64 or
# openssl-1.0.0d-1.fc14.x86 64)
[ policy_match ]
countryName = optional
stateOrProvinceName = optional
organizationName = optional
organizationalUnitName = optional
commonName = supplied
emailAddress = optional
[ req ]
default_bits = 1024 # Size of keys
default_keyfile = key.pem # name of generated keys
default md = md5 # message digest algorithm
string_mask = nombstr # permitted characters
distinguished_name = req_distinguished_name
req_extensions = v3_req
x509_extensions = v3_ca
[ req_distinguished_name ]
# Variable name Prompt string
#----
0.organizationName = Organization Name (company)
organizationalUnitName = Organizational Unit Name (department, division)
emailAddress = Email Address
emailAddress_max = 40
localityName = Locality Name (city, district)
stateOrProvinceName = State or Province Name (full name)
countryName = Country Name (2 letter code)
countryName_min = 2
countryName max = 2
commonName = Common Name (hostname, IP, or your name)
commonName_max = 64
# Default values for the above, for consistency and less typing.
```

### Sign and create a server certificate

Now you can sign the server certificate with your CA. Copy the commands below to the command line and execute. This will cause the one file, example-internal-cert.crt, to be created:

```
export EXAMPLE_SERVER_CERT_FILE='example-internal-cert.crt'
export EXAMPLE_SERVER_CSR_FILE='example-internal-cert.csr'
export EXAMPLE_CA_KEY_FILE='example-CA.key'
export EXAMPLE_CA_CERT_FILE='example-CA.crt'

touch index.txt
openssl rand -hex -out serial 6

openssl ca -batch -notext -md sha256 -in "$EXAMPLE_SERVER_CSR_FILE" \
-cert "$EXAMPLE_CA_CERT_FILE" \
-keyfile "$EXAMPLE_CA_KEY_FILE" \
-out "$EXAMPLE_SERVER_CERT_FILE" \
-out "$EXAMPLE_SERVER_CERT_FILE" \
-config openssl.cnf -extensions v3_req
```

Finally, concatenate both the server key and certificate in preparation for uploading to the lifecycle manager.

```
cat example-internal-cert.key example-internal-cert.crt > example-internal-cert
```

Note that you have only created the internal-cert in this example. Repeat the above sequence for example-public-cert. Make sure you use the appropriate certificate request generated by the configuration processor.

## 25.7 Upload to the lifecycle manager

The following two files created from the example run above will need to be uploaded to the lifecycle manager and copied into config/tls:

- example-internal-cert
- example-CA.crt

Once on the lifecycle manager, execute the following two copy commands to copy to their respective directories. Note if you had created an external cert, you can copy that in a similar manner, specifying its name using the copy command as well.

```
cp example-internal-cert ~/helion/my_cloud/config/tls/certs/
cp example-CA.crt ~/helion/my_cloud/config/tls/cacerts/
```

#### Continue with the deployment

Next, log into the lifecycle manager node, and save and commit the changes to the local git repository:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "updated certificate and CA"
```

Next, rerun the config-processor-run playbook, and run ready-deployment.yml:

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
ansible-playbook -i hosts/localhost ready-deployment.yml
```

If you receive any prompts, enter the required information.



#### Note

For automated installation (e.g. CI) you can specify the required passwords on the Ansible command line. For example, the command below will disable encryption by the configuration processor:

```
ansible-playbook -i hosts/localhost config-processor-run.yml -e encrypt="" -e
rekey=""
```

Run this series of runbooks to complete the deployment:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts tls-reconfigure.yml
ansible-playbook -i hosts/verb_hosts FND-CLU-stop.yml
ansible-playbook -i hosts/verb_hosts FND-CLU-start.yml
ansible-playbook -i hosts/verb_hosts monasca-stop.yml
ansible-playbook -i hosts/verb_hosts monasca-start.yml
ansible-playbook -i hosts/verb_hosts hlm-reconfigure.yml
```

## 25.8 Configuring the cipher suite

By default, the cipher suite is set to: <a href="https://docs.openstack.org/security-guide/secure-communication/introduction-to-ssl-and-tls.html">https://docs.openstack.org/security-guide/secure-communication/introduction-to-ssl-and-tls.html</a>. You may override this. To do so, open <a href="mailto:config/haproxy/defaults.yml">config/haproxy/defaults.yml</a> and edit it. The parameters can be found under the <a href="haprox-y\_globals">haprox-y\_globals</a> list.

```
- "ssl-default-bind-ciphers HIGH:!aNULL:!eNULL:!DES:!3DES"
- "ssl-default-server-ciphers HIGH:!aNULL:!eNULL:!DES:!3DES"
```

Make the changes as needed. It's best to keep the two options identical.

## 25.9 Testing

You can easily determine if an endpoint is behind TLS. To do so, run the following command, which probes a Keystone identity service endpoint that's behind TLS:

```
echo | openssl s_client -connect 192.168.245.5:5000 | openssl x509 -fingerprint -noout
depth=0 CN = helion-vip
verify error:num=20:unable to get local issuer certificate
verify return:1
depth=0 CN = helion-vip
verify error:num=27:certificate not trusted
verify return:1
depth=0 CN = helion-vip
verify error:num=21:unable to verify the first certificate
verify return:1
DONE
SHA1 Fingerprint=C6:46:1E:59:C6:11:BF:72:5E:DD:FC:FF:B0:66:A7:A2:CC:32:1C:B8
```

The next command probes a MySQL endpoint that is not behind TLS:

```
echo | openssl s_client -connect 192.168.245.5:3306 | openssl x509 -fingerprint -noout 140448358213264:error:140770FC:SSL routines:SSL23_GET_SERVER_HELLO:unknown protocol:s23_clnt.c:795: unable to load certificate 140454148159120:error:0906D06C:PEM routines:PEM_read_bio:no start line:pem_lib.c:703:Expecting: TRUSTED CERTIFICATE
```

# 25.10 Verifying that the trust chain is correctly deployed

You can determine if the trust chain is correctly deployed by running the following commands:

```
echo | openssl s_client -connect 192.168.245.9:5000 2>/dev/null | grep code

Verify return code: 21 (unable to verify the first certificate)

echo | openssl s_client -connect 192.168.245.9:5000 -CAfile /usr/local/share/ca-

certificates/helion_frontend_cacert.crt 2>/dev/null | grep code

Verify return code: 0 (ok)
```

Here, the first command produces error 21, which is then fixed by providing the CA certificate file. This verifies that the CA certificate matches the server certificate.

## 25.11 Turning TLS on or off

You should leave TLS enabled in production. However, if you need to disable it for any reason, you must change "tls-components" to "components" in <a href="network\_groups.yml">network\_groups.yml</a> (as shown earlier) and comment out the cert-file. Additionally, if you have a <a href="network\_groups.yml">network\_groups.yml</a> file from a previous installation, you won't have TLS enabled unless you change "components" to "tls-components" in that file. By default, Horizon is configured with TLS in the input model. Note that you should not disable TLS in the input model for Horizon as that is a public endpoint and is required. Additionally, you should keep all services behind TLS, but using the input model file <a href="network\_groups.yml">network\_groups.yml</a> you may turn TLS off for a service for troubleshooting or debugging. TLS should always be enabled for production environments.

If you are using an example input model on a clean install, all supported TLS services will be enabled before deployment of your cloud. If you want to change this setting later, for example, when upgrading, you can change the input model and reconfigure the system. The process is as follows:

Edit the input model <a href="network\_groups.yml">network\_groups.yml</a> file appropriately, as described above. Then, commit the changes to the git repository:

```
cd ~/helion/hos/ansible/
git add -A
git commit -m "TLS change"
```

Change directories again and run the configuration processor and ready deployment playbooks:

```
ansible-playbook -i hosts/localhost config-processor-run.yml
ansible-playbook -i hosts/localhost ready-deployment.yml
```

Change directories again and run the reconfigure playbook:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts hlm-reconfigure.yml
```

**421** Turning TLS on or off

# 26 Configuring Availability Zones

The lifecycle manager only creates a default availability zone during installation. If your system has multiple failure/availability zones defined in your input model, these zones will not get created automatically.

Once the installation has finished, you can run the <u>nova-cloud-configure.yml</u> playbook to configure availability zones and assign compute nodes to those zones based on the configuration specified in the model.

You can run the playbook <u>nova-cloud-configure.yml</u> any time you make changes to the configuration of availability zones in your input model. Alternatively, you can use Horizon or the command line to perform the configuration.

For more details, see .

## 27 Configuring Load Balancer as a Service

The HPE Helion OpenStack Neutron LBaaS service supports several load balancing providers. By default, both Octavia and the namespace haproxy driver are configured to be used. We describe this in more detail here.



## Warning

If you are planning to upgrade from HPE Helion OpenStack 3.0 or HPE Helion OpenStack 4.0, please contact F5 and HPE PointNext to determine which F5 drivers have been certified for use with Helion OpenStack. Loading drivers not certified by HPE may result in catastrophic failure of your cloud deployment. The last tested versions are 8.0.8 for HPE Helion OpenStack 3.x and 9.0.3 for HPE Helion OpenStack 4.x. More information is expected in 4th quarter 2017, including the correct drivers for HPE Helion OpenStack 5.x.

The HPE Helion OpenStack Neutron LBaaS service supports several load balancing providers. By default, both Octavia and the namespace haproxy driver are configured to be used. A user can specify which provider to use with the <a href="https://example.com/en-provider">--provider</a> flag upon load balancer creation. Example:

```
neutron lbaas-loadbalancer-create --name <name> --provider [octavia|haproxy] <subnet>
```

If you don't specify the <u>--provider</u> option it will default to Octavia. The Octavia driver provides more functionality than the haproxy namespace driver which is deprecated. The haproxy namespace driver will be retired in a future version of HPE Helion OpenStack.

There are additional drivers for 3rd party hardware load balancers. Please refer to the vendor directly. The <a href="neutron service-provider-list">neutron service-provider-list</a> command displays not only the currently installed load balancer drivers but also other installed services such as VPN. You can see a list of available services as follows:

```
| LOADBALANCERV2 | octavia | True |
| VPN | openswan | True |
| LOADBALANCERV2 | haproxy | False |
+-----+
```



#### Note

In the example above, the providers are listed twice. This is a limitation in HPE Helion OpenStack 3.0. Also note that the Octavia load balancer provider is listed as the default.

## 27.1 Prerequisites

You will need to create an external network and create an image to test LBaaS functionality. If you have already created an external network and registered an image, this step can be skipped.

Creating an external network: Section 23.4, "Creating an External Network".

Creating and uploading a Glance image:

## 27.2 Octavia Load Balancing Provider

The Octavia Load balancing provider bundled with HPE Helion OpenStack 5.0 is an operator grade load balancer for OpenStack. It is based on the Mitaka version of Octavia. It differs from the namespace driver by starting a new nova virtual machine to house the haproxy load balancer software, called an *amphora*, that provides the load balancer function. A virtual machine for each load balancer requested provides a better separation of load balancers between tenants and makes it easier to grow load balancing capacity alongside compute node growth. Additionally, if the virtual machine fails for any reason Octavia will replace it with a replacement VM from a pool of spare VM's, assuming that the feature is configured.



#### Note

The Health Monitor will not create or replace failed amphorae. If the pool of spare VM's is exhausted there will be no additional virtual machines to handle load balancing requests.

**424** Prerequisites

Octavia uses two-way SSL encryption to communicate with the amphora. There are demo Certificate Authority (CA) certificates included with HPE Helion OpenStack 5.0 in /home/stack/scratch/ansible/next/hos/ansible/roles/octavia-common/files on the lifecycle manager. For additional security in production deployments, all certificate authorities should be replaced with ones you generated yourself by running the following commands:

```
openssl genrsa -passout pass:foobar -des3 -out cakey.pem 2048
openssl req -x509 -passin pass:foobar -new -nodes -key cakey.pem -out ca_01.pem
openssl genrsa -passout pass:foobar -des3 -out servercakey.pem 2048
openssl req -x509 -passin pass:foobar -new -nodes -key cakey.pem -out serverca_01.pem
```

For more details refer to the openssl man page (https://www.openssl.org/docs/manmaster/apps/openssl.html) . ♣.



#### Note

If you change the certificate authority and have amphora running with an old CA you won't be able to control the amphora. The amphora's will need to be failed over so they can utilize the new certificate. If you change the CA password for the server certificate you need to change that in the Octavia configuration files as well. See the guide for more information.

# 27.3 Setup of prerequisites

#### Octavia Network and Management Network Ports

The Octavia management network and Management network must have access to each other. If you have a configured firewall between the Octavia management network and Management network, you must open up the following ports to allow network traffic between the networks.

- From Management network to Octavia network
  - TCP 9443 (amphora API)
- From Octavia network to Management network
  - TCP 9876 (Octavia API)
  - UDP 5555 (Octavia Health Manager)

#### **Installing the Amphora Image**

Octavia is utilizing Nova VMs for its load balancing function and HPE provides images used to boot those VM's called octavia-amphora-haproxy.



## Warning

Without these images the Octavia load balancer will not work.

You can download the **Octavia Amphora HA Proxy Guest Image** from HPE Software Entitlement Portal.

- Sign in and download the product and signature files from the following link: Software Entitlement Portal (http://www.hpe.com/software/entitlements) 

  ✓
- Once the image is downloaded it needs to be placed on the lifecycle manager node and the image registered.

#### Register the image.

1. Switch to the ansible directory and register the image by giving the full path and name (e.g. /tmp/octavia-amphora-haproxy-guest-image.tgz) as argument to service\_package:

```
$ cd ~/scratch/ansible/next/hos/ansible/
$ ansible-playbook -i hosts/verb_hosts -e service_package=<image path/name> service-
guest-image.yml
```

2. Source the service user (this can be done on a different computer)

```
$ source service.osrc
```

3. Verify that the image was uploaded and registered (this can be done on a computer with access to the glance CLI client)



In rare circumstances the image won't be registered and you will have to run the whole registration again. If you run the registration by accident, the system will only upload a new image if the underlying image has been changed.

4. Check the status of the image by running glance image-show <image-id>.

Property	Value	-
	094a553d38fb5bfdb43f4a662d84ec2e	
container_format	bare	-
created_at	2016-04-14T23:05:21Z	-
disk_format	qcow2	
id	e64cb914-15d2-4ad8-a63c-b7c60a6c232e	
min_disk	0	
min_ram	0	-
name	octavia-amphora-x64-haproxy_hos-3.0.0	)
owner	0671a8d4d71c44ffb210c11cb5d11f7f	
protected	False	-
size	434379264	
status	active	-
tags	[]	-
updated_at	2016-04-14T23:05:36Z	-
virtual_size	None	-
visibility	private	- 1



## **Important**

In the example above, the status of the image is *active* which means the image was successfully registered. If a status of the images is *queued*, you must run the image registration again.

Please be aware that if you have already created load balancers they won't receive the new image. Only load balancers created after the image has been successfully installed will use the new image. If existing load balancers need to be switched to the new image please follow the instructions in the guide.

#### Setup network, subnet, router, security and IP's

If you have already created a network, subnet, router, security settings and IP's you can skip the following steps and go directly to creating the load balancers.

#### 1. Create a network.

eutron net-create lb_net1	.+	+
Field	Value	
	-+	+
admin_state_up	True	1
id	71a1ac88-30a3-48a3-a18b-d98509fbef5c	1
mtu	0	1
name	lb_net1	1
provider:network_type	vxlan	1
<pre>provider:physical_network</pre>		1
<pre>provider:segmentation_id</pre>	1061	1
router:external	False	1
shared	False	1
status	ACTIVE	1
subnets		1
tenant_id	4b31d0508f83437e83d8f4d520cda22f	1

#### 2. Create a subnet.

```
neutron subnet-create --name lb_subnet1 lb_net1 10.247.94.128/26 --gateway
10.247.94.129
+-----+
           | Value
+-----
| allocation_pools | {"start": "10.247.94.130", "end": "10.247.94.190"} |
       | 10.247.94.128/26
| cidr
| dns_nameservers |
| enable_dhcp | True
| gateway_ip
           | 10.247.94.129
| host_routes
| id
           | 6fc2572c-53b3-41d0-ab63-342d9515f514
| ip_version | 4
| ipv6_address_mode |
| ipv6_ra_mode
name
           | lb_subnet1
```

+-----+

3. Create a router.

```
neutron router-create --distributed False lb_router1
+-----+
+-----+
| admin_state_up
              | True
| distributed | False
| external_gateway_info |
               | False
              | 6aafc9a9-93f6-4d7e-94f2-3068b034b823 |
| id
| name
              | lb_router1
| routes
              | ACTIVE
| status
| tenant_id
              | 4b31d0508f83437e83d8f4d520cda22f
```

4. Add interface to router. In this example, <a href="neutron router-interface-add lb\_router1">neutron router-interface-add lb\_router1</a>
<a href="lb\_subnet1">lb\_subnet1</a> will add interface 426c5898-f851-4f49-b01f-7a6fe490410c to router lb router1.

```
neutron router-interface-add lb_router1 lb_subnet1
```

5. Set gateway for router.

```
neutron router-gateway-set lb_router1 ext-net
```

6. Check networks.

+	+	+	+
id	name	subnets	I
+	+	+	+
d3cb12a6-a000-4e3e-	ext-net	f4152001-2500-4ebe-ba9d-	I
82c4-ee04aa169291	I	a8d6149a50df 10.247.96.0/23	
8306282a-3627-445a-a588-c18	OCTAVIA-MGMT-NET	f00299f8-3403-45ae-ac4b-	
8b6a13163	I	58af41d57bdc	
	I	10.247.94.128/26	
71a1ac88-30a3-48a3-a18b-	lb_net1	6fc2572c-	1
d98509fbef5c		53b3-41d0-ab63-342d9515f514	1
		10.247.94.128/26	1

7. Create security group.

neutron security-grou	up-create lb_secgroup1 	
+		
+		
Field	Value	
Tietu	vacue	
+	·	
+		
+		
description	T.	
465611961011		
id	75343a54-83c3-464c-8773-802598afaee9	
1 = 4		
name	lb_secgroup1	
T Traine	cb_5eeg.eap1	
l security aroup rule	es   {"remote_group_id": null, "direction": "egress",	
	null, "protocol": null, "tenant_id":	
I I I I I I I I I I I I I I I I I I I	"4b31d0508f83437e83d8f4d520cda22f", "port_range_max": r	
"security aroun id"	: "75343a54-83c3-464c-8773-802598afaee9",	iucc,
I	"port range min": null, "ethertype": "IPv4", "id":	
"20ae3723-050c-44ed		
1	{"remote group id": null, "direction": "egress",	
"remote in nrefix":	null, "protocol": null, "tenant_id":	
	"4b31d0508f83437e83d8f4d520cda22f", "port range max": r	nul 1
"security aroun id"	: "75343a54-83c3-464c-8773-802598afaee9",	14 6 6 7
	"port range min": null, "ethertype": "IPv6", "id":	
"563c5ca3-3hhd-4c96.	-9324-1ca9bc3aaef9"}	
tenant id	4b31d0508f83437e83d8f4d520cda22f	
relialit_tu	4DJ10030010343760300140320C00221	
1		
+		
+		
+		

## 8. Create icmp security group rule.

neutron security-g	roup-rule-create lb_secgroup1protocol	. icmp
+	+	· ·-+
Field	Value	
+	+	-+
direction	ingress	1
ethertype	IPv4	1
id	16d74150-a5b2-4cf6-82eb-a6c49a972d93	3
port_range_max		1
port_range_min		1
protocol	icmp	1
remote_group_id		

9. Create TCP port 22 rule.

10. Create TCP port 80 rule.

11. If you haven't already created a keypair, create one now with <a href="nova">nova</a> keypair-add. You'll use the keypair to boot images.

```
nova keypair-add lb kp1 > lb kp1.pem
chmod 400 lb_kp1.pem
cat lb_kp1.pem
----BEGIN RSA PRIVATE KEY-----
MIIEqAIBAAKCAQEAkbW5W/XWGRGC0LAJI7lttR7EdDfiTDeFJ7A9b9Cff+OMXjhx
WL26eKIr+jp8DR64YjV2mNnQLsDyCxekFpkyjnGRId3KVAeV5sRQgXqtaCXI+Rvd
IyUtd8p1cp3DRgTd1dx00oL6bBmwrZatNrrRn4HgKc2c7ErekeXrwLHyE0Pia/pz
C6qs0coRdfIeXxsmS3kXExP0YfsswRS/OyDl8QhRAF0ZW/zV+DQIi8+HpLZT+RW1
8sTTYZ6b0kXoH9wLER4IUBj1I11IyrYdxlAhe2VIn+tF0Ec4nDBn1py9iwEfGmn0+
N2jHCJAkrK/QhWdX0408zeXfL4mCZ9FybW4nzQIDAQABAoIBACe0PvgB+v8FuIGp
FjR32J8b7ShF+hI0pufzrCoFzRCKLruV4bzuphstBZK/0QG6Nz/7lX99Cq9SwCGp
pXrK7+3EoGl8CB/xmTUylVA4gRb6BNNsdkuXW9ZigrJirs0rkk8uIwRV0GsYbP5A
Kp7ZNTmjqDN75aC1ngRfhGgTlQU0dxBH+4xSb7GukekD13V8V5MF1Qft089asdWp
l/TpvhYeW9092xEnZ3qXQYpXYQgEFQoM2PKa3VW7FGLgfw9gdS/MSqpHuHGyKmjl
uT6upUX+Lofbe7V+9kfxuV32sLL/S5YFvkBy2q8VpuEV1sXI707Sc411WX4cqmlb
YoFwhrkCggCBALkYE70MTtdCAGcMotJhTiiS5l8d4U/fn1x0zus43XV5Y7wCnMuU
r5vCoK+a+TR9Ekzc/GjccAx7Wz/YYKp6G8FXW114dLcADXZjqjIlX7ifUud4sLCS
y+x3KAJa7LqyzH53I6F0ts9RaB5xx4gZ2WjcJquCTbATZWj7j1yGeNgvAoIAgQDJ
h0r0Te5IliYbCRq+ES9YRZzH/PSLuIn00bbLvp0PNEoKe2Pxs+KI8Fqp6ZIDAB3c
4EPOK5QrJvAny9Z58ZArrNZi15t84KEVAkWUATl+c4SmHc8sW/atgmUoqIzgDQXe
AlwadHLY7JCdg7EYDuUxuTKLLOdqfpf6fKkhNxtEwwKCAIAMxi+d5aIPUxvKA0I/
2L1XKYRCrkI9i/ZooBsjusH1+JG8iQWf0zy/aDhExlJKoBMiQ0IerpABHIZYqqtJ
OLIvrsK8ebK8aoGDWS+G1HN9v2kuVnMDTK5MPJEDUJkj7XEVjU1lNZSCTGD+MOYP
a5FInmEA1zZbX4tRKoNjZFh0uwKCAIEAiLs7drAdOLBu4C72fL4KIljwu5t7jATD
zRAwduIxmZq/lYcMU2RaEdEJonivsUt193NNbeeRWwnLLSUWupvT1l4pAt0ISNzb
TbbB4F5IVOwpls9ozc8DecubuM9K7YTIc02kkepqNZWjtMsx74HDrU3a5iSsSkvj
73Z/BeMupCMCggCAS48BsrcsDsHSHE3t04D8pAIr1r+6WPaQn49pT3GIrdQNc7a0
d9PfXmPoe/PxUlqaXoNAvT99+nNEadp+GTId21VM0Y28pn3EkIGE1Cqoeyl3BE08
f9SUiRNruDnH4F4OclsDBlmqWXImuXRfeiDHxM8X03UDZoqyHmGD3RqA53I=
----END RSA PRIVATE KEY-----
```

#### 12. Check and boot images.

```
+-----+
+-----+
```

#### Boot first VM.

```
nova boot --flavor 1 --image 04b1528b-b1e2-45d4-96d1-fbe04c6b2efd --key-name lb_kp1
--security-groups lb secgroup1 --nic net-id=71a1ac88-30a3-48a3-a18b-d98509fbef5c
lb_vm1 --poll
| Property
                            | Value
                              | MANUAL
| OS-DCF:diskConfig
| OS-EXT-AZ:availability_zone |
| OS-EXT-SRV-ATTR:host
                       | -
| OS-EXT-SRV-ATTR:hypervisor_hostname | -
| OS-EXT-STS:power_state
                       | 0
| OS-EXT-STS:task_state | scheduling
                       | building
| OS-EXT-STS:vm_state
| OS-SRV-USG:launched_at
| OS-SRV-USG:terminated_at
                              | | -
| accessIPv4
| accessIPv6
| adminPass
                               | NeVvhP5E8iCy
| config_drive
| created
                               | 2016-06-15T16:53:00Z
| flavor
                               | m1.tiny (1)
```

```
| hostId
| id
                                   | dfdfe15b-ce8d-469c-a9d8-2cea0e7ca287
                                   | cirros-0.3.3-x86_64 (04b1528b-
| image
b1e2-45d4-96d1-fbe04c6b2efd) |
                                   | lb_kp1
| key_name
| metadata
                                   | {}
| name
                                   | lb_vm1
| os-extended-volumes:volumes_attached | []
                                   | 0
| progress
| security_groups
                                   | lb_secgroup1
                                   | BUILD
| status
| tenant_id
                                   | 4b31d0508f83437e83d8f4d520cda22f
| updated
                                   | 2016-06-15T16:53:00Z
| user_id
                                   | fd471475faa84680b97f18e55847ec0a
+----+
          Server building... 100% complete
          Finished
```

#### Boot second VM.

```
| OS-EXT-SRV-ATTR:host
| OS-EXT-SRV-ATTR:hypervisor_hostname | -
| OS-EXT-SRV-ATTR:instance_name | instance-00000034
| OS-EXT-STS:power_state
                                   | 0
| OS-EXT-STS:task_state
                                   | scheduling
| OS-EXT-STS:vm_state
                           | building
| OS-SRV-USG:launched_at
| OS-SRV-USG:terminated_at | -
| accessIPv4
| accessIPv6
| adminPass
                                     | 3nFXjNrTrmNm
| config_drive
| created
                                     | 2016-06-15T16:55:10Z
| flavor
                                     | m1.tiny (1)
| hostId
                                     | 3844bb10-2c61-4327-a0d4-0c043c674344
| id
                                     | cirros-0.3.3-x86_64 (04b1528b-
| image
b1e2-45d4-96d1-fbe04c6b2efd) |
                                     | lb_kp1
| key_name
                                     | {}
| metadata
| name
                                     | lb_vm2
| os-extended-volumes:volumes_attached | []
| progress
                                     | 0
| security_groups
                                     | lb_secgroup1
```

#### 13. List the running VM with nova list

#### 14. Check ports.

15. Create the first floating IP.

```
neutron floatingip-create ext-net --port-id 7e5e0038-88cf-4f97-a366-b58cd836450e
| Field
                | Value
+-----+
| fixed_ip_address | 10.247.94.132
| floating_ip_address | 10.247.96.26
| floating_network_id | d3cb12a6-a000-4e3e-82c4-ee04aa169291 |
               | 3ce608bf-8835-4638-871d-0efe8ebf55ef |
| port id
               | 7e5e0038-88cf-4f97-a366-b58cd836450e |
              | 6aafc9a9-93f6-4d7e-94f2-3068b034b823 |
| router_id
| status
               | DOWN
| tenant_id
                | 4b31d0508f83437e83d8f4d520cda22f
+-----
```

16. Create the second floating IP.

```
neutron floatingip-create ext-net --port-id ca95cc24-4e8f-4415-9156-7b519eb36854
+----+
          | Value
| Field
+-----+
| fixed_ip_address | 10.247.94.133
| floating_ip_address | 10.247.96.27
| floating_network_id | d3cb12a6-a000-4e3e-82c4-ee04aa169291 |
| id
               | 680c0375-a179-47cb-a8c5-02b836247444 |
| port_id
               | ca95cc24-4e8f-4415-9156-7b519eb36854 |
| router_id
              | 6aafc9a9-93f6-4d7e-94f2-3068b034b823 |
| status
               | DOWN
| tenant_id
                | 4b31d0508f83437e83d8f4d520cda22f
```

17. List the floating IP's.

#### 18. Show first Floating IP.

```
neutron floatingip-show 3ce608bf-8835-4638-871d-0efe8ebf55ef
+-----
l Field
                | Value
+-----+
| fixed_ip_address | 10.247.94.132
| floating_ip_address | 10.247.96.26
| floating_network_id | d3cb12a6-a000-4e3e-82c4-ee04aa169291 |
              | 3ce608bf-8835-4638-871d-0efe8ebf55ef |
| 7e5e0038-88cf-4f97-a366-b58cd836450e |
| id
| port_id
              | 6aafc9a9-93f6-4d7e-94f2-3068b034b823 |
| router_id
                | ACTIVE
| status
               | 4b31d0508f83437e83d8f4d520cda22f
| tenant_id
```

#### 19. Show second Floating IP.

```
neutron floatingip-show 680c0375-a179-47cb-a8c5-02b836247444
+----+
| Field
              | Value
+-----+
| fixed_ip_address | 10.247.94.133
| floating_ip_address | 10.247.96.27
| floating_network_id | d3cb12a6-a000-4e3e-82c4-ee04aa169291 |
     | 680c0375-a179-47cb-a8c5-02b836247444 |
| id
| port_id
             | ca95cc24-4e8f-4415-9156-7b519eb36854 |
| router_id
            | 6aafc9a9-93f6-4d7e-94f2-3068b034b823 |
| status
             | ACTIVE
```

#### 20. Ping first Floating IP.

```
ping -c 1 10.247.96.26
PING 10.247.96.26 (10.247.96.26) 56(84) bytes of data.
64 bytes from 10.247.96.26: icmp_seq=1 ttl=62 time=3.50 ms

--- 10.247.96.26 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
```

```
rtt min/avg/max/mdev = 3.505/3.505/3.505/0.000 ms
```

21. Ping second Floating IP.

```
ping -c 1 10.247.96.27
PING 10.247.96.27 (10.247.96.27) 56(84) bytes of data.
64 bytes from 10.247.96.27: icmp_seq=1 ttl=62 time=3.47 ms
--- 10.247.96.27 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 3.473/3.473/0.000 ms
```

22. Listing the VM's will give you both the fixed and floating IP's for each virtual machine.

23. List Floating IP's.

## 27.4 Create Load Balancers

The following steps will setup new Octavia Load Balancers.



#### Note

The following examples assume names and values from the previous section.

1. Create load balancer for subnet

2. List load balancers. You will need to wait until the load balancer <a href="provisioning\_status">provisioning\_status</a> is ACTIVE before proceeding to the next step.

3. Once the load balancer is created, create the listener. This may take some time.

```
neutron lbaas-listener-create --loadbalancer lb1 --protocol HTTP --protocol-port=80
--name lb1 listener
+-----+
| Field
               | Value
+-----+
| default_pool_id
| default_tls_container_ref |
| description
name
              | lb1_listener
| protocol
              | HTTP
| protocol
| protocol_port
               | 80
| sni_container_refs
| tenant_id
              | 4b31d0508f83437e83d8f4d520cda22f
```

4. Create the load balancing pool. During the creation of the load balancing pool, the status for the load balancer goes to <a href="PENDING\_UPDATE">PENDING\_UPDATE</a>. Use <a href="neutron lbaas-loadbalancer-list">neutron lbaas-loadbalancer-list</a> to watch for the change to <a href="ACTIVE">ACTIVE</a>. Once the load balancer returns to <a href="ACTIVE">ACTIVE</a>, proceed with the next step.

```
neutron lbaas-pool-create --lb-algorithm ROUND ROBIN --listener lb1 listener --
protocol HTTP --name lb1_pool
+-----
| Field
        | Value
+-----
| healthmonitor_id |
members
      | lb1_pool
name
| protocol | HTTP
| session persistence |
+-----
```

5. Create first member of the load balancing pool.

6. Create the second member.

7. You should check to make sure the load balancer is active and check the pool members.

subnet id		. proceed _port	weight	-
			+	_
61da1e21-e0ae-4158-935a-c909a81470e1 6fc2572c-53b3-41d0-ab63-342d9515f514	10.247.94.132	·	1	I
459c7f21-46f7-49e8-9d10-dc7da09f8d5a 6fc2572c-53b3-41d0-ab63-342d9515f514	•	80 	1	I

8. You can view the details of the load balancer, listener and pool.

Field	Value	I
+   admin_state_up	+    True	+ 
description	I	I
id	3d9170a1-8605-43e6-9255-e14a8b4aae53	I
listeners	{"id": "c723b5c8-e2df-48d5-a54c-fc240ac7b539"}	I
name	lb1	I
operating_status	ONLINE	I
provider	octavia	I
provisioning_status	ACTIVE	I
tenant_id	4b31d0508f83437e83d8f4d520cda22f	I
vip_address	10.247.94.134	I
vip_port_id	da28aed3-0eb4-4139-afcf-2d8fd3fc3c51	
vip_subnet_id	6fc2572c-53b3-41d0-ab63-342d9515f514	
+   id	list 	   name
+	a54c-fc240ac7b539   0f5951ee-c2a0-4e62-ae44-e1491a	·
	++	

## 27.5 Create Floating IPs for Load Balancer

To create the floating IP's for the load balancer, you will need to list the current ports to get the load balancer id. Once you have the id, you can then create the floating IP.

1. List the current ports.

```
neutron port-list
| id
| mac_address | fixed_ips
| 7e5e0038-88cf-4f97-a366-b58cd836450e |
| fa:16:3e:66:fd:2e | {"subnet_id": "6fc2572c-
                                     | 53b3-41d0-ab63-342d9515f514",
| "ip_address": "10.247.94.132"}
| a3d0b0fe-e7ff-4b00-a033-44e833b55efe |
| fa:16:3e:91:a2:5b | {"subnet_id": "f00299f8-3403-45ae-ac4b-
| 58af41d57bdc", "ip_address":
| "10.247.94.142"}
| ca95cc24-4e8f-4415-9156-7b519eb36854 |
| fa:16:3e:e0:37:c4 | {"subnet_id": "6fc2572c-
                    | 53b3-41d0-ab63-342d9515f514",
| "ip_address": "10.247.94.133"}
| da28aed3-0eb4-4139-afcf-2d8fd3fc3c51 | loadbalancer-
| fa:16:3e:1d:a2:1c | {"subnet_id": "6fc2572c-
                                      | 3d9170a1-8605-43e6-9255-e14a8b4aae53
ı
                    | 53b3-41d0-ab63-342d9515f514",
| "ip_address": "10.247.94.134"}
```

2. Create the floating IP for the load balancer.

# 27.6 Testing the Octavia Load Balancer

To test the load balancers, create the following web server script so you can run it on each virtual machine. You will use curl <ip address> to test if the load balance services are responding properly.

1. Start running web servers on both of the virtual machines. Create the webserver.sh script with below contents. In this example, the port is 80.

```
vi webserver.sh

#!/bin/bash

MYIP=$(/sbin/ifconfig eth0|grep 'inet addr'|awk -F: '{print $2}'| awk '{print $1}');
while true; do
    echo -e "HTTP/1.0 200 OK

Welcome to $MYIP" | sudo nc -l -p 80
done
```

2. Deploy the web server and run it on the first virtual machine.

```
ssh-keygen -R 10.247.96.26
```

```
/home/stack/.ssh/known_hosts updated.
Original contents retained as /home/stack/.ssh/known_hosts.old

scp -o StrictHostKeyChecking=no -i lb_kp1.pem webserver.sh cirros@10.247.96.26:
webserver.sh 100% 263 0.3KB/s 00:00

ssh -o StrictHostKeyChecking=no -i lb_kp1.pem cirros@10.247.96.26 'chmod +x ./
webserver.sh'
ssh -o StrictHostKeyChecking=no -i lb_kp1.pem cirros@10.247.96.26 ./webserver.sh
```

3. Test the first web server.

```
curl 10.247.96.26
Welcome to 10.247.94.132
```

4. Deploy and start the web server on the second virtual machine like you did in the previous steps. Once the second web server is running, list the floating IP's.

5. Display the floating IP for the load balancer.

status	ACTIVE	
tenant_id	4b31d0508f83437e83d8f4d520cda22f	1
+	-+	+

## 6. Finally, test the load balancing.

curl 10.247.96.28
Welcome to 10.247.94.132

curl 10.247.96.28
Welcome to 10.247.94.133

curl 10.247.96.28
Welcome to 10.247.94.132

curl 10.247.96.28
Welcome to 10.247.94.133

curl 10.247.96.28
Welcome to 10.247.94.132

curl 10.247.96.28
Welcome to 10.247.94.132

## 28 Other Common Post-Installation Tasks

## 28.1 Determining Your User Credentials

On your lifecycle manager, in the <u>~/scratch/ansible/next/hos/ansible/group\_vars/</u> directory you will find several files. In the one labeled as first control plane node you can locate the user credentials for both the Administrator user (<u>admin</u>) and your Demo user (<u>demo</u>) which you will use to perform many other actions on your cloud.

For example, if you are using the Entry-scale KVM with VSA model and used the default naming scheme given in the example configuration files, you can use these commands on your lifecycle manager to GREP for your user credentials:

#### Administrator

```
grep keystone_admin_pwd entry-scale-kvm-vsa-control-plane-1
```

#### Demo

```
grep keystone_demo_pwd entry-scale-kvm-vsa-control-plane-1
```

# 28.2 Configure your lifecycle manager to use the command-line tools

This playbook will do a series of steps to update your environment variables for your cloud so you can use command-line clients.

Run the following command, which will replace /etc/hosts on the lifecycle manager:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts cloud-client-setup.yml
```

As the /etc/hosts file no longer has entries for Helion lifecycle management, sudo commands may become a bit slower. To fix this issue, once this step is complete, add "hlm" after "127.0.0.1 localhost". The result will look like this:

```
# Localhost Information
```

## 28.3 Protect home directory

The home directory of the user that owns the HPE Helion OpenStack 5.0 scripts should not be world readable. Change the permissions so that they are only readable by the owner:

chmod 0700  $\sim$ 

## 28.4 Back up Your SSH Keys

As part of the cloud deployment setup process, SSH keys to access the systems are generated and stored in ~/.ssh on your lifecycle manager.

These SSH keys allow access to the subsequently deployed systems and should be included in the list of content to be archived in any backup strategy.

# 28.5 Retrieving Service Endpoints

- 1. Log in to your lifecycle manager.
- 2. Source the keystone admin credentials:

```
unset OS_TENANT_NAME
source ~/keystone.osrc
```

3. Using the OpenStack command-line tool you can then query the Keystone service for your endpoints:

openstack endpoint list



Tip

You can use openstack -h to access the client help file and a full list of commands.

To learn more about Keystone, see .

449 Protect home directory

# 28.6 Other Common Post-Installation Tasks

Here are the links to other common post-installation tasks that either the Administrator or Demo users can perform:

- •
- Section 23.4, "Creating an External Network"
- •
- •
- •